

# Global Illumination II



70001 – Advanced Computer Graphics: Photographic Image Synthesis

Abhijeet Ghosh

Lecture 15, Feb. 23<sup>rd</sup> 2024

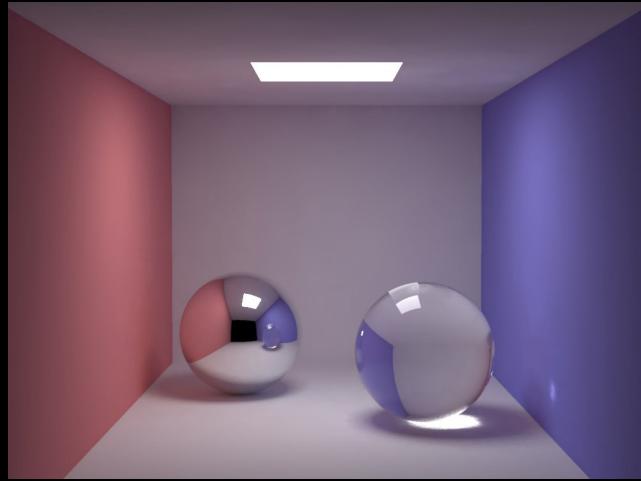
1

## Biased approaches

- Reuse previously computed results
  - Assuming slow change in radiance
- Efficient computation of smooth noise-free images
  - Price: unpredictable behavior!
- Irradiance & radiance caching
- Photon mapping

2

## Global Illumination



Henrik Wann Jensen

3

## Indirect Irradiance



Henrik Wann Jensen

4

2

## Irradiance caching – Ward et al. 98

- Indirect lighting changes **slowly** compared to direct lighting
  - especially for diffuse surfaces
- Compute at sparse points and interpolate!
- Two issues:
  - How to store an irradiance estimate?
  - When to sample new estimate?

5

## Irradiance caching

- Irradiance
$$E(x, n) = \int_{\Omega} L_i(x, \omega_i) \cos\theta_i d\omega_i.$$
- Radiance
$$L_r(x, \omega_r) = \int_{\Omega} f_r(x, \omega_r, \omega_i) L_i(x, \omega_i) \cos\theta_i d\omega_i.$$
- Diffuse surface: BRDF constant  $c$ 
$$L_r(x, \omega_r) = c \int_{\Omega} L_i(x, \omega_i) \cos\theta_i d\omega_i = c \cdot E(x, n)$$

6

## Irradiance caching

- First compute  $E(x,n)$  with MC path tracing

$$E(x,n) = \frac{1}{N} \sum_j \frac{L_i(x, \omega_j) \cos \theta_j}{p(\omega_j)}$$

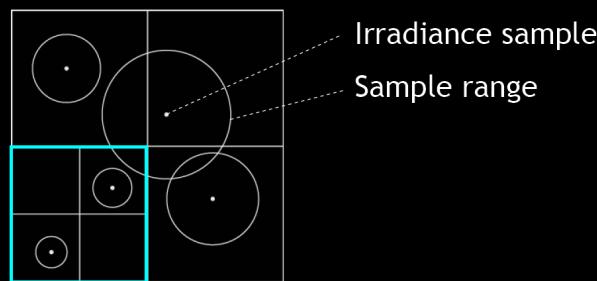
- Sample from cosine weighted distribution  $p(\omega_j) = \cos \theta / \pi$

$$E(x,n) = \frac{\pi}{N} \sum_j L_i(x, \omega_j)$$

- Irradiance estimates stored in an octree data-structure

7

## Octree example



- Adaptive subdivision such that each cell contains < 3 samples

8

## Irradiance sampling

- Each sample is assigned a range for interpolation
- Rate of change of irradiance decides interpolation distance
  - inversely proportional
- Depends on distance to visible surfaces
- Harmonic Mean 
$$\frac{N}{\sum_{i=1}^N 1/d_i}$$
  - $d_i$  is the distance to first intersection along path  $i$

9

## Irradiance sample

```
struct irradiance_sample {  
    vector3 E // irradiance  
    vector3 n // normal  
    vector3 p // position  
    float r // range  
}
```

10

## Irradiance interpolation

- Compute the error in estimate when using a cached sample

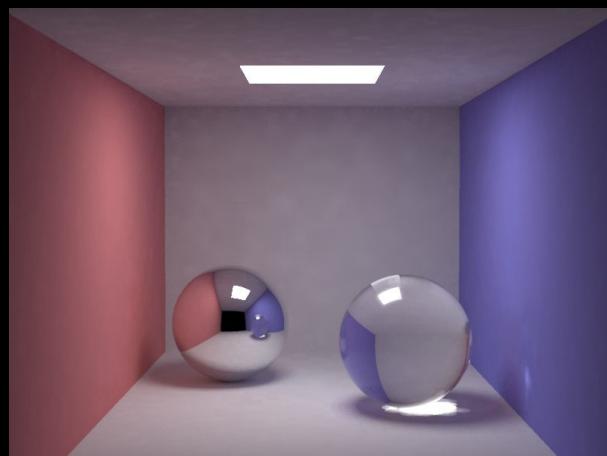
$$w_j(\mathbf{x}) = \frac{1}{\frac{\|\mathbf{x}-\mathbf{x}_j\|}{r_j} + \sqrt{1 - \mathbf{n}(\mathbf{x}) \cdot \mathbf{n}(\mathbf{x}_j)}} \approx \epsilon_j(\mathbf{x})$$

- If error is acceptable, use a weighted sum of interpolated irradiance

$$E(\mathbf{x}) = \frac{\sum_i w_i(\mathbf{x}) E(\mathbf{x}_i)}{\sum_i w_i(\mathbf{x})}$$

11

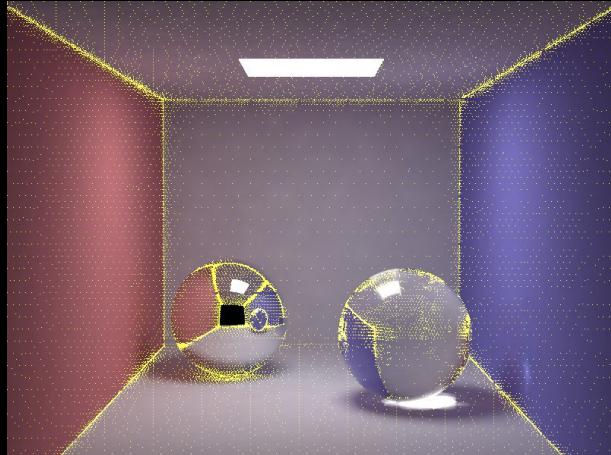
## Irradiance gradients



1000 sample rays,  $w > 10$

12

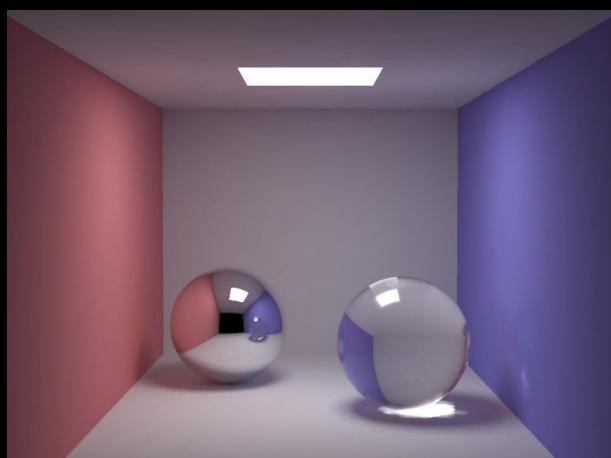
## Irradiance cache positions



1000 sample rays,  $w > 10$

13

## Irradiance gradients



5000 sample rays,  $w > 10$

14

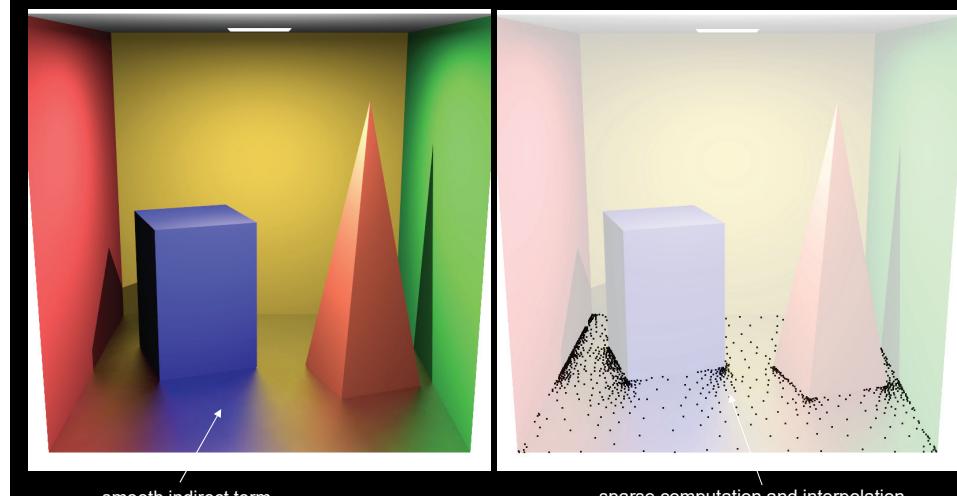
## Glossy surfaces?

- Irradiance caching works well for diffuse surfaces
- For glossy surfaces we need to account for the directional dependence of the BRDF!

$$\begin{aligned} L_r(x, \omega_r) &\approx (\int_{\Omega} f_r(x, \omega_r, \omega_i) d\omega_i) (\int_{\Omega} L_i(x, \omega_i) \cos\theta_i d\omega_i) \\ &\approx \rho_{hd}(\omega_r) \cdot E(x, n) \end{aligned}$$

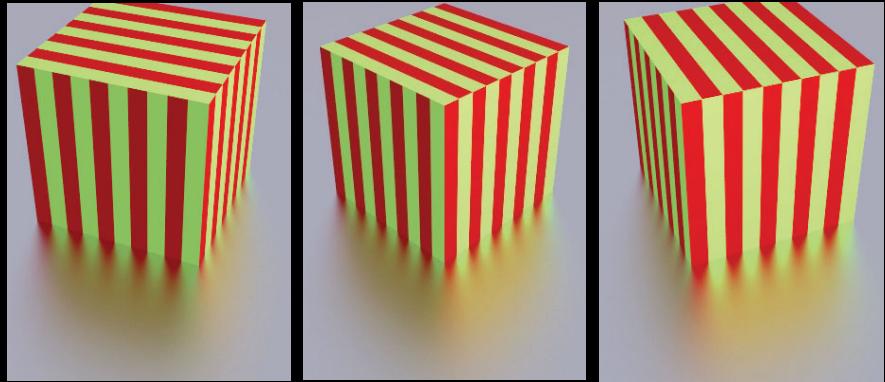
15

## Radiance caching – Křivánek et al. 05



16

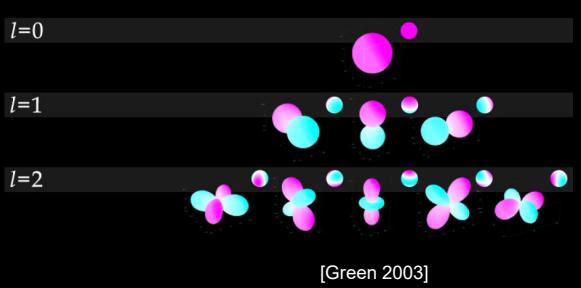
## View dependence



17

## Incoming radiance representation

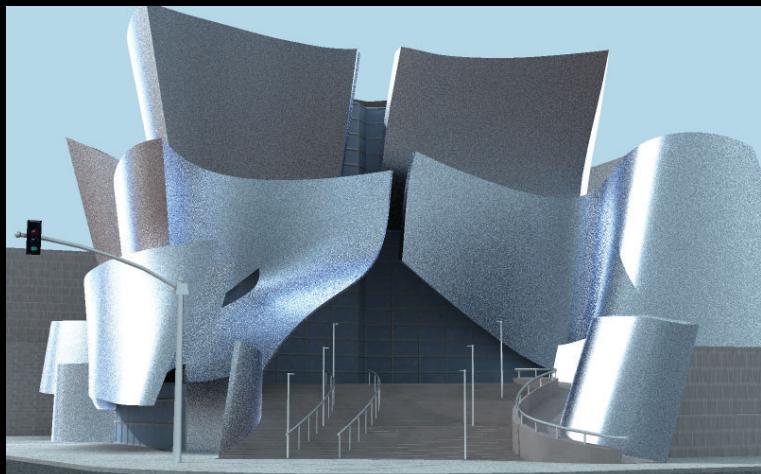
- Spherical Harmonics
  - basis functions on a sphere
- Approximate  $\rho_{hd}(\omega_r)$  in the SH basis
  - 3D rotations of SH coefficients for change in viewpoint
- Dot with SH irradiance coefficient



[Green 2003]

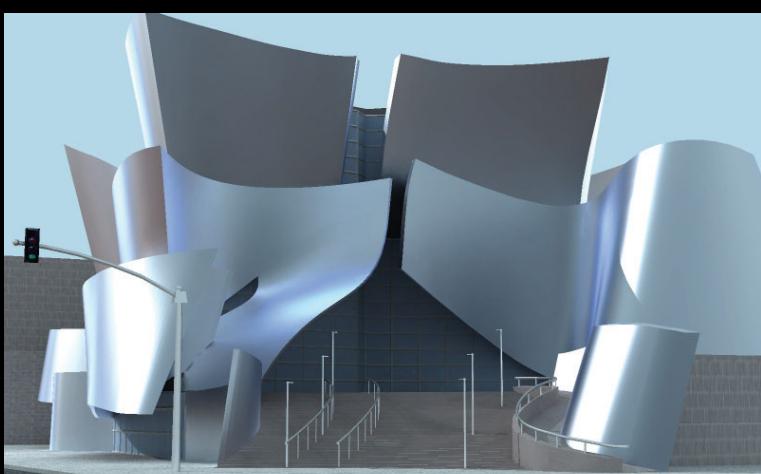
18

## MC Importance Sampling



19

## Radiance caching



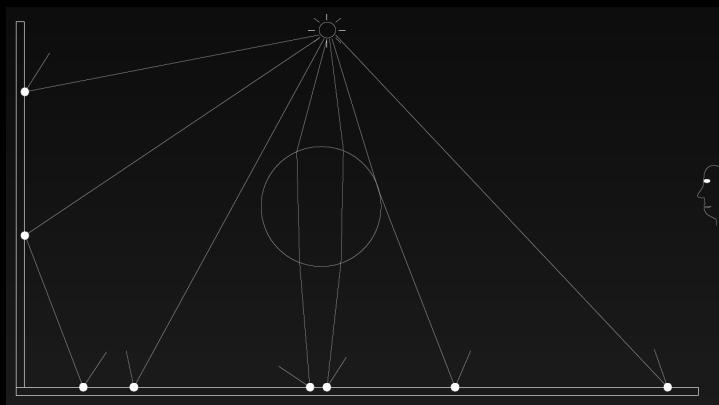
20

## Indirect specular?

- Caustics are still hard to simulate with radiance caching ☺
- Solution: **Photon Mapping** [Jensen 2001]
- 2 pass algorithm
  - trace photon into the scene from light source
  - gather photon density from camera

21

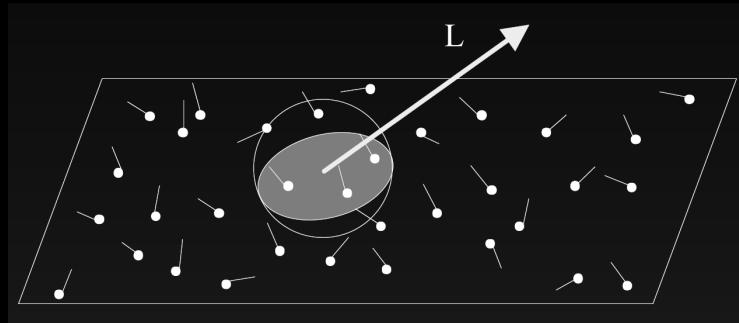
## Photon tracing



Henrik Wann Jensen

22

## Photon density



Henrik Wann Jensen

23

## Radiance estimate

$$\begin{aligned} L(x, \vec{\omega}) &= \int_{\Omega} f_r(x, \vec{\omega}', \vec{\omega}) L'(x, \vec{\omega}') \cos \theta' d\omega \\ &= \int_{\Omega} f_r(x, \vec{\omega}', \vec{\omega}) \frac{d\Phi^2(x, \vec{\omega}')}{\cos \theta' d\omega dA} \cos \theta' d\omega \\ &= \int_{\Omega} f_r(x, \vec{\omega}', \vec{\omega}) \frac{d\Phi(x, \vec{\omega}')}{dA} \\ &\approx \sum_{p=1}^n f_r(x, \vec{\omega}'_p, \vec{\omega}) \frac{\Delta\Phi_p(x, \vec{\omega}'_p)}{\pi r^2} \end{aligned}$$

24

12

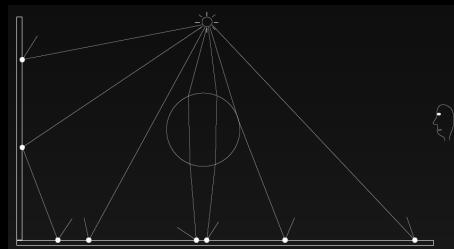
## Photon map data-structure

- Balanced kd-tree

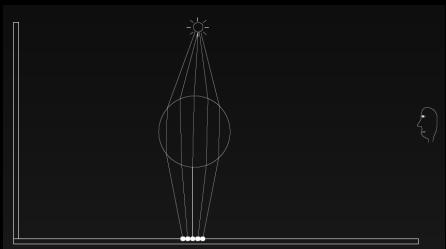
```
struct photon = {  
    float position[3];  
    rgbe power;          // power packed as 4 bytes  
    char phi, theta;    // incoming direction  
    short flags;  
}
```

25

## Two photon maps



Global photon map

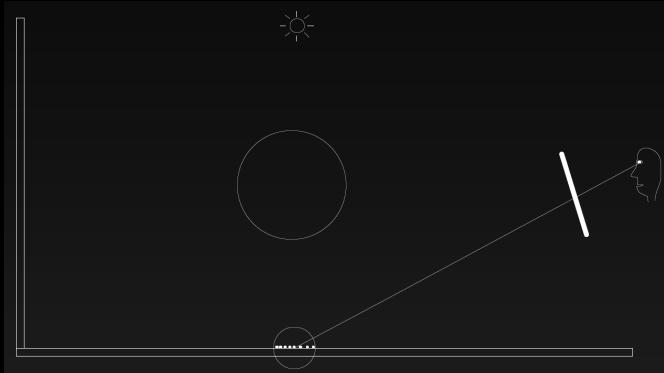


Caustic photon map

26

13

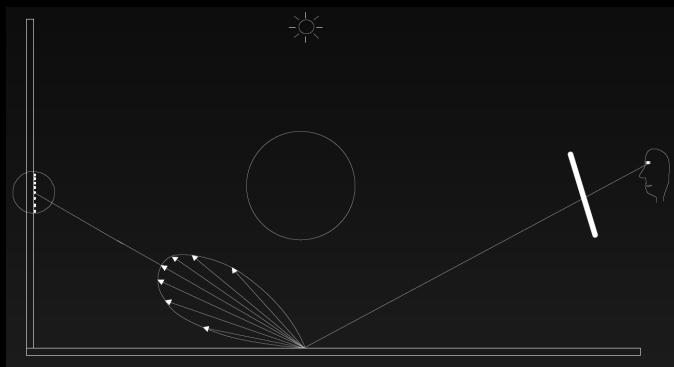
## Rendering caustics



1<sup>st</sup> diffuse bounce

27

## Rendering indirect illumination

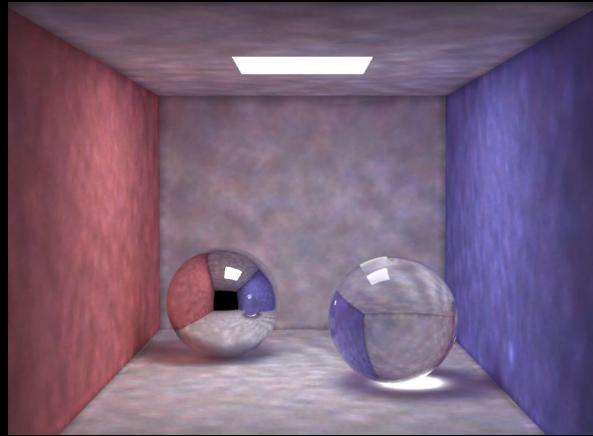


2<sup>nd</sup> diffuse bounce

- Also known as final gather!

28

## Photon map visualization



Henrik Wann Jensen

100,000 photons, 50 photons for density estimate  
(1<sup>st</sup> bounce density evaluation)

29

## Photon map visualization



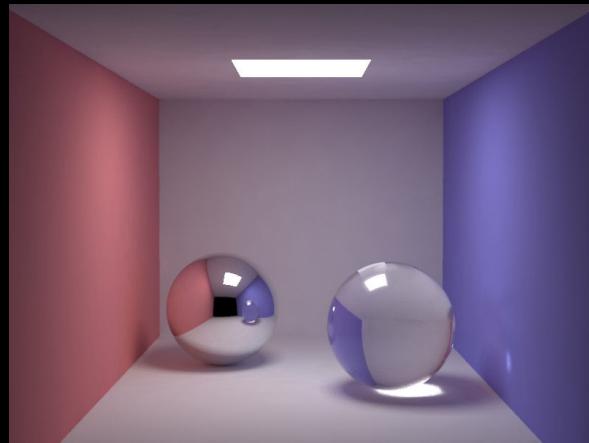
Henrik Wann Jensen

100,000 photons, 50 photons for density estimate

30

15

## Rendering with final gather



Henrik Wann Jensen

- 2<sup>nd</sup> bounce density evaluation for indirect illumination

31

## Photon Mapping - Cognac glass



HENRIK WANN JENSEN 1995

32

## Photon Mapping – Mies house



Henrik Wann Jensen

33

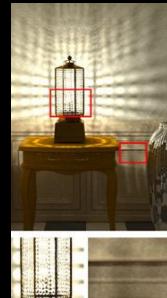
## Photon Mapping – subsurface scattering



34

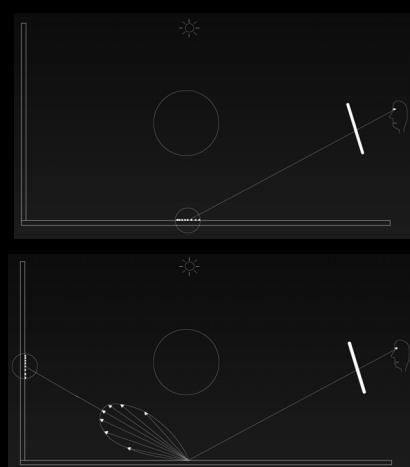
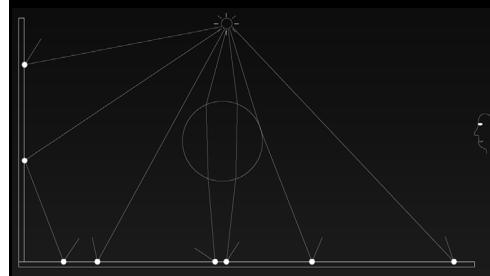
## Caustics dominated scene

- Photon mapping still too expensive
  - too much memory consumption
- Solution: **Progressive Photon Mapping** [Hachisuka et al. 08]
  - resolve SDS paths with arbitrary accuracy with finite memory!
- Multiple photon tracing steps for incremental estimate
  - no storage of previously traced photons



35

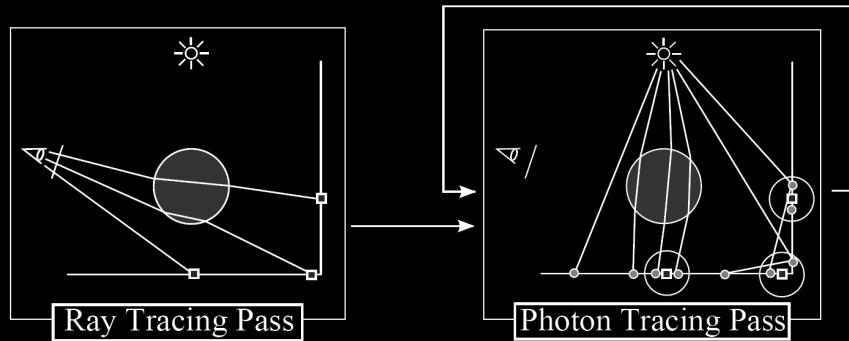
## Regular Photon Mapping



- 1<sup>st</sup> pass photon tracing from light source to diffuse surfaces
- 2<sup>nd</sup> pass of ray tracing from camera to scene hit points for evaluation of photon density estimate

36

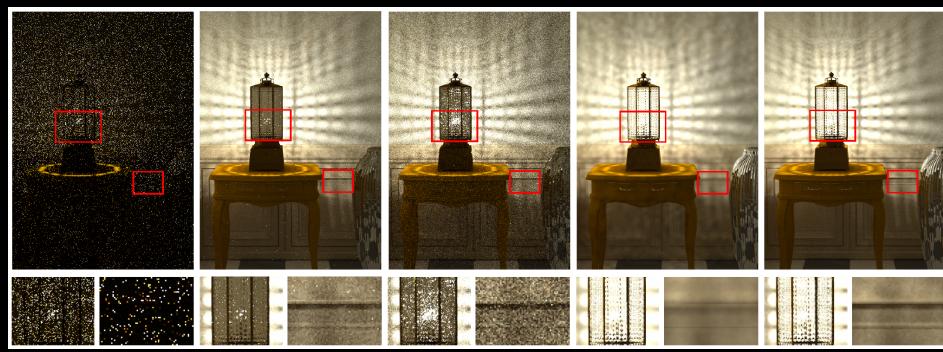
## Progressive Photon Mapping



- 1<sup>st</sup> pass ray trace from camera to scene hit points
- 2<sup>nd</sup> pass (multiple) of photon tracing from light source to evaluate photon density at camera hit points. No need to store older photons!  
Update of density estimate after each pass.

37

## Progressive photon mapping



[Hachisuka et al. 08]

38

19

## Photon density estimation

- Source of bias:

$$L(x, \vec{\omega}) \approx \sum_{p=1}^n \frac{f_r(x, \vec{\omega}, \vec{\omega}_p) \phi_p(x_p, \vec{\omega}_p)}{\pi r^2},$$

- radiance estimate converges with increase in density
- infinite density in the limit

- Trace N photons, estimate with  $N^\beta$  photons,  $\beta$  in  $[0, 1]$ :

$$L(x, \vec{\omega}) = \lim_{N \rightarrow \infty} \sum_{p=1}^{\lfloor N^\beta \rfloor} \frac{f_r(x, \vec{\omega}, \vec{\omega}_p) \phi_p(x_p, \vec{\omega}_p)}{\pi r^2},$$

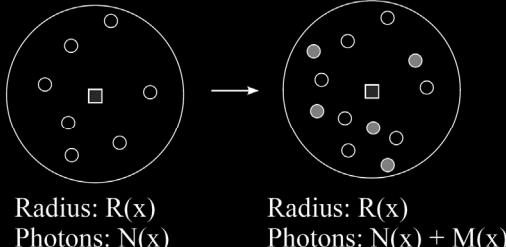
39

## Progressive density estimate

- Density  $d(x) = n/\pi r^2$
- If radius  $r$  remains constant, then density estimate is averaged
  - suboptimal
- Key idea is to progressively **reduce**  $r$  with each photon tracing pass
  - while increasing photon accumulation!

40

## Radius reduction

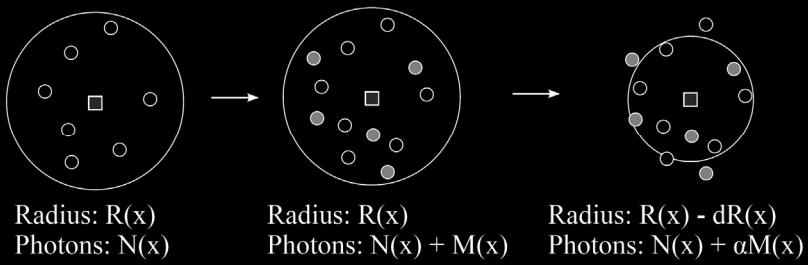


- After second round of photon tracing, local density with constant radius  $R(x)$

$$d'(x) = (N(x) + M(x)) / (\pi * R(x)^2)$$

41

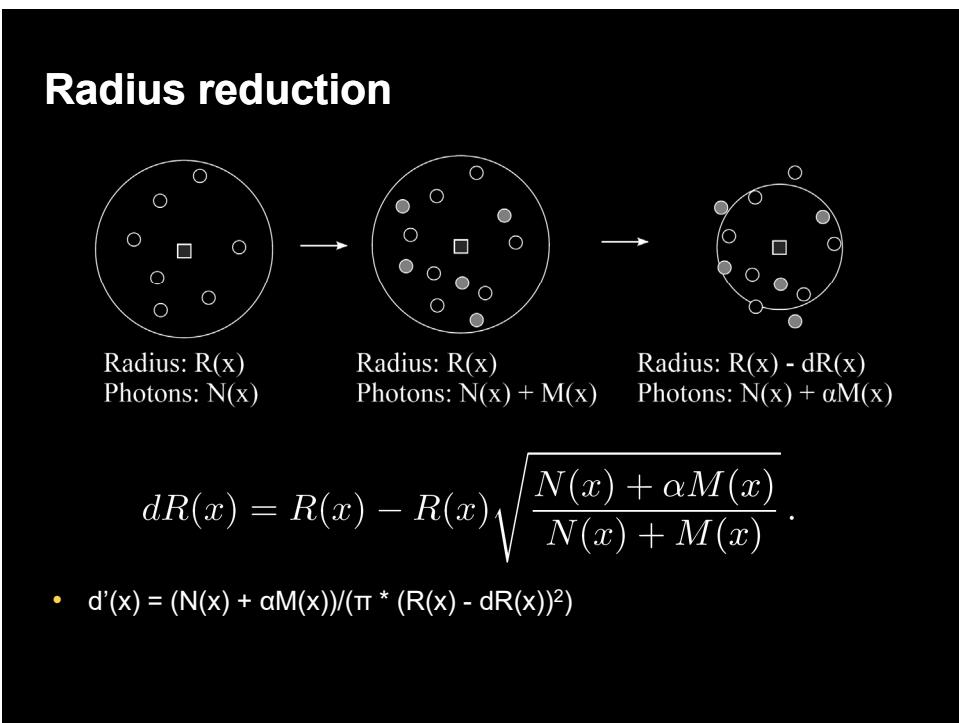
## Radius reduction



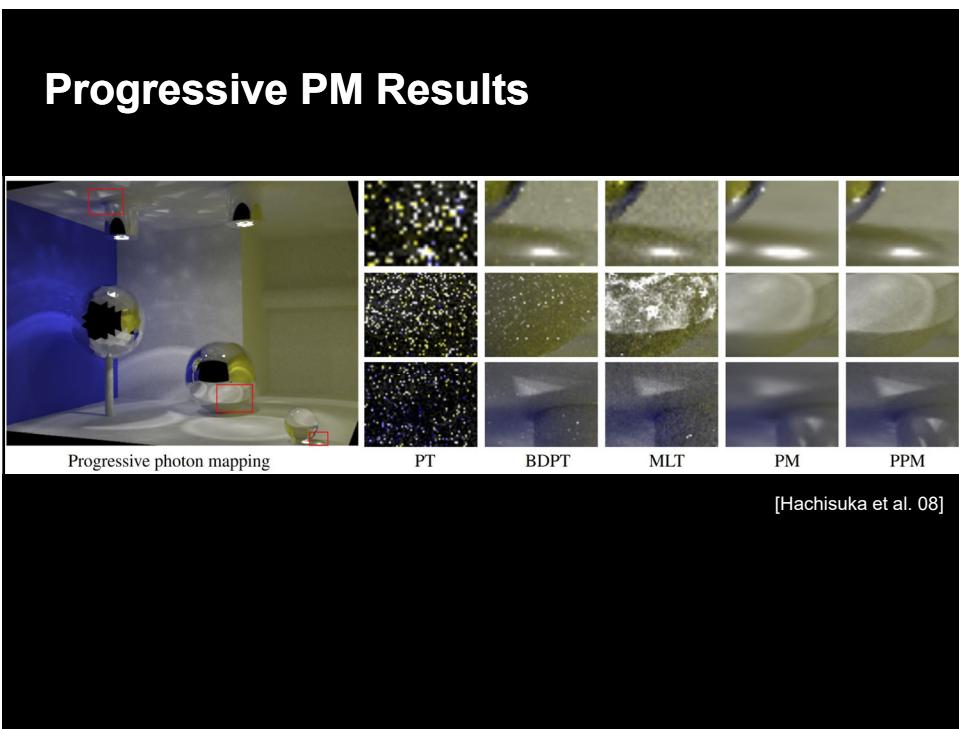
$$dR(x) = R(x) - R(x) \sqrt{\frac{N(x) + \alpha M(x)}{N(x) + M(x)}}.$$

- $\alpha \rightarrow [0, 1]$

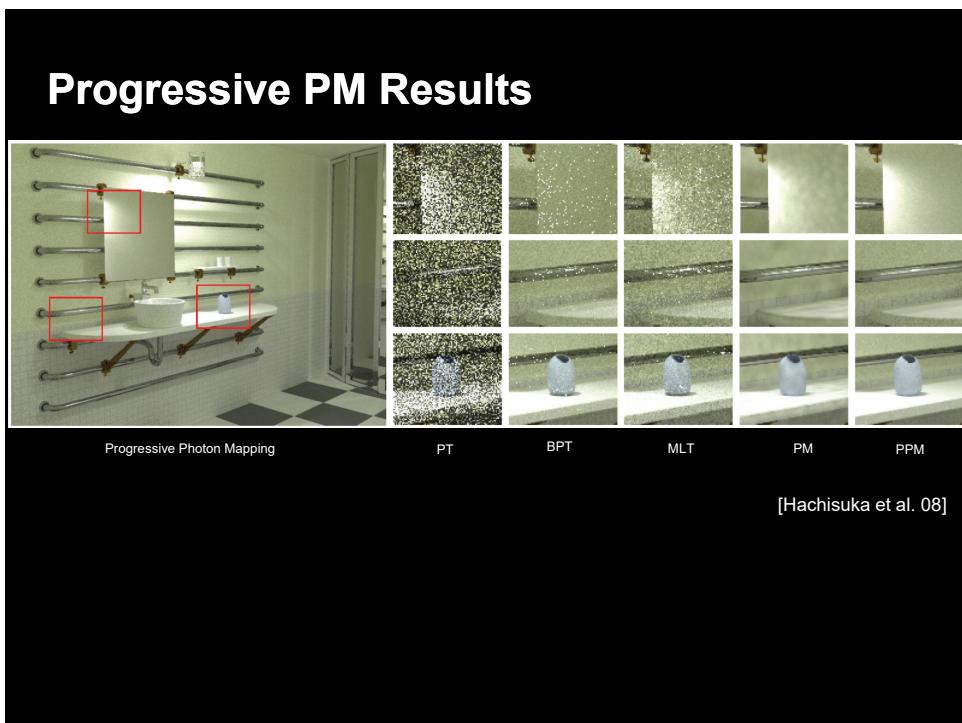
42



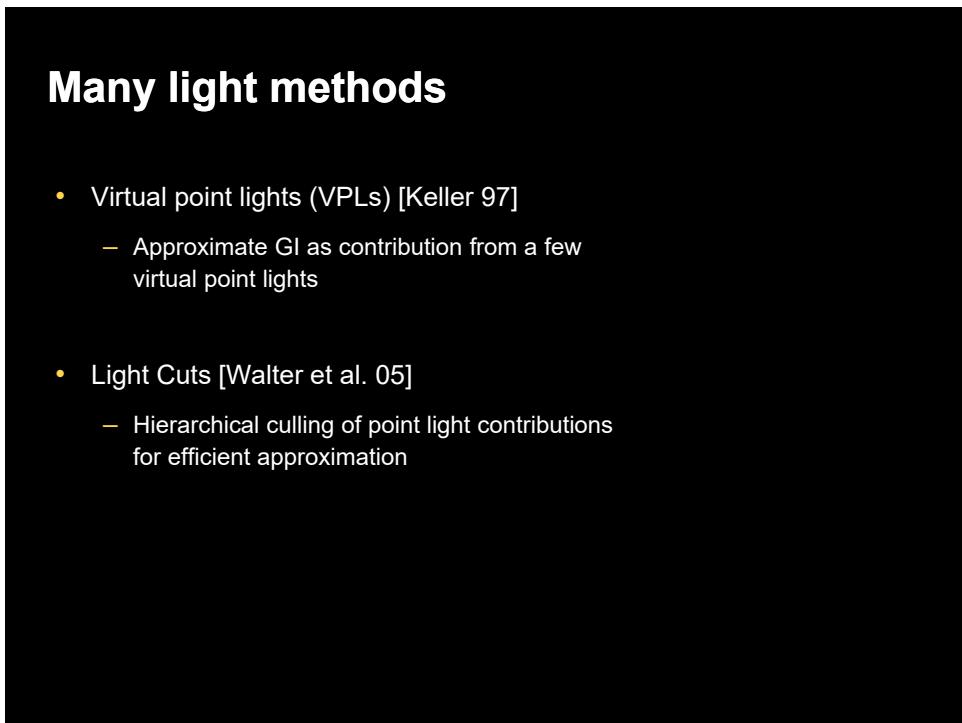
43



44



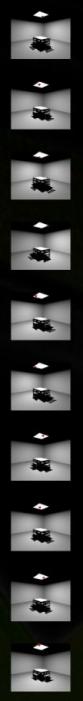
45



46

## Instant Radiosity with VPLs [Keller 97]

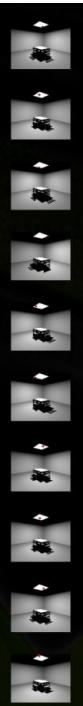
- Approximate GI as contribution from a few virtual point lights
- Hardware (OpenGL) implementation using the accumulation buffer!
  - Soft shadowing computed as average of many images rendered with a point light placed at different positions on the area light



47

## Instant Radiosity with VPLs [Keller 97]

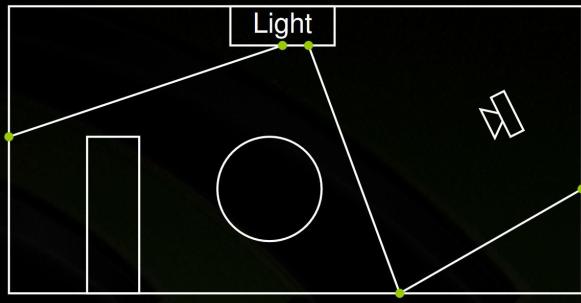
- Approximate GI as contribution from a few virtual point lights
- Hardware (OpenGL) implementation using the accumulation buffer!
  - Soft shadowing computed as average of many images rendered with a point light placed at different positions on the area light



48

## Instant Radiosity with VPLs [Keller 97]

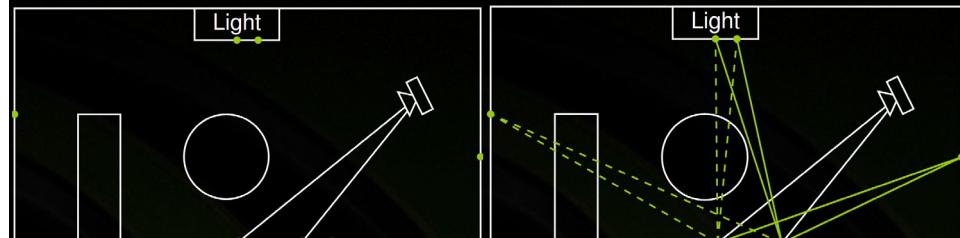
- Light path vertices stored as point lights
  - similar to a small photon map!



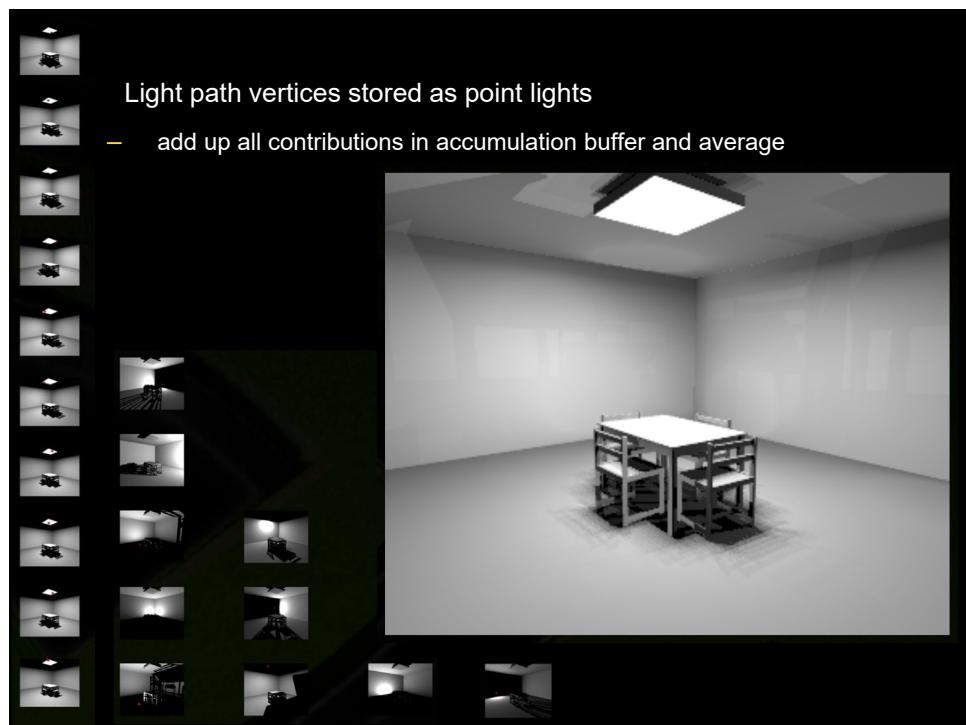
49

## Instant Radiosity with VPLs [Keller 97]

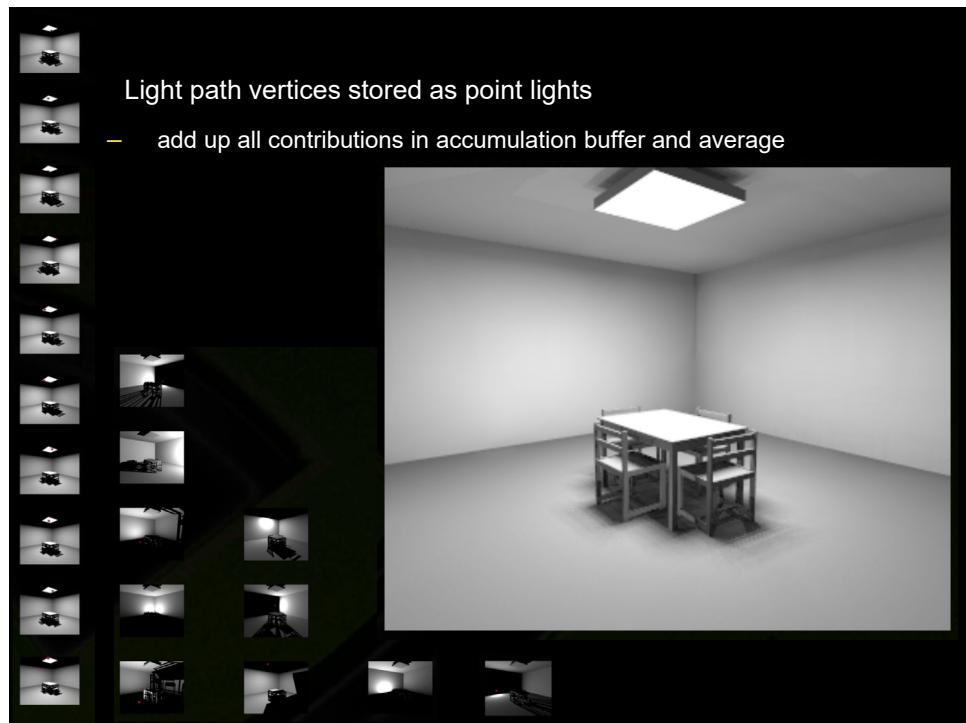
- Light path vertices stored as point lights
  - trace camera rays and then connect to VPLs with shadow rays



50



51



52

## Scalability with many lights – Lightcuts

[Walter et al. 05]

- Simulate complex illumination using point lights
  - Area lights
  - HDR environment maps
  - Sun & sky light
  - Indirect illumination
- More lights → more accurate
  - And more expensive
  - Naive cost linear in lights



Area lights + Sun/sky + Indirect

53

## Accurate approximation with few lights

[Walter et al. 05]

- Do we really need to evaluate all lights?
  - Thousands to millions of lights
  - Average contribution minuscule
  - Below human perceptual limits
    - Weber's Law (roughly 2%)
- Evaluate a small subset
  - Chosen adaptively
  - Perceptually accurate approximation

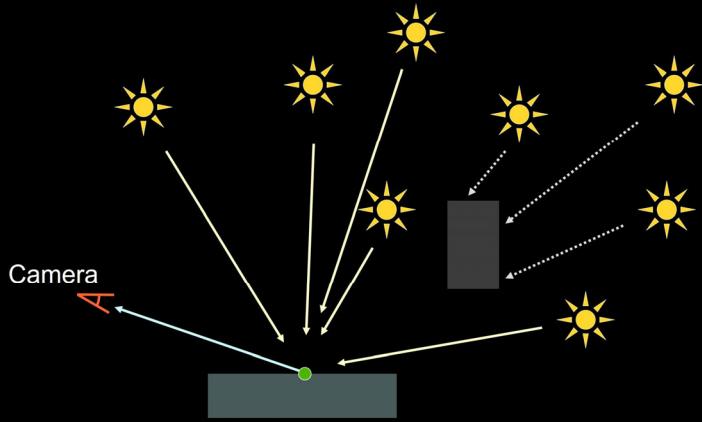


Environment lighting + Indirect

54

## Lightcuts problem

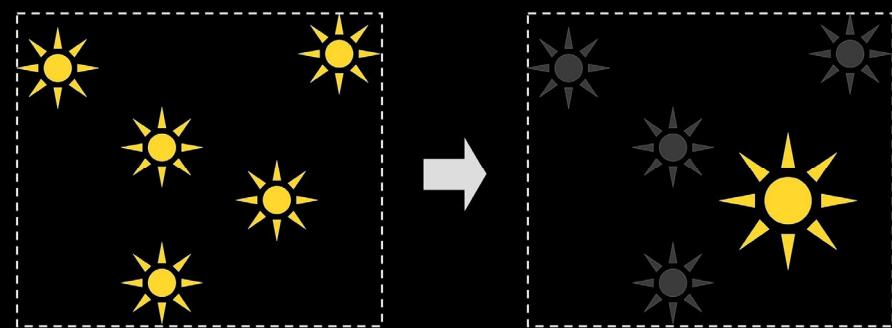
- All lights do not contribute equally to a surface
  - Some may be blocked or affected by BRDF



55

## Lightcuts key idea - clustering

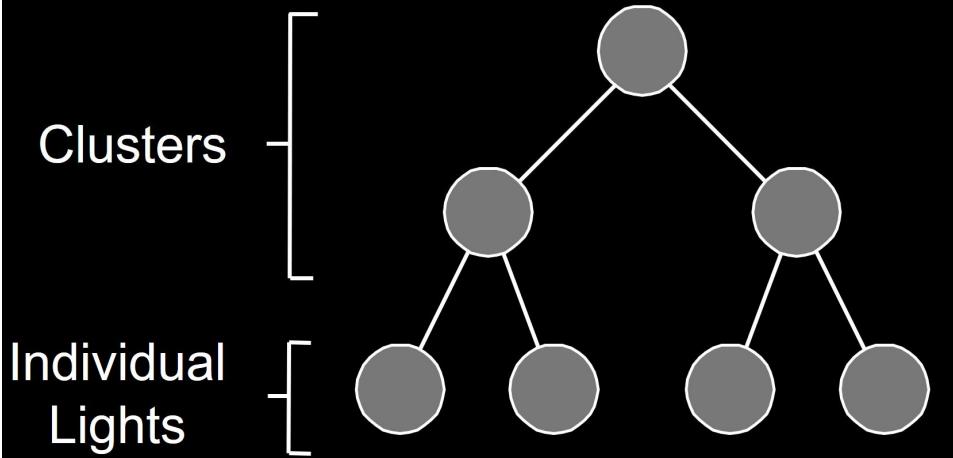
- Cluster a group of lights and represent them with a representative light



56

## Light tree

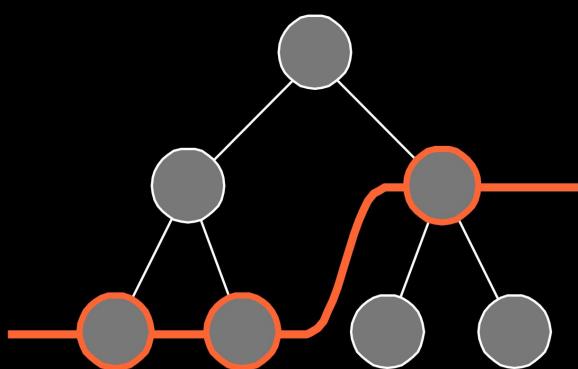
- A binary tree of lights and light clusters



57

## A cut

- A set of nodes that partitions the lights into clusters
  - Leaf nodes are individual lights

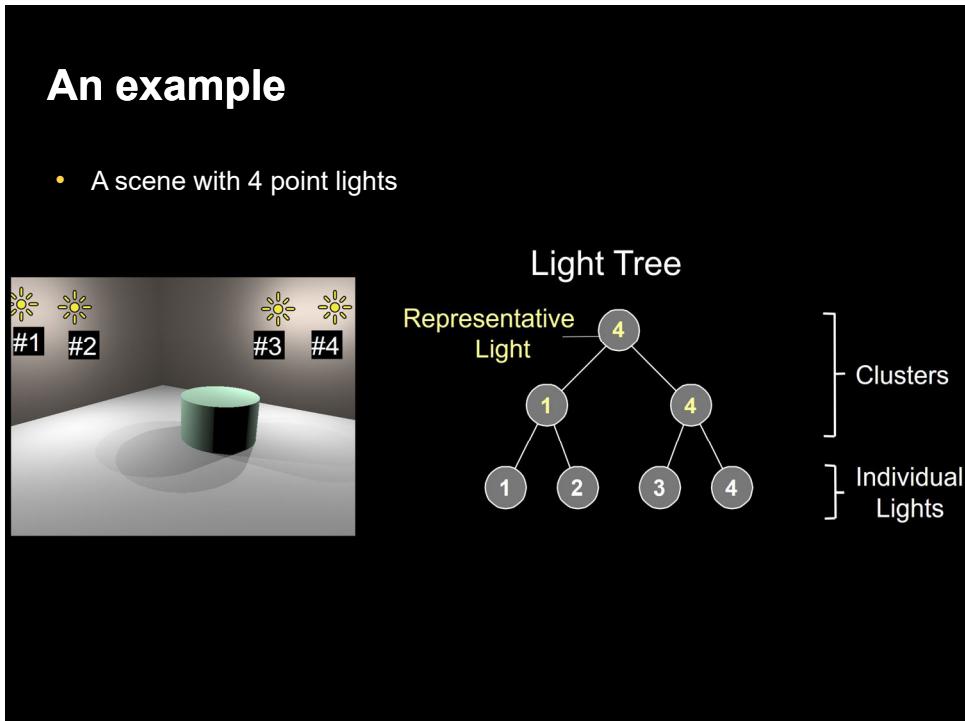


58

29

## An example

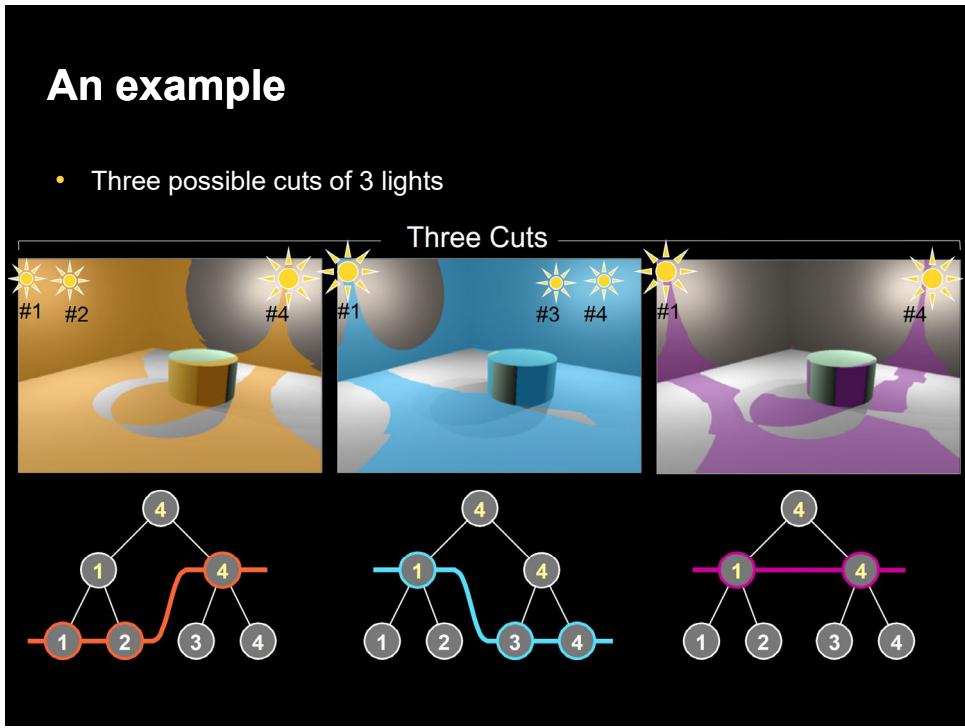
- A scene with 4 point lights



59

## An example

- Three possible cuts of 3 lights

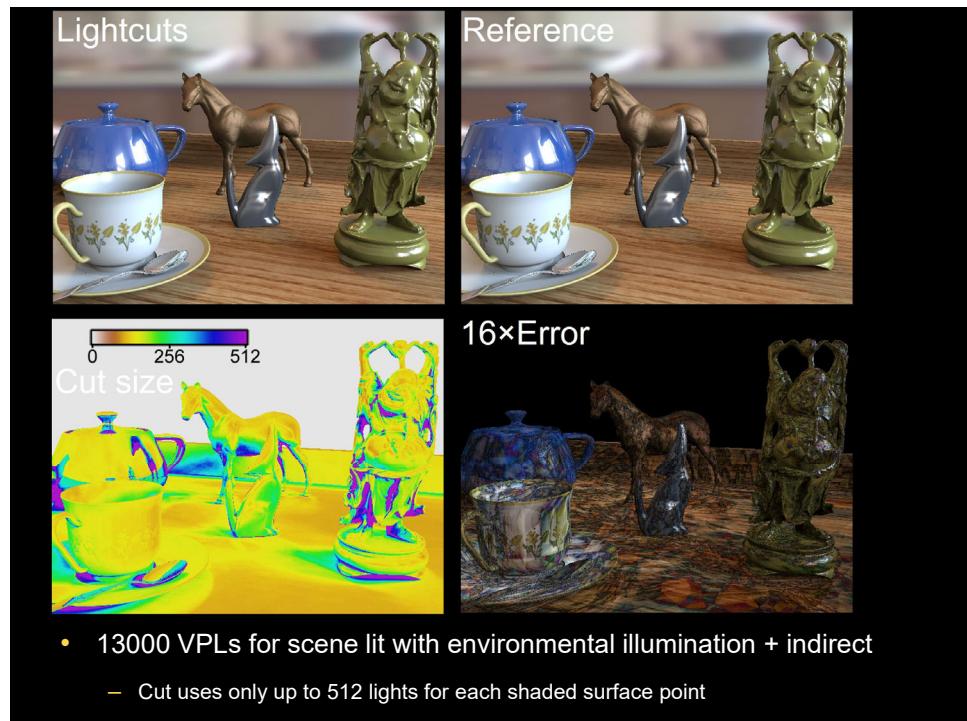


60

## Lightcuts algorithm

- Pre-process
  - Convert illumination to point lights
  - Build light tree
- For each visible point
  - Choose a cut to approximate the local illumination
    - Bound maximum error of cluster approximation
    - Refine cluster if error bound is too large
- Ensure each cluster's error < perceptual threshold
  - Webber's Law 2% of signal

61



62