

Machine Learning for Imaging

Image Segmentation

Ben Glocker

b.glocker@imperial.ac.uk

Common image analysis tasks

Image Classification



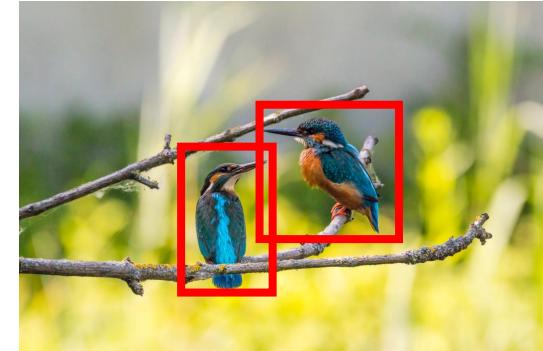
Output: Category (e.g., “bird”)

Object Detection



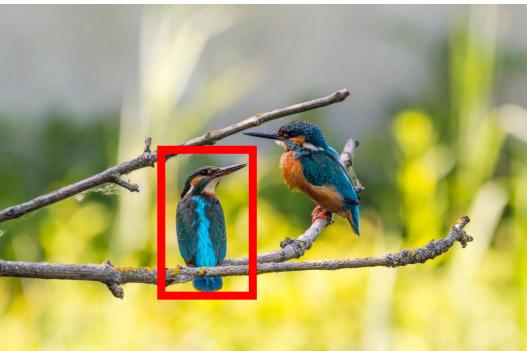
Output: Coordinates (e.g., centroid)

Object Localisation



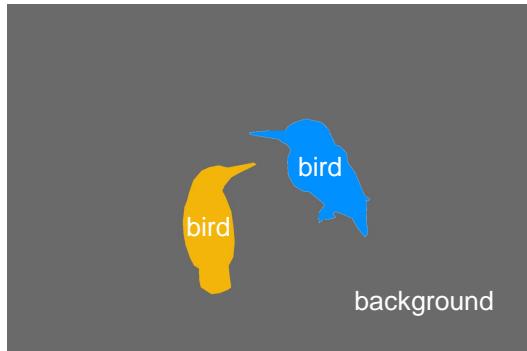
Output: Coordinates (e.g., bounding box)

Object Recognition



Output: Category (e.g., “kingfisher”)

Semantic Segmentation



Output: Labelmap

Image Captioning

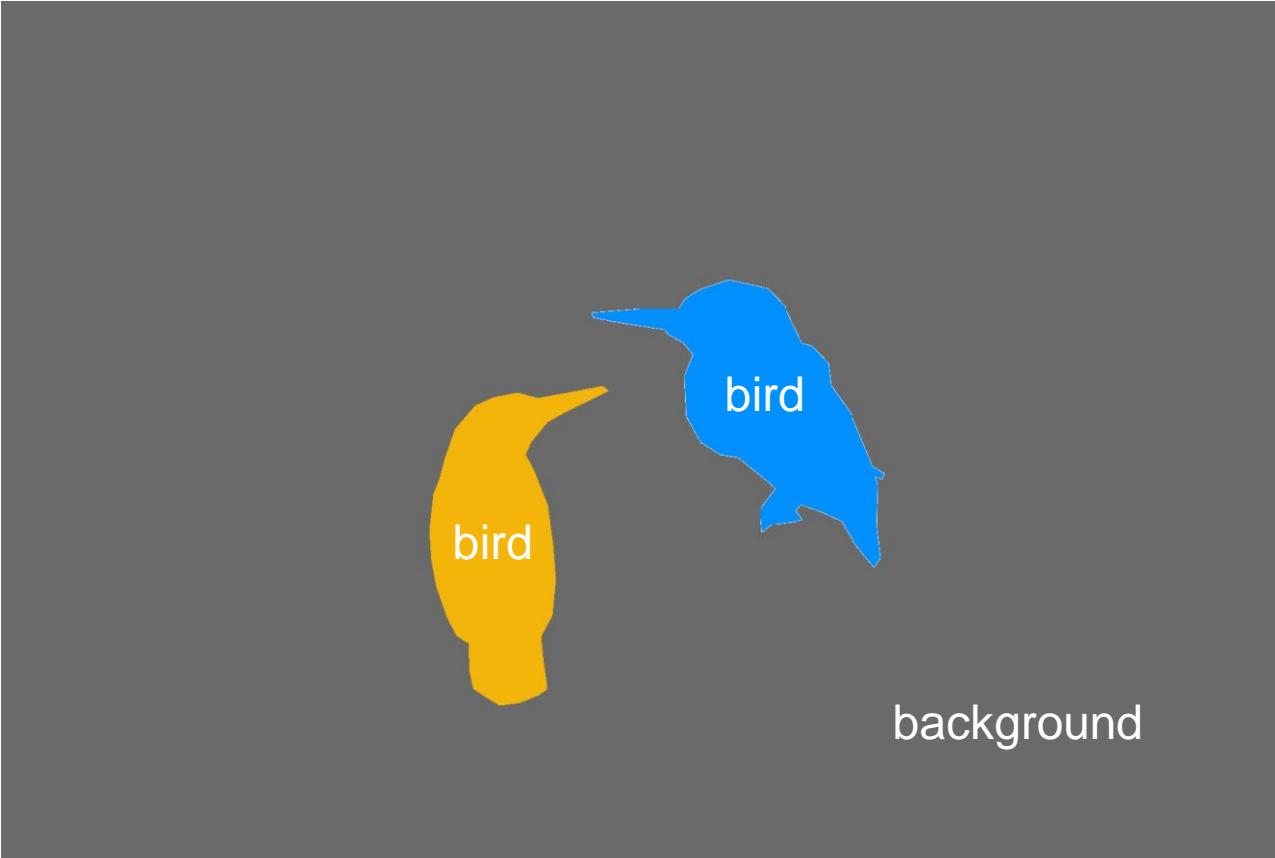


Output: Text
(e.g., “two birds sitting on a branch”)

Semantic segmentation

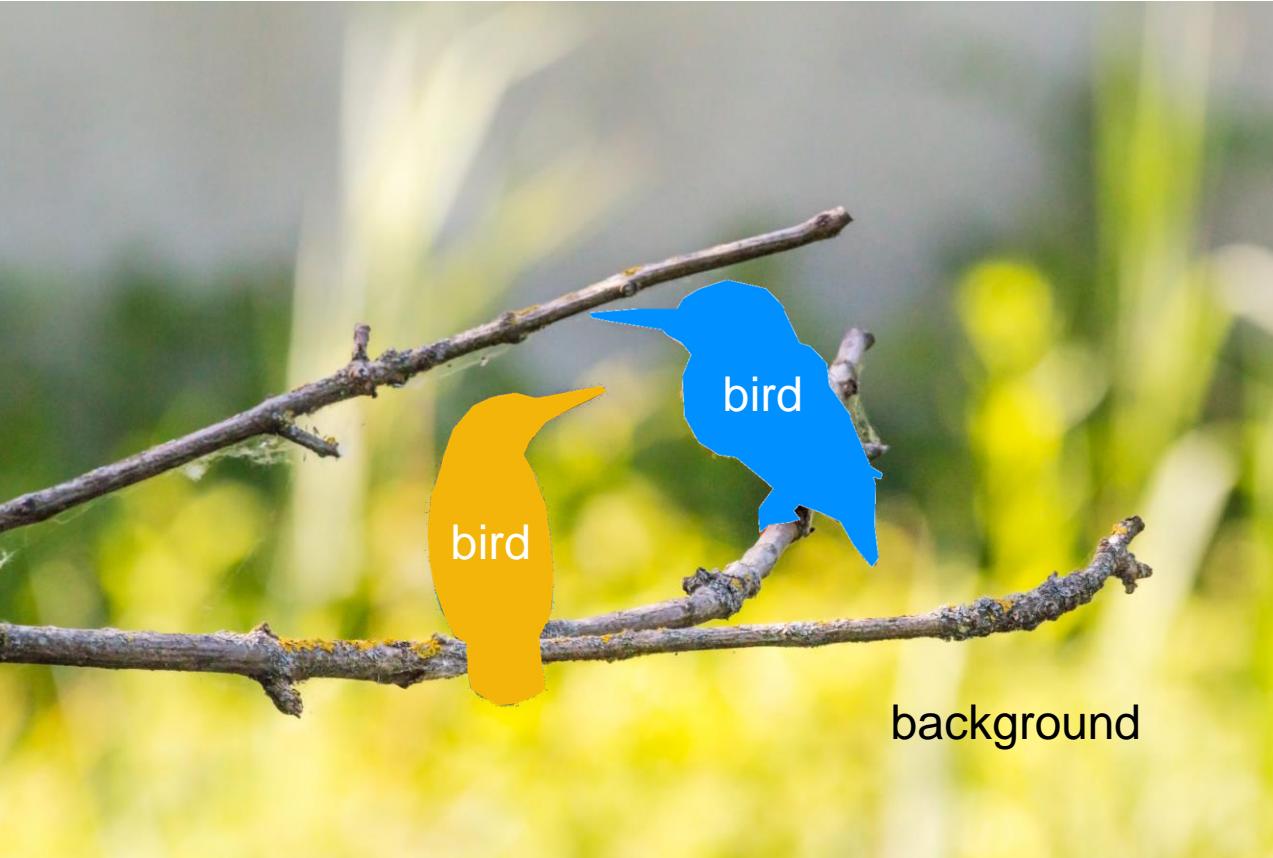


Semantic segmentation



Output: Labelmap

Semantic segmentation



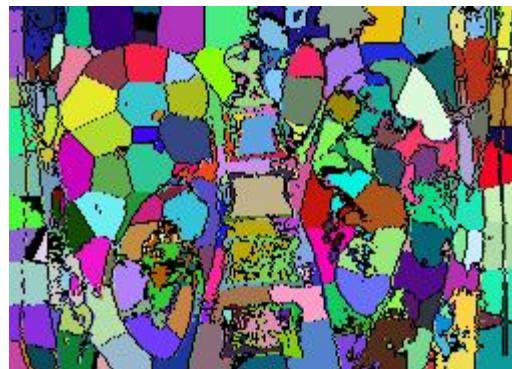
Output: Labelmap

Semantic segmentation

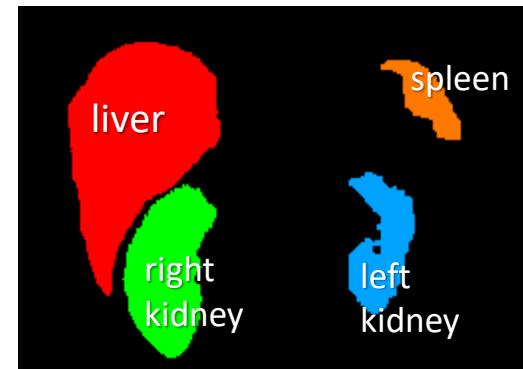
- In semantic segmentation, a segmented region is also assigned a **semantic meaning**
- This contrasts with segmentation based on ‘pure’ clustering of the image into coherent regions
- Here, we are mostly interested in semantic segmentation



raw image



segmentation/clustering

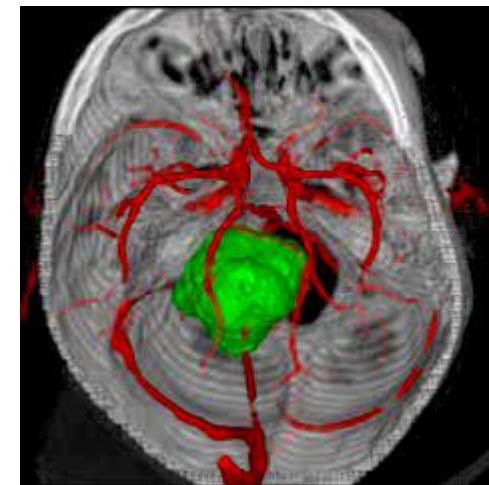
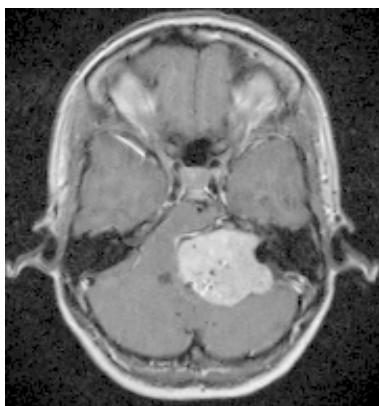
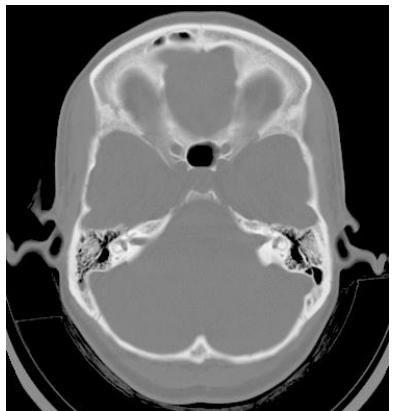


semantic segmentation

Segmentation in medical imaging

- **Image segmentation is useful for...**
 - conducting **quantitative analyses**, e.g. measuring the volume of the ventricular cavity.
 - determining the precise **location and extent** of an organ or a certain type of tissue, e.g. a tumour, for treatment such as radiation therapy.
 - creating **3D models** used for **simulation**, e.g. generating a model of an abdominal aortic aneurysm for simulating stress/strain distributions.

Example



C. Ruff, 1995

Challenges for image segmentation

- There are several effects in imaging that may hamper the application of segmentation algorithms
 - noise,
 - partial volume effects,
 - intensity inhomogeneities, anisotropic resolution,
 - imaging artifacts,
 - limited contrast,
 - morphological variability,
 - and many more ...

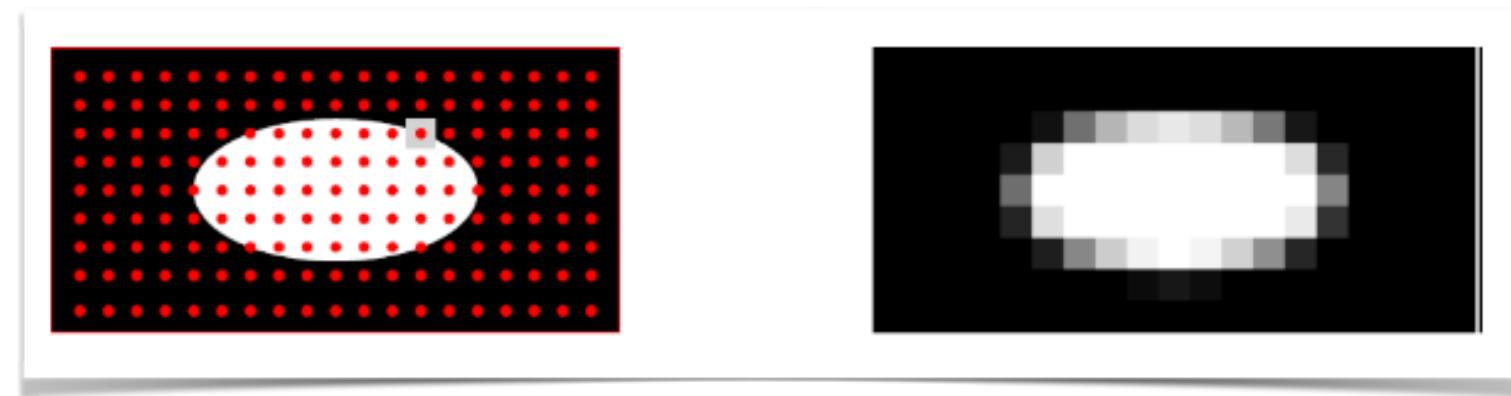
Noise

- Example: Transcranial 2D ultrasound image of the brain

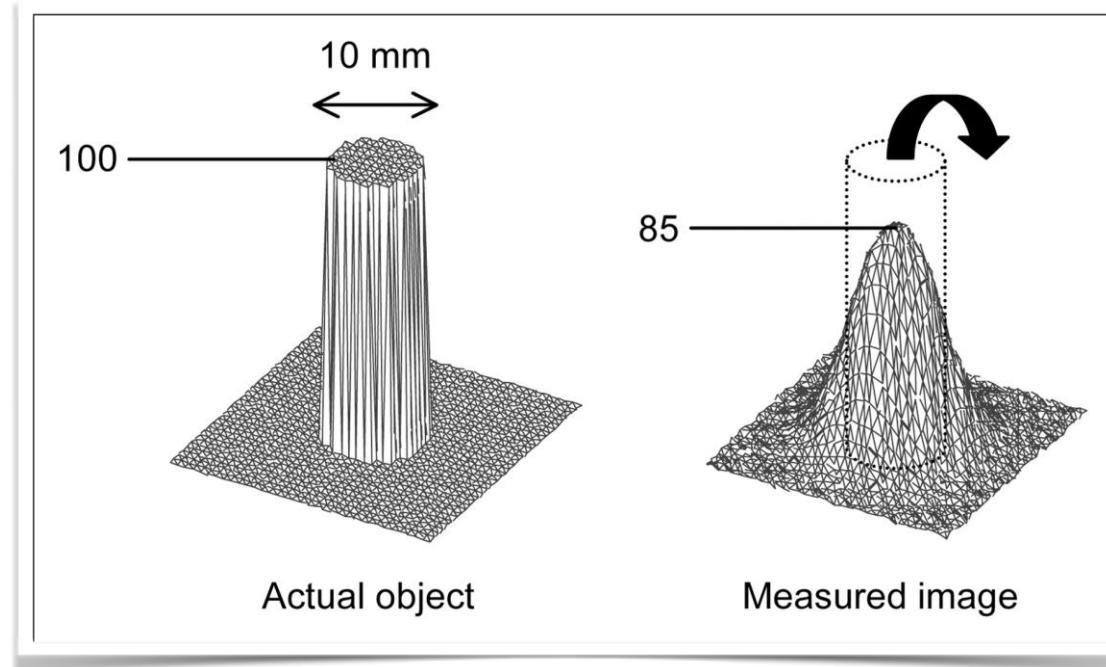


Partial volume effects

- Due to the coarse sampling (red dots) the resulting image shows partial volume effects at the boundary of the white “organ”
- Both tissue types (black and white) contribute to the intensity value of the generated image (right) due to the relatively large influence area of each pixel (gray square in left image)

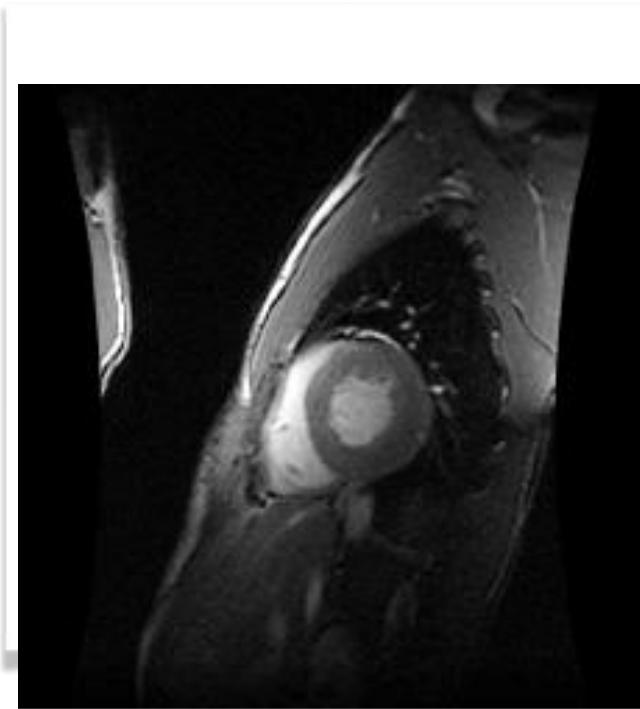


Partial volume effects



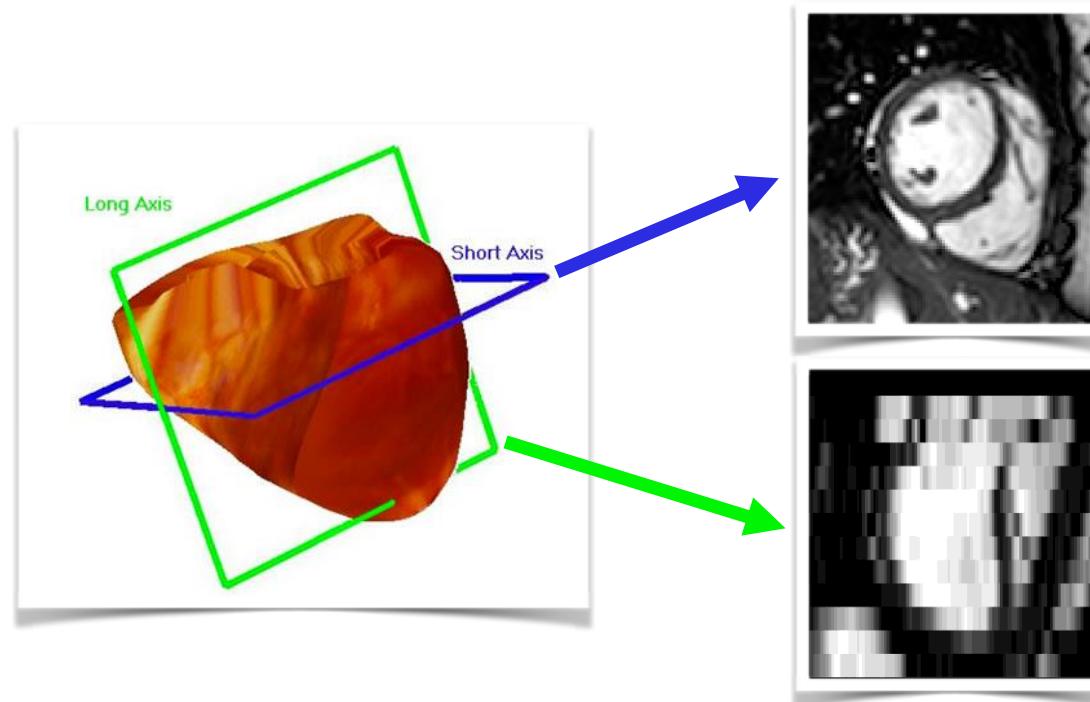
Intensity inhomogeneities

- Typical for MRI and ultrasound imaging data

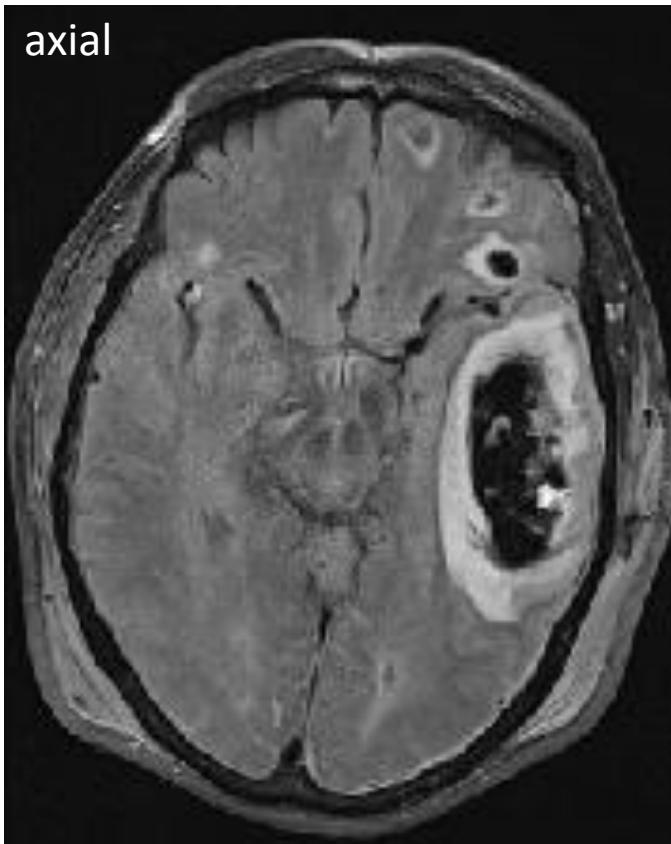


Anisotropic resolution

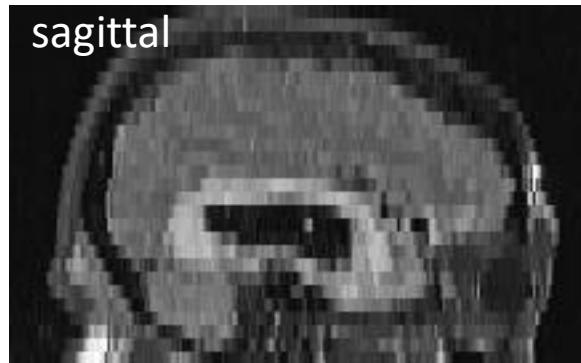
- Almost all 3D imaging modalities, such as CT or MRI suffer from this problem
 - For example, (dynamic) MRI acquisition of short axis slices of the left ventricle may have intra-slice resolution of 1.3mm and an inter-slice resolution of 8mm



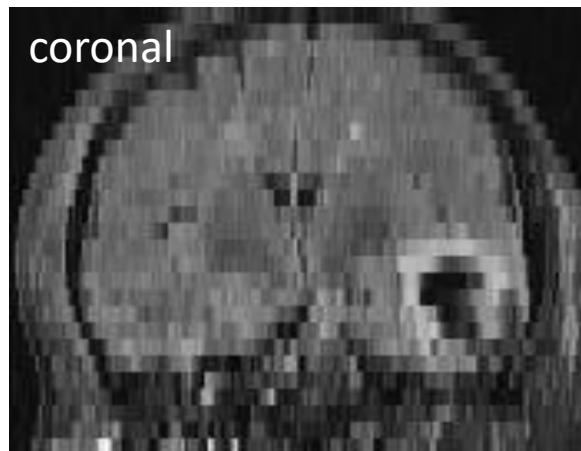
Anisotropic resolution



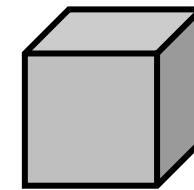
axial



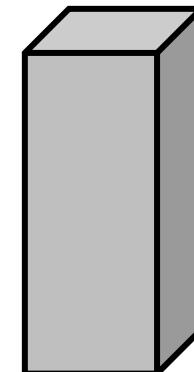
sagittal



coronal



isotropic



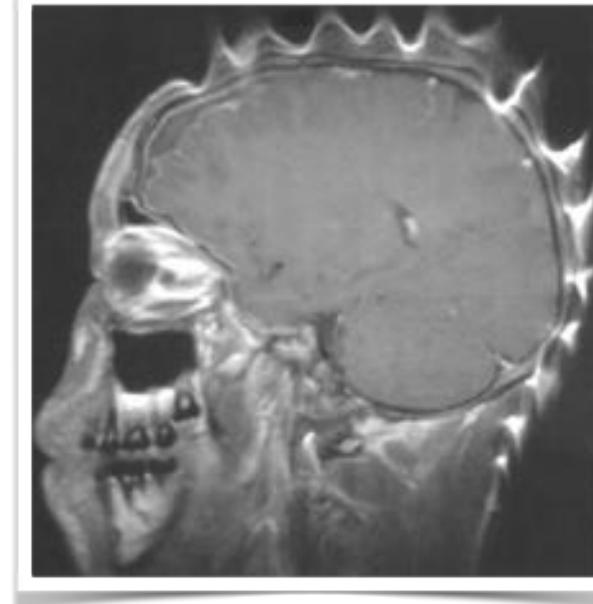
anisotropic

Voxel size: 0.7 x 0.7 x 5mm

Imaging artifacts



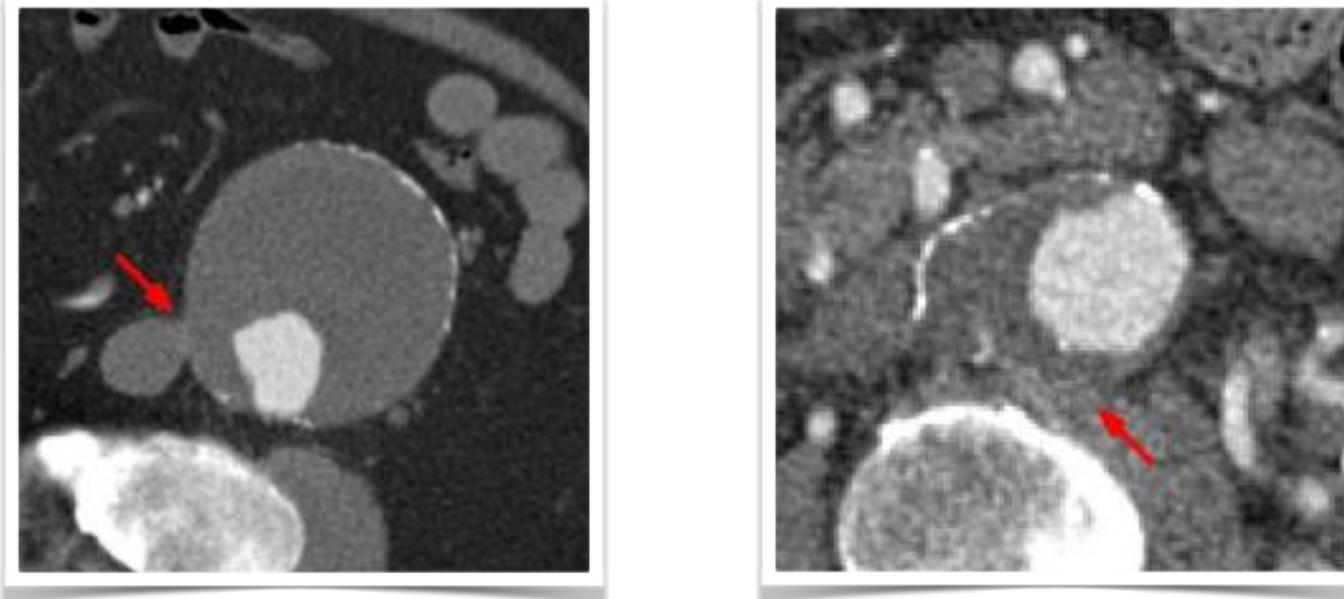
CT image showing streak artifacts caused by metal implants.



T1-weighted MRI image showing susceptibility artifacts caused by iron oxide particles suspended in the beeswax dressing in the hair of the patient.

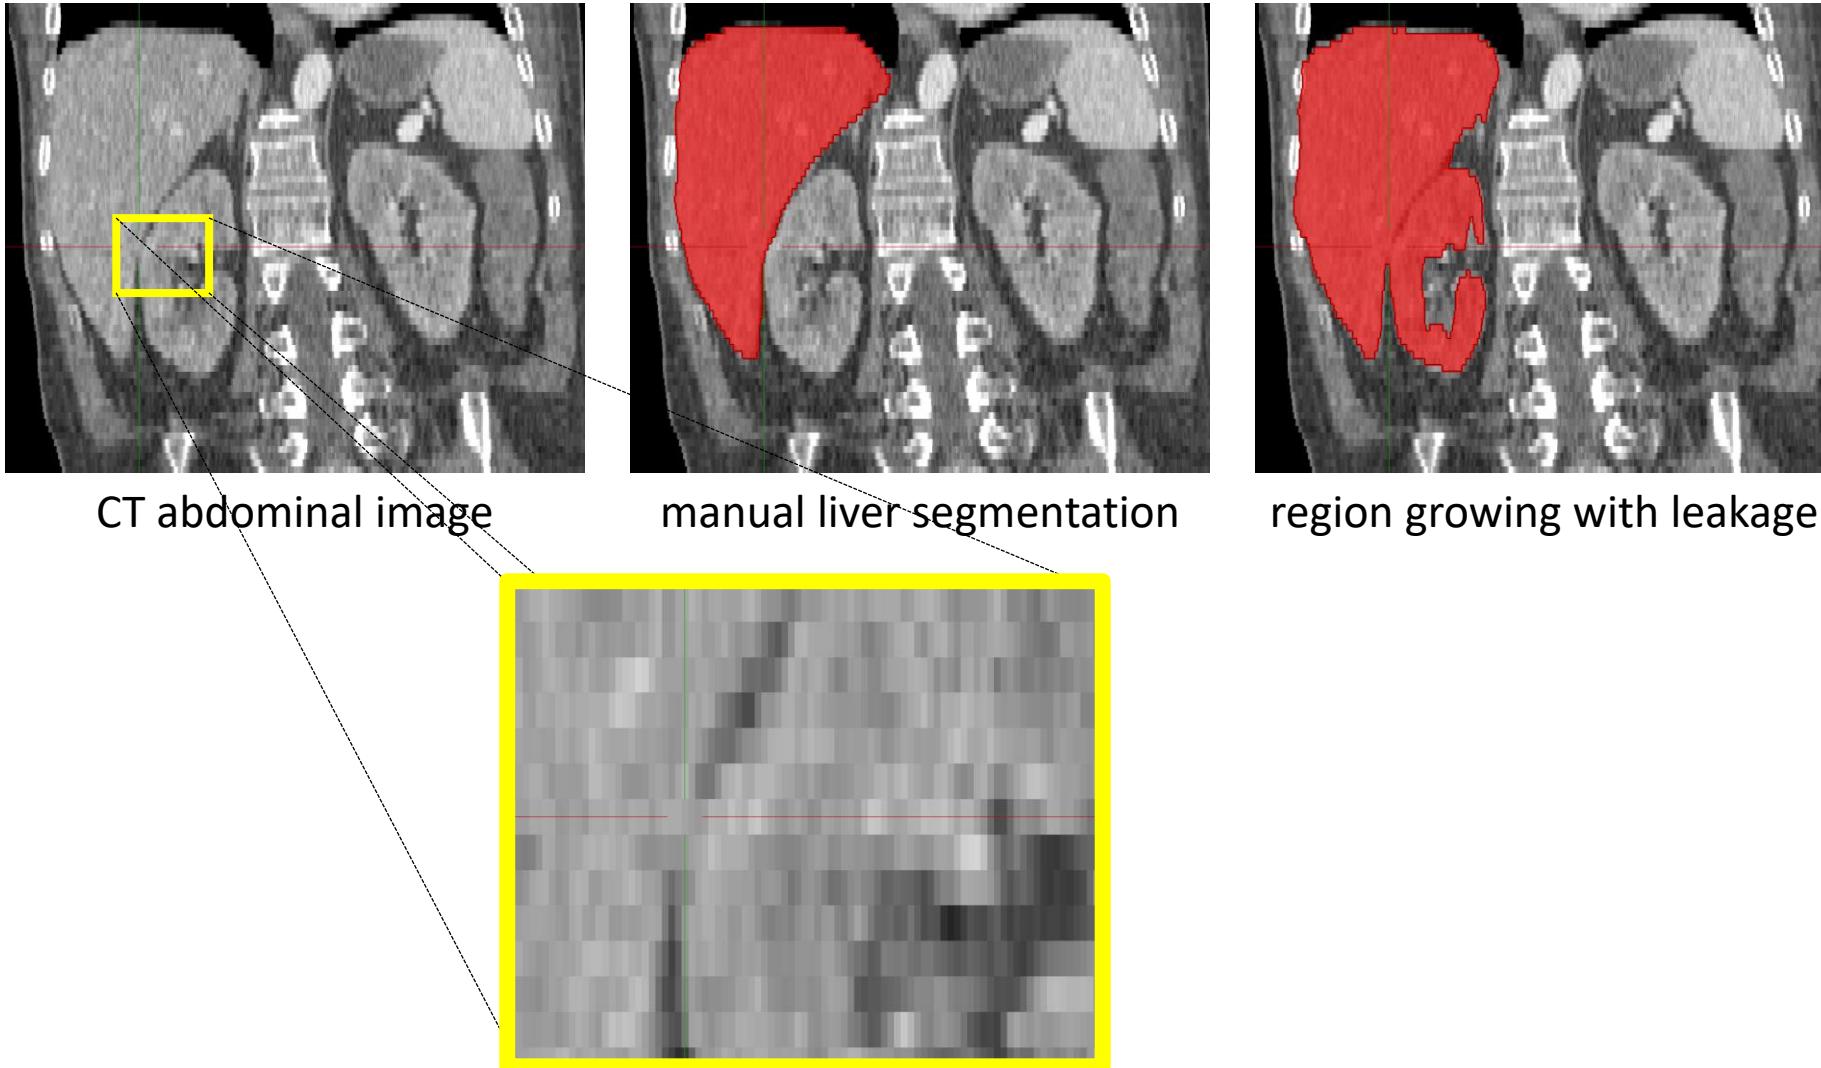
Limited contrast

- Different tissues can have similar physical properties and thus similar intensity values
- Purely intensity-based algorithms are prone to fail or “leak” into adjacent tissue



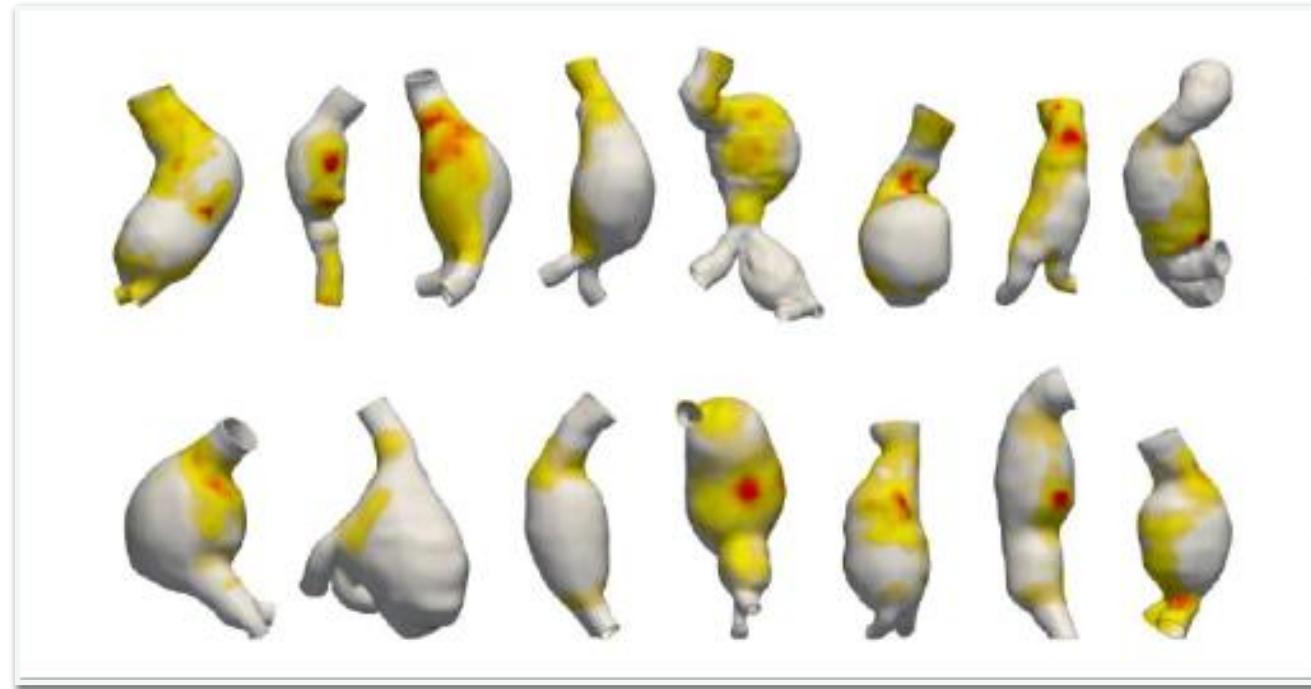
Example: Thrombus in the wall of abdominal aortic aneurysms is often hard to distinguish from the surrounding tissue (modality: CTA).

Limited contrast



Morphological variability

- Morphological variability makes it hard to incorporate meaningful prior information or useful shape models



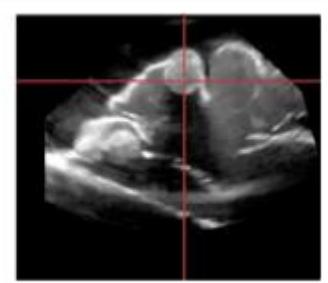
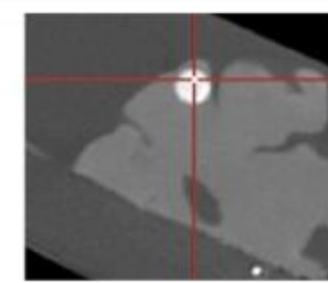
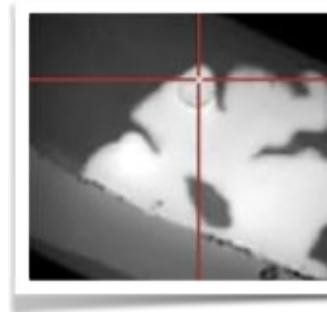
Example: A collection of abdominal aortic aneurysms acquired with PET-CT and colored by FDG-18 uptake values.

Segmentation evaluation

- Ground truth
- Gold standard
- Performance measures
 - precision/recall
 - sensitivity/specificity
 - accuracy
 - overlap
 - distance measures

Ground truth

- Reference or standard against a method can be compared, e.g. the optimal transformation, or a true segmentation boundary
- GT is usually only available for:
 - Synthetic or simulated phantoms (e.g., computer phantoms)
 - Physical phantom (e.g., gel phantoms)



PVA mould of a brain model, scanned with MR, CT and US.

Gold standard

- Expert (often referred to as *gold* standard)
 - Manual segmentation by human observer
(e.g. experienced clinician)
- Disadvantage
 - Requires training and is tedious and time-consuming
 - Intra-observer variability (disagreement between same observer on different occasions)
 - Inter-observer variability (disagreement between observers)
- Remedy
 - Human observer can perform segmentation repeatedly
 - Multiple experts can perform segmentations
 - Agreement or disagreement can be quantified

How to assess performance?

Precision

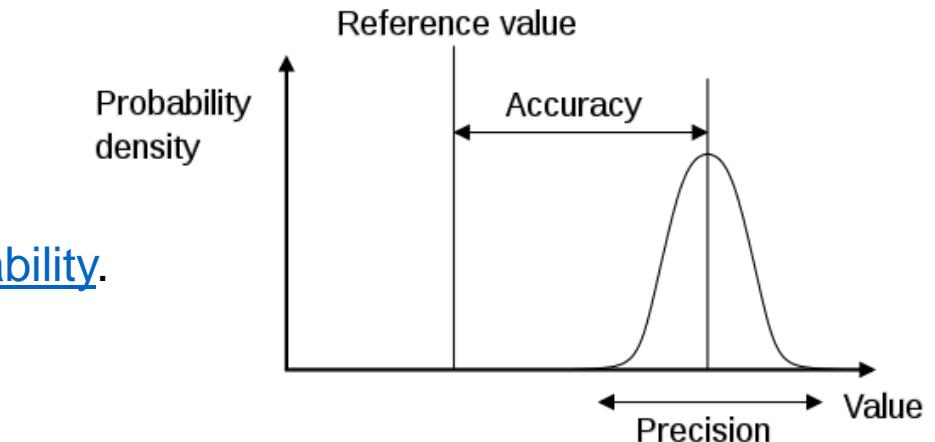
- is a description of [random errors](#), a measure of [statistical variability](#).
- the repeatability, or reproducibility of the measurement

Accuracy has two definitions

- More commonly, it is a description of [systematic errors](#), a measure of [statistical bias](#); as these cause a difference between a result and a "true" value, [ISO](#) calls this *trueness*.
- Alternatively, ISO defines accuracy as describing a combination of both types of [observational error](#) above (random and systematic), so high accuracy requires both high precision and high trueness

Robustness

- refers to the degradation in performance with respect to varying noise levels or other imaging artefacts

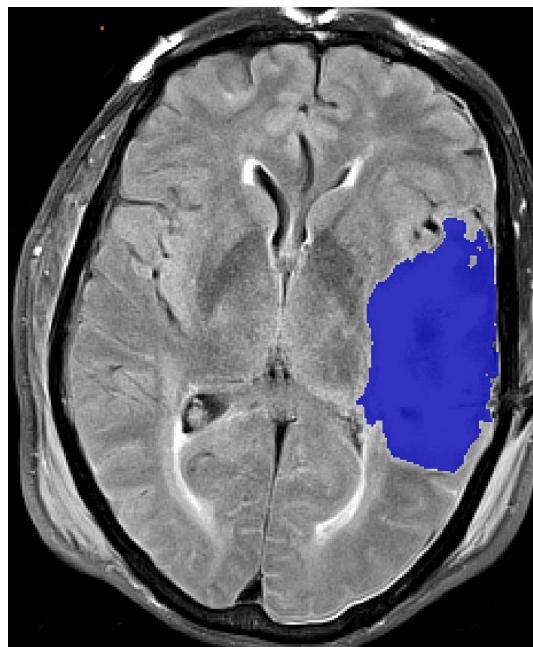


Performance measures

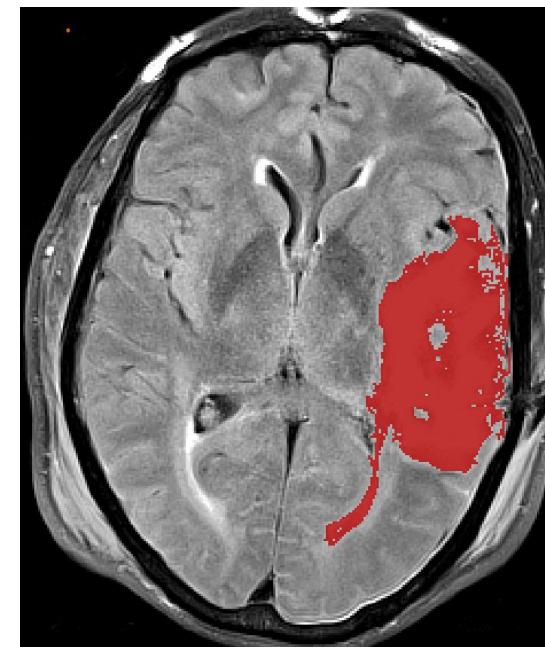
- Assume we have an automated segmentation algorithm: how do we assess 'how good it is'?
- There are several ways of quantifying segmentation performance



MR image



Gold standard



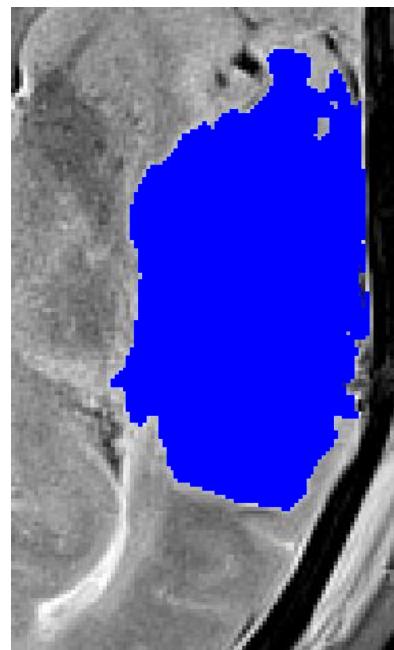
Auto Segmentation

Performance measures

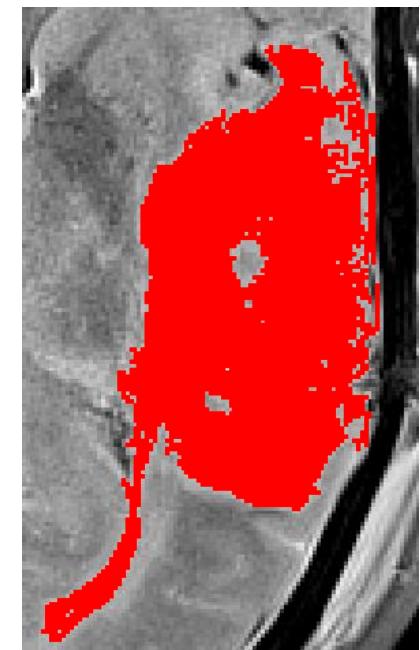
- Assume we have an automated segmentation algorithm, how do we assess 'how good it is'?
- There are several ways of quantifying segmentation performance



MR image



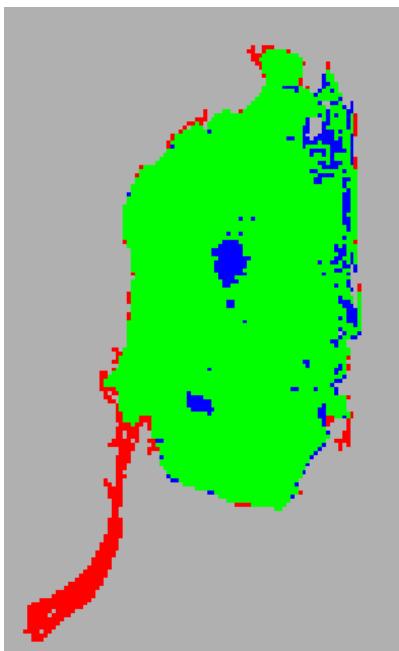
Gold standard



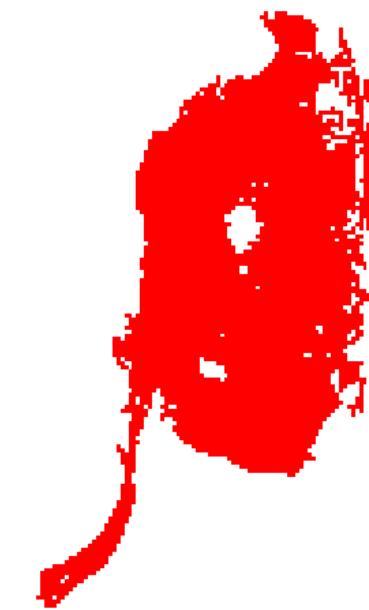
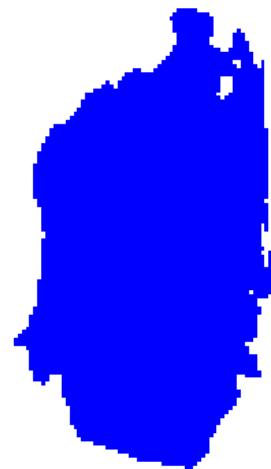
Auto Segmentation

Performance measures

- Measure performance of an automated method in terms of **agreement** of its result with a reference gold standard



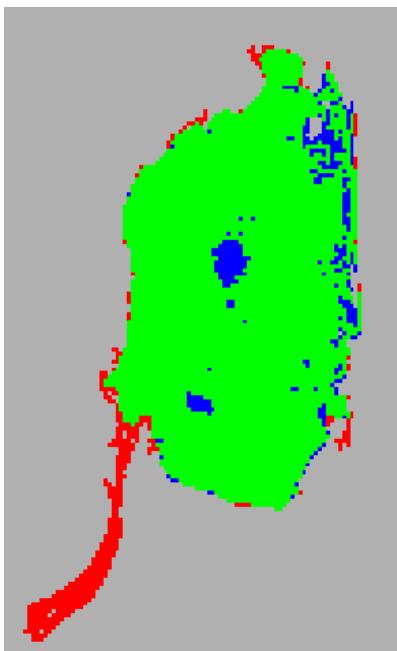
Gold standard



Auto Segmentation

Performance measures

- Measure performance of an automated method in terms of **agreement** of its result with a reference gold standard



agreement

disagreement

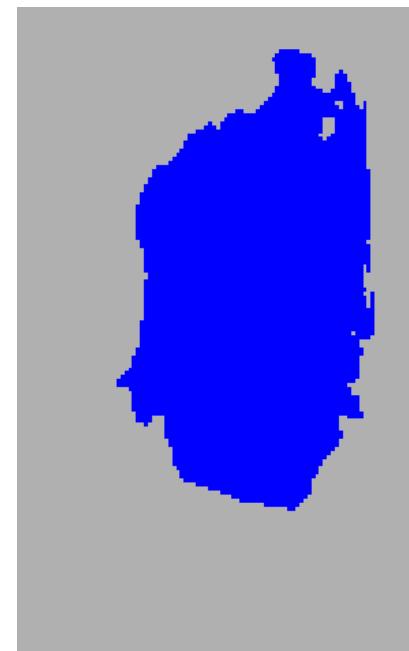
disagreement

agreement

positives

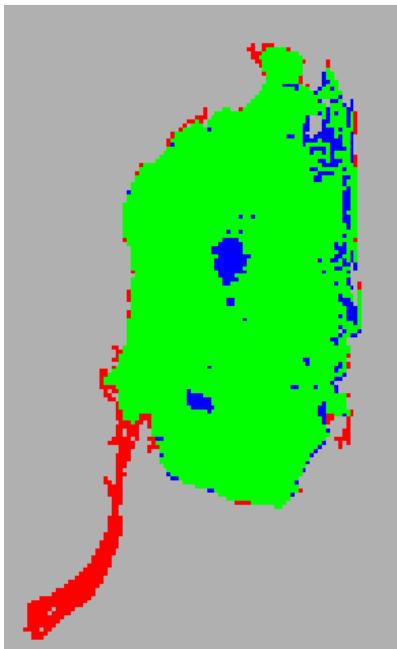


negatives



Performance measures

- Measure performance of an automated method in terms of **agreement** of its result with a reference gold standard



■ true positives

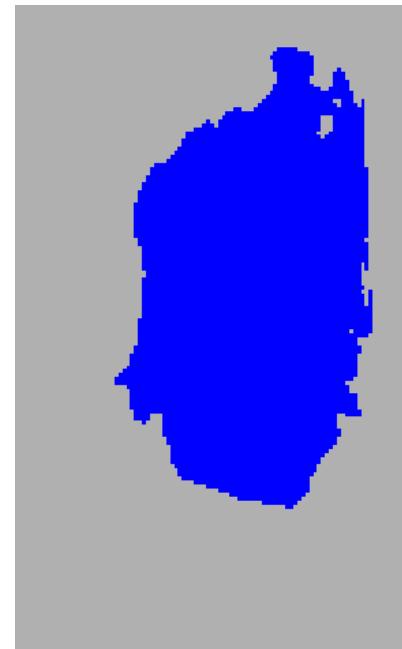
■ false positives

■ false negatives

■ true negatives

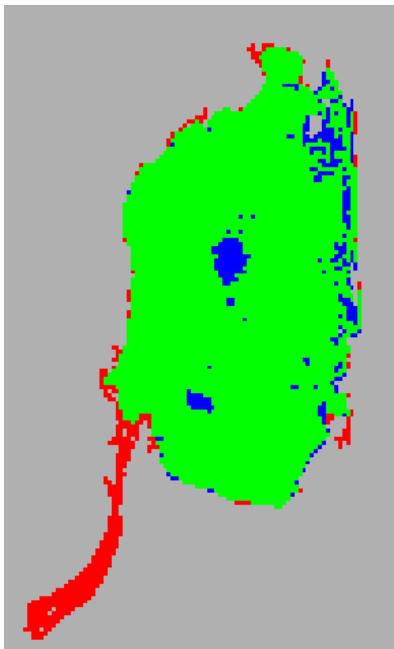
positives ■

negatives ■

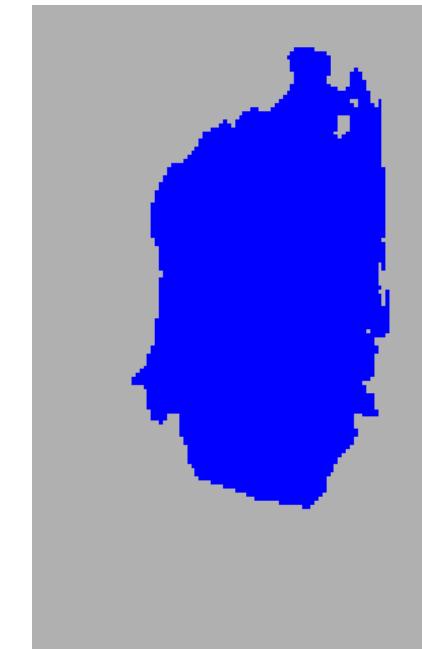


Performance measures

- Measure performance of an automated method in terms of **agreement** of its result with a reference gold standard



TP
FP
FN
TN



P
N

Confusion matrix

condition positive (P)

the number of real positive cases in the data

condition negatives (N)

the number of real negative cases in the data

true positive (TP)

eqv. with hit

true negative (TN)

eqv. with correct rejection

false positive (FP)

eqv. with false alarm, Type I error

false negative (FN)

eqv. with miss, Type II error

		True condition	
		Total population	Condition positive
Predicted condition	Predicted condition positive	True positive	False positive, Type I error
	Predicted condition negative	False negative, Type II error	True negative

Accuracy, precision, recall, ...

accuracy

$$ACC = \frac{TP+TN}{P+N}, P = TP + FN, N = TN + FP$$

precision or positive predictive value

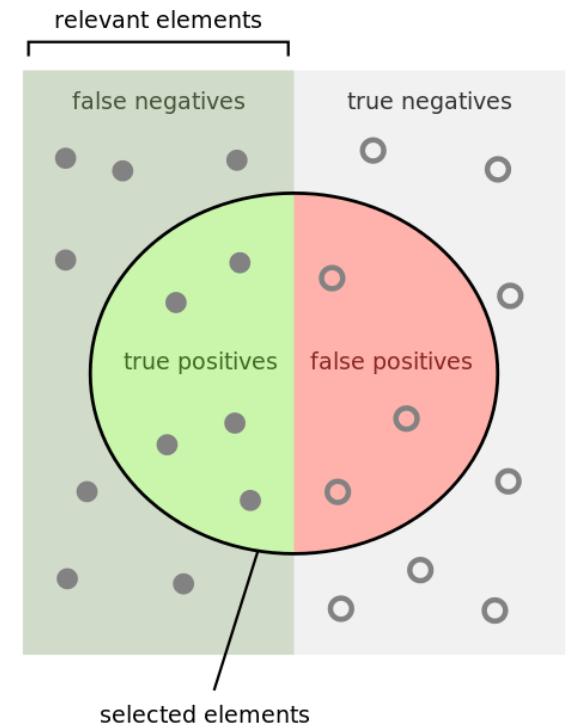
$$PPV = \frac{TP}{TP + FP}$$

recall, sensitivity, hit rate or true positive rate

$$TPR = \frac{TP}{P} = \frac{TP}{TP + FN}$$

specificity or true negative rate

$$TNR = \frac{TN}{N} = \frac{TN}{TN + FP}$$



How many selected items are relevant?

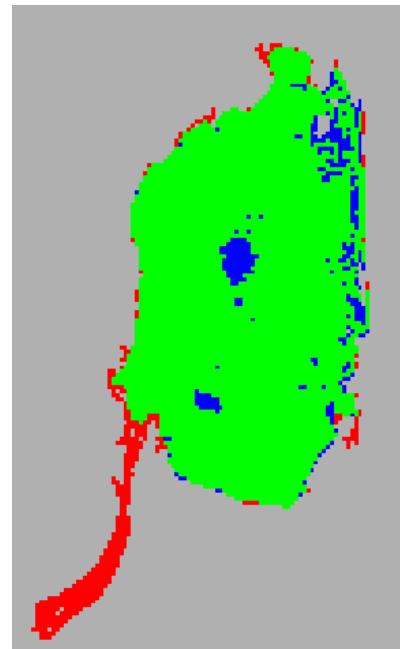
$$\text{Precision} = \frac{\text{green}}{\text{green} + \text{red}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{green}}{\text{green} + \text{blue}}$$

F1 score

precision $PPV = \frac{TP}{TP+FP}$



recall $TPR = \frac{TP}{TP+FN}$

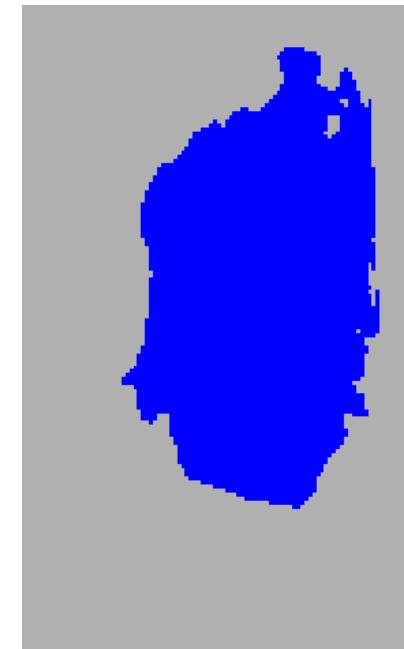
- TP
- FP
- FN
- TN

F1 score

is the [harmonic mean](#) of precision and recall

$$F_1 = 2 \frac{PPV \cdot TPR}{PPV + TPR} = \frac{2TP}{2TP + FP + FN}$$

- P ■
- N ■

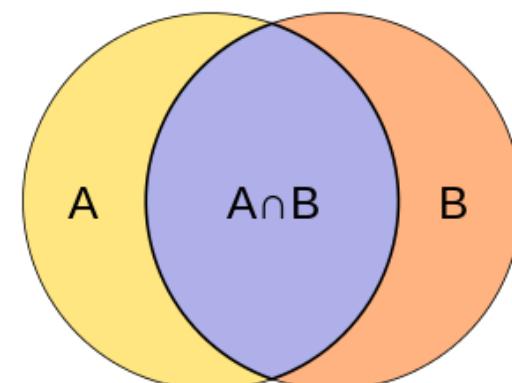


Overlap measures

- Measure agreement by looking at overlap of reference and predicted segmentation (e.g., auto segmentation)
- **Jaccard Index**

aka **Intersection over Union (IoU)**
aka **Jaccard Similarity Coefficient**

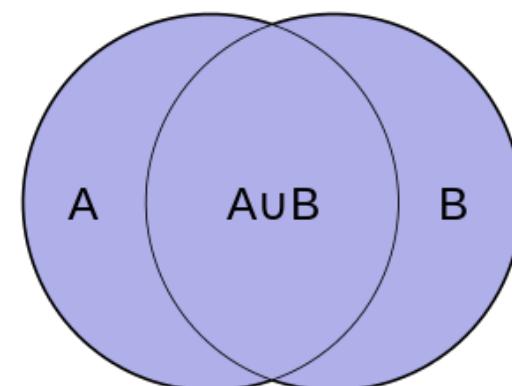
$$JI = \frac{|A \cap B|}{|A \cup B|}$$



- **Dice's Coefficient**

aka **Sørensen Index**
aka **Dice Similarity Coefficient (DSC)**

$$DSC = \frac{2|A \cap B|}{|A| + |B|}$$



Dice similarity coefficient

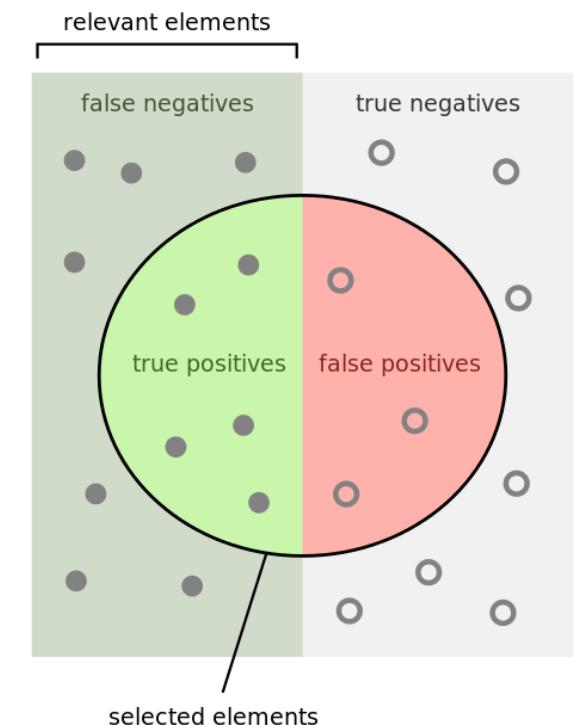
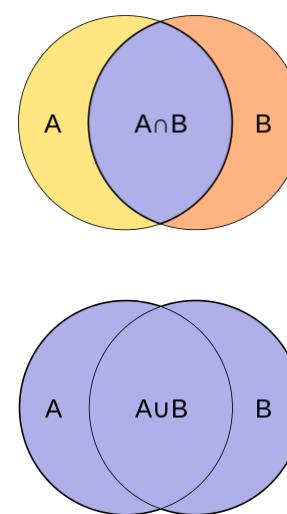
- Most widely used measure for evaluating segmentation
- Assume A is the reference, and B the prediction

$$DSC = \frac{2|A \cap B|}{|A| + |B|}$$

- $|A| = TP + FN$
- $|B| = TP + FP$
- $|A \cap B| = TP$

$$DSC = \frac{2TP}{2TP + FP + FN} = F_1$$

DSC is equivalent to F1 score!



Other measures

volume similarity

$$VS = 1 - \frac{||A| - |B||}{|A| + |B|} = 1 - \frac{|FN - FP|}{2TP + FP + FN}$$

surface distance measures

- Hausdorff distance

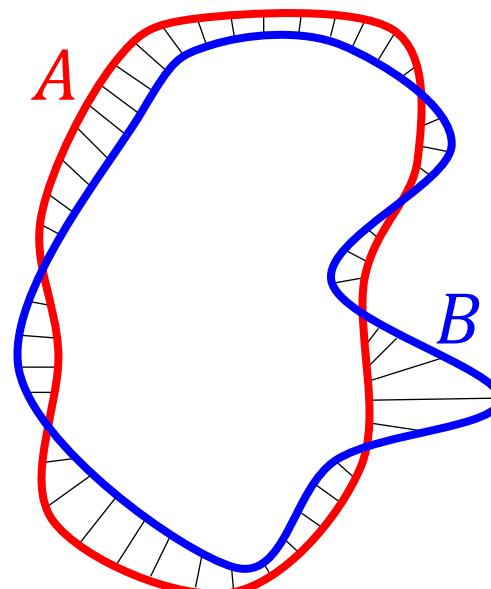
$$HD = \max(h(A, B), h(B, A))$$

$$h(A, B) = \max_{a \in A} \min_{b \in B} \|a - b\|$$

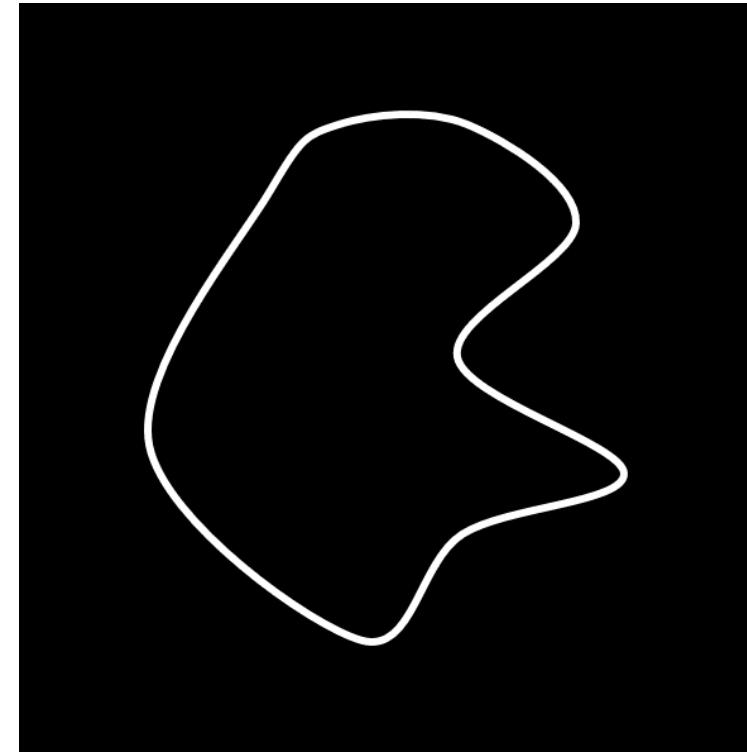
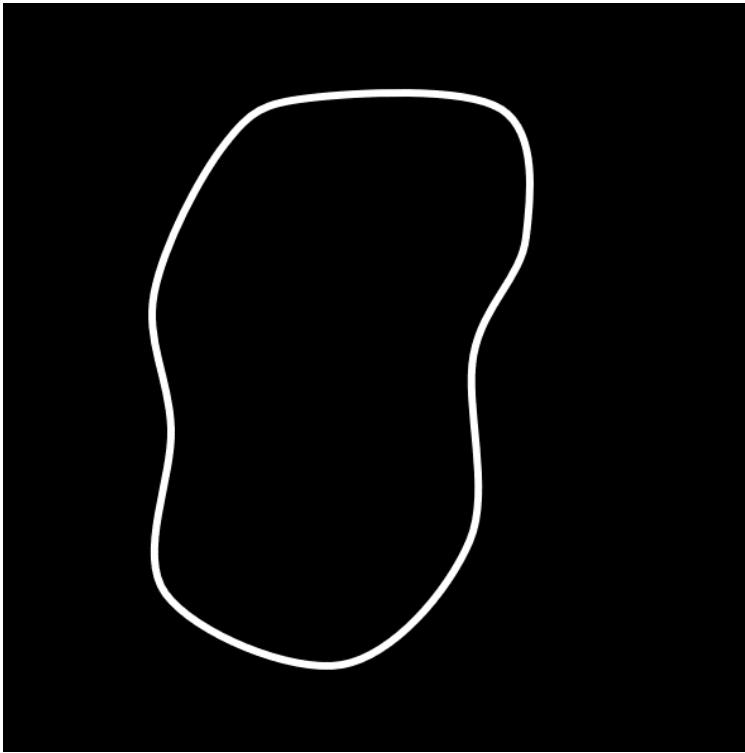
- (symmetric) average surface distance

$$ASD = \frac{d(A, B) + d(B, A)}{2}$$

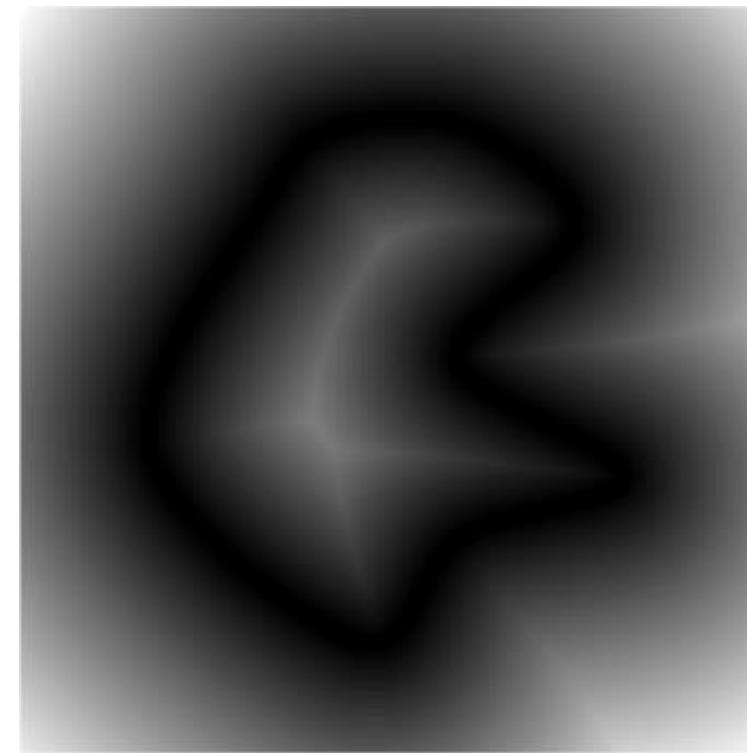
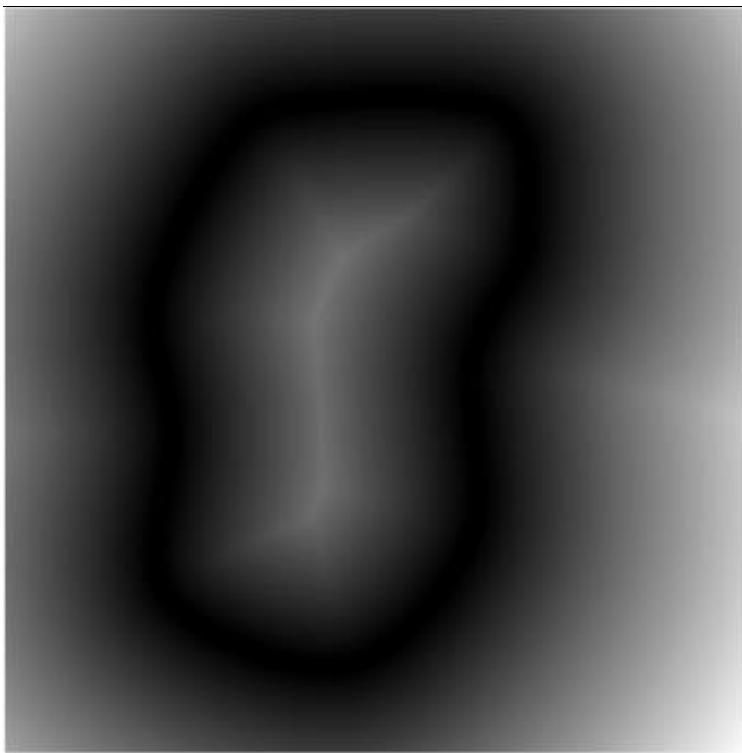
$$d(A, B) = \frac{1}{N} \sum_{a \in A} \min_{b \in B} \|a - b\|$$



Surface distance

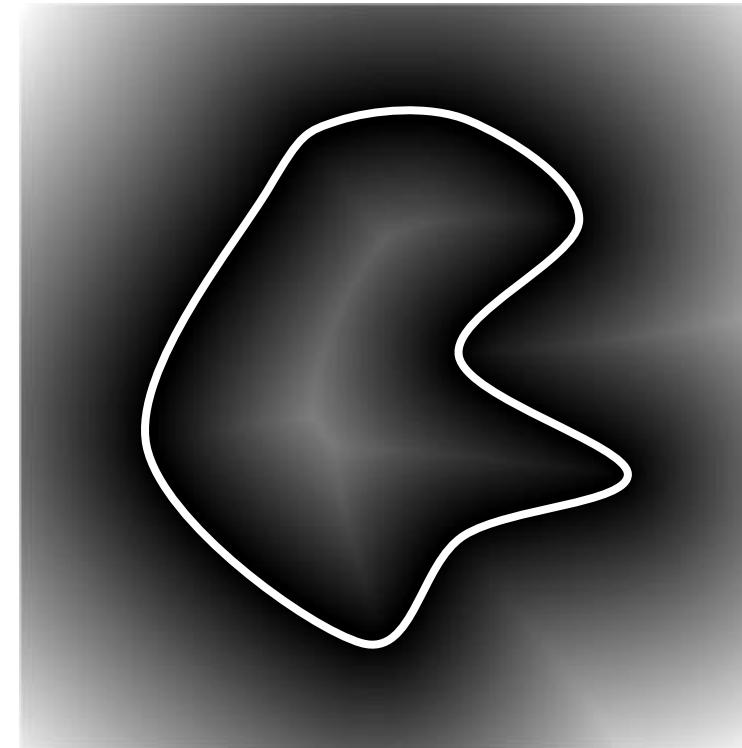
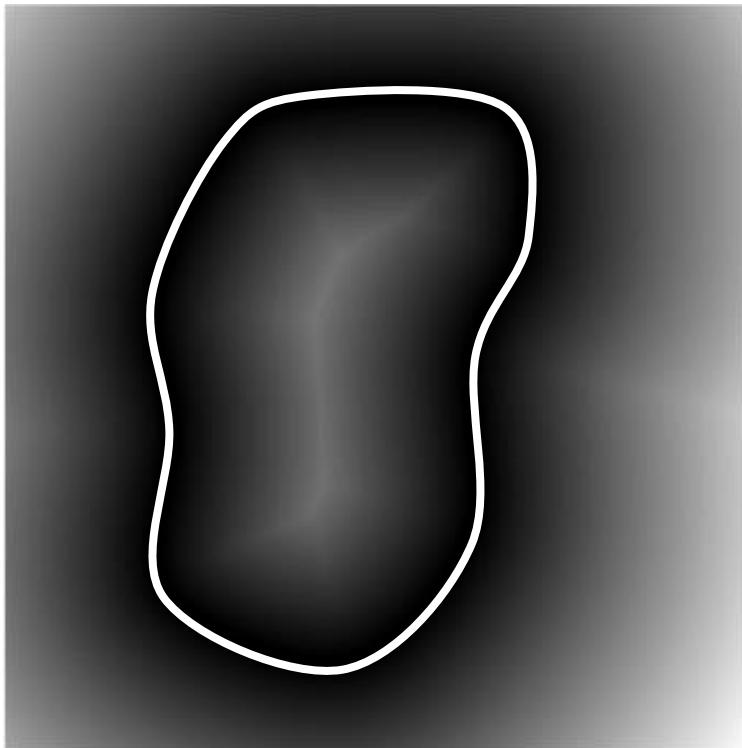


Surface distance



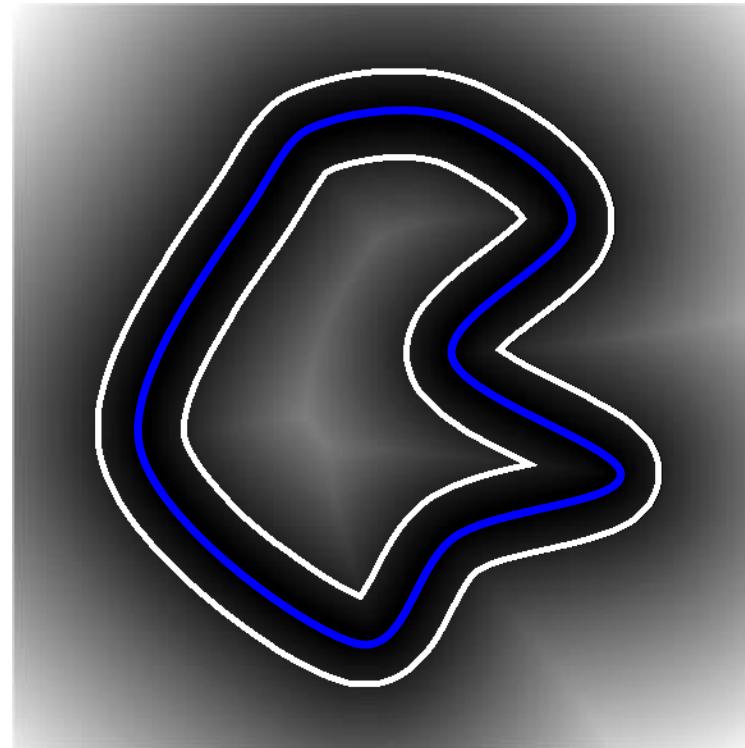
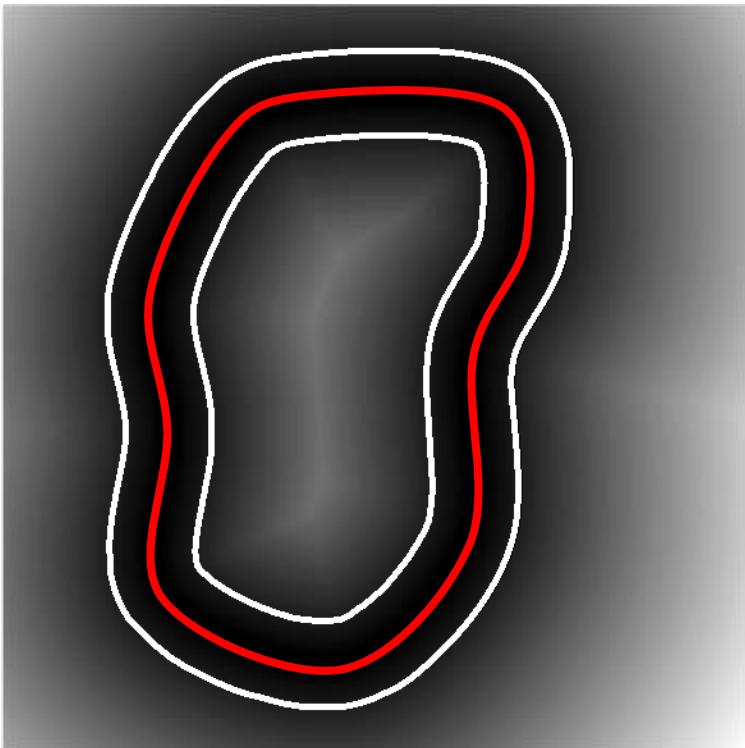
Euclidean Distance Maps

Surface distance



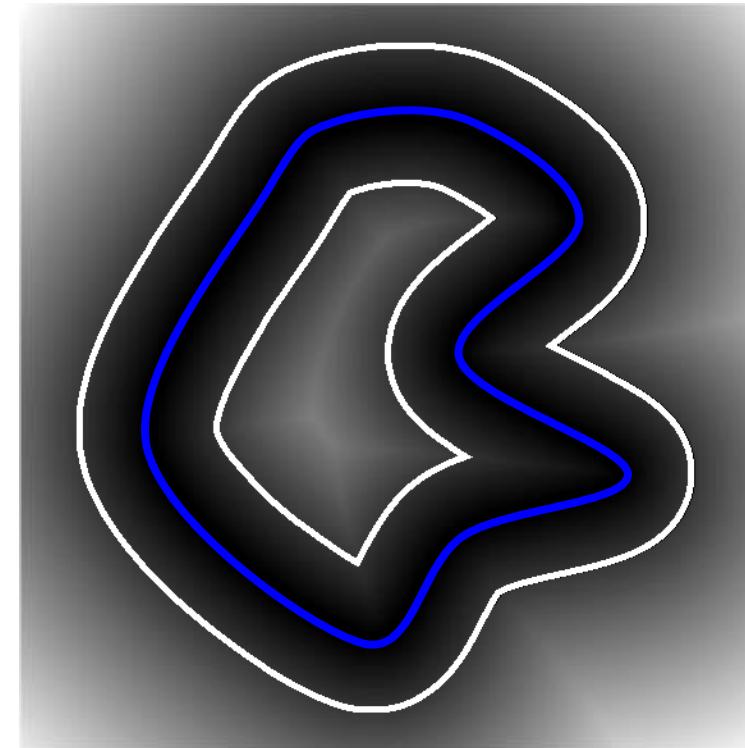
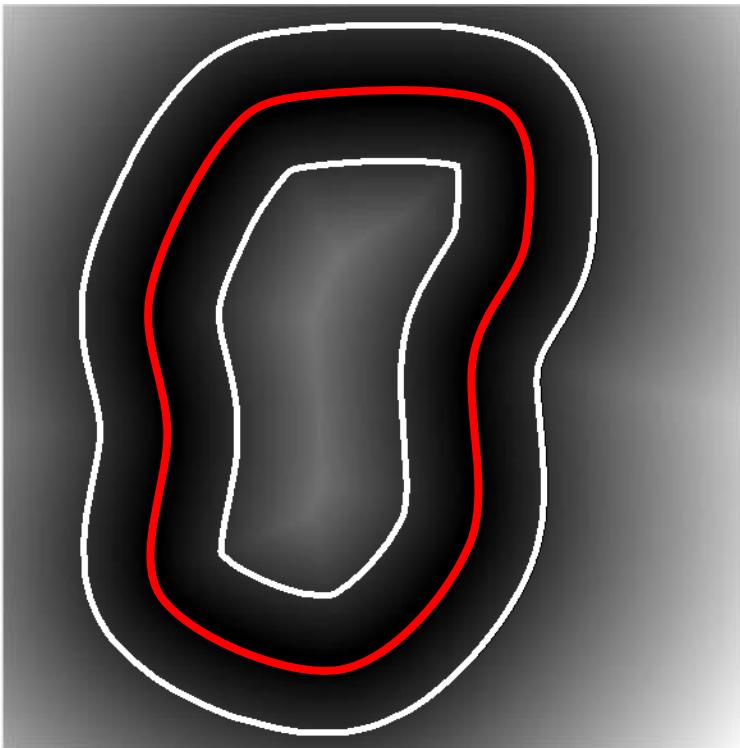
Zero-Level Surface

Surface distance



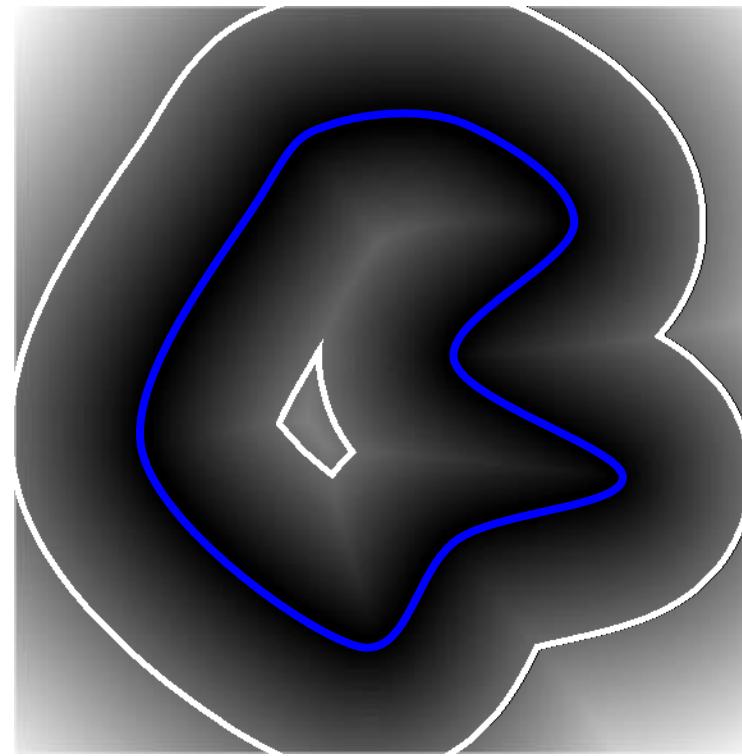
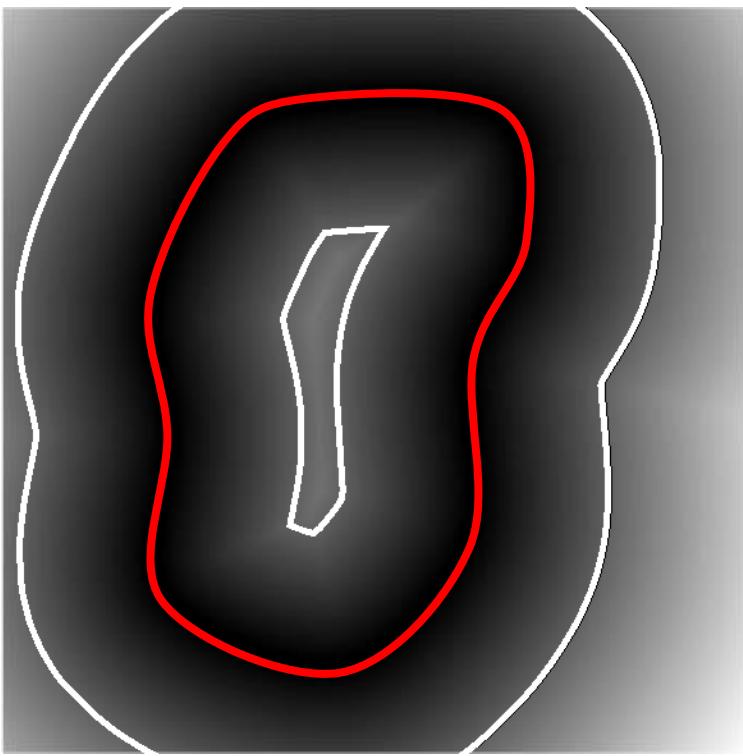
30-Pixels Distance Surface

Surface distance



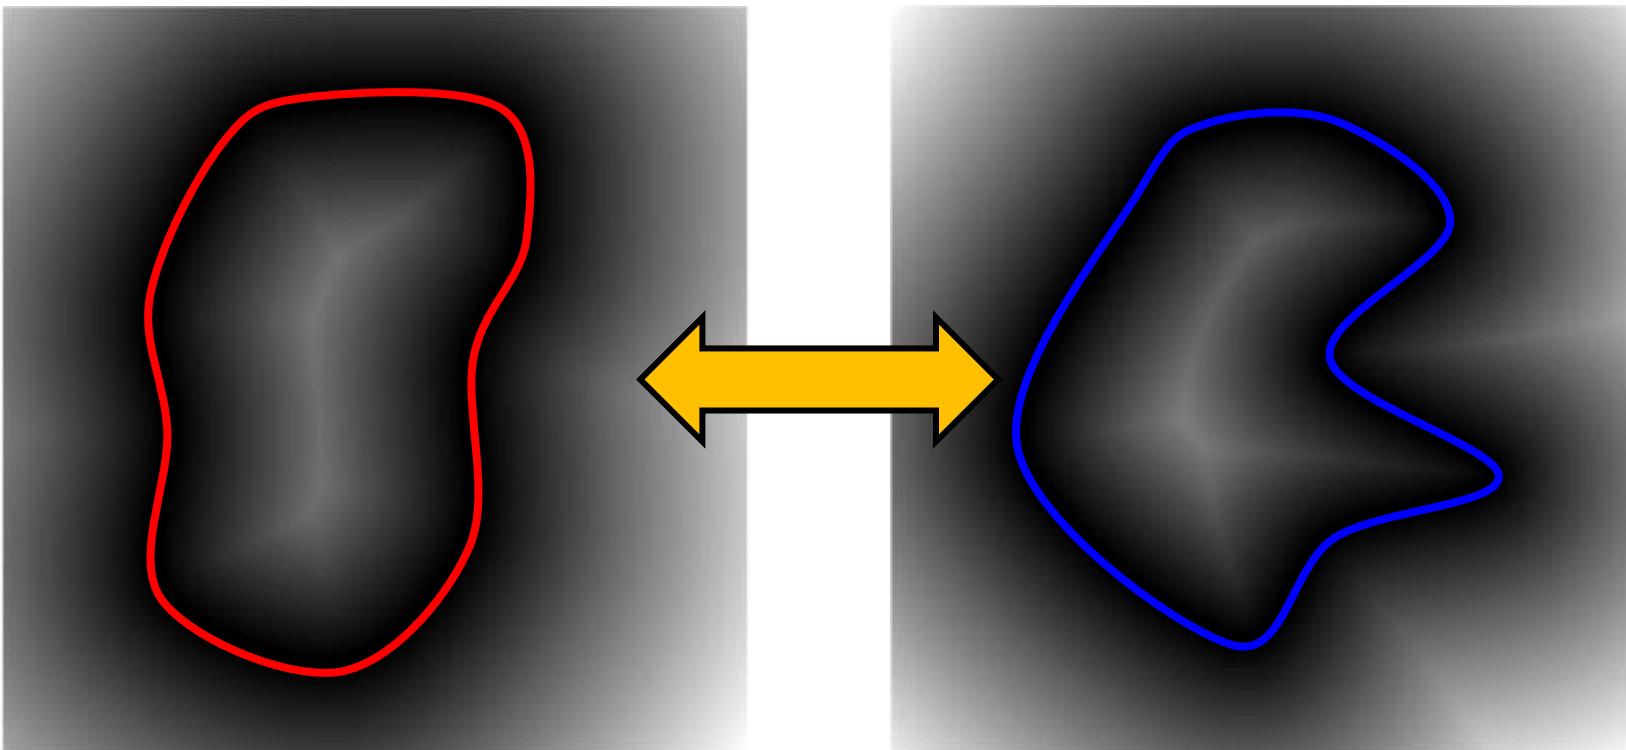
50-Pixels Distance Surface

Surface distance



100-Pixels Distance Surface

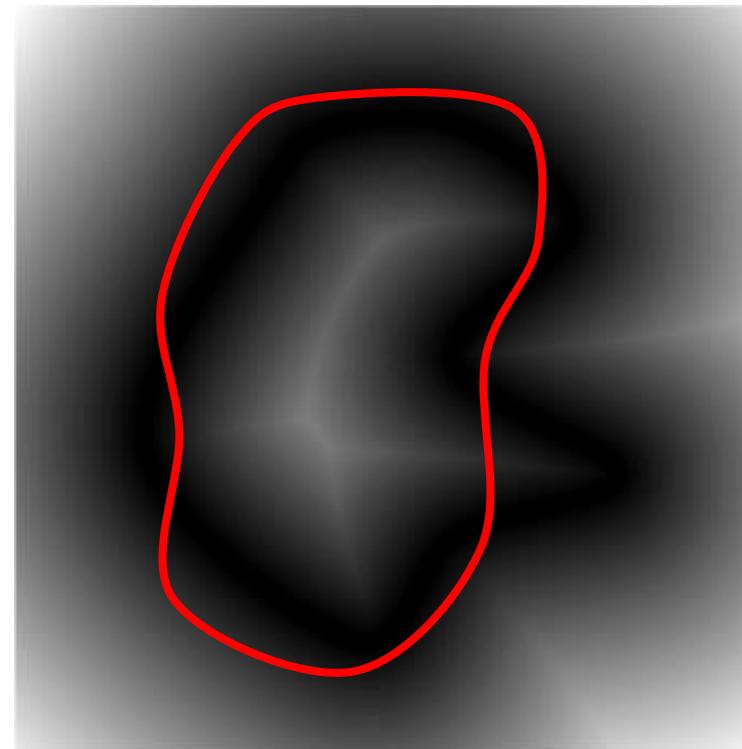
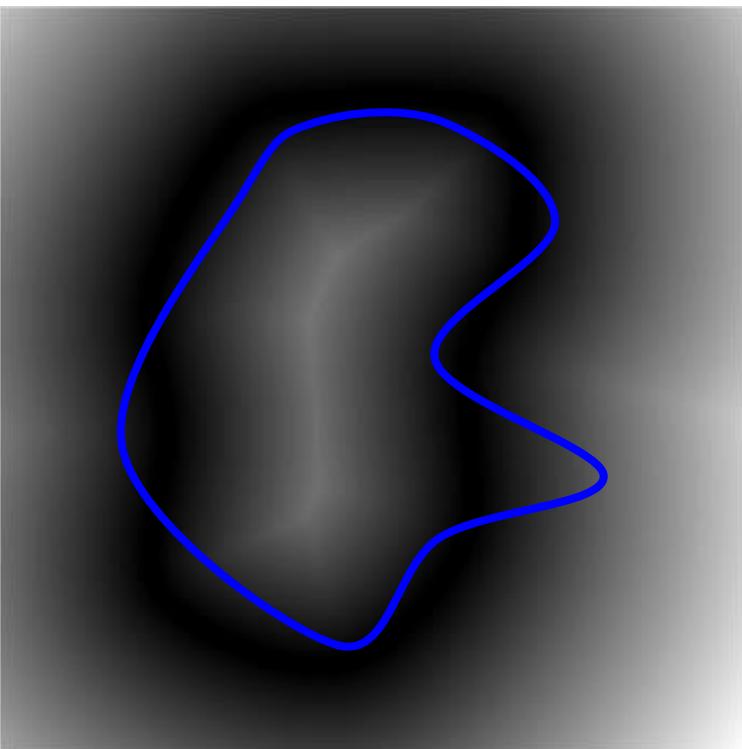
Surface distance



Euclidean Distance Maps

Surface distance

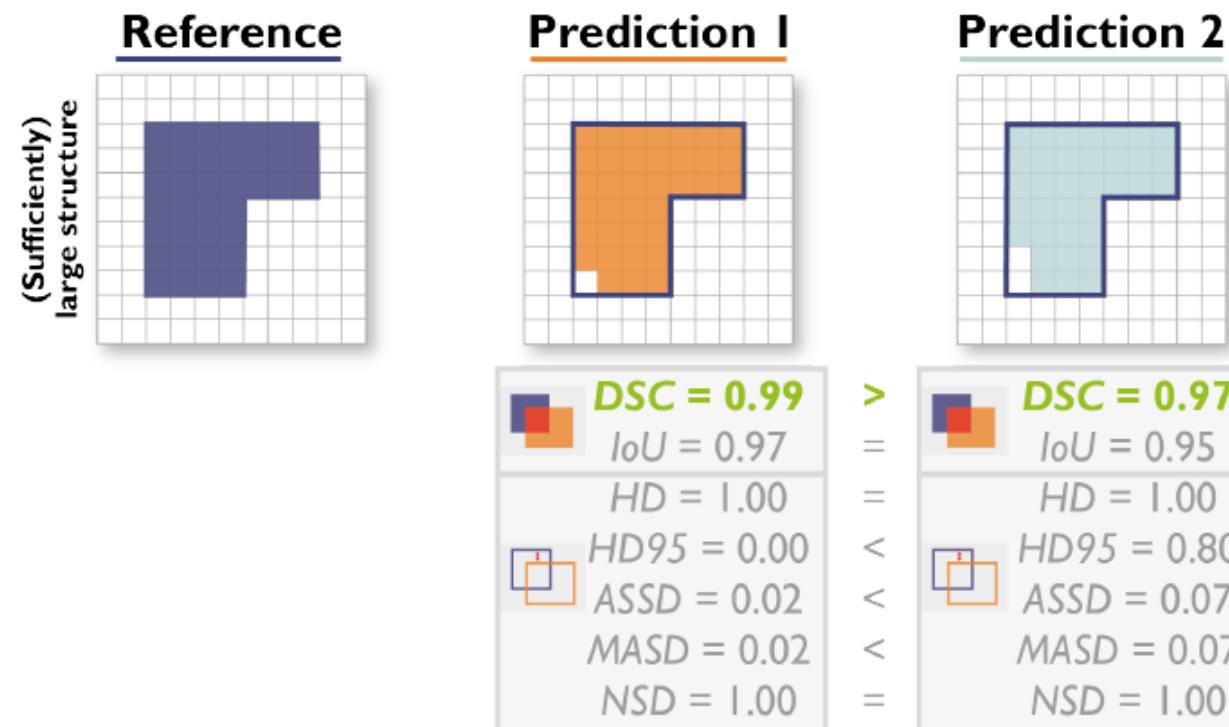
Sum up distances along pixels on the boundaries of one contour overlaid on the distance map of the other



Euclidean Distance Maps

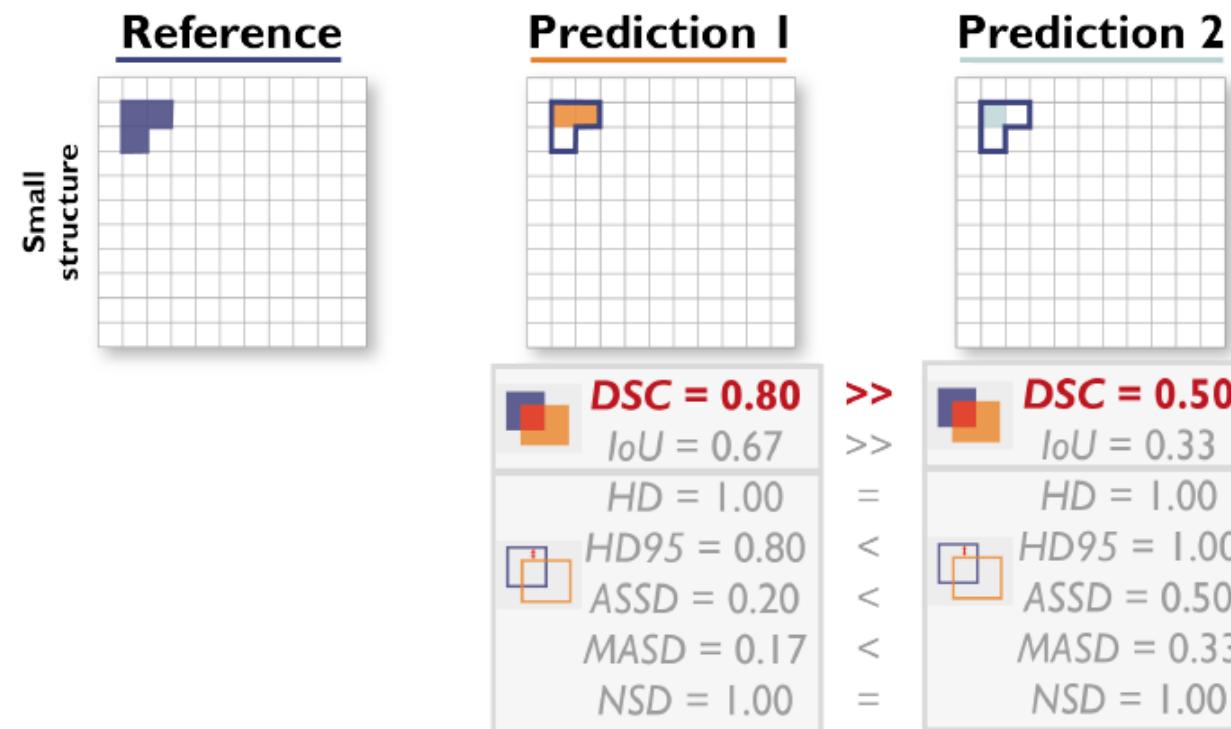
Pitfalls in segmentation evaluation

Effect of structure size



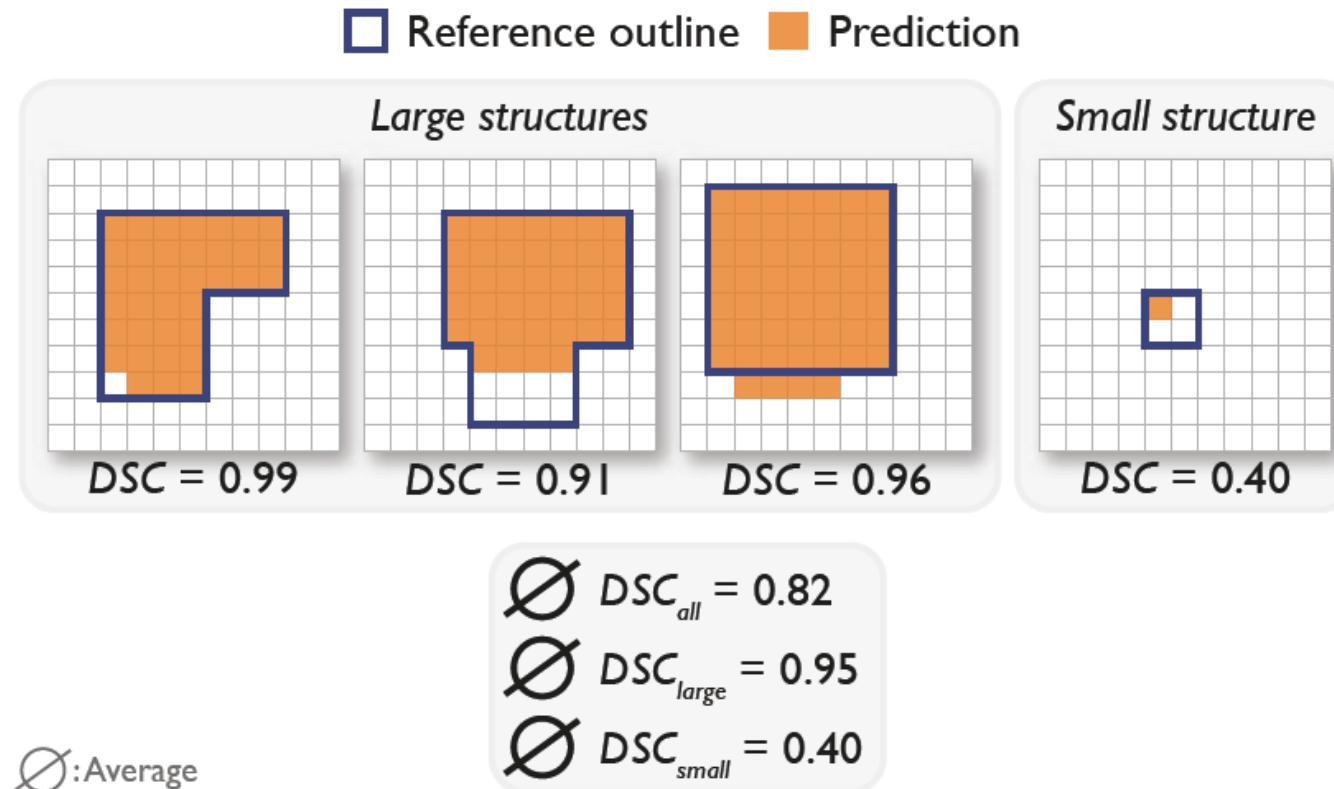
Pitfalls in segmentation evaluation

Effect of structure size



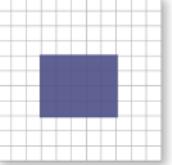
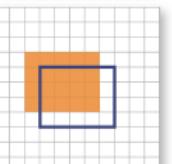
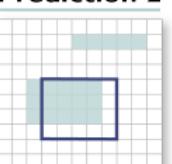
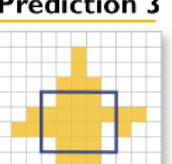
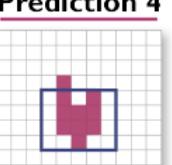
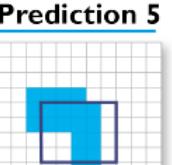
Pitfalls in segmentation evaluation

Effect of structure size



Pitfalls in segmentation

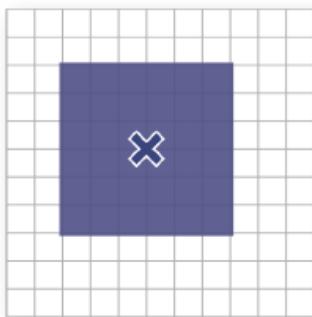
Effect of structure shape

Reference	DSC	IoU	HD	HD95	ASSD	MASD	NSD
							
Prediction 1	DSC = 0.6	IoU = 0.4	HD = 1.4	HD95 = 1.3	ASSD = 0.9	MASD = 0.9	NSD = 1.0
							
Prediction 2	DSC = 0.6	IoU = 0.4	HD = 3.6	HD95 = 3.1	ASSD = 1.0	MASD = 1.0	NSD = 0.7
							
Prediction 3	DSC = 0.6	IoU = 0.4	HD = 3.0	HD95 = 2.0	ASSD = 0.8	MASD = 0.7	NSD = 0.8
							
Prediction 4	DSC = 0.6	IoU = 0.4	HD = 2.2	HD95 = 2.0	ASSD = 0.8	MASD = 0.7	NSD = 0.8
							
Prediction 5	DSC = 0.6	IoU = 0.4	HD = 2.0	HD95 = 1.2	ASSD = 0.8	MASD = 0.8	NSD = 0.9
							

Pitfalls in segmentation evaluation

Effect of spatial alignment

Reference



Center = (4.5, 4.5) = **Center = (4.5, 4.5)** ≠ **Center = (3.5, 3.5)**

DSC = 0.62

IoU = 0.44

HD = 1.41

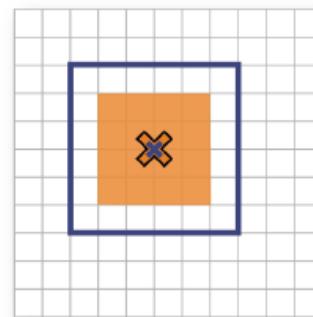
HD95 = 1.41

ASSD = 1.05

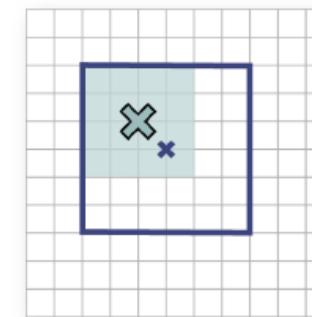
MASD = 1.04

NSD = 0.92

Prediction I

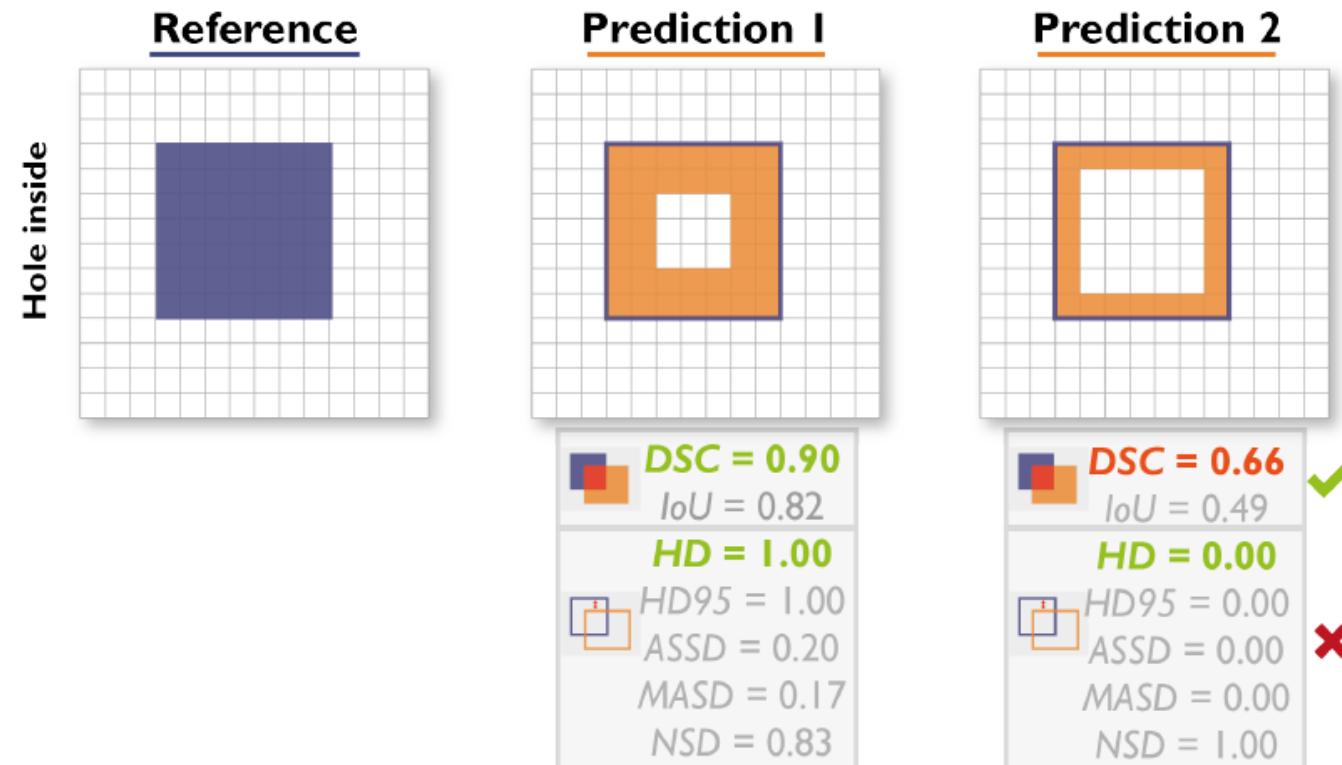


Prediction 2



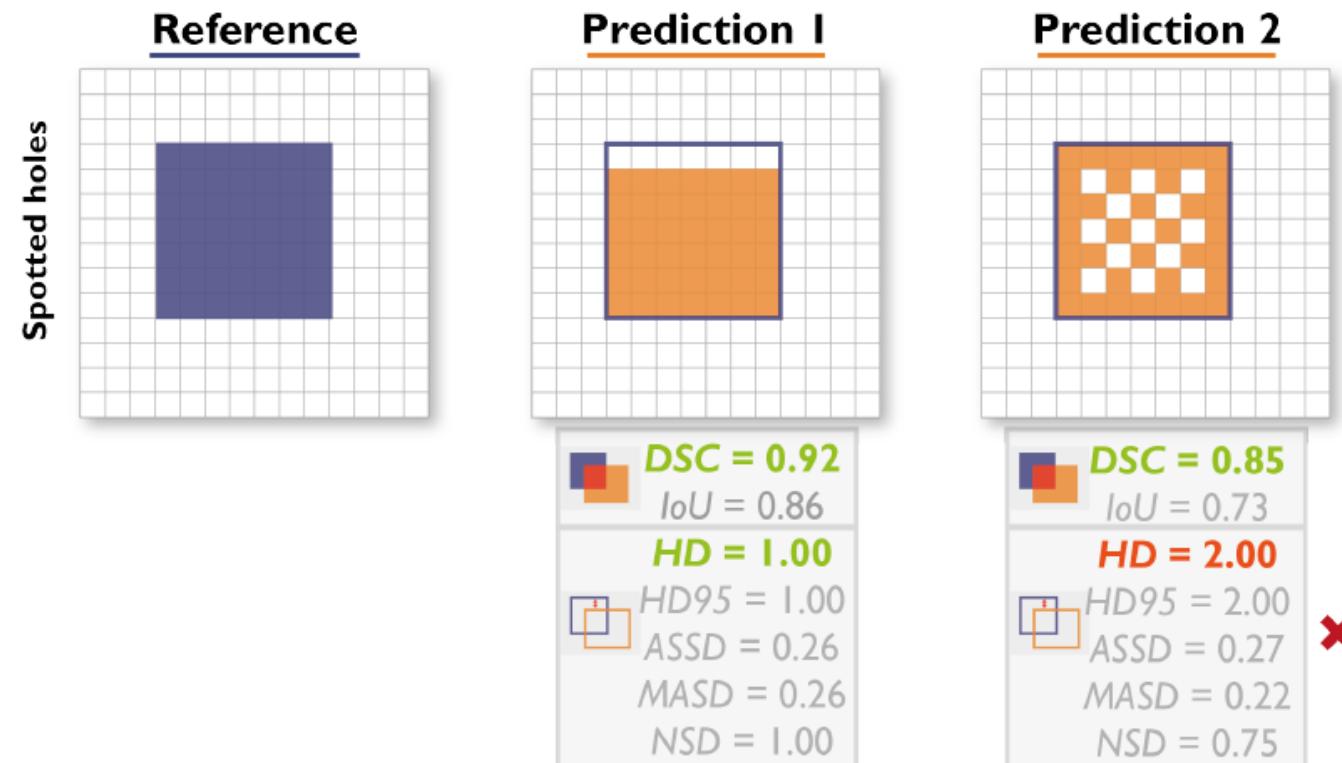
Pitfalls in segmentation evaluation

Effect of “holes”



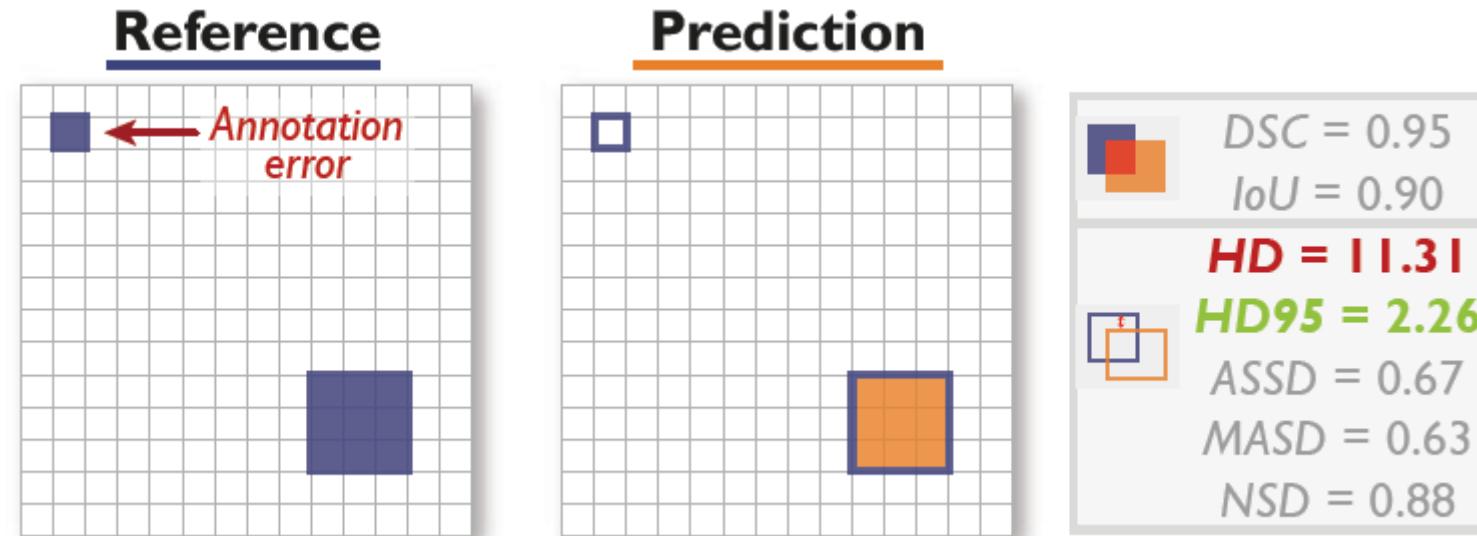
Pitfalls in segmentation evaluation

Effect of “holes”



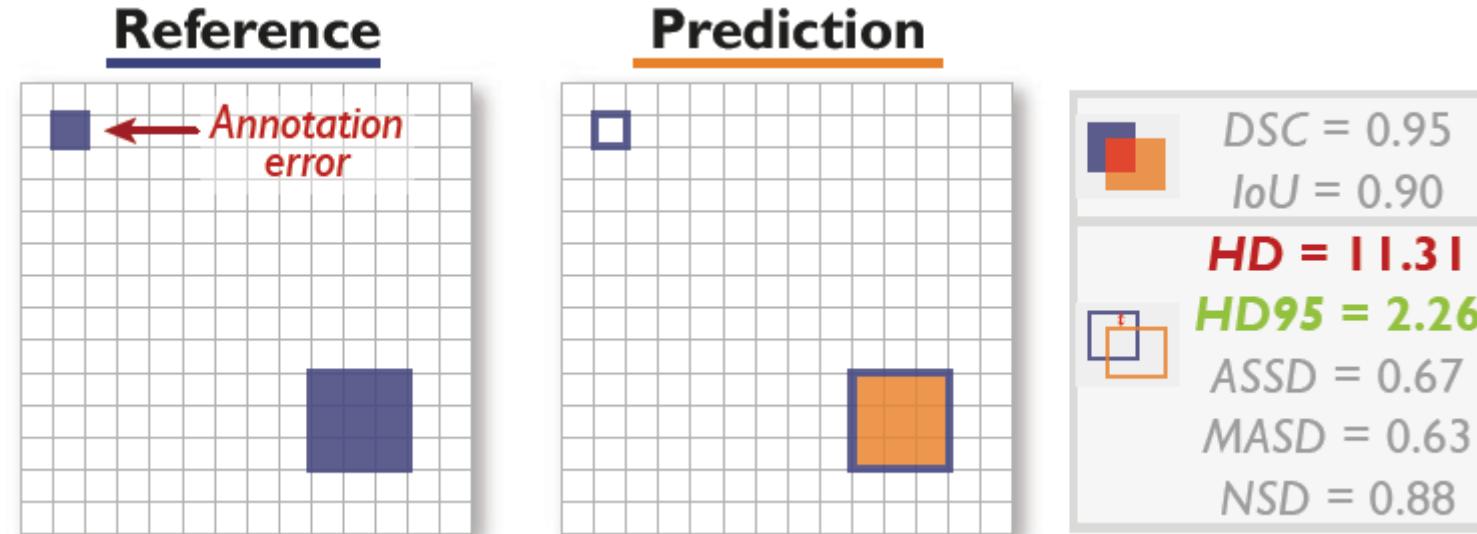
Pitfalls in segmentation evaluation

Effect of annotation noise



Pitfalls in segmentation evaluation

Effect of annotation noise



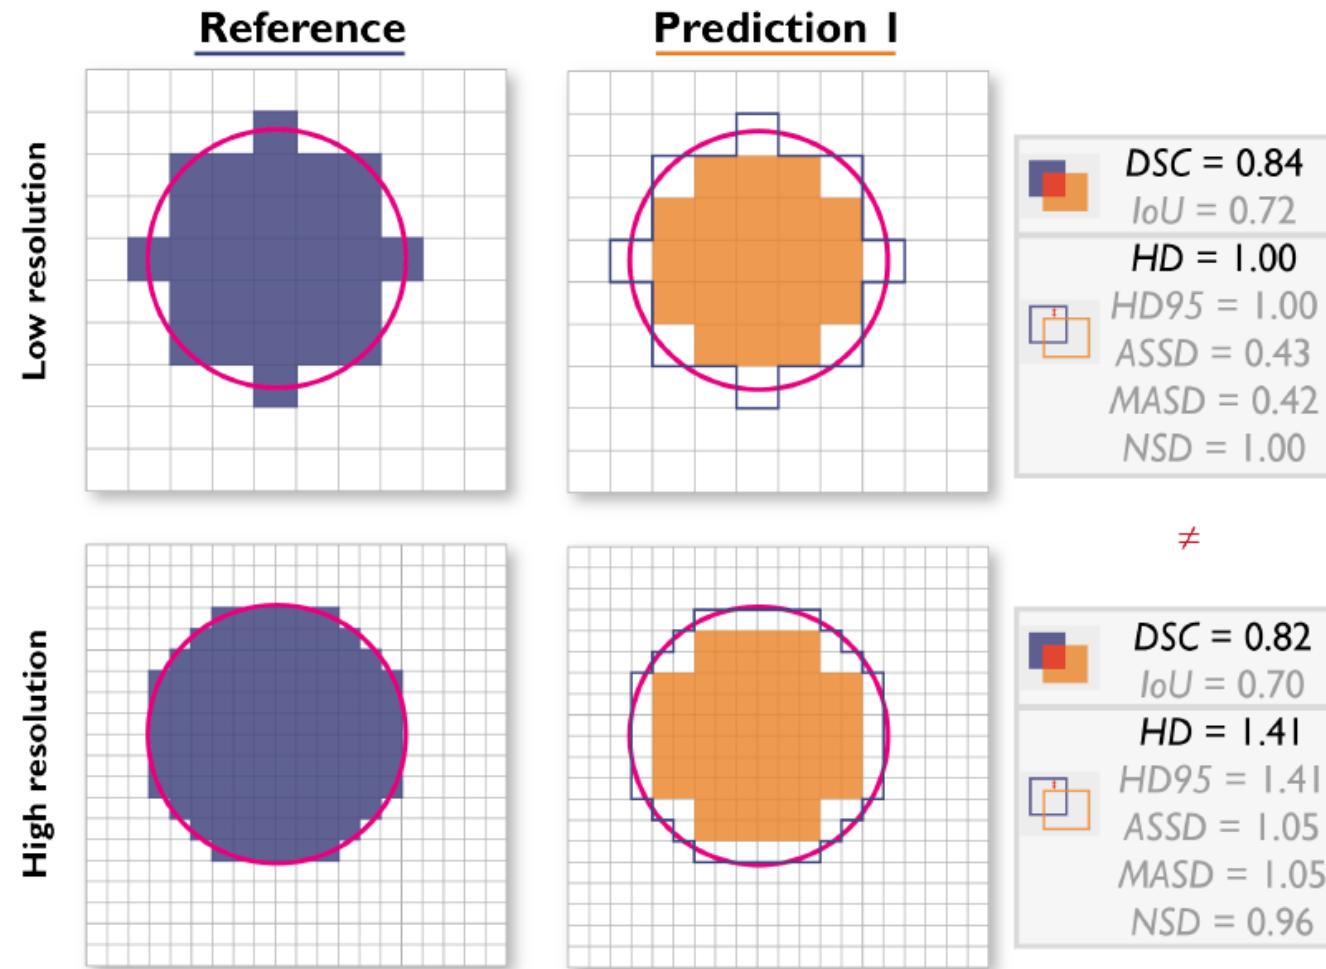
Pitfalls in segmentation evaluation

Effect of “empty” labelmaps

	<u>Reference</u>	<u>Prediction</u>	
Reference empty			TP = 0, FN = 0, FP = 4 $DSC = IoU = Precision = NSD = 0$ Sensitivity = NaN (0/0) $HD = HD95 = ASSD = MASD = \text{NaN}$
Prediction empty			TP = 0, FN = 4, FP = 0 $DSC = IoU = Sensitivity = NSD = 0$ Precision = NaN (0/0) $HD = HD95 = ASSD = MASD = \text{NaN}$
Reference and prediction empty			TP = 0, FN = 0, FP = 0 $DSC = IoU = \text{NaN (0/0)}$ Sensitivity = Precision = NaN (0/0) $HD = HD95 = \text{NaN}$ $ASSD = MASD = NSD = \text{NaN}$

Pitfalls in segmentation evaluation

Effect of resolution

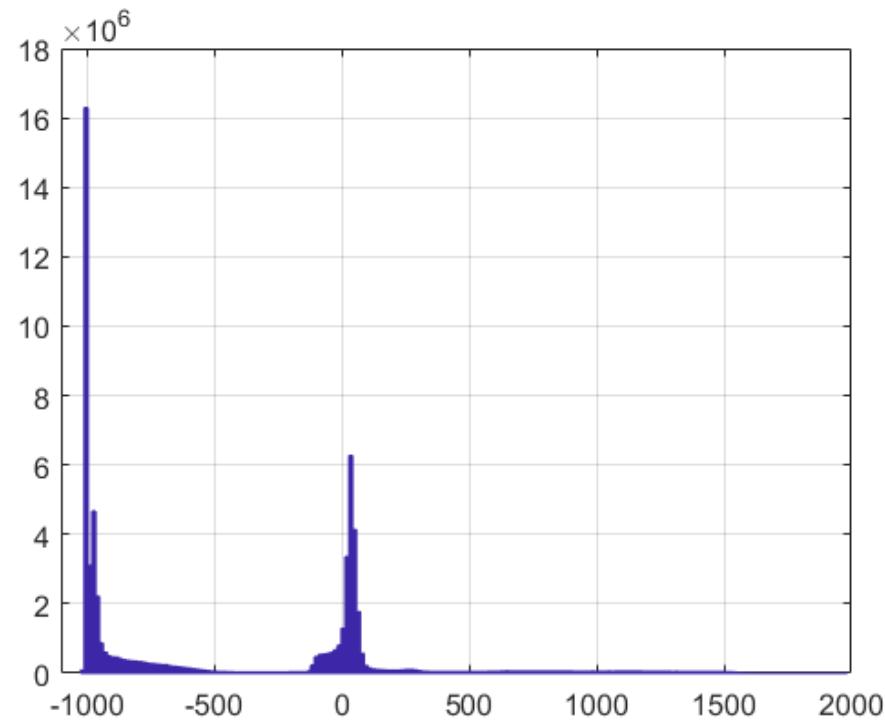


Segmentation methods

- Intensity-based segmentation
 - e.g., thresholding
- Region-based
 - e.g., region growing
- Graph-based segmentation
 - e.g., graph cuts
- Active contours
 - e.g., level sets
- Atlas-based segmentation
 - e.g., multi-atlas label propagation
- Learning-based segmentation
 - e.g., random forests, **convolutional neural networks**

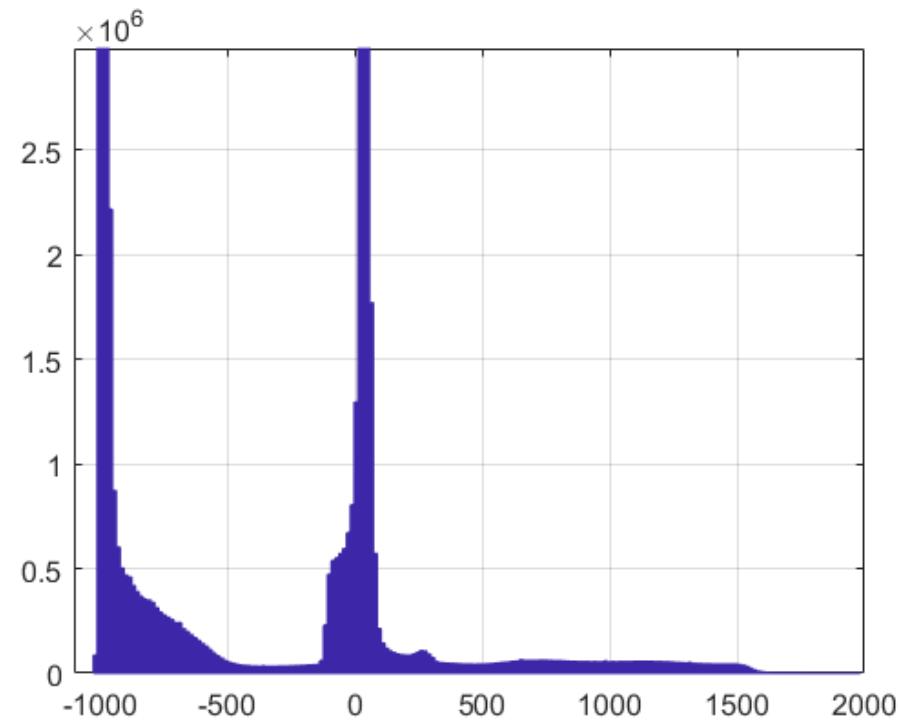
Simple thresholding

- Select a threshold on the intensity range



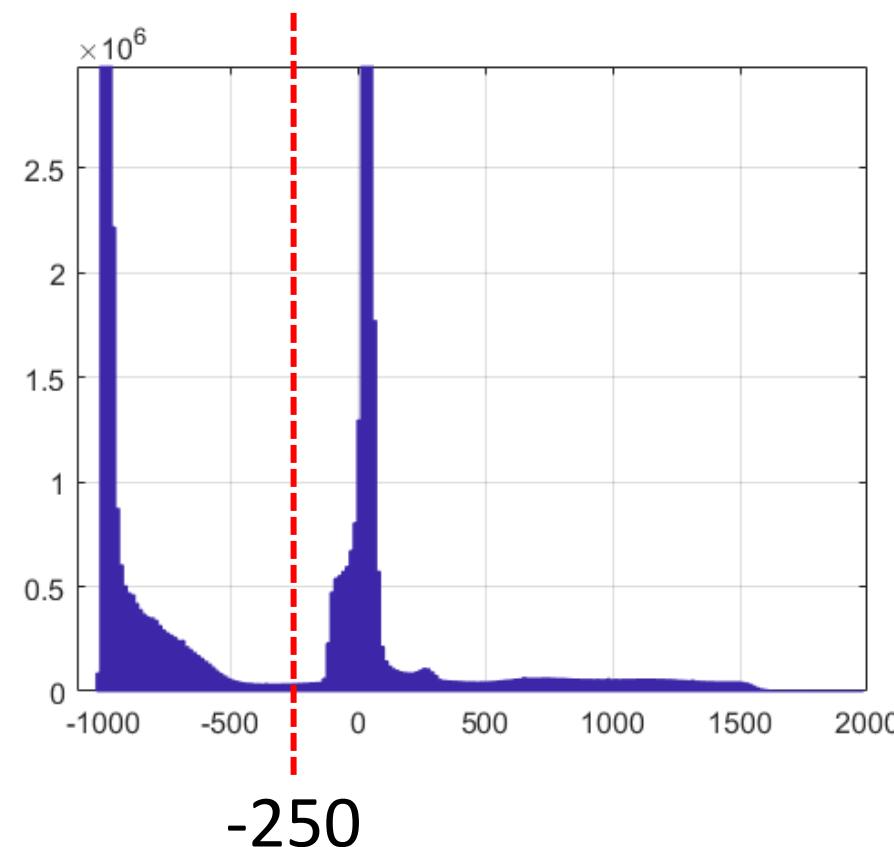
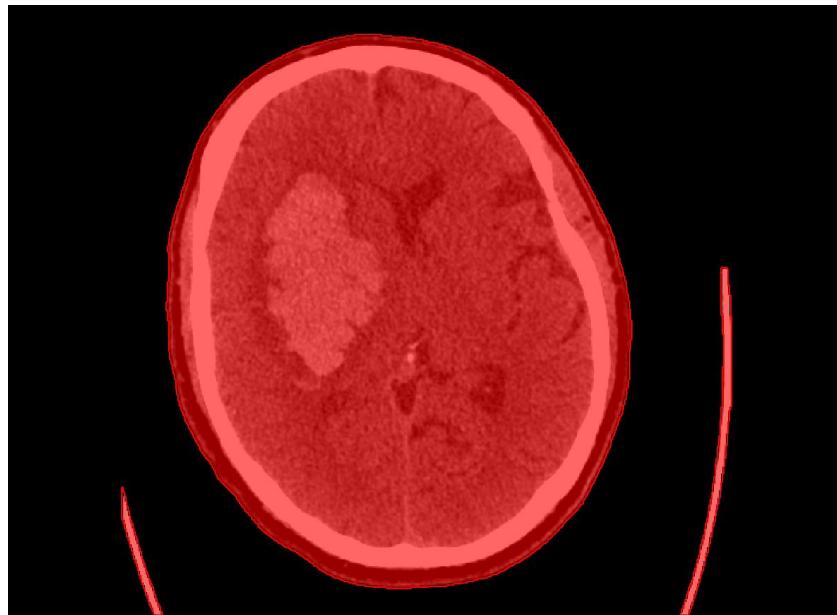
Simple thresholding

- Select a threshold on the intensity range



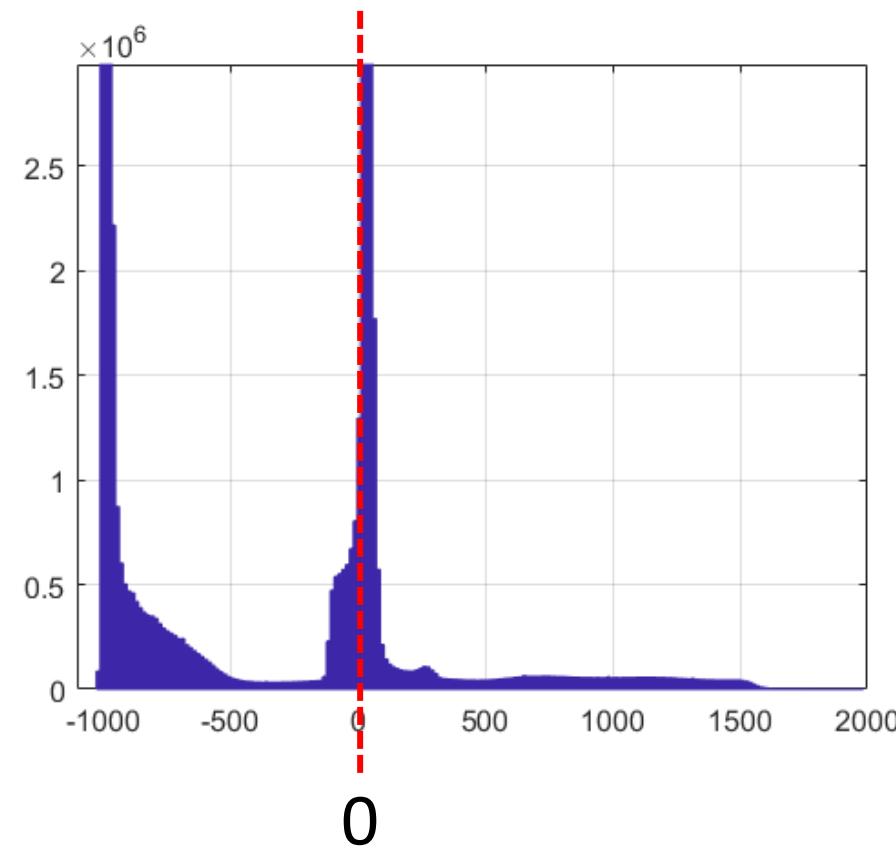
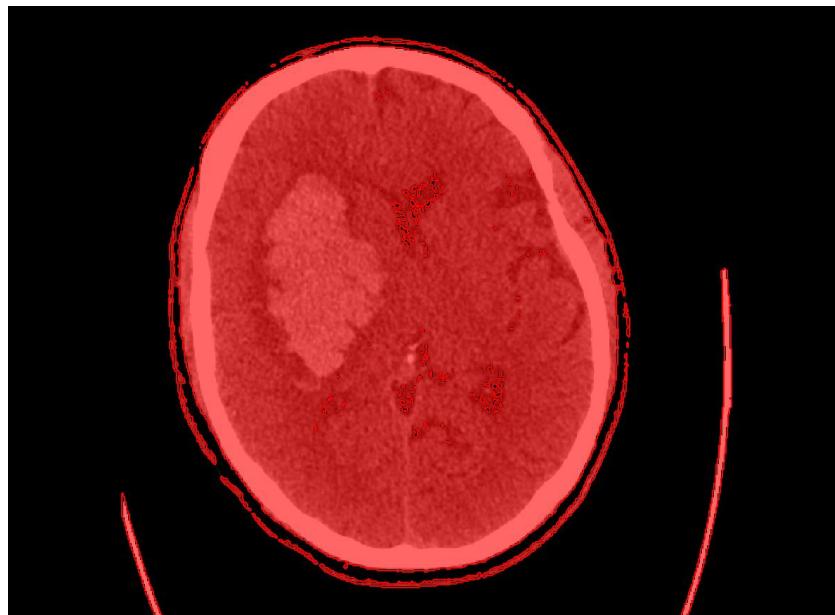
Simple thresholding

- Select a threshold on the intensity range



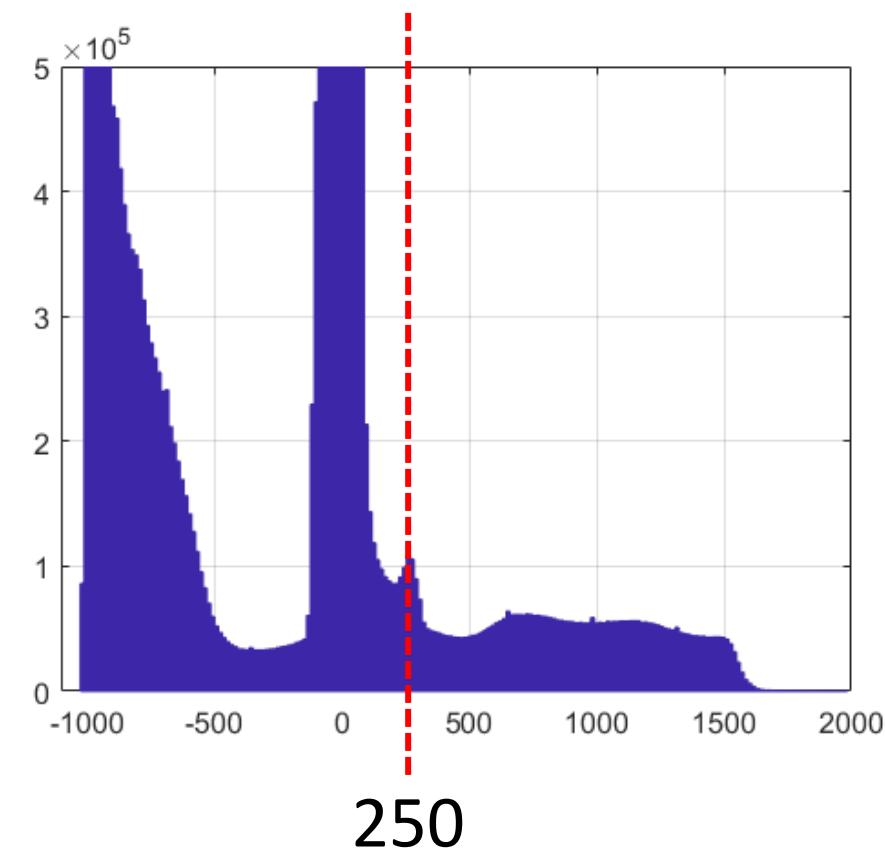
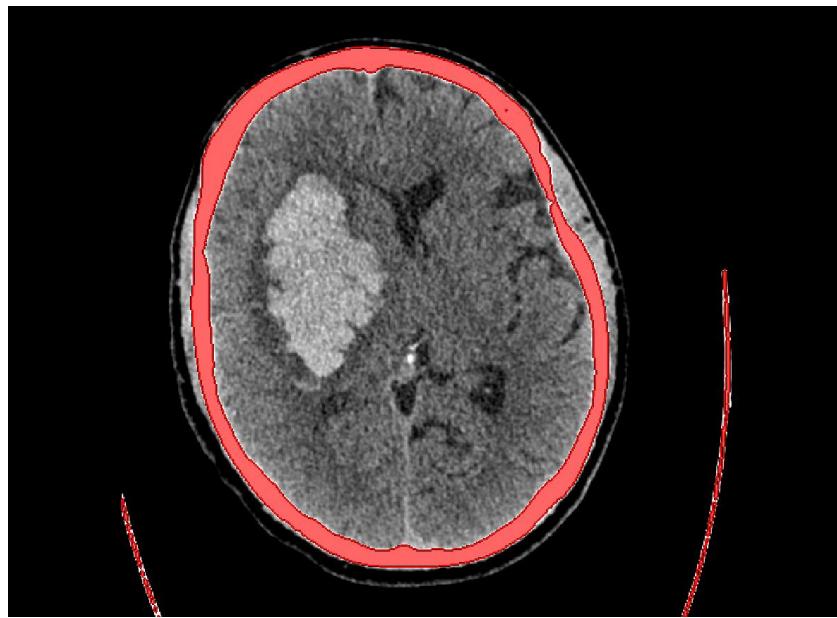
Simple thresholding

- Select a threshold on the intensity range



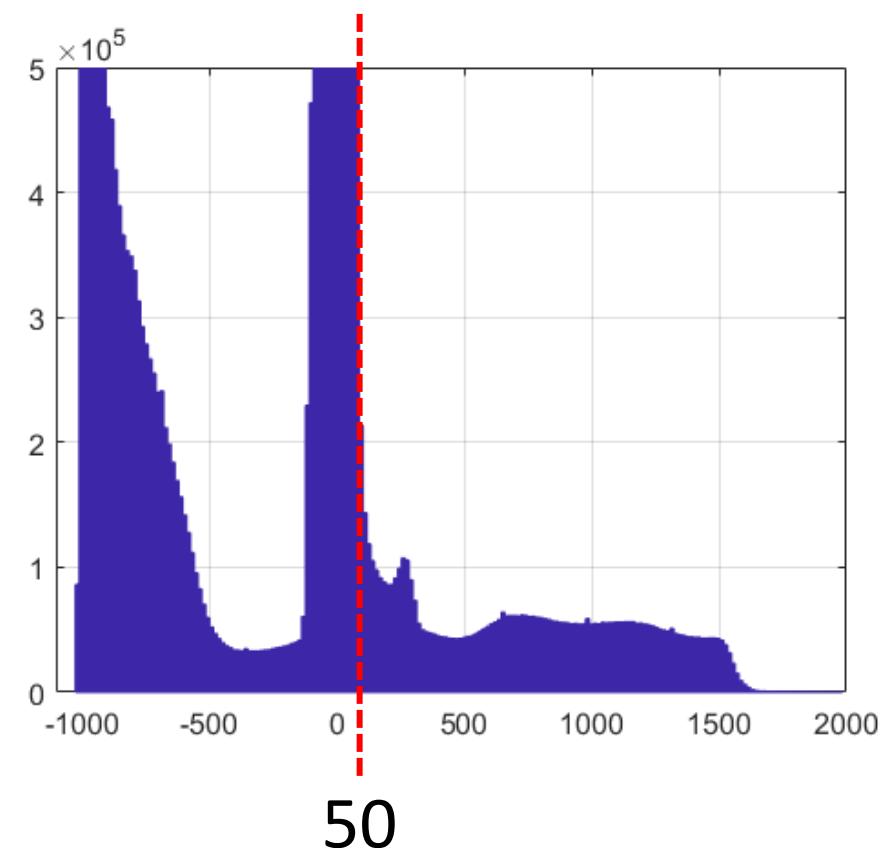
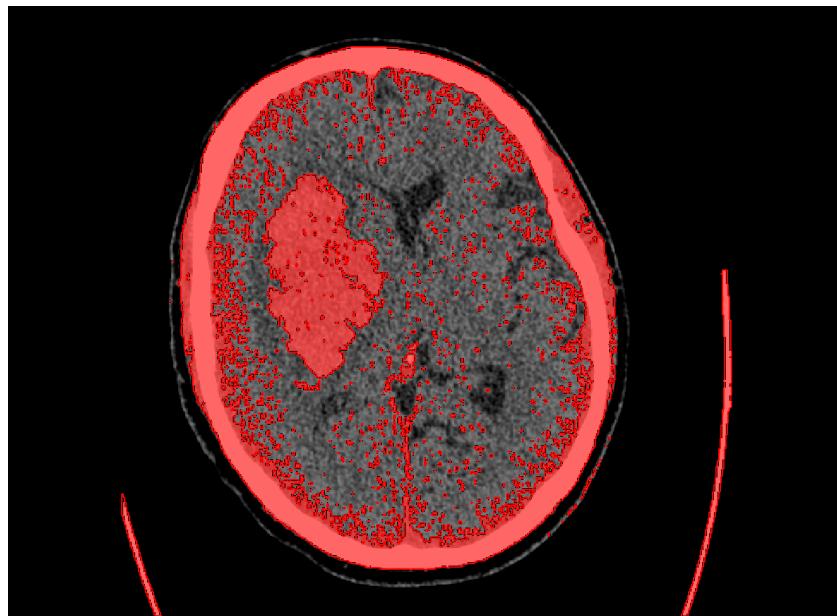
Simple thresholding

- Select a threshold on the intensity range



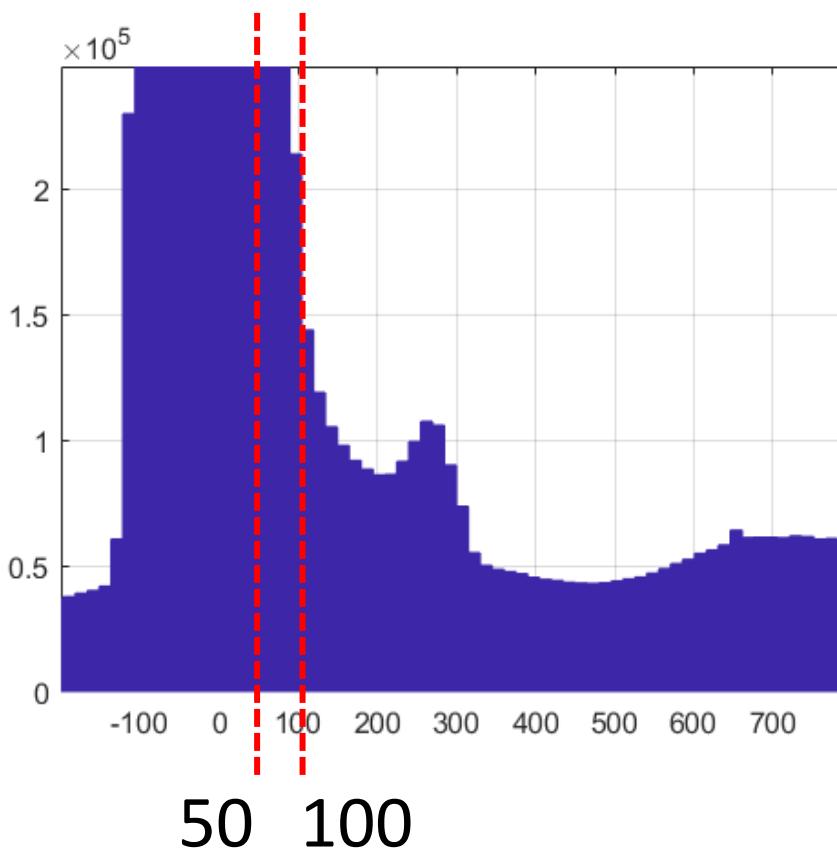
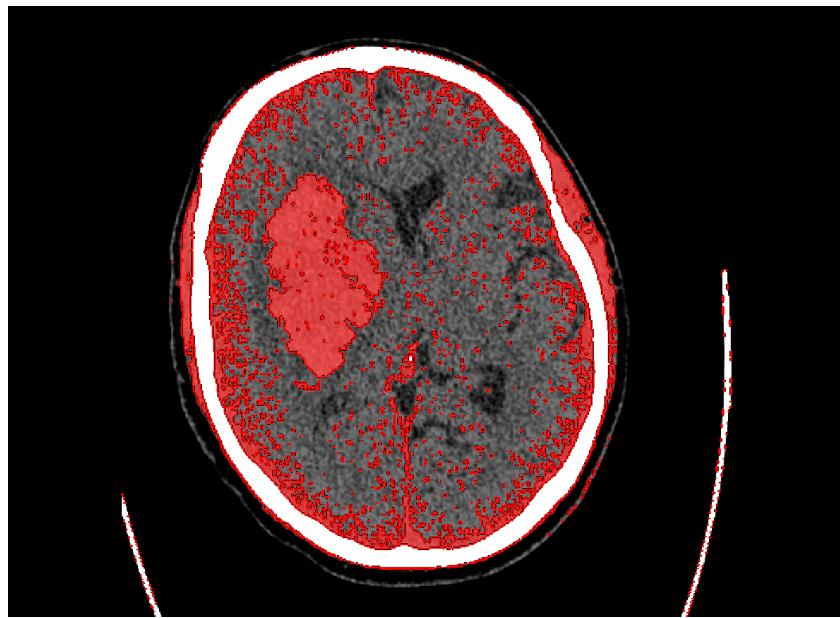
Simple thresholding

- Select a threshold on the intensity range



UL thresholding

- Select a lower and upper threshold



Thresholding

Advantages

- simple
- fast

Disadvantages

- regions must be homogeneous and distinct
- difficulty in finding consistent thresholds across images
- leakages, isolated pixels and ‘rough’ boundaries likely

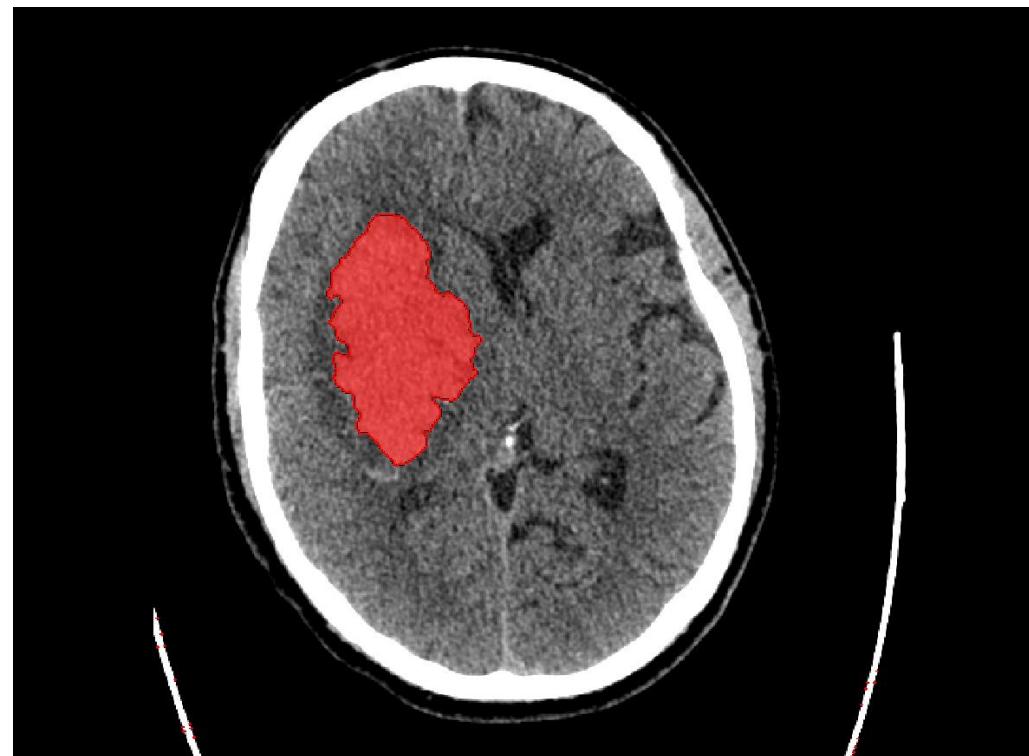
Region growing

- Start from (user selected) seed point(s), and grow a region according to an intensity threshold



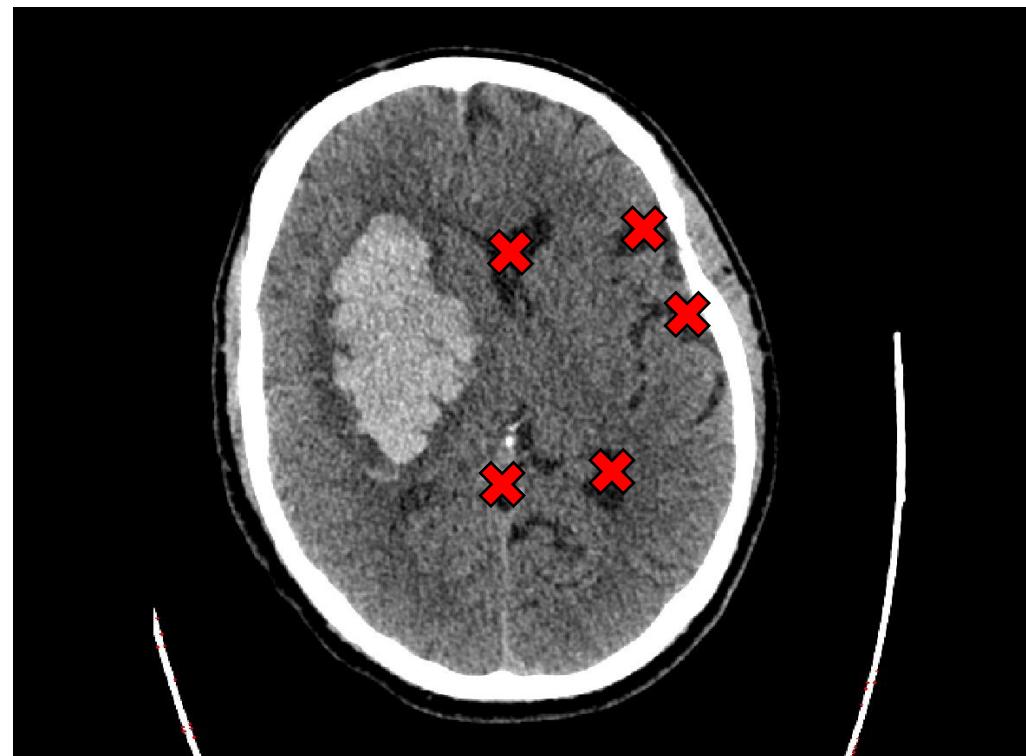
Region growing

- Start from (user selected) seed point(s), and grow a region according to an intensity threshold



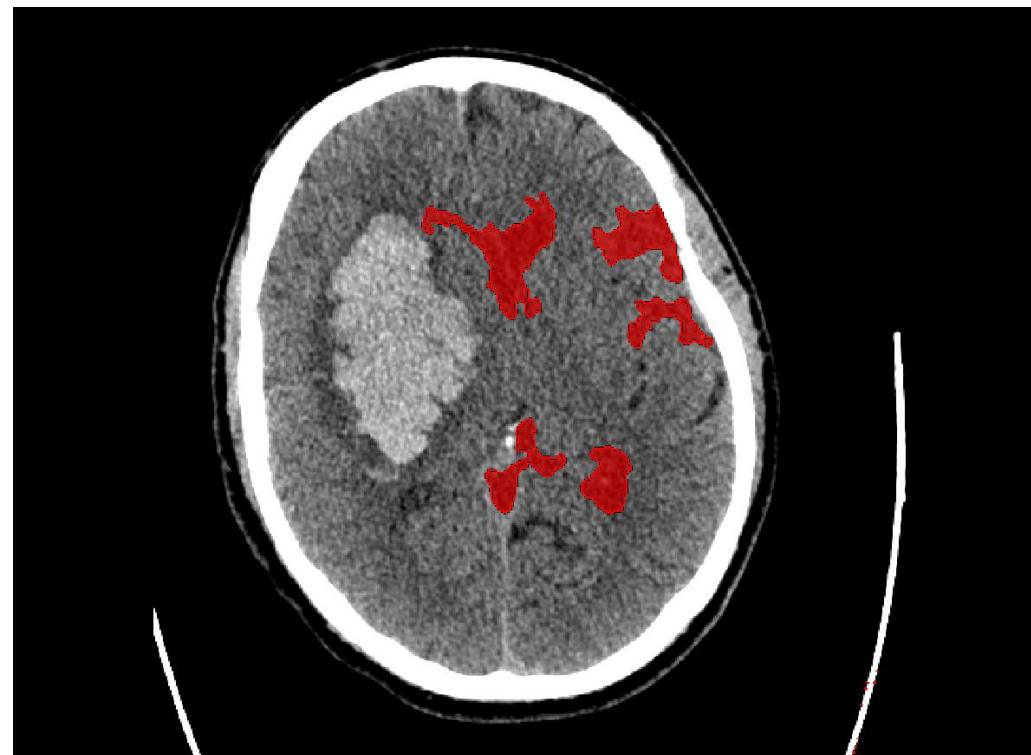
Region growing

- Start from (user selected) seed point(s), and grow a region according to an intensity threshold



Region growing

- Start from (user selected) seed point(s), and grow a region according to an intensity threshold



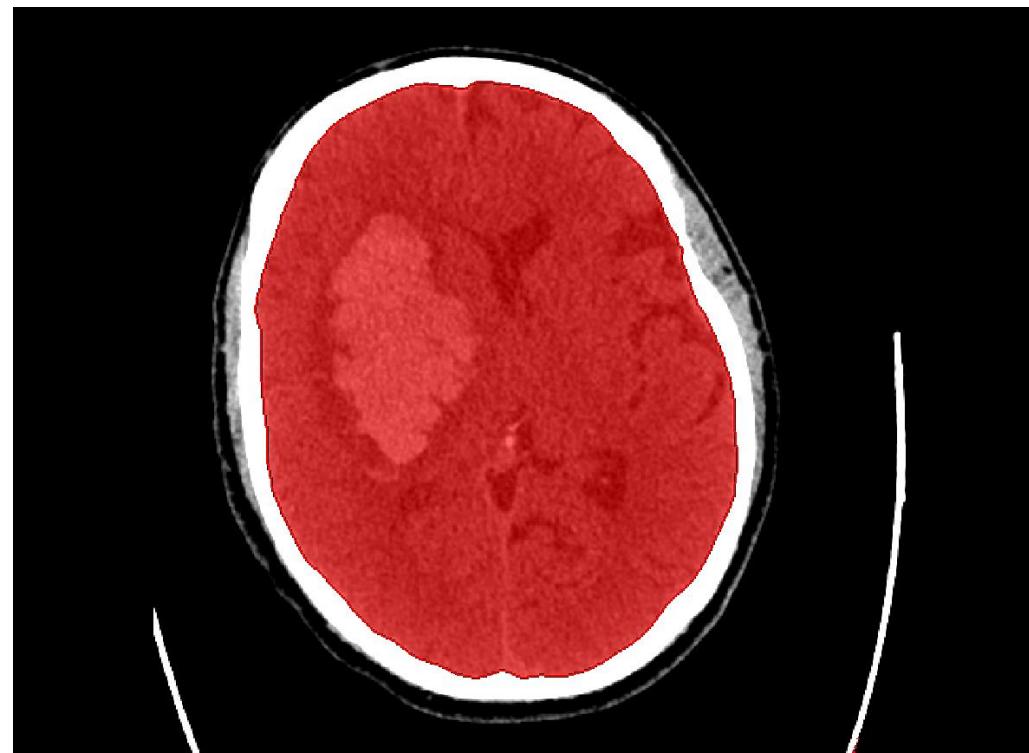
Region growing

- Start from (user selected) seed point(s), and grow a region according to an intensity threshold



Region growing

- Start from (user selected) seed point(s), and grow a region according to an intensity threshold



Region growing

Advantages

- relatively fast
- yields connected region (from a seed point)

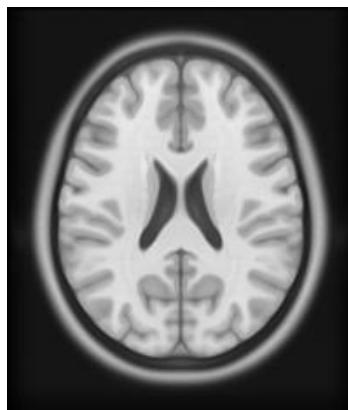
Disadvantages

- regions must be homogeneous
- leakages and ‘rough’ boundaries likely
- requires (user) input for seed points

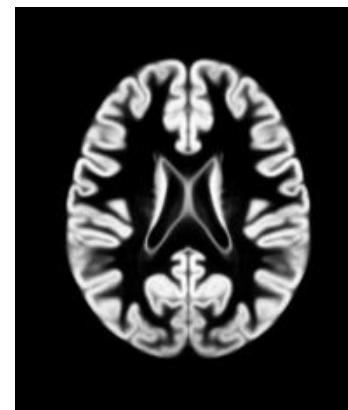
Atlas-based segmentation

What is an atlas?

MNI ICBM 152 Nonlinear atlases (2009)



T1 MRI



grey matter



white matter



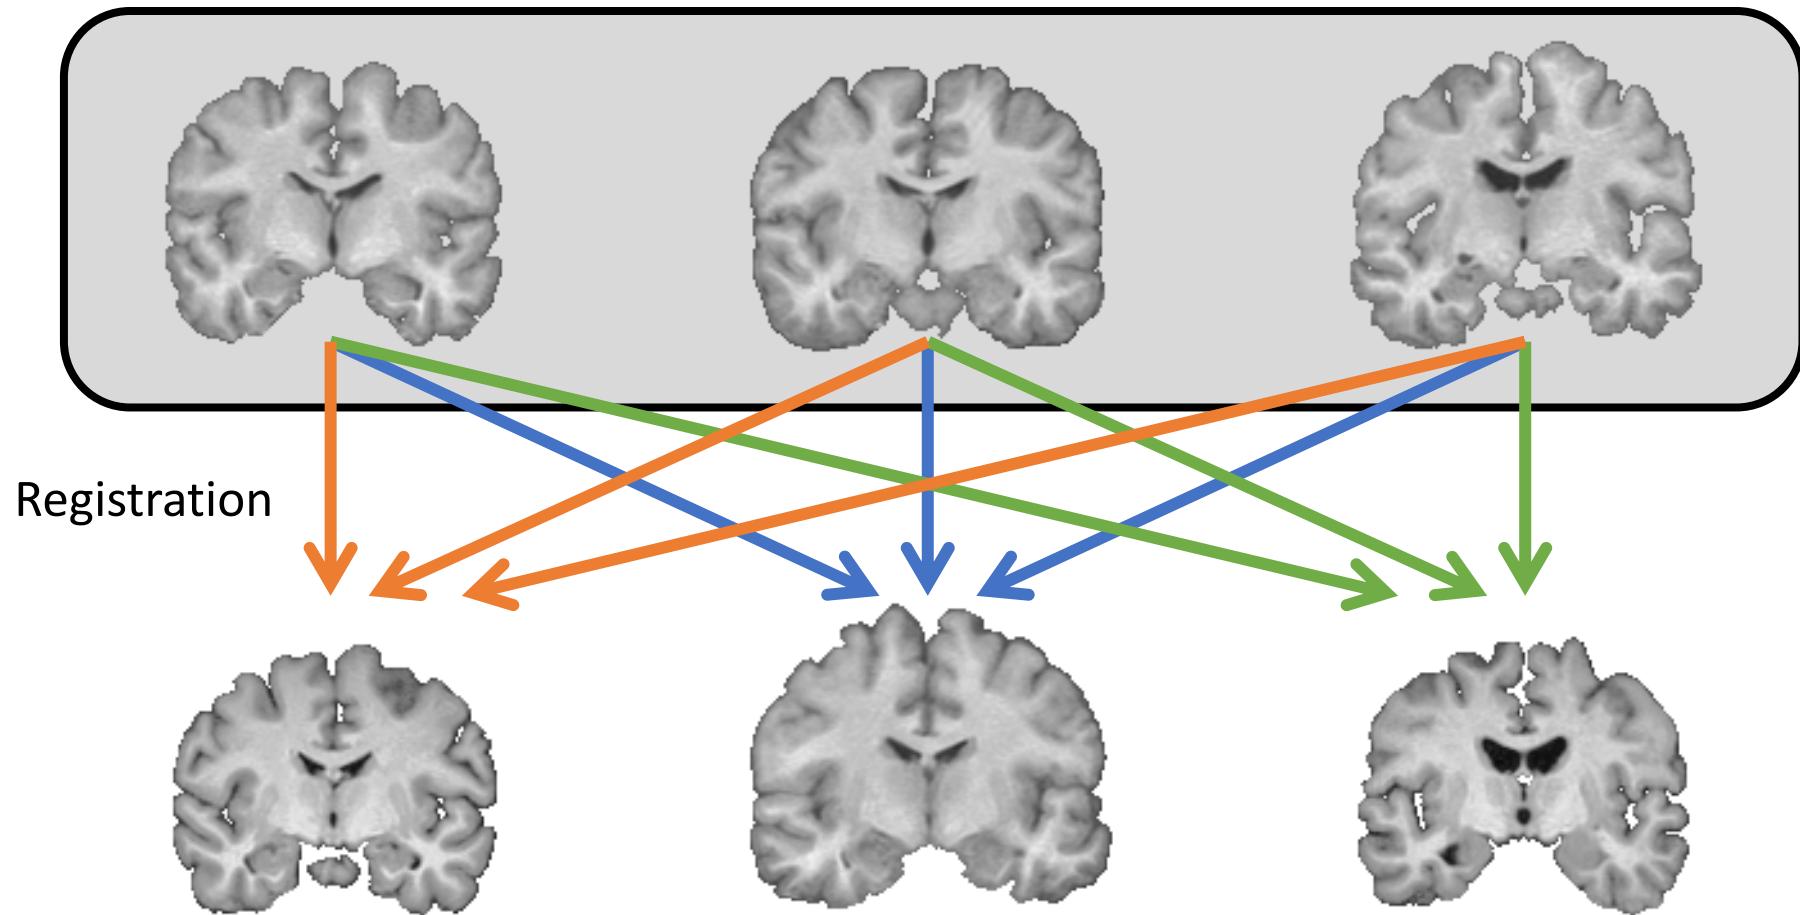
CSF

- A map or chart of the anatomy
- Atlases usually have
 - geometric information about points, curves or surfaces, or
 - label information about voxels (anatomical regions or function)
- Atlases are usually constructed from example data
 - single subjects
 - populations of subjects, e.g. by averaging to produce probabilistic atlases

Segmentation via registration

Label propagation

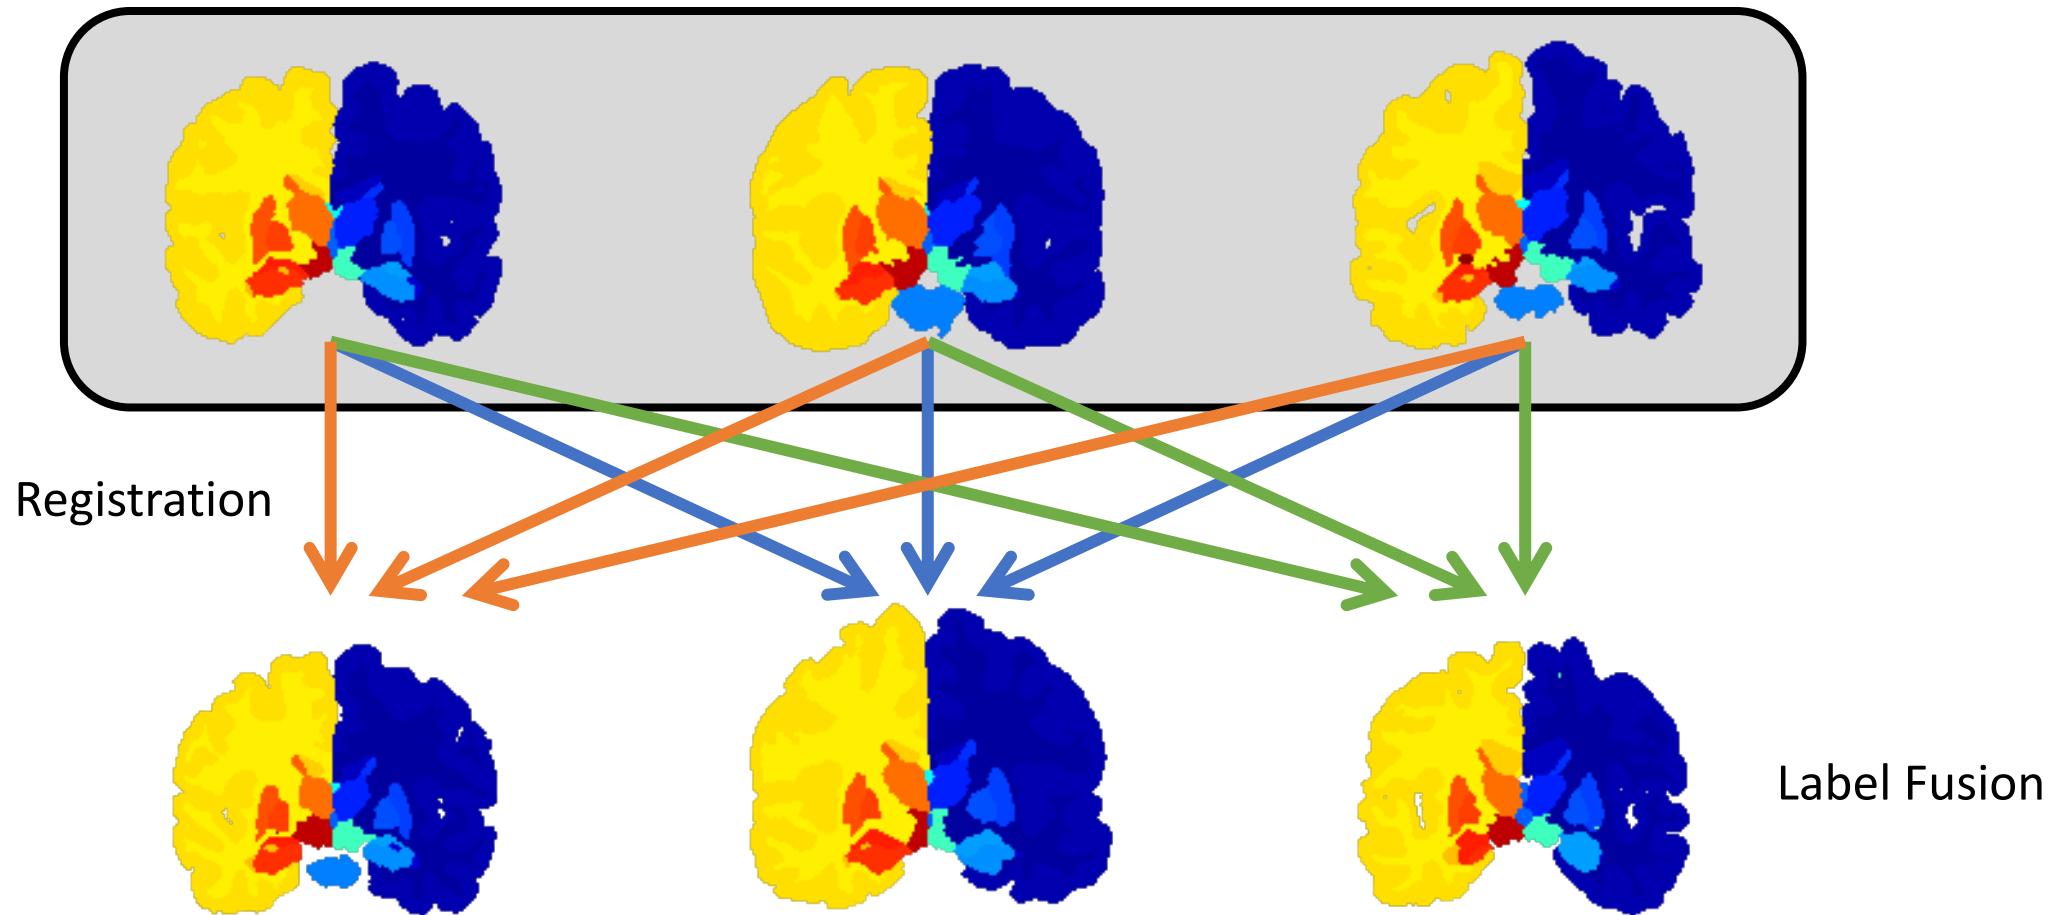
Database



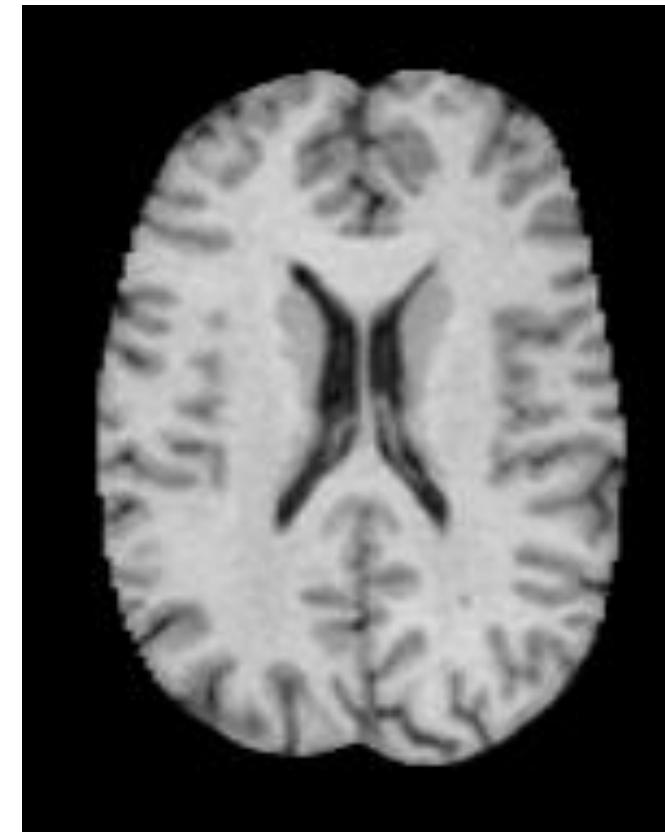
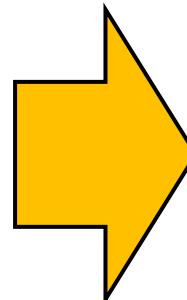
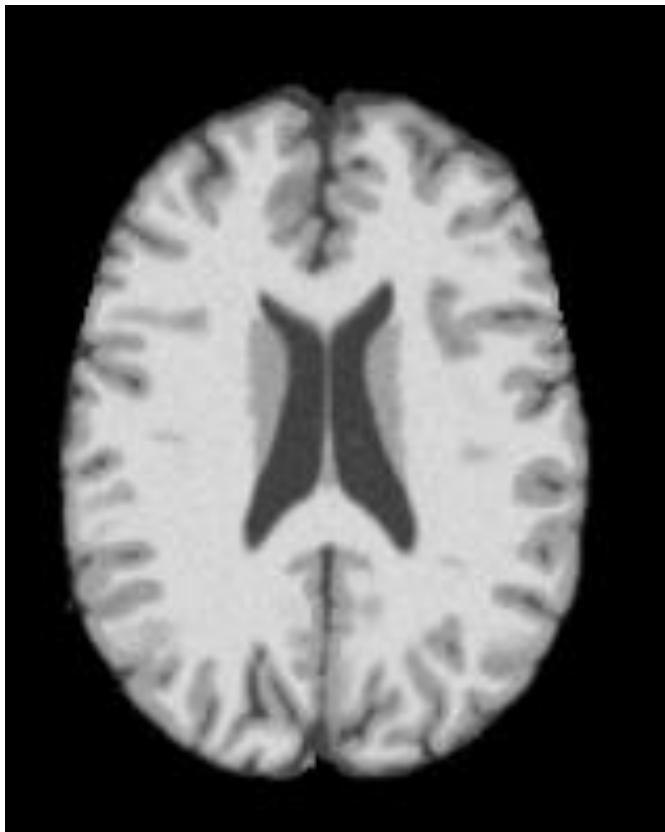
Segmentation via registration

Label propagation

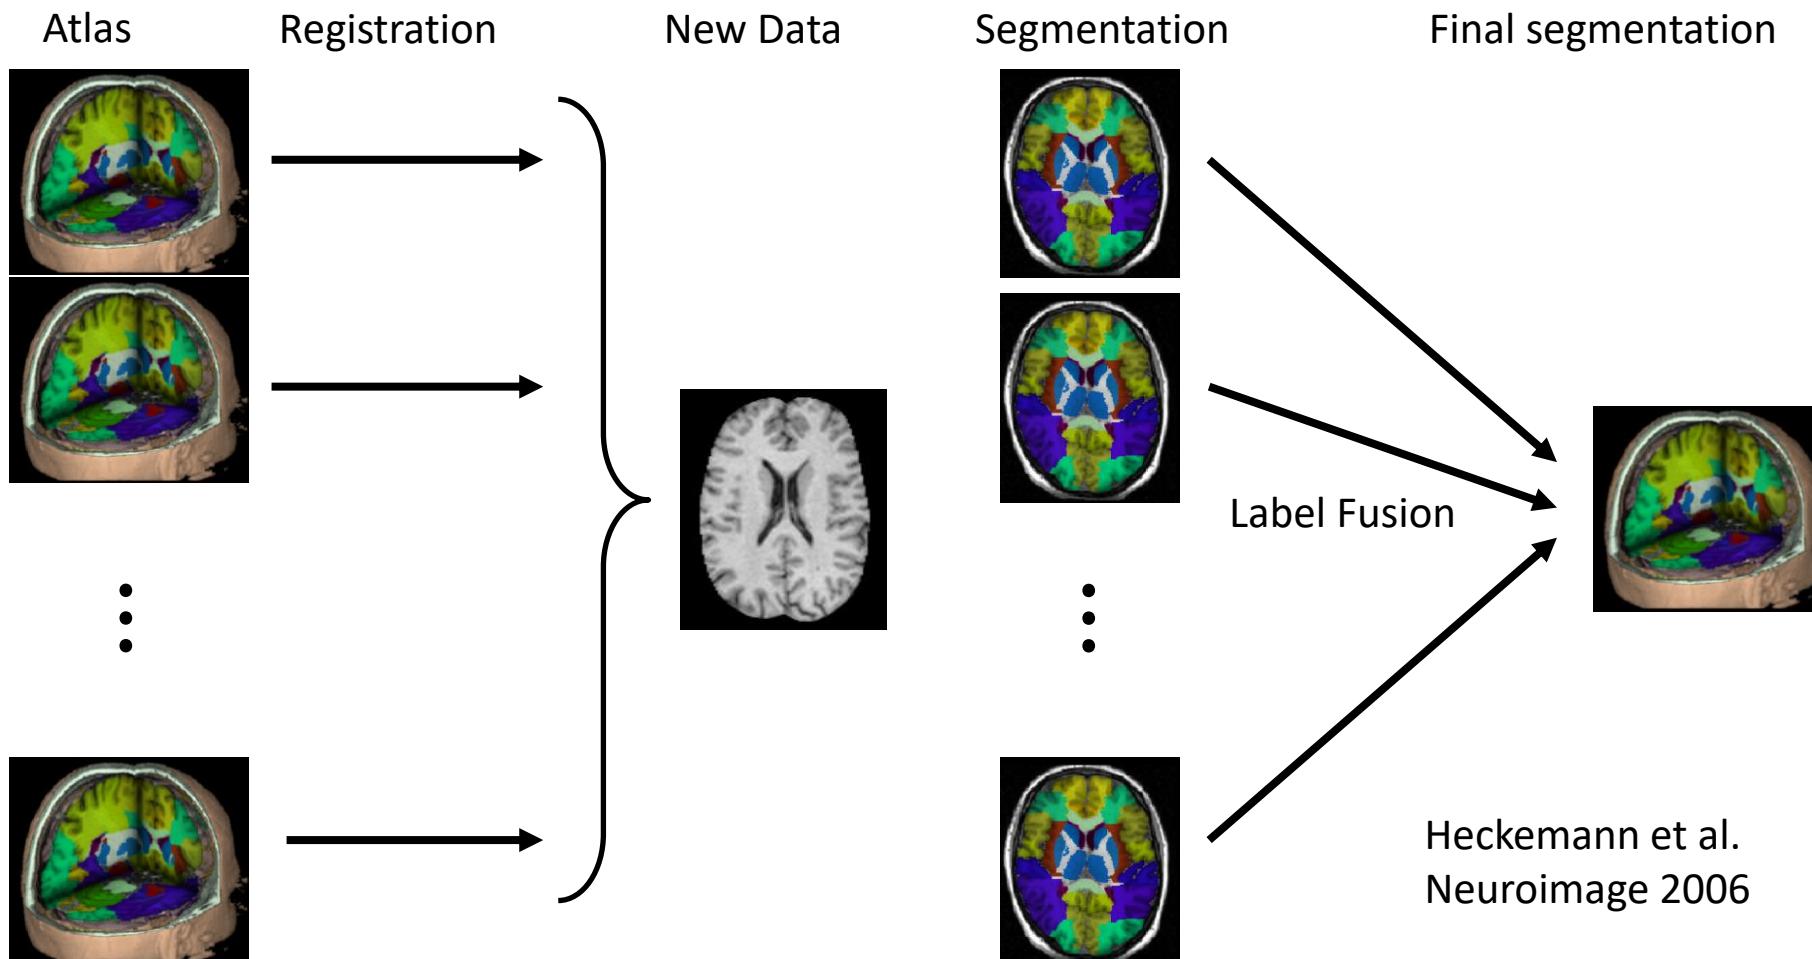
Database



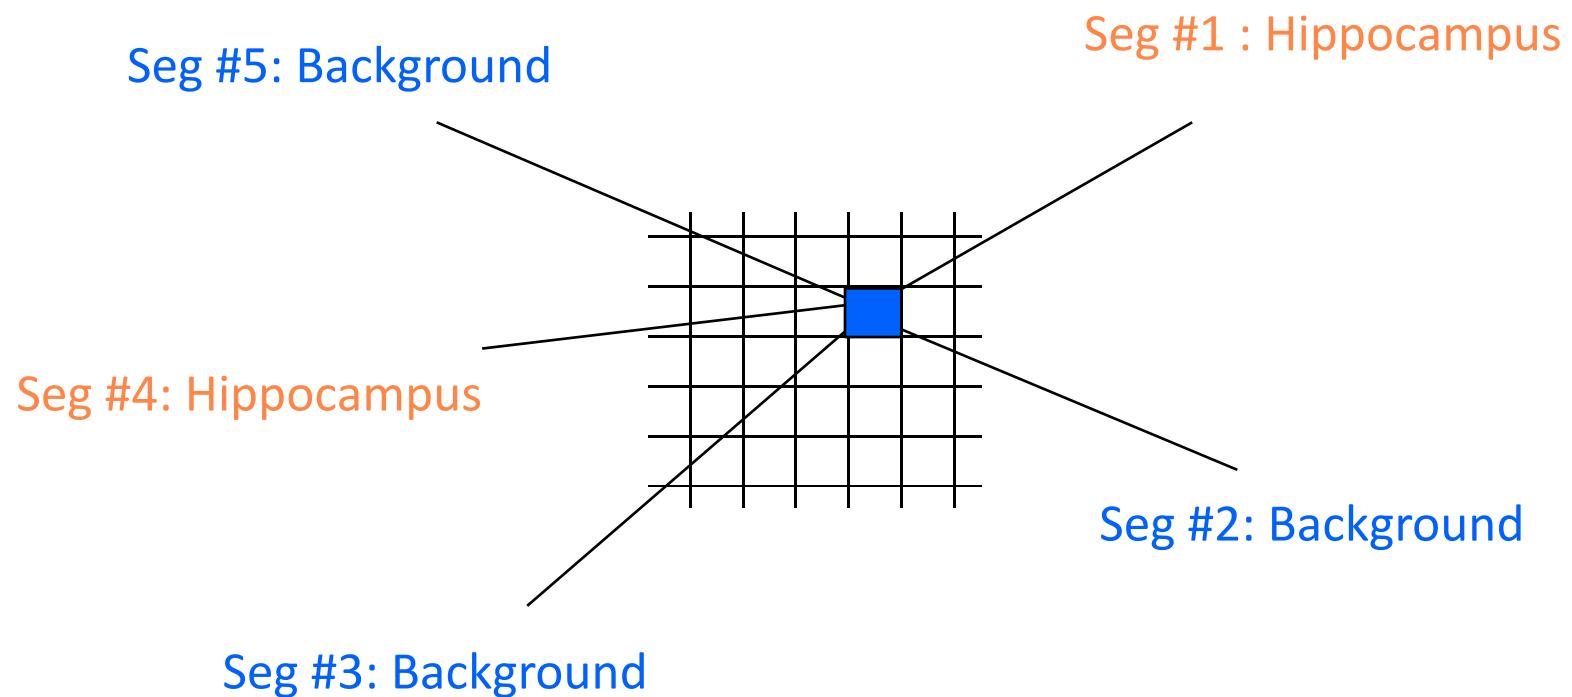
Segmentation via registration



Multi-atlas label propagation



Label fusion via majority voting



Multi-atlas label propagation

Advantages

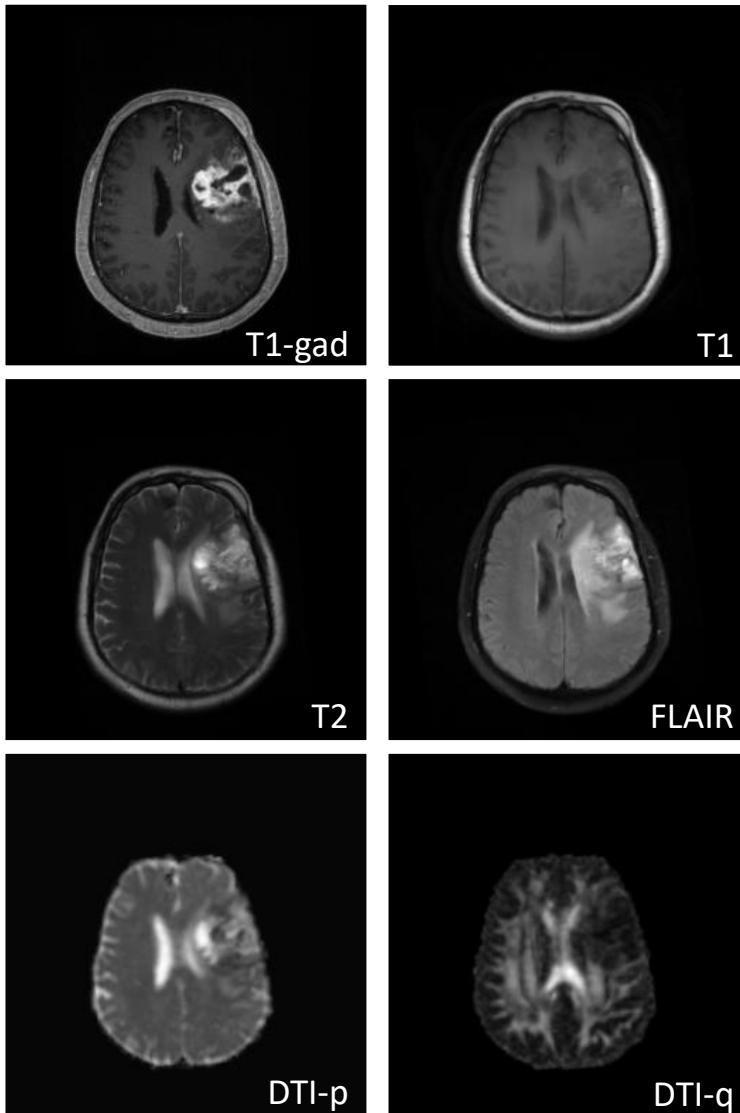
- robust and accurate (like ensembles)
- yields plausible segmentations
- fully automatic

Disadvantages

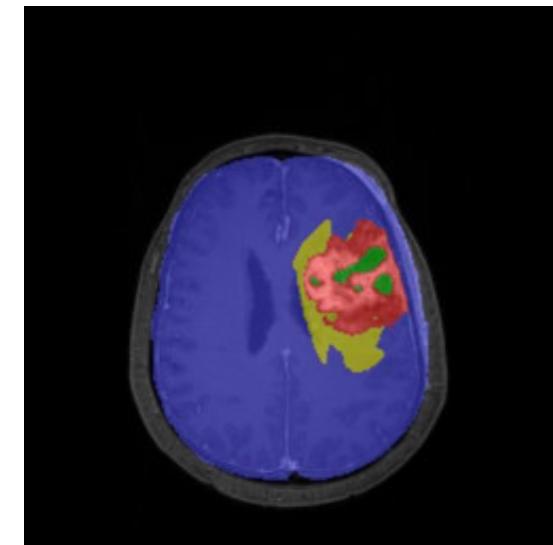
- computationally expensive
- cannot deal well with abnormalities
- not suitable for tumour segmentation

Random forests

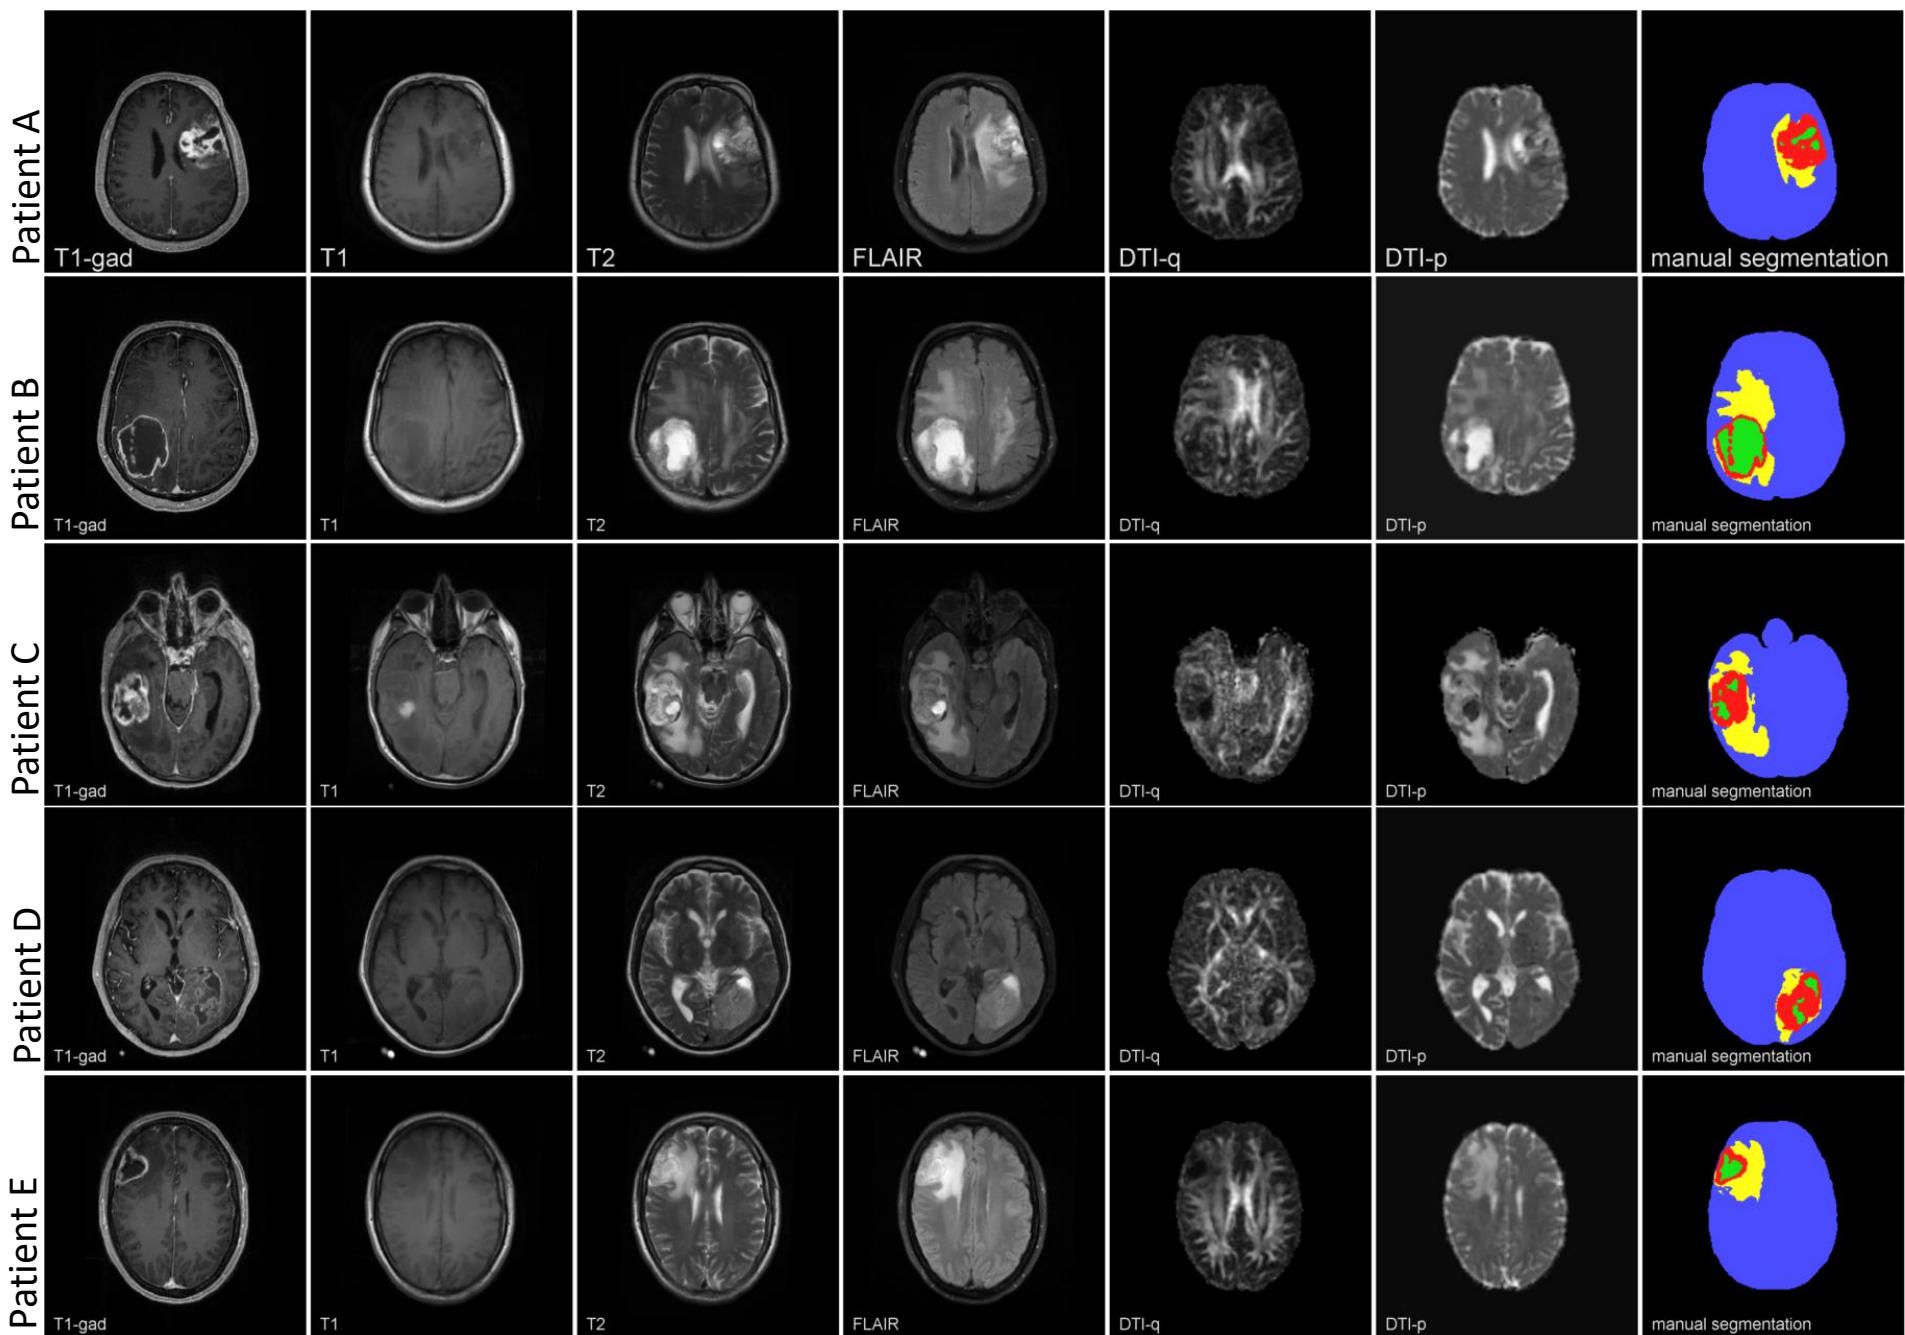
Multi-channel 3D MRI input data



Segmentation of
tumorous tissues:

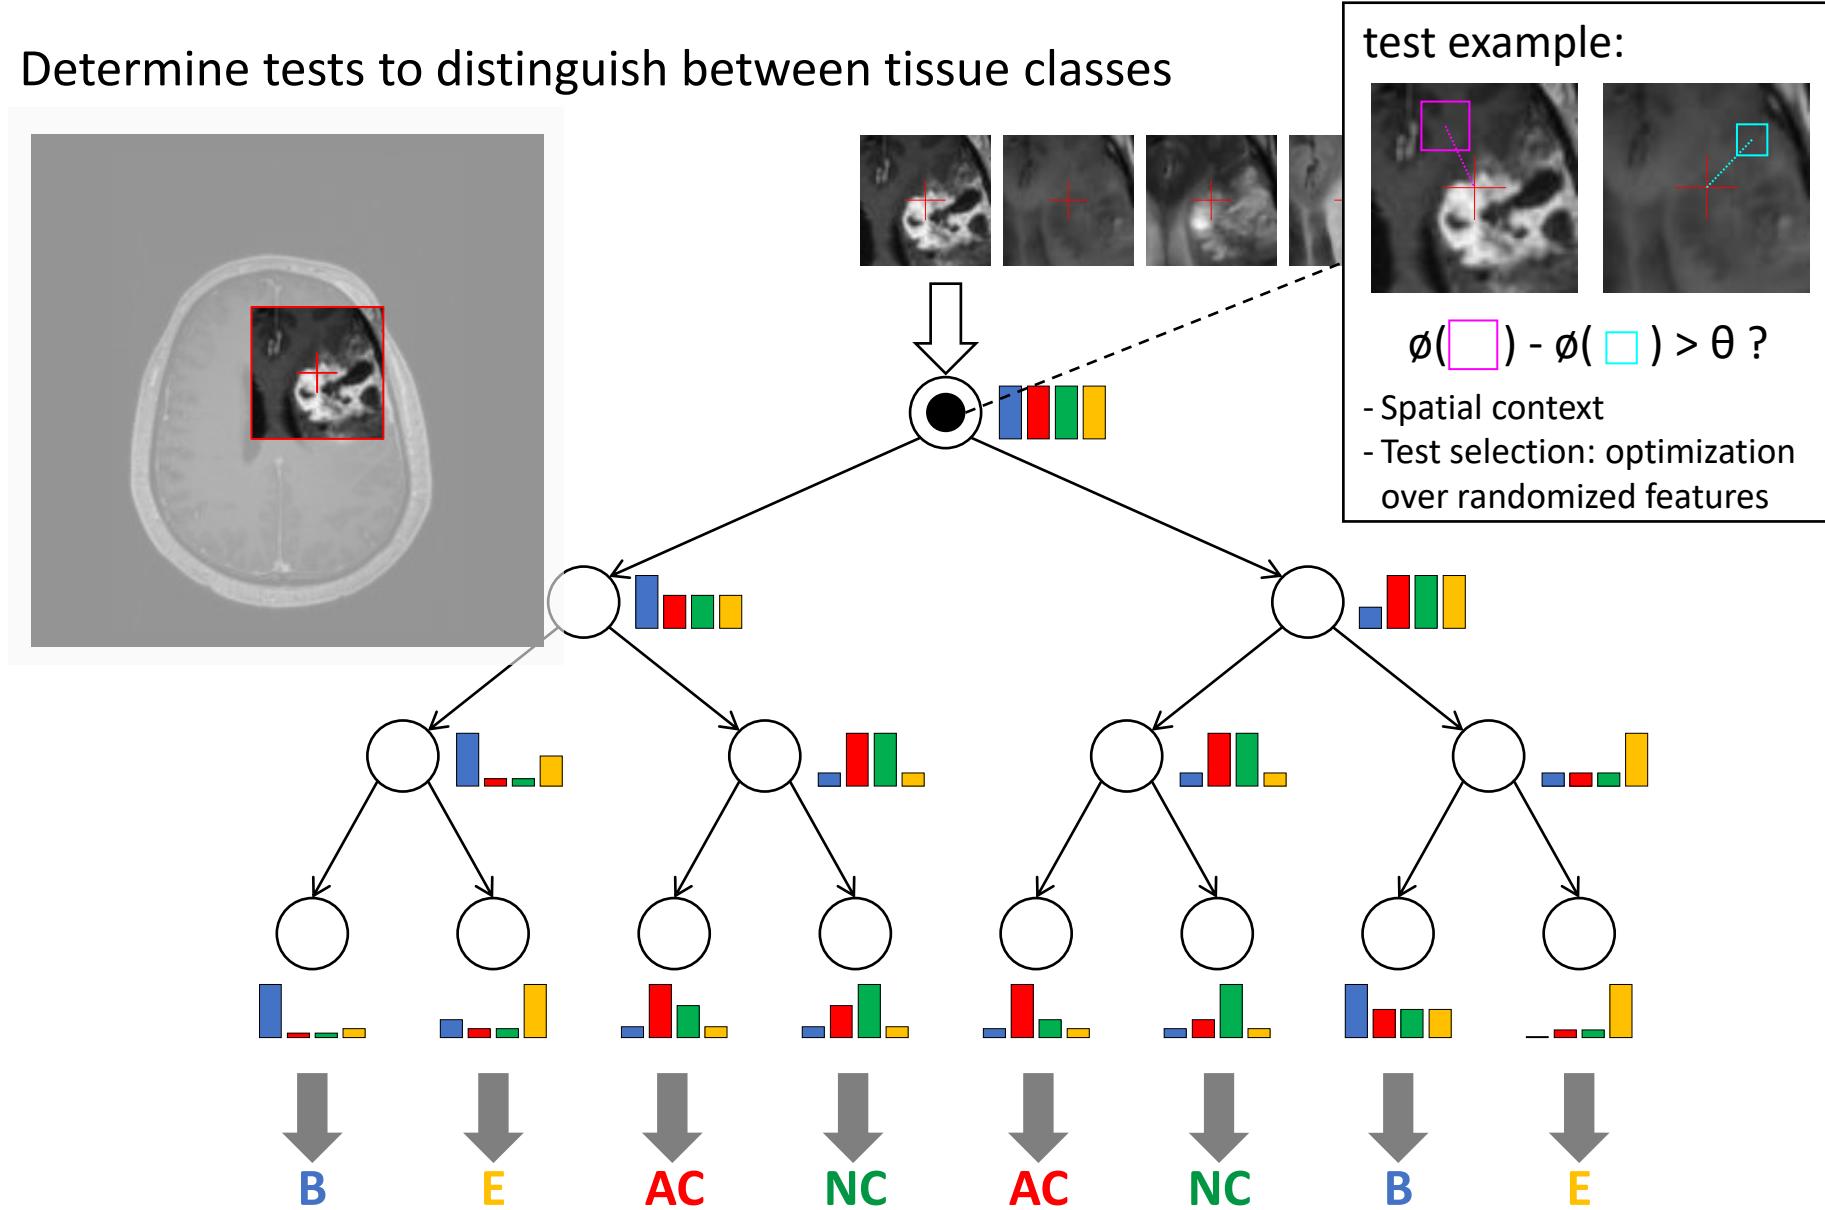


- Active cells
- - - Necrotic core
- Edema
- Background



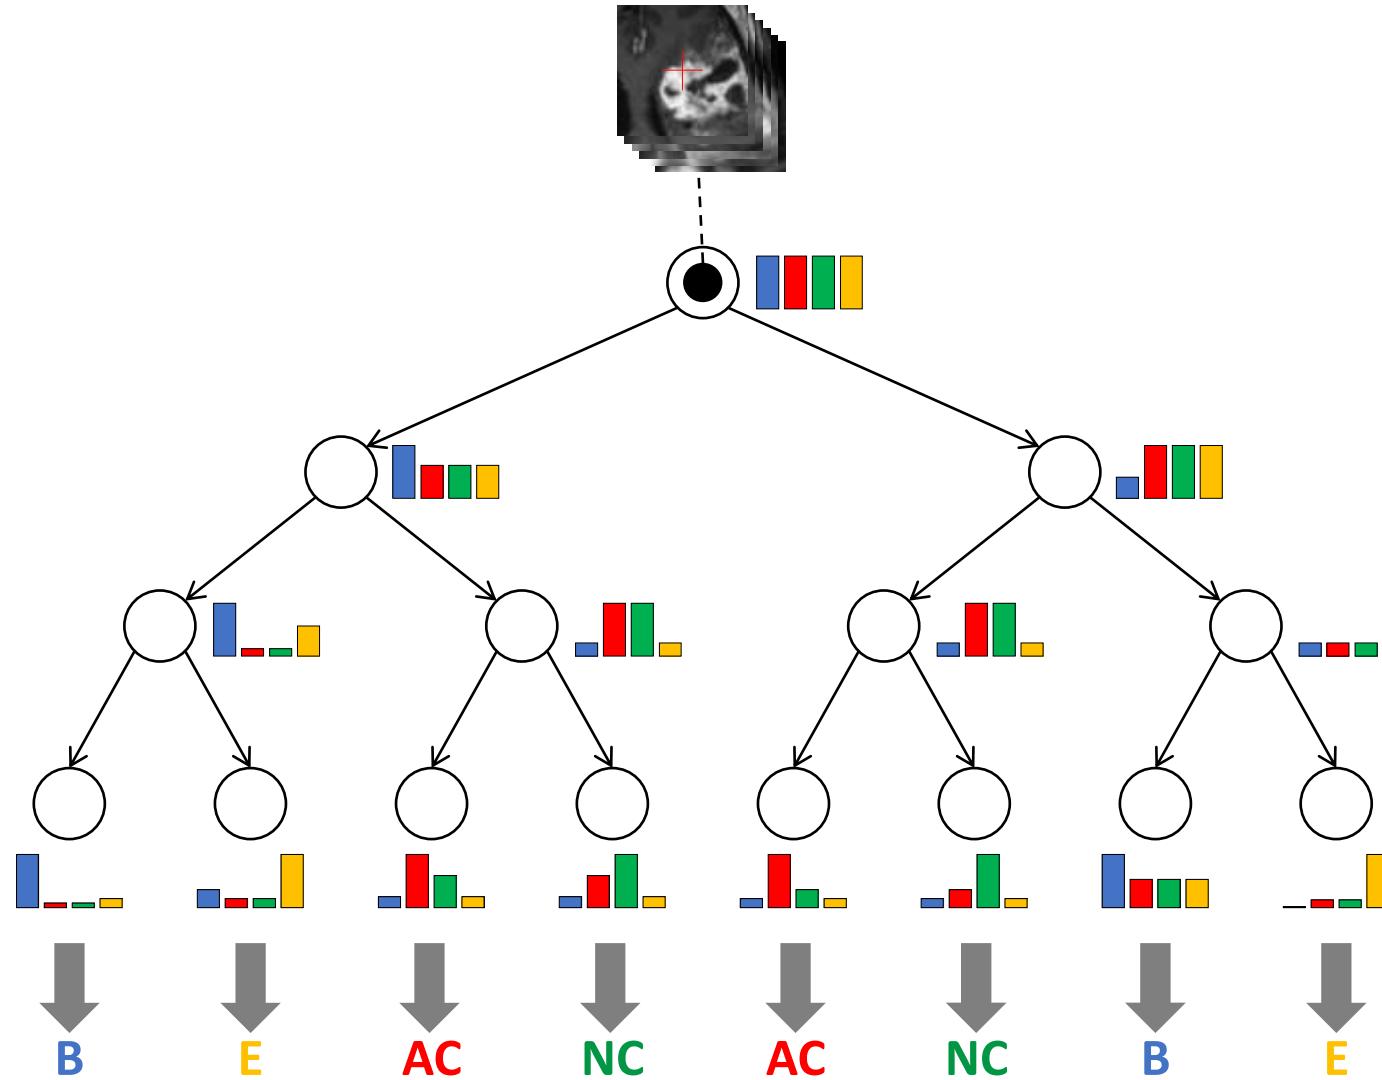
Training: Growing the trees

Determine tests to distinguish between tissue classes

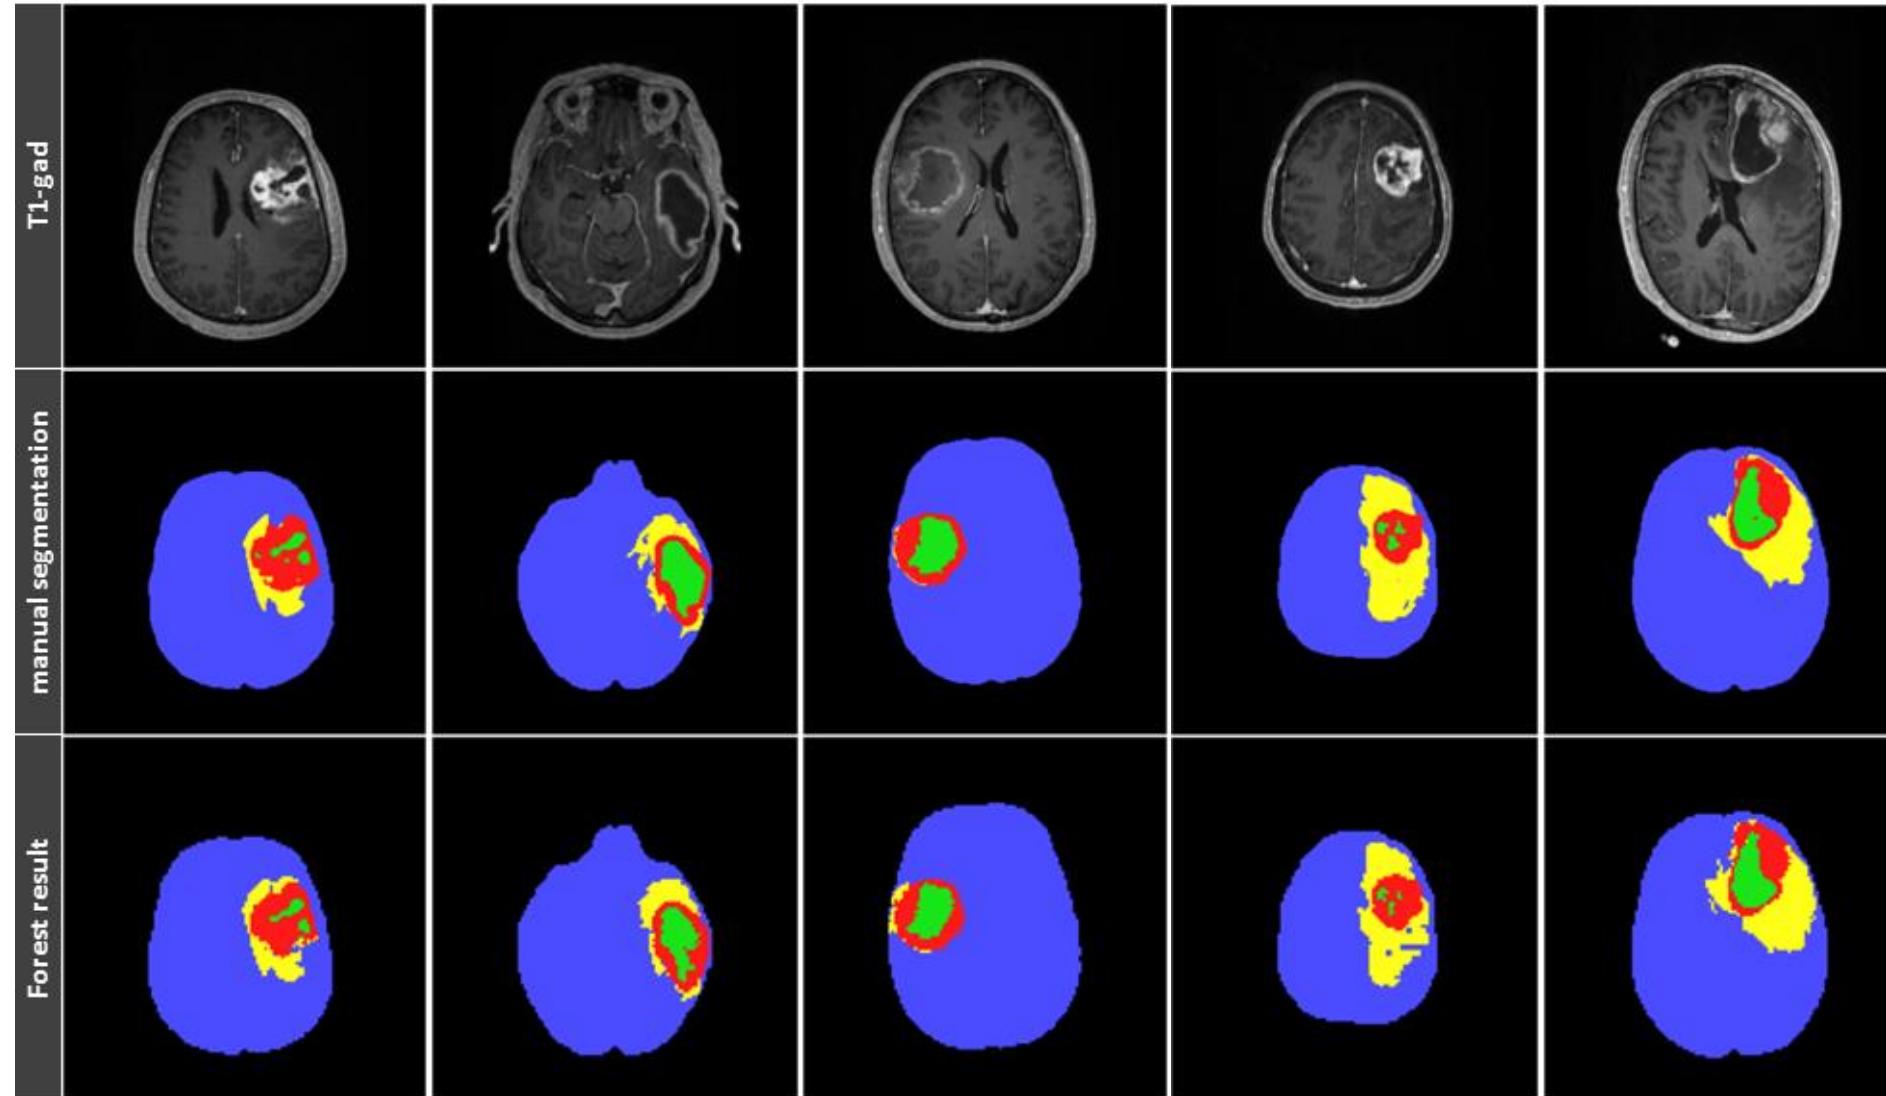


Testing: Traversing the trees

Apply the learned tests to classify tissue points



Visual results



Random forests

Advantages

- ensemble classifiers are robust and accurate
- computationally efficient
- fully automatic

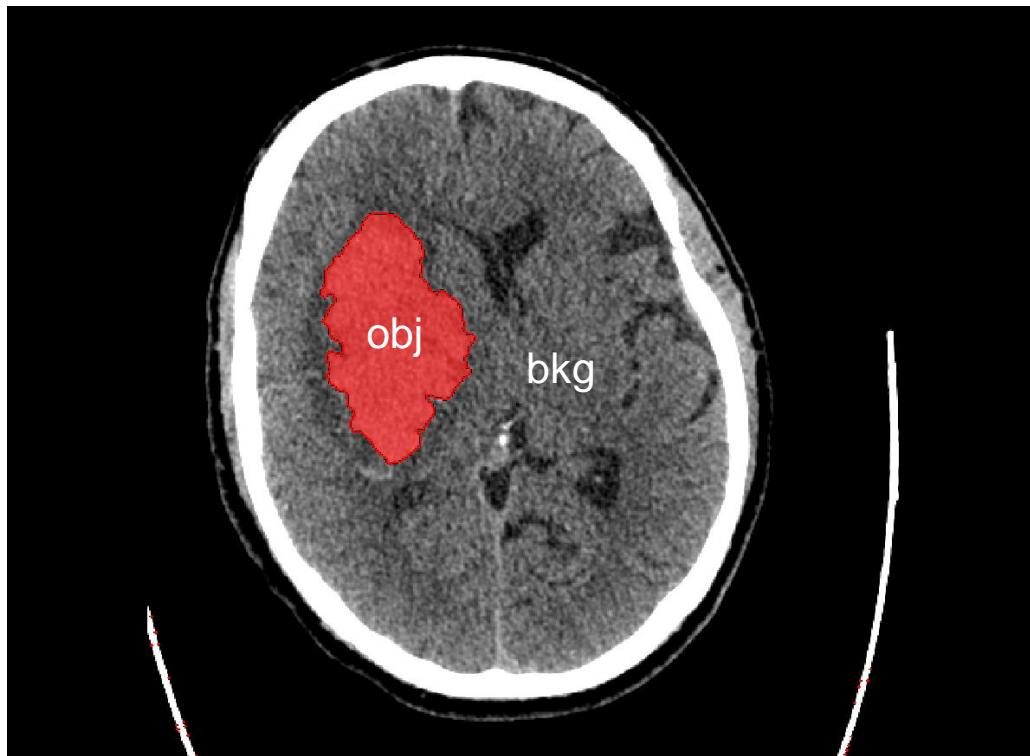
Disadvantages

- shallow model, no hierarchical features
- no guarantees on connectedness

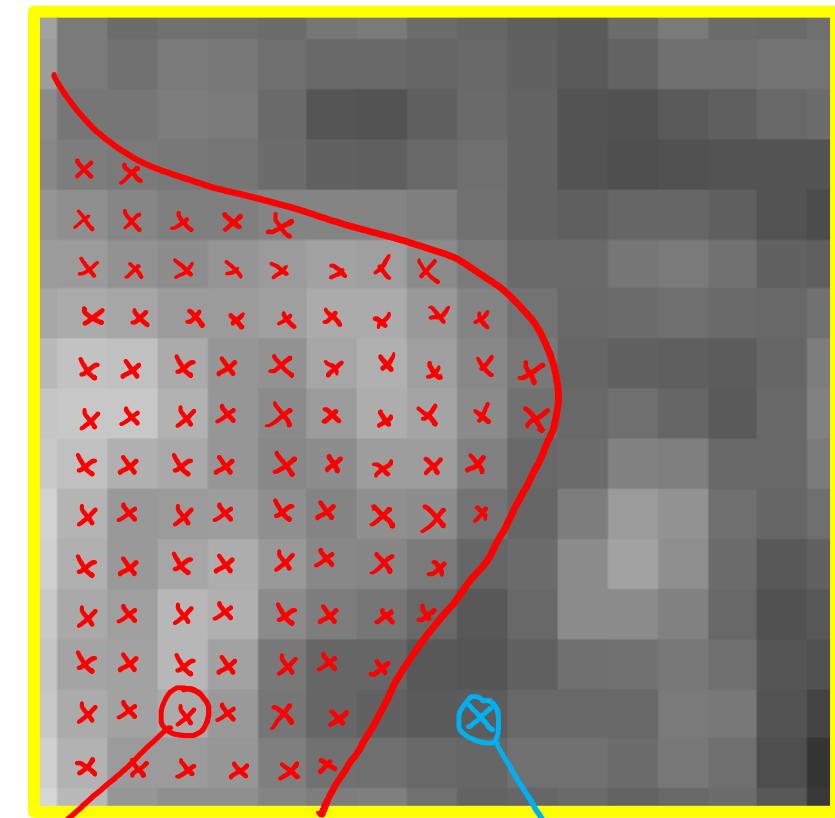
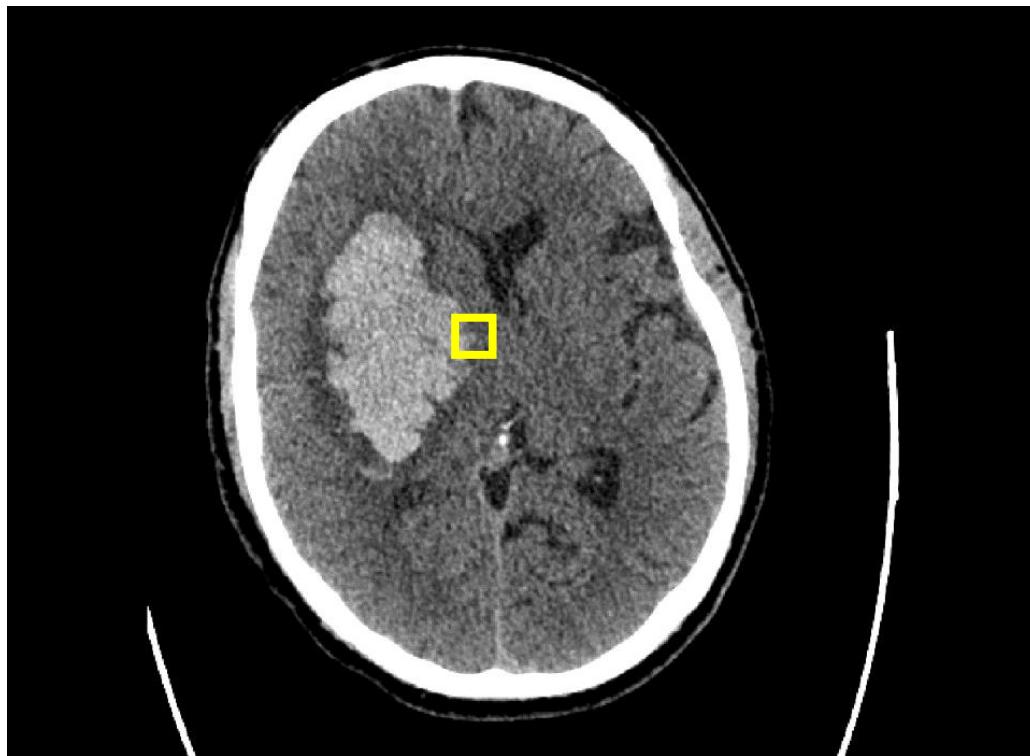
...and then...



Segmentation via dense classification



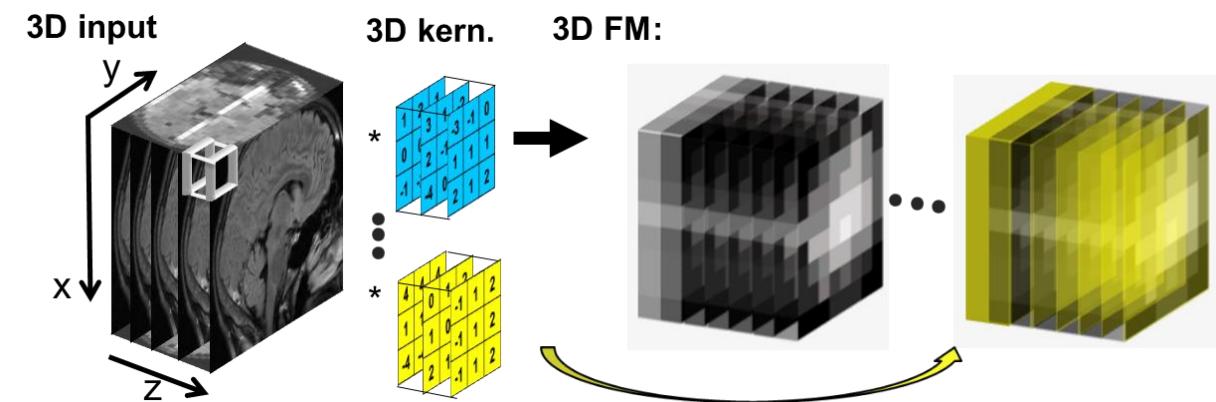
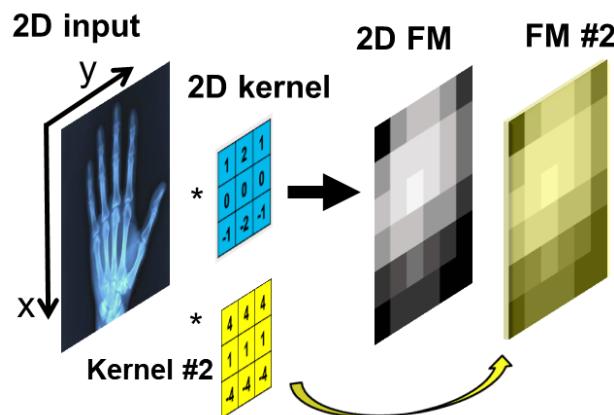
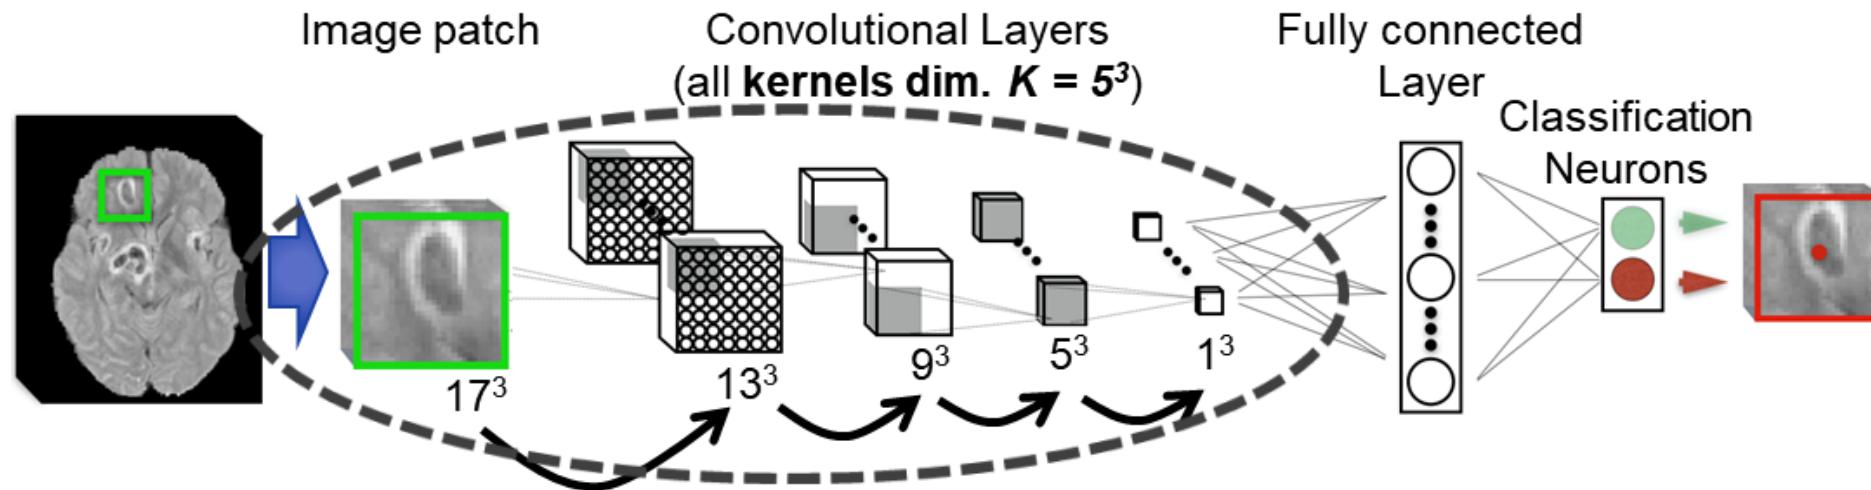
Segmentation via dense classification



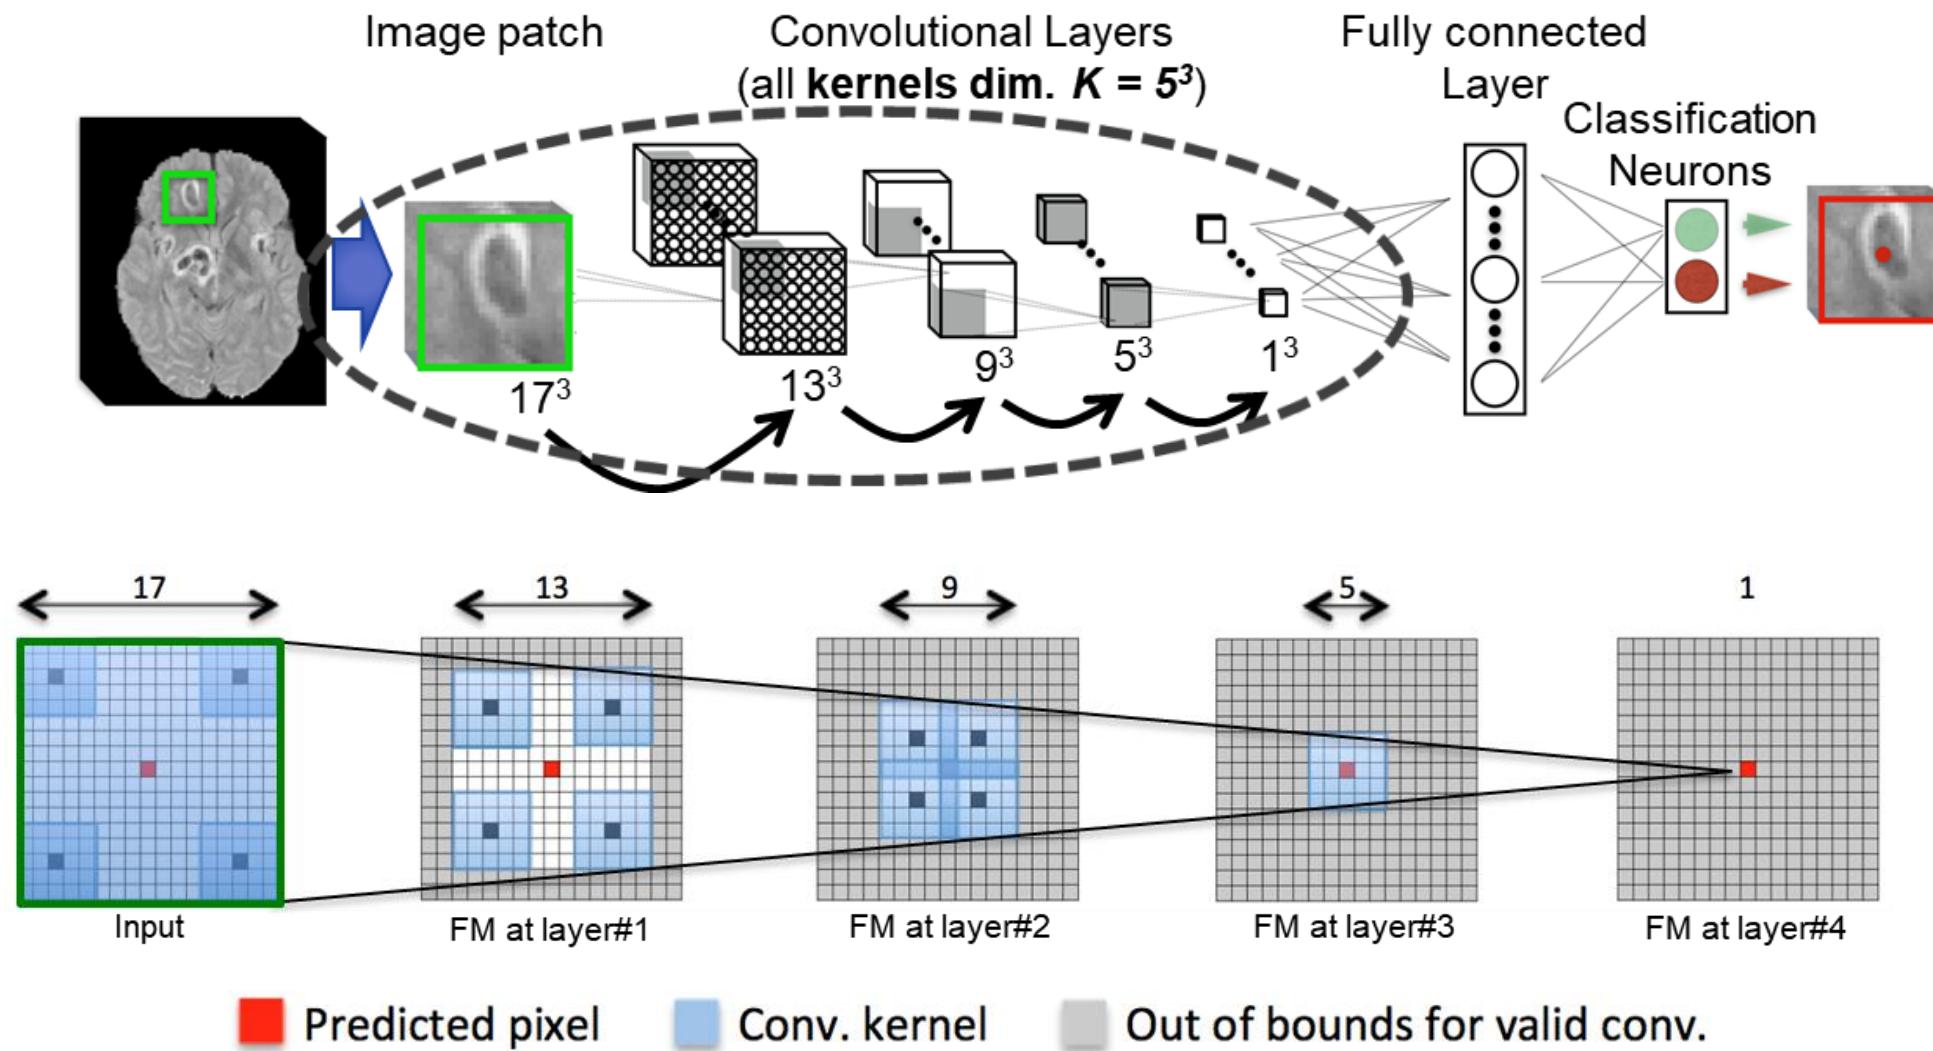
object

background

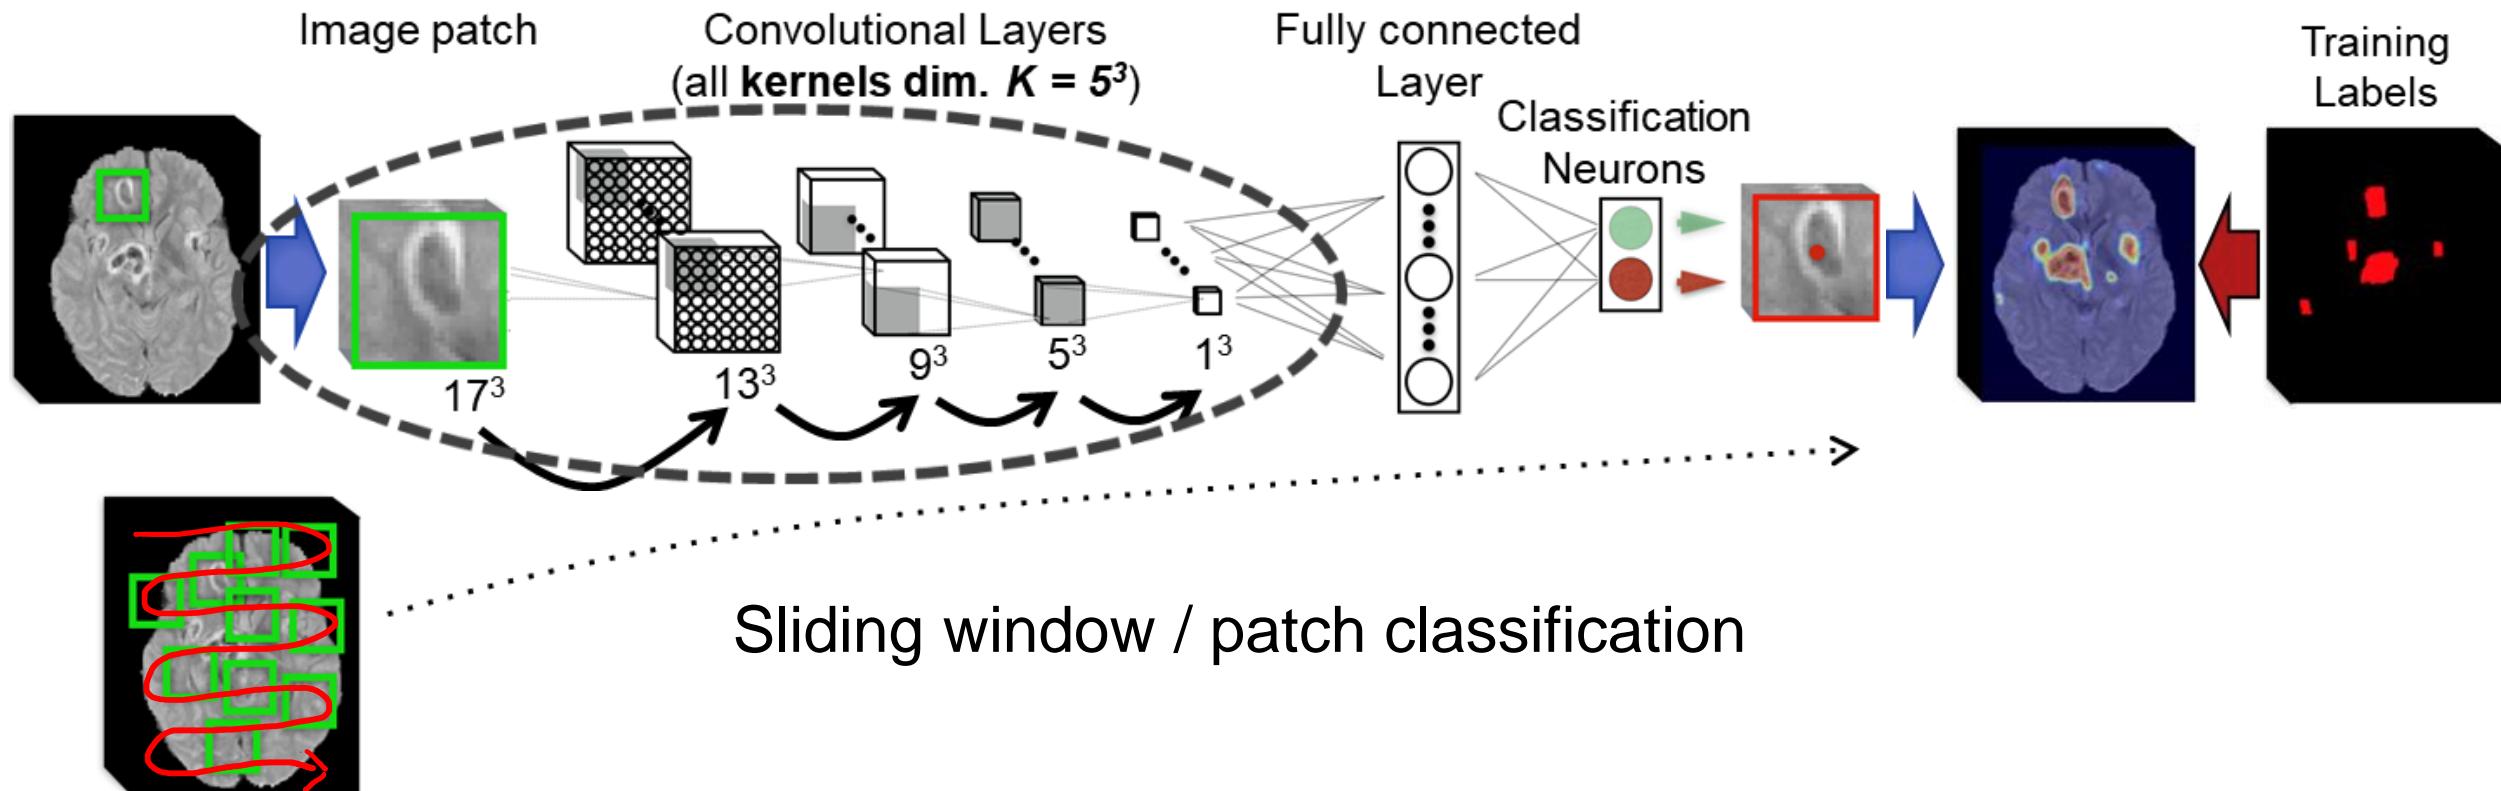
Segmentation via dense classification



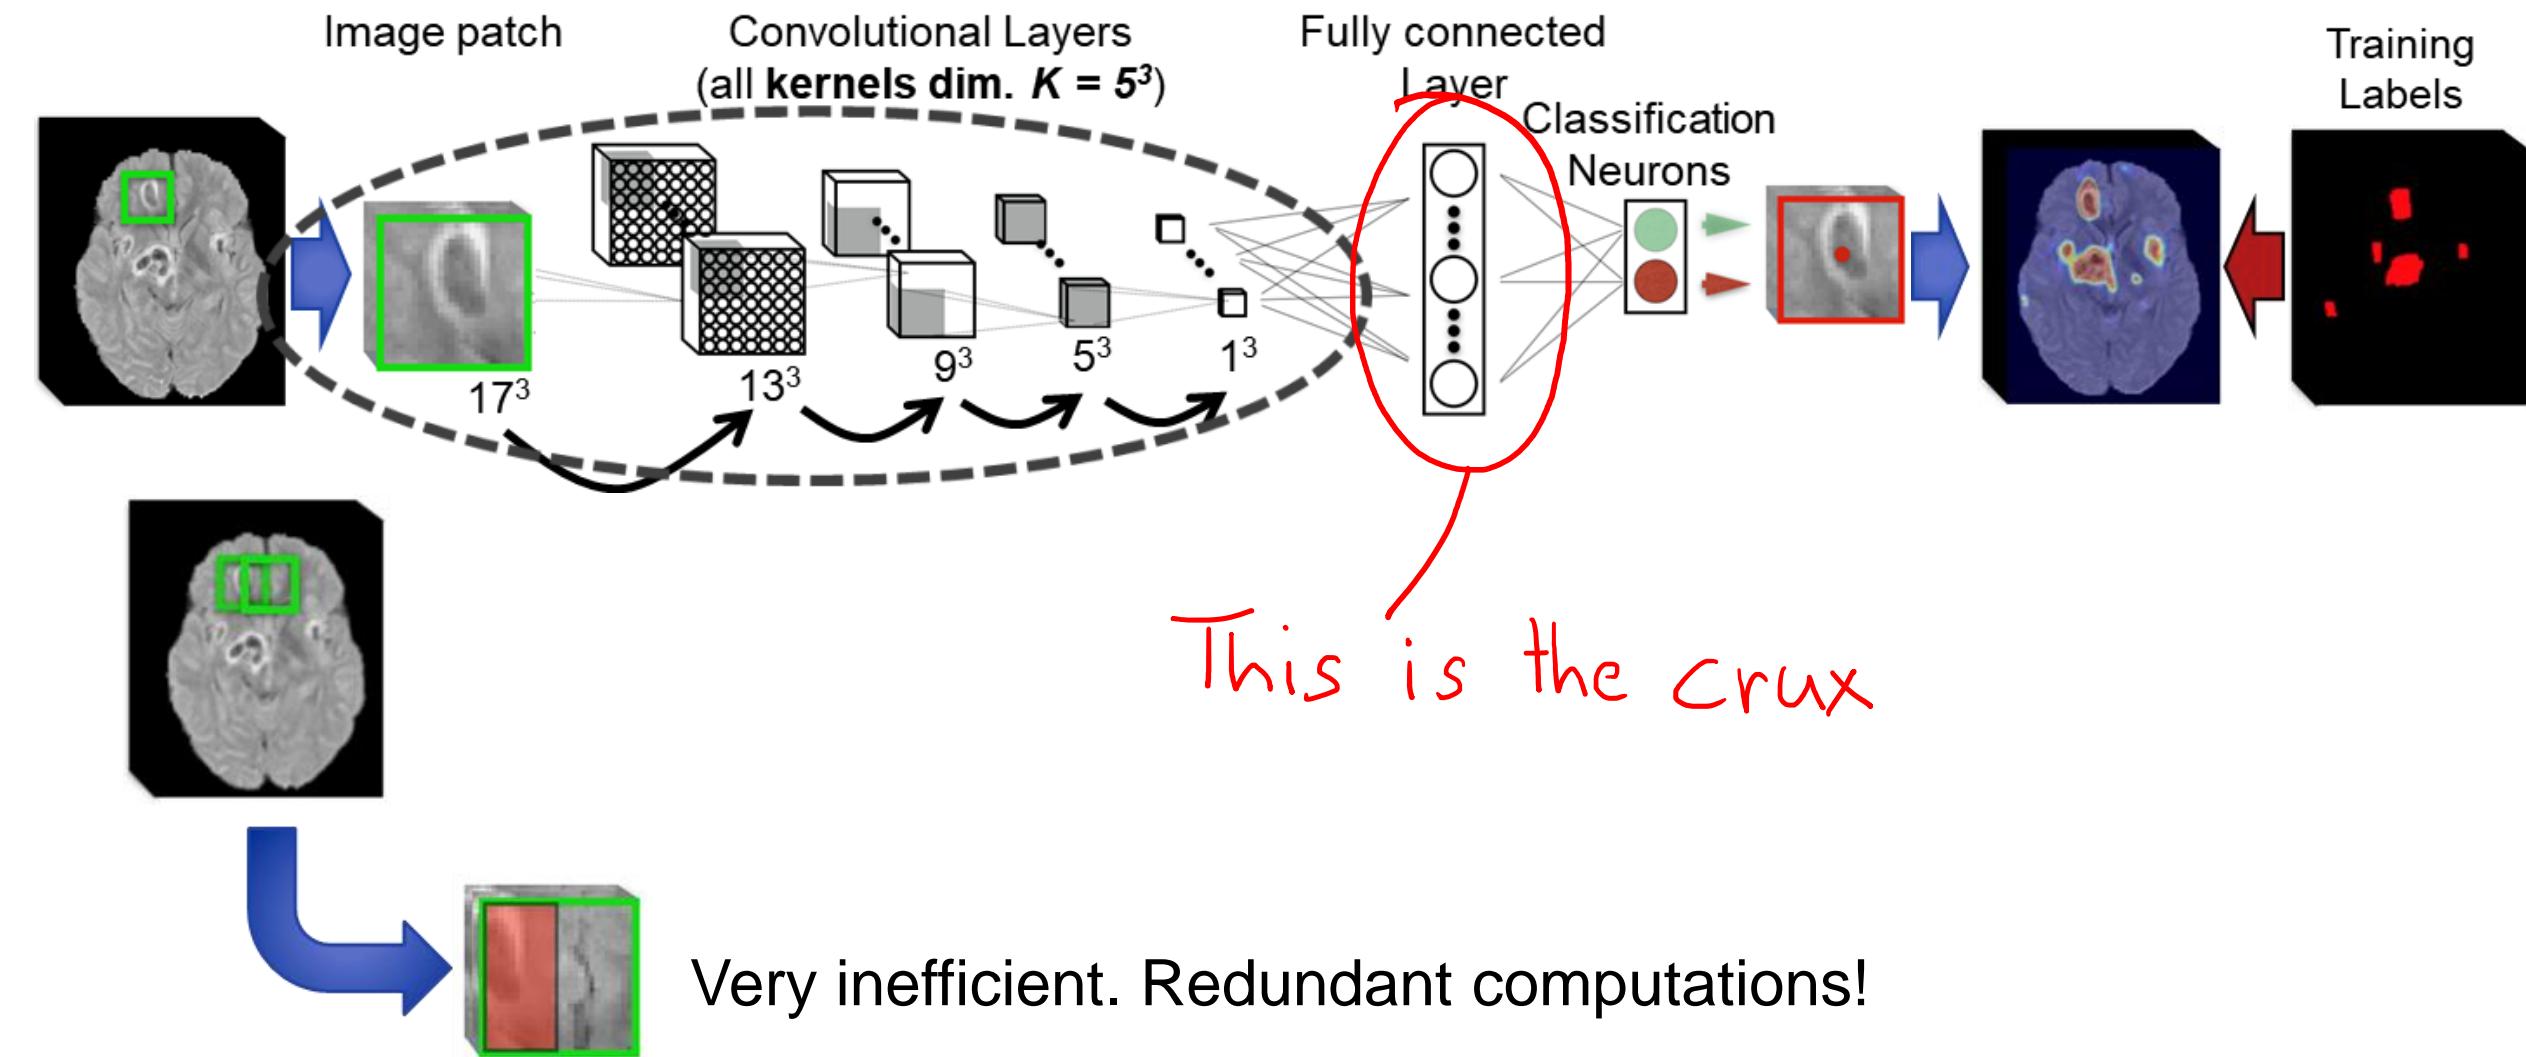
Segmentation via dense classification



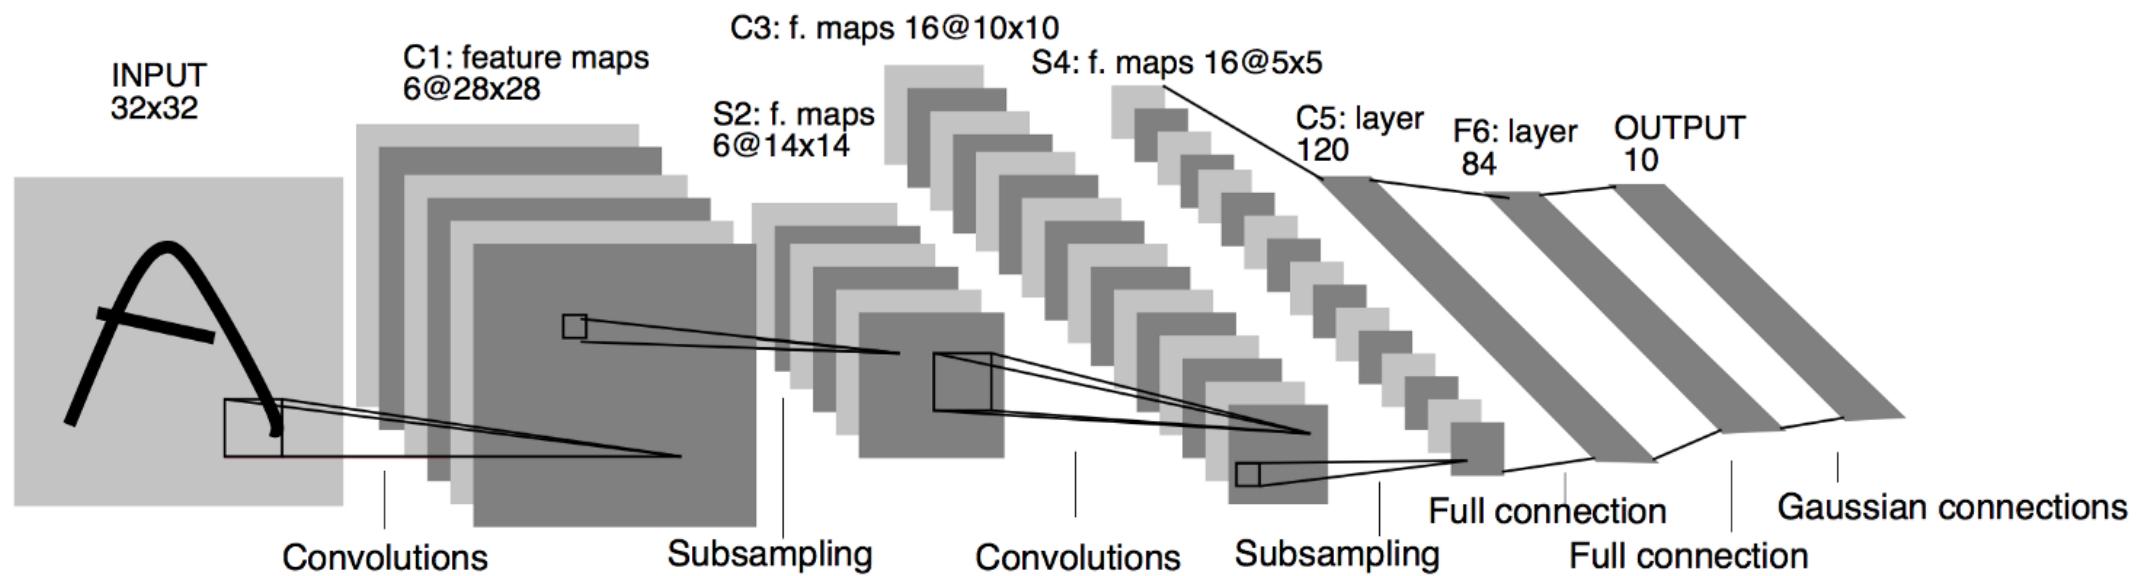
Segmentation via dense classification



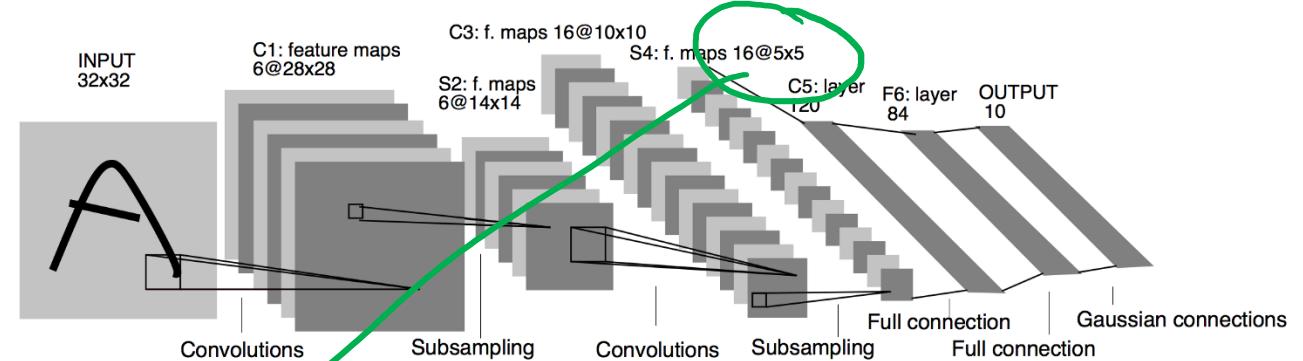
Segmentation via dense classification



Example: LeNet



LeNet in PyTorch

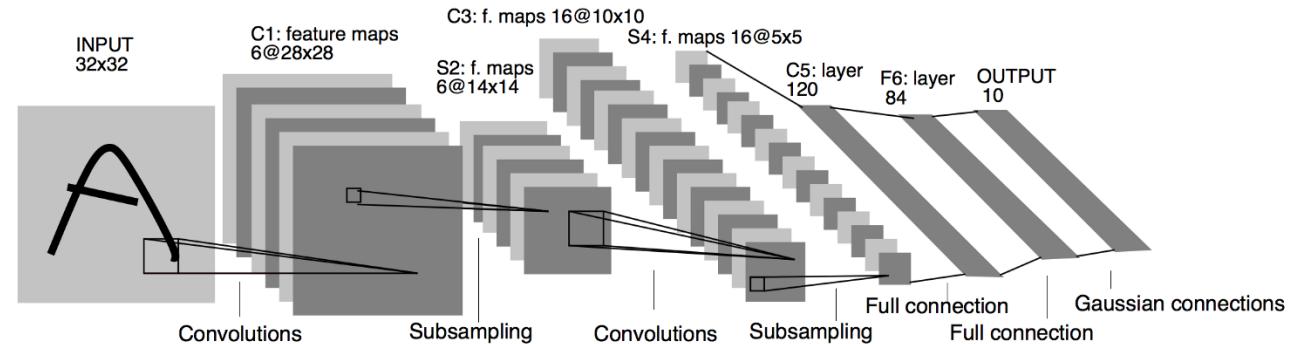


```
class LeNet(nn.Module):

    def __init__(self):
        super(LeNet, self).__init__()
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)
        self.fc1 = nn.Linear(16 * ? * ?, 120)          # Question marks here
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(x.size(0), -1)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return F.log_softmax(x, dim=1)
```

LeNet in PyTorch



```
class LeNet(nn.Module):

    def __init__(self):
        super(LeNet, self).__init__()
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

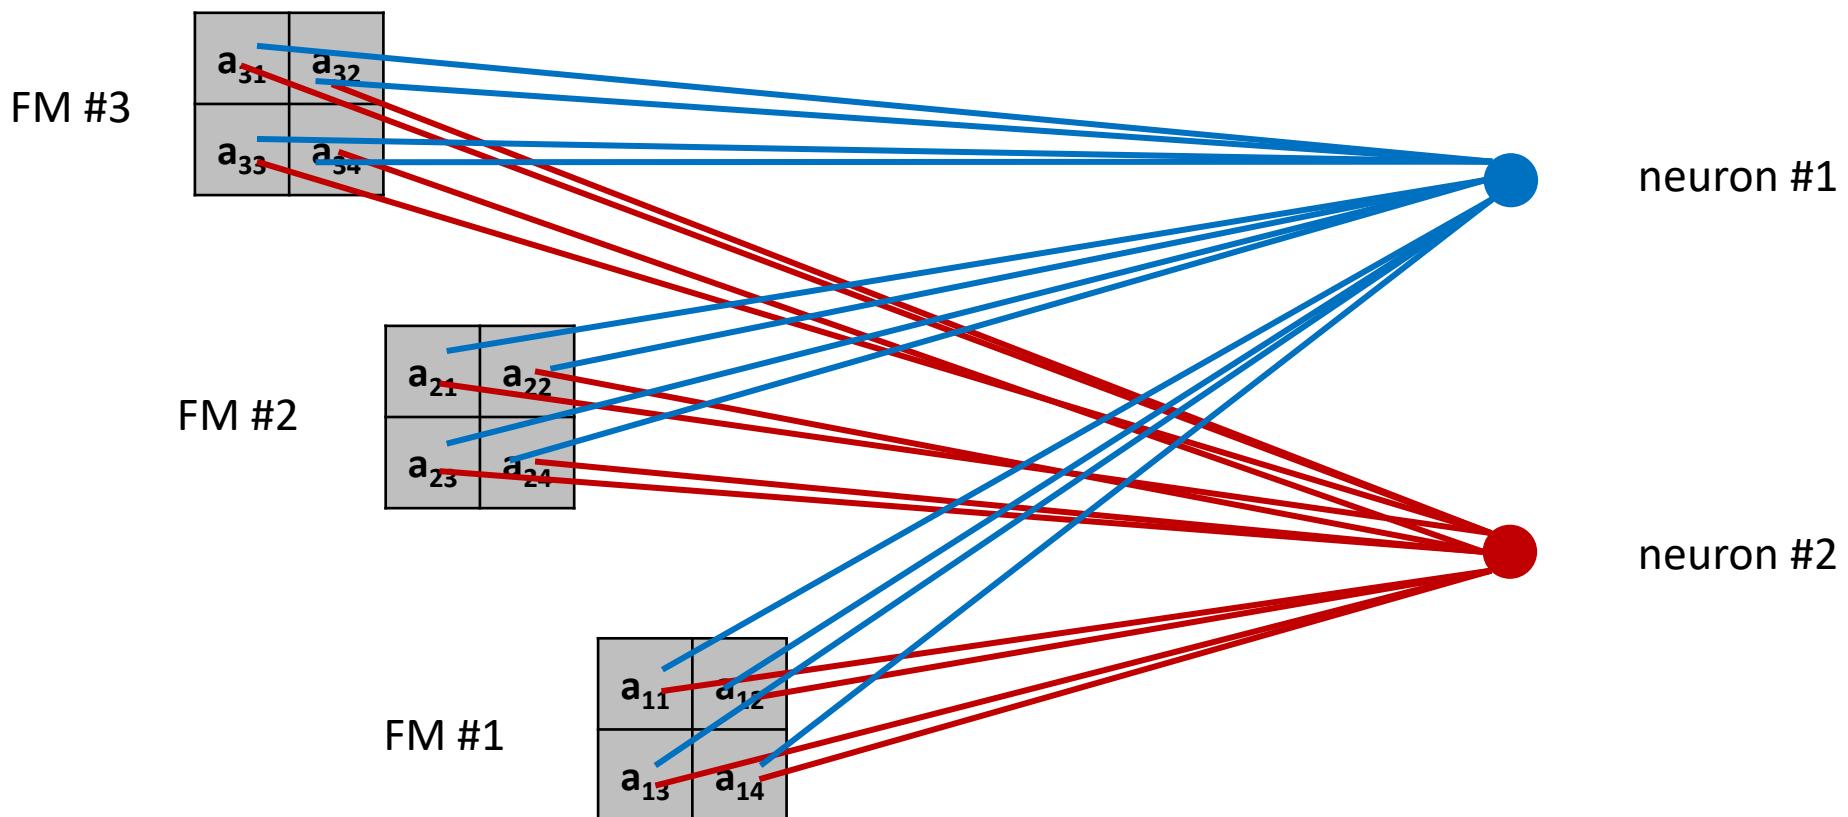
    def forward(self, x):
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(x.size(0), -1)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return F.log_softmax(x, dim=1)
```

Question:
What happens when we feed in
an image of size 64×64 ?

Fully connected to fully convolutional

3 x 2D feature maps (FM) of previous convolutional layer

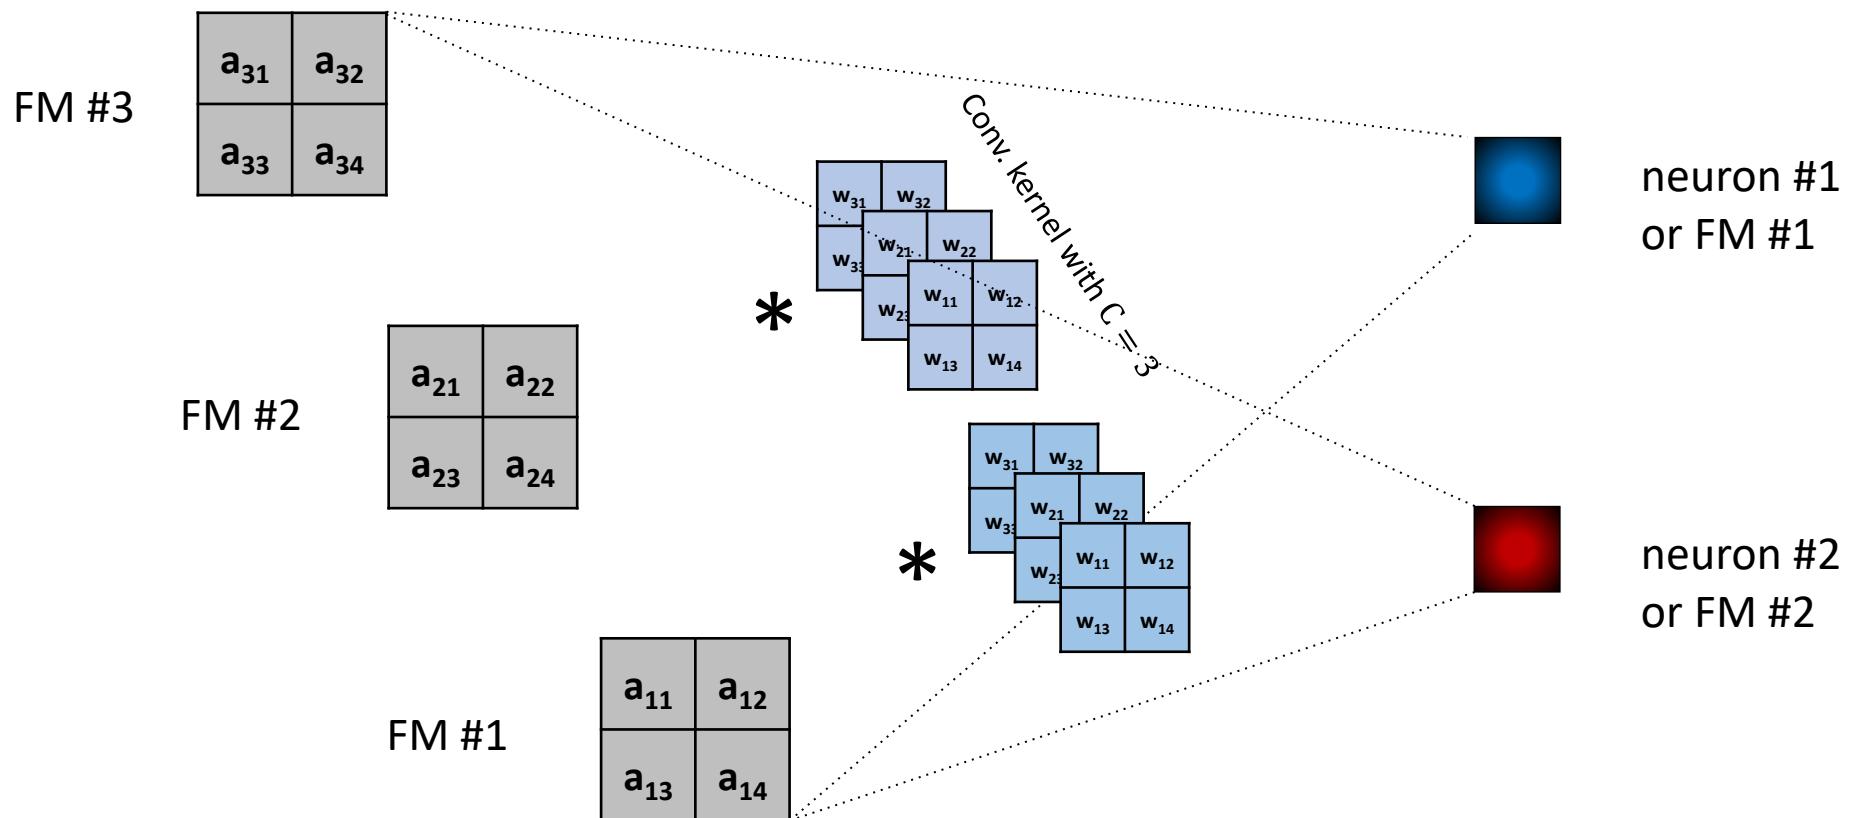
Fully connected neurons of next fully connected layer



Fully connected to fully convolutional

3 x 2D feature maps (FM) of previous convolutional layer

Fully connected neurons of next fully connected layer



Fully convolutional LeNet

```
class LeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.fc1 = nn.Linear(16 * 5 * 5, 120)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```

```
class FCLeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.fc1 = nn.Linear(16 * 5 * 5, 120)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```



Fully convolutional LeNet

```
class LeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.fc1 = nn.Linear(16 * 5 * 5, 120)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```

```
class FCLeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.conv3 = nn.Conv2d(??, ???, kernel_size=?)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```

Fully convolutional LeNet

```
class LeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.fc1 = nn.Linear(16 * 5 * 5, 120)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```

```
class FCLeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.conv3 = nn.Conv2d(16, 120, kernel_size=5)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```



Fully convolutional LeNet

```
class LeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.fc1 = nn.Linear(16 * 5 * 5, 120)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```

```
class FCLeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.conv3 = nn.Conv2d(16, 120, kernel_size=5)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```



Fully convolutional LeNet

```
class LeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.fc1 = nn.Linear(16 * 5 * 5, 120)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```

```
class FCLeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.conv3 = nn.Conv2d(16, 120, kernel_size=5)  
        self.conv4 = nn.Conv2d(120, ??, kernel_size=?)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```

Fully convolutional LeNet

```
class LeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.fc1 = nn.Linear(16 * 5 * 5, 120)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```

```
class FCLeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.conv3 = nn.Conv2d(16, 120, kernel_size=5)  
        self.conv4 = nn.Conv2d(120, 84, kernel_size=1)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```

Fully convolutional LeNet

```
class LeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.fc1 = nn.Linear(16 * 5 * 5, 120)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```

```
class FCLeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.conv3 = nn.Conv2d(16, 120, kernel_size=5)  
        self.conv4 = nn.Conv2d(120, 84, kernel_size=1)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```



Fully convolutional LeNet

```
class LeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.fc1 = nn.Linear(16 * 5 * 5, 120)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```

```
class FCLeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.conv3 = nn.Conv2d(16, 120, kernel_size=5)  
        self.conv4 = nn.Conv2d(120, 84, kernel_size=1)  
        self.conv5 = nn.Conv2d(84, ??, kernel_size=?)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```

Fully convolutional LeNet

```
class LeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.fc1 = nn.Linear(16 * 5 * 5, 120)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)   
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```

```
class FCLeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.conv3 = nn.Conv2d(16, 120, kernel_size=5)  
        self.conv4 = nn.Conv2d(120, 84, kernel_size=1)  
        self.conv5 = nn.Conv2d(84, 10, kernel_size=1)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)   
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```

Fully convolutional LeNet

```
class LeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.fc1 = nn.Linear(16 * 5 * 5, 120)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```

```
class FCLeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.conv3 = nn.Conv2d(16, 120, kernel_size=5)  
        self.conv4 = nn.Conv2d(120, 84, kernel_size=1)  
        self.conv5 = nn.Conv2d(84, 10, kernel_size=1)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```

Fully convolutional LeNet

```
class LeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.fc1 = nn.Linear(16 * 5 * 5, 120)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return F.log_softmax(x, dim=1)
```

```
class FCLeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.conv3 = nn.Conv2d(16, 120, kernel_size=5)  
        self.conv4 = nn.Conv2d(120, 84, kernel_size=1)  
        self.conv5 = nn.Conv2d(84, 10, kernel_size=1)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = x.view(x.size(0), -1)  
        x = F.relu(self.conv3(x))  
        x = F.relu(self.conv4(x))  
        x = self.conv5(x)  
        return F.log_softmax(x, dim=1)
```

Fully convolutional LeNet

```
class LeNet(nn.Module):

    def __init__(self):
        super(LeNet, self).__init__()
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(x.size(0), -1)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return F.log_softmax(x, dim=1)
```

```
class FCLeNet(nn.Module):

    def __init__(self):
        super(LeNet, self).__init__()
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)
        self.conv3 = nn.Conv2d(16, 120, kernel_size=5)
        self.conv4 = nn.Conv2d(120, 84, kernel_size=1)
        self.conv5 = nn.Conv2d(84, 10, kernel_size=1)

    def forward(self, x):
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
x = x.view(x.size(0), -1)
        x = F.relu(self.conv3(x))
        x = F.relu(self.conv4(x))
        x = self.conv5(x)
        return F.log_softmax(x, dim=1)
```

Fully convolutional LeNet

```
class LeNet(nn.Module):

    def __init__(self):
        super(LeNet, self).__init__()
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(x.size(0), -1)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return F.log_softmax(x, dim=1)
```

```
class FCLeNet(nn.Module):

    def __init__(self):
        super(LeNet, self).__init__()
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)
        self.conv3 = nn.Conv2d(16, 120, kernel_size=5)
        self.conv4 = nn.Conv2d(120, 84, kernel_size=1)
        self.conv5 = nn.Conv2d(84, 10, kernel_size=1)

    def forward(self, x):
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
x = x.view(x.size(0), -1)
        x = F.relu(self.conv3(x))
        x = F.relu(self.conv4(x))
        x = self.conv5(x)
        return F.log_softmax(x, dim=1)
```

Fully convolutional LeNet

Question:

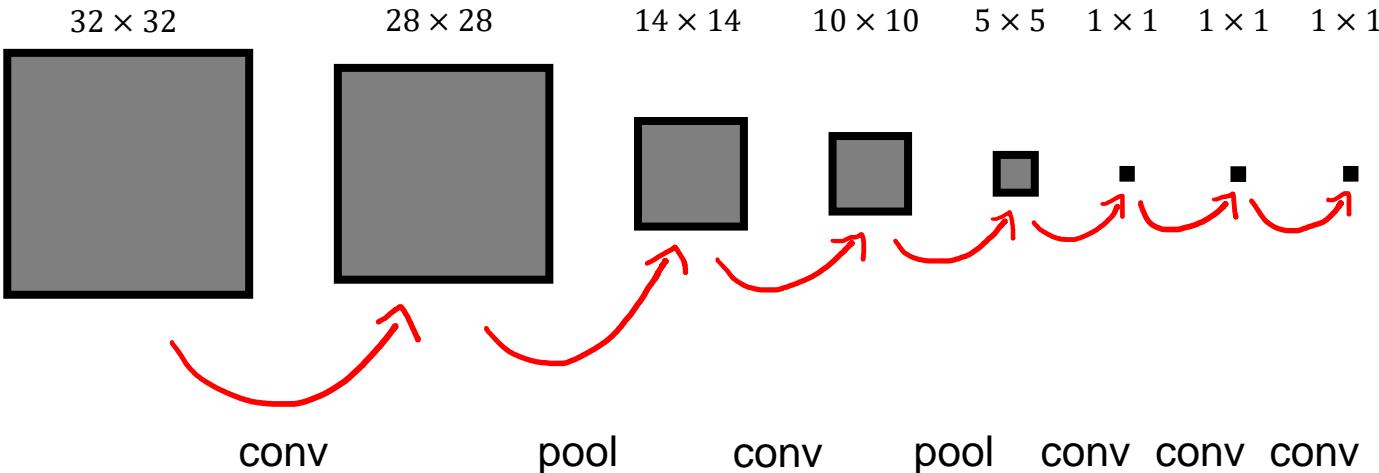
What happens when we feed in
an image of size 64×64 ?

```
class FCLeNet(nn.Module):

    def __init__(self):
        super(LeNet, self).__init__()
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)
        self.conv3 = nn.Conv2d(16, 120, kernel_size=5)
        self.conv4 = nn.Conv2d(120, 84, kernel_size=1)
        self.conv5 = nn.Conv2d(84, 10, kernel_size=1)

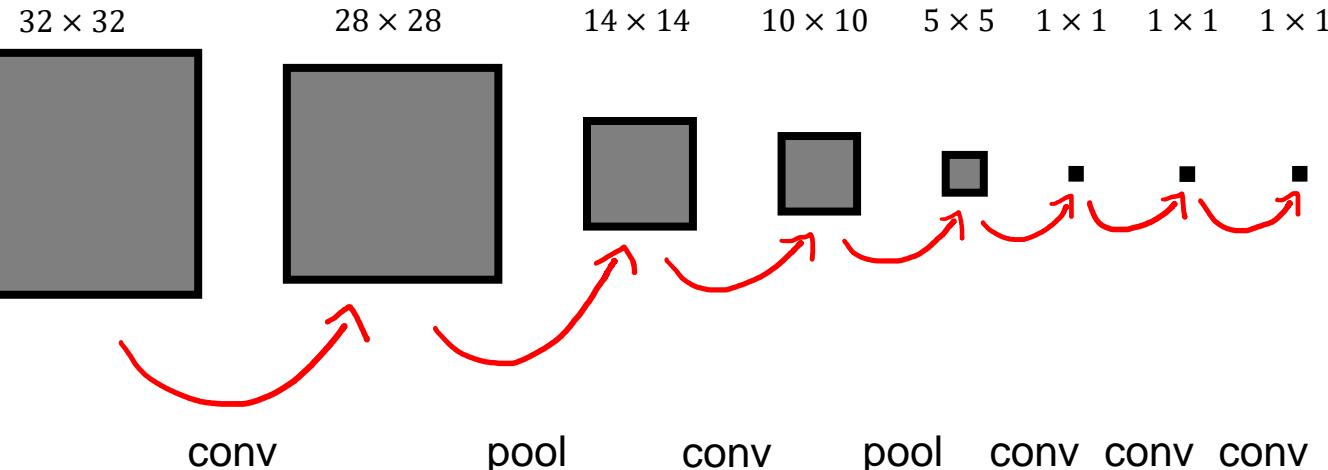
    def forward(self, x):
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = F.relu(self.conv3(x))
        x = F.relu(self.conv4(x))
        x = self.conv5(x)
        return F.log_softmax(x, dim=1)
```

Fully convolutional LeNet



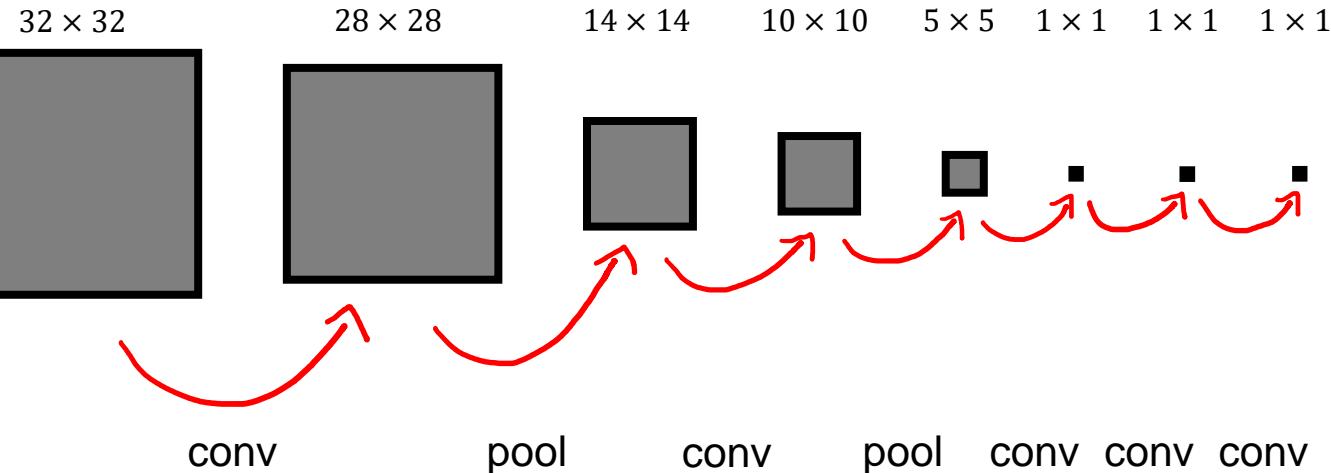
```
class FCLeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.conv3 = nn.Conv2d(16, 120, kernel_size=5)  
        self.conv4 = nn.Conv2d(120, 84, kernel_size=1)  
        self.conv5 = nn.Conv2d(84, 10, kernel_size=1)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = F.relu(self.conv3(x))  
        x = F.relu(self.conv4(x))  
        x = self.conv5(x)  
        return F.log_softmax(x, dim=1)
```

Fully convolutional LeNet

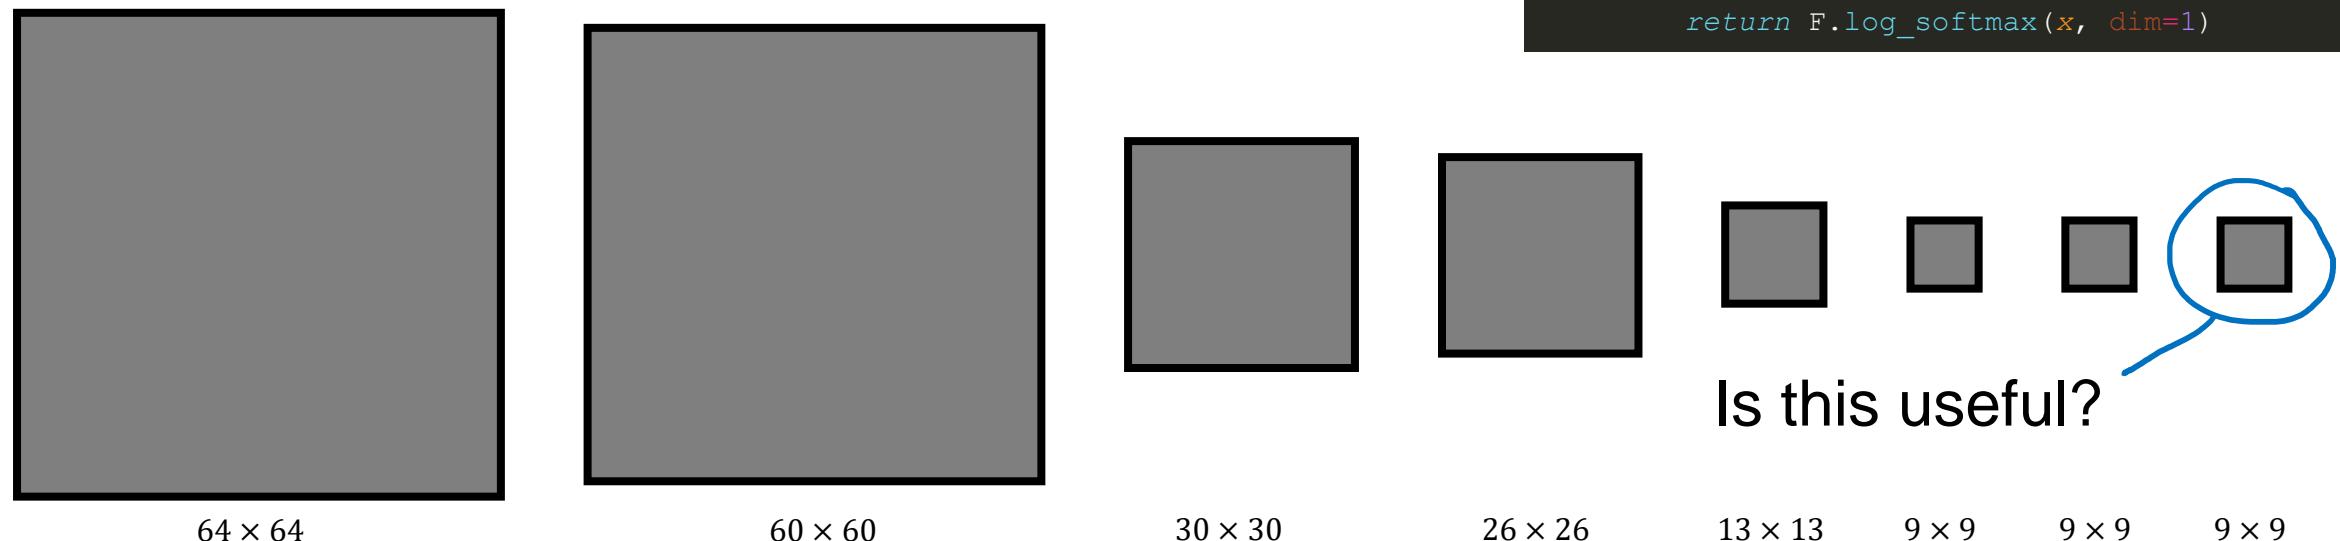


```
class FCLeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.conv3 = nn.Conv2d(16, 120, kernel_size=5)  
        self.conv4 = nn.Conv2d(120, 84, kernel_size=1)  
        self.conv5 = nn.Conv2d(84, 10, kernel_size=1)  
  
    def forward(self, x):  
        x = F.relu(self.conv1(x), 2)  
        x = F.relu(self.conv2(x), 2)  
        x = F.relu(self.conv3(x))  
        x = F.relu(self.conv4(x))  
        x = self.conv5(x)  
        return F.log_softmax(x, dim=1)
```

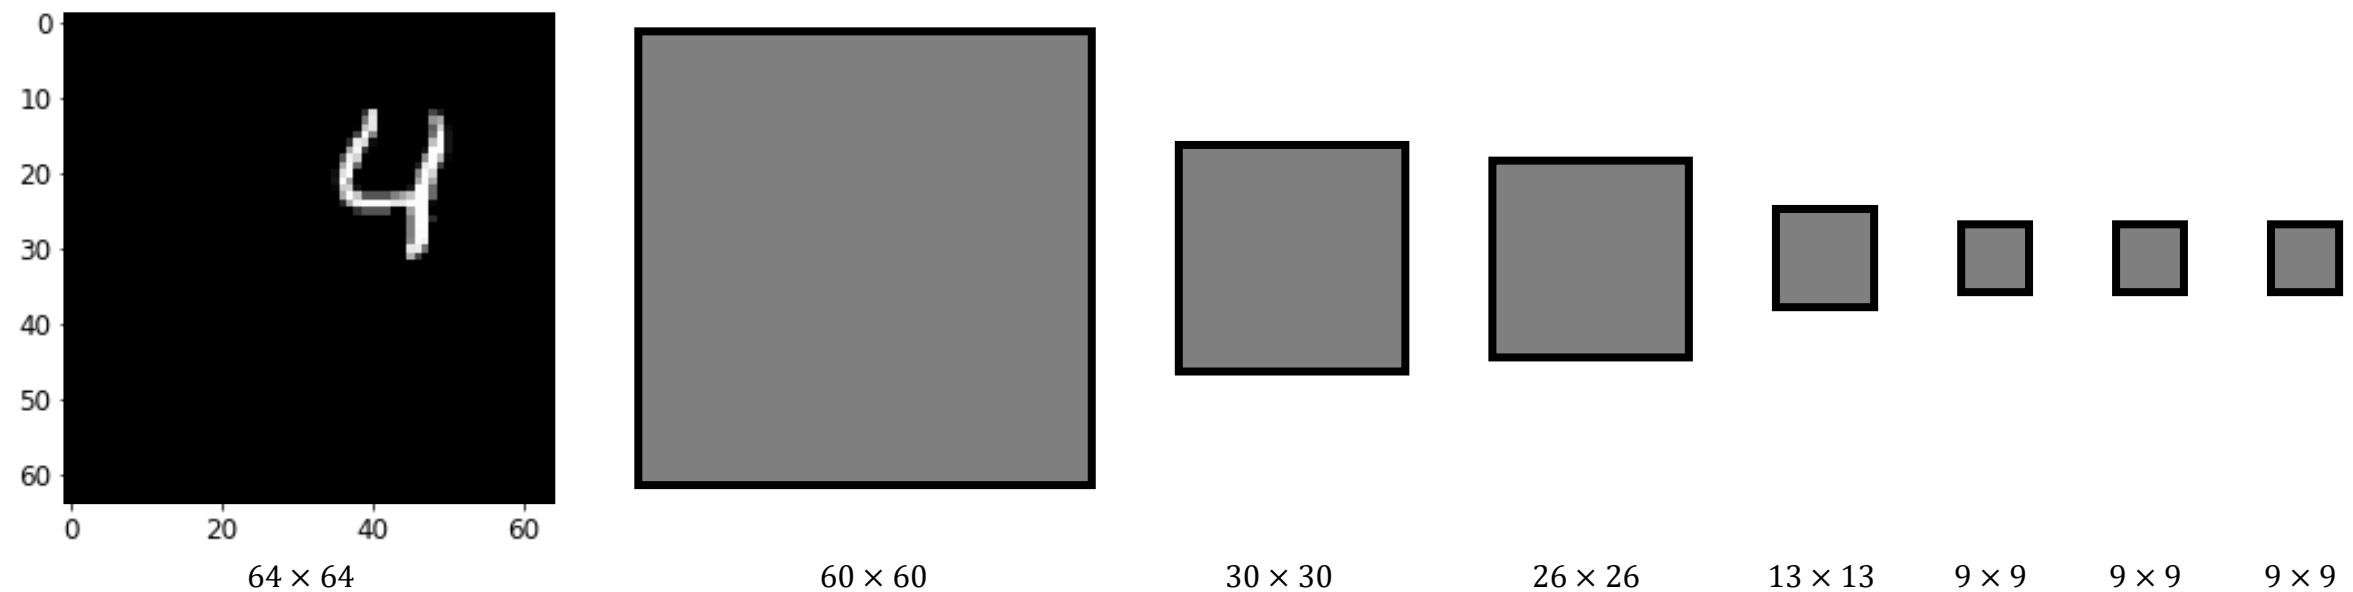
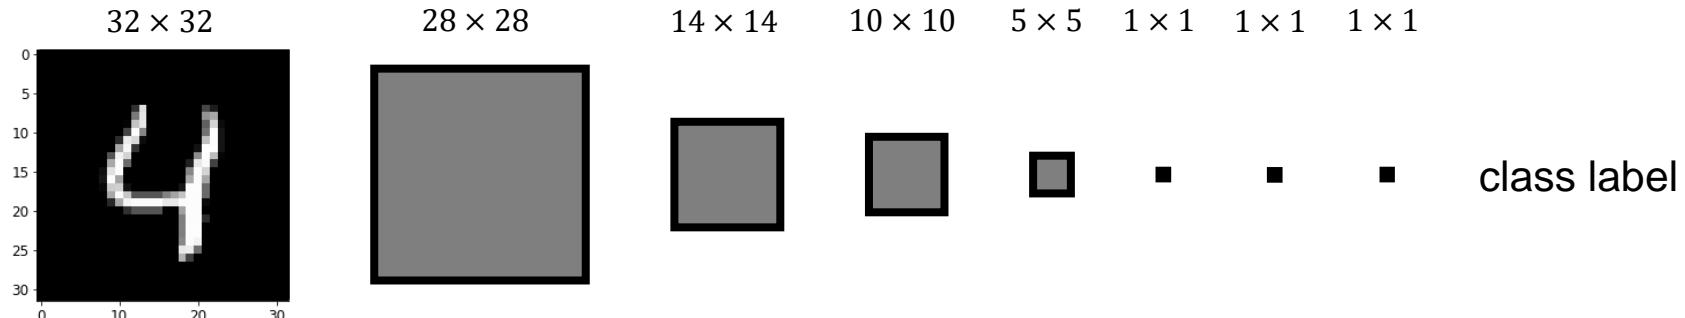
Fully convolutional LeNet



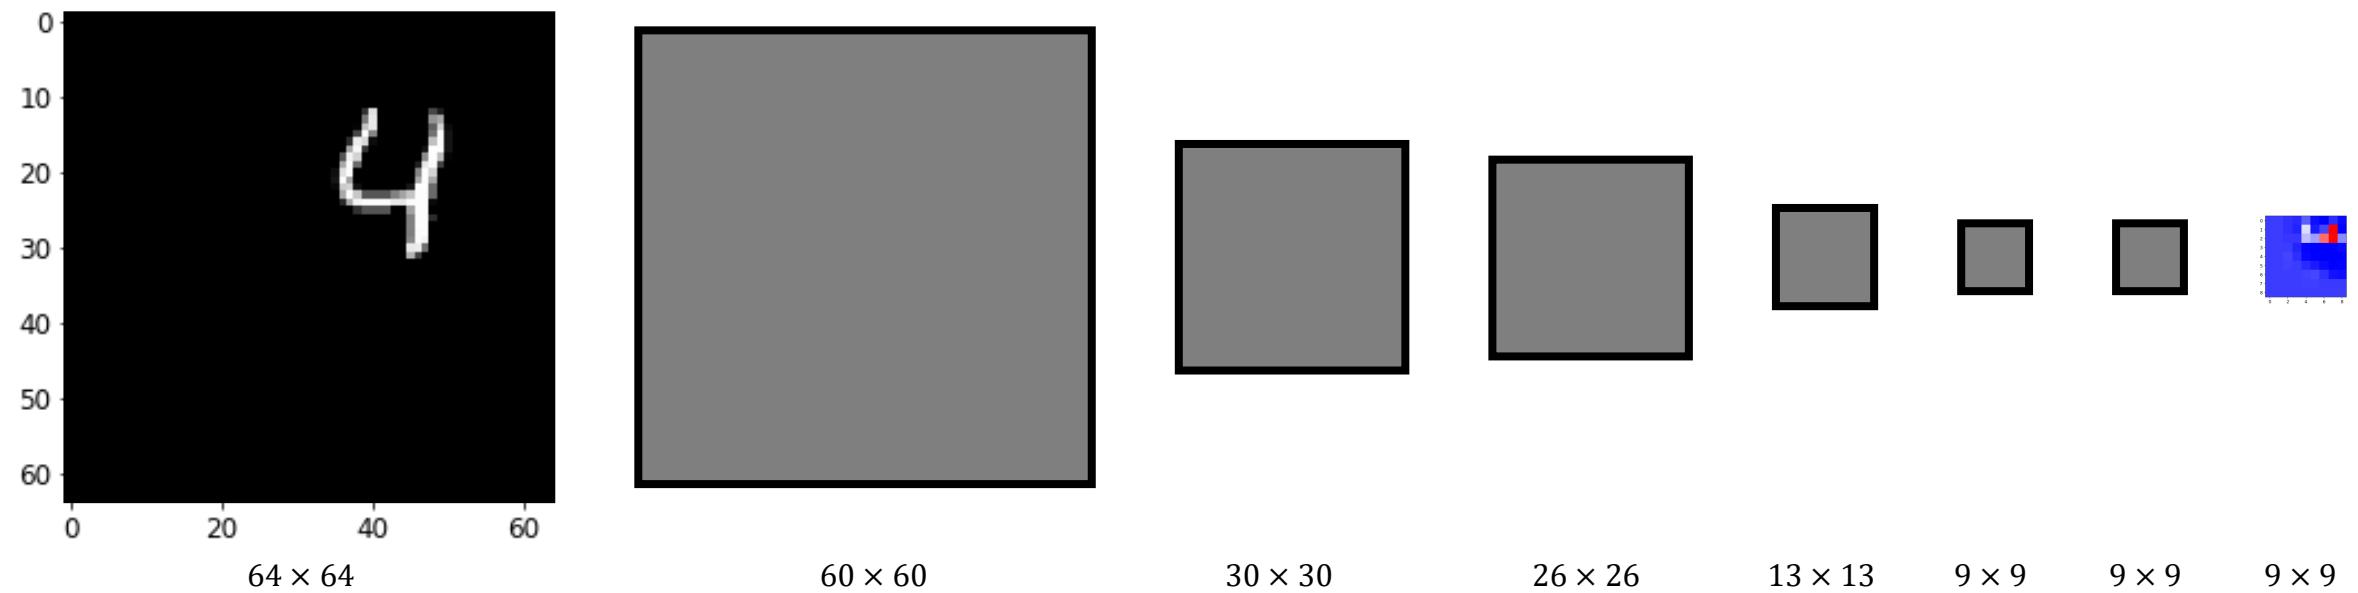
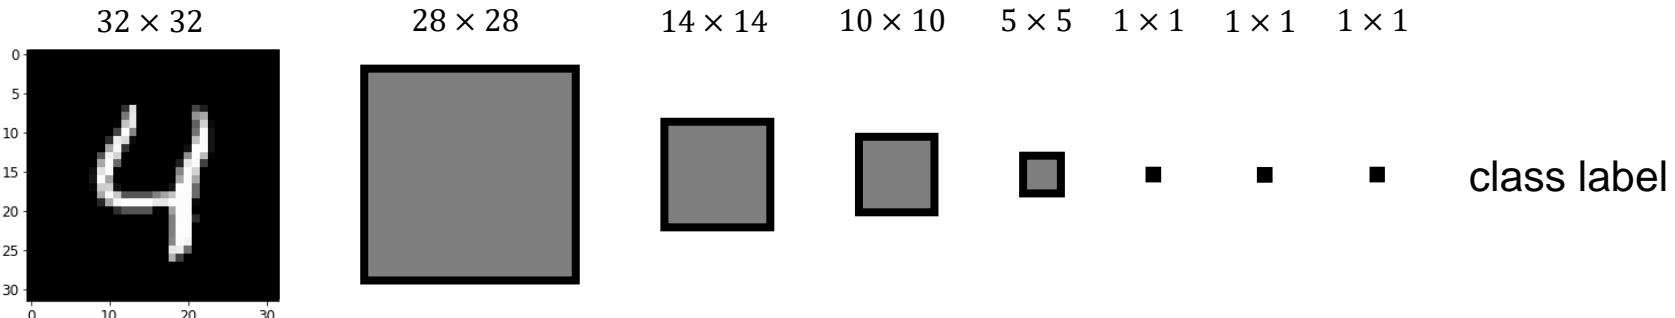
```
class FCLeNet(nn.Module):  
  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)  
        self.conv3 = nn.Conv2d(16, 120, kernel_size=5)  
        self.conv4 = nn.Conv2d(120, 84, kernel_size=1)  
        self.conv5 = nn.Conv2d(84, 10, kernel_size=1)  
  
    def forward(self, x):  
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)  
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)  
        x = F.relu(self.conv3(x))  
        x = F.relu(self.conv4(x))  
        x = self.conv5(x)  
        return F.log_softmax(x, dim=1)
```



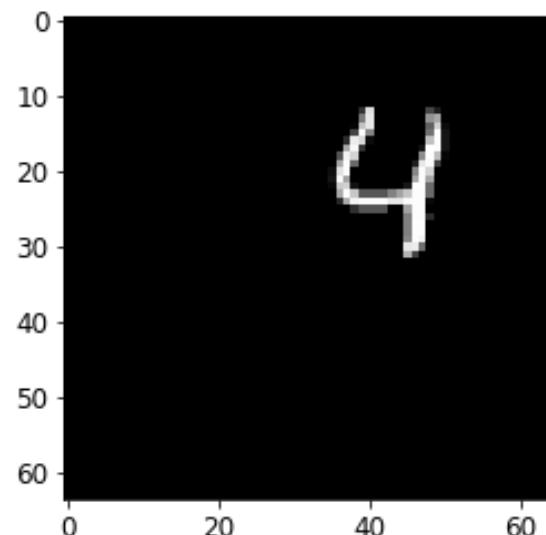
Beyond classification



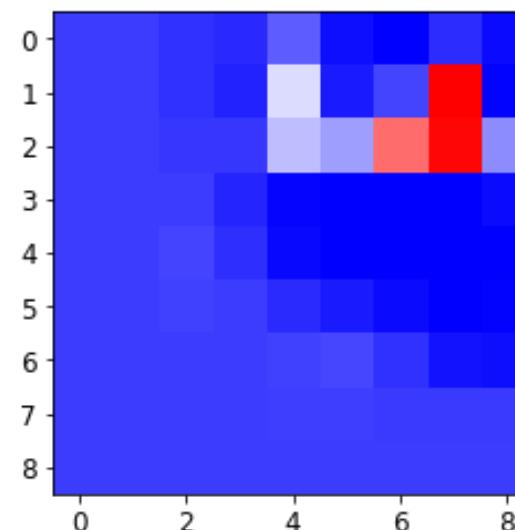
Beyond classification



Beyond classification

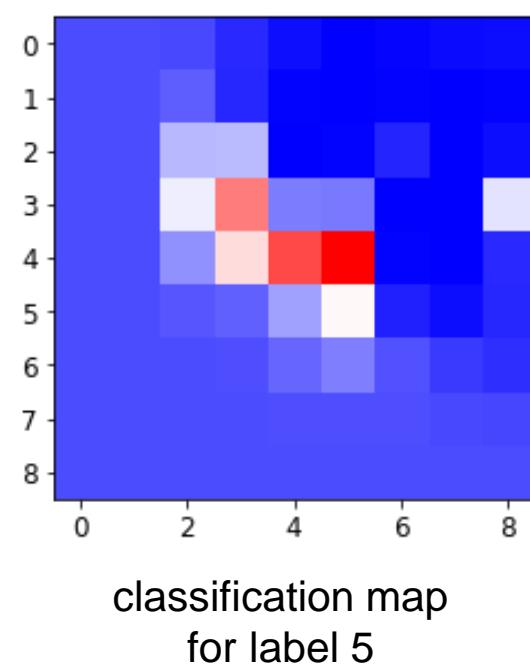
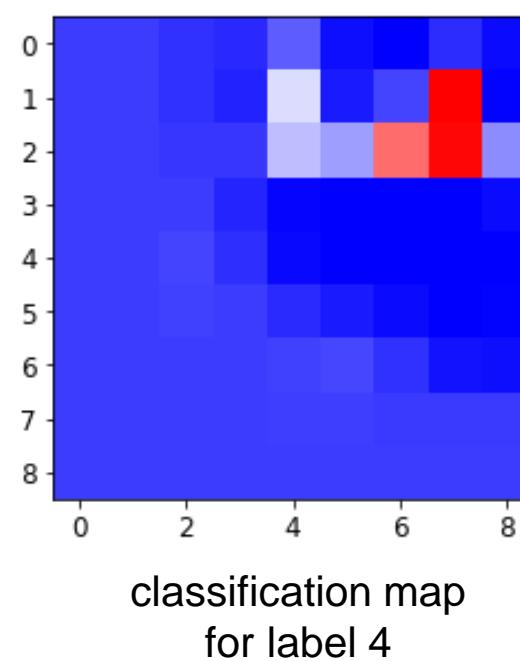
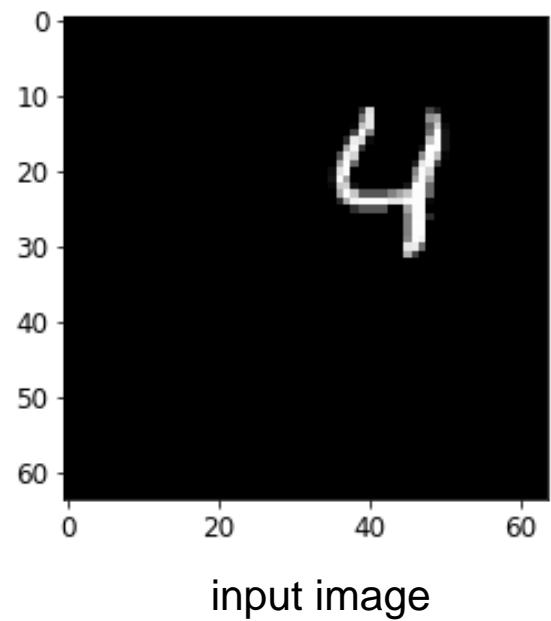


input image

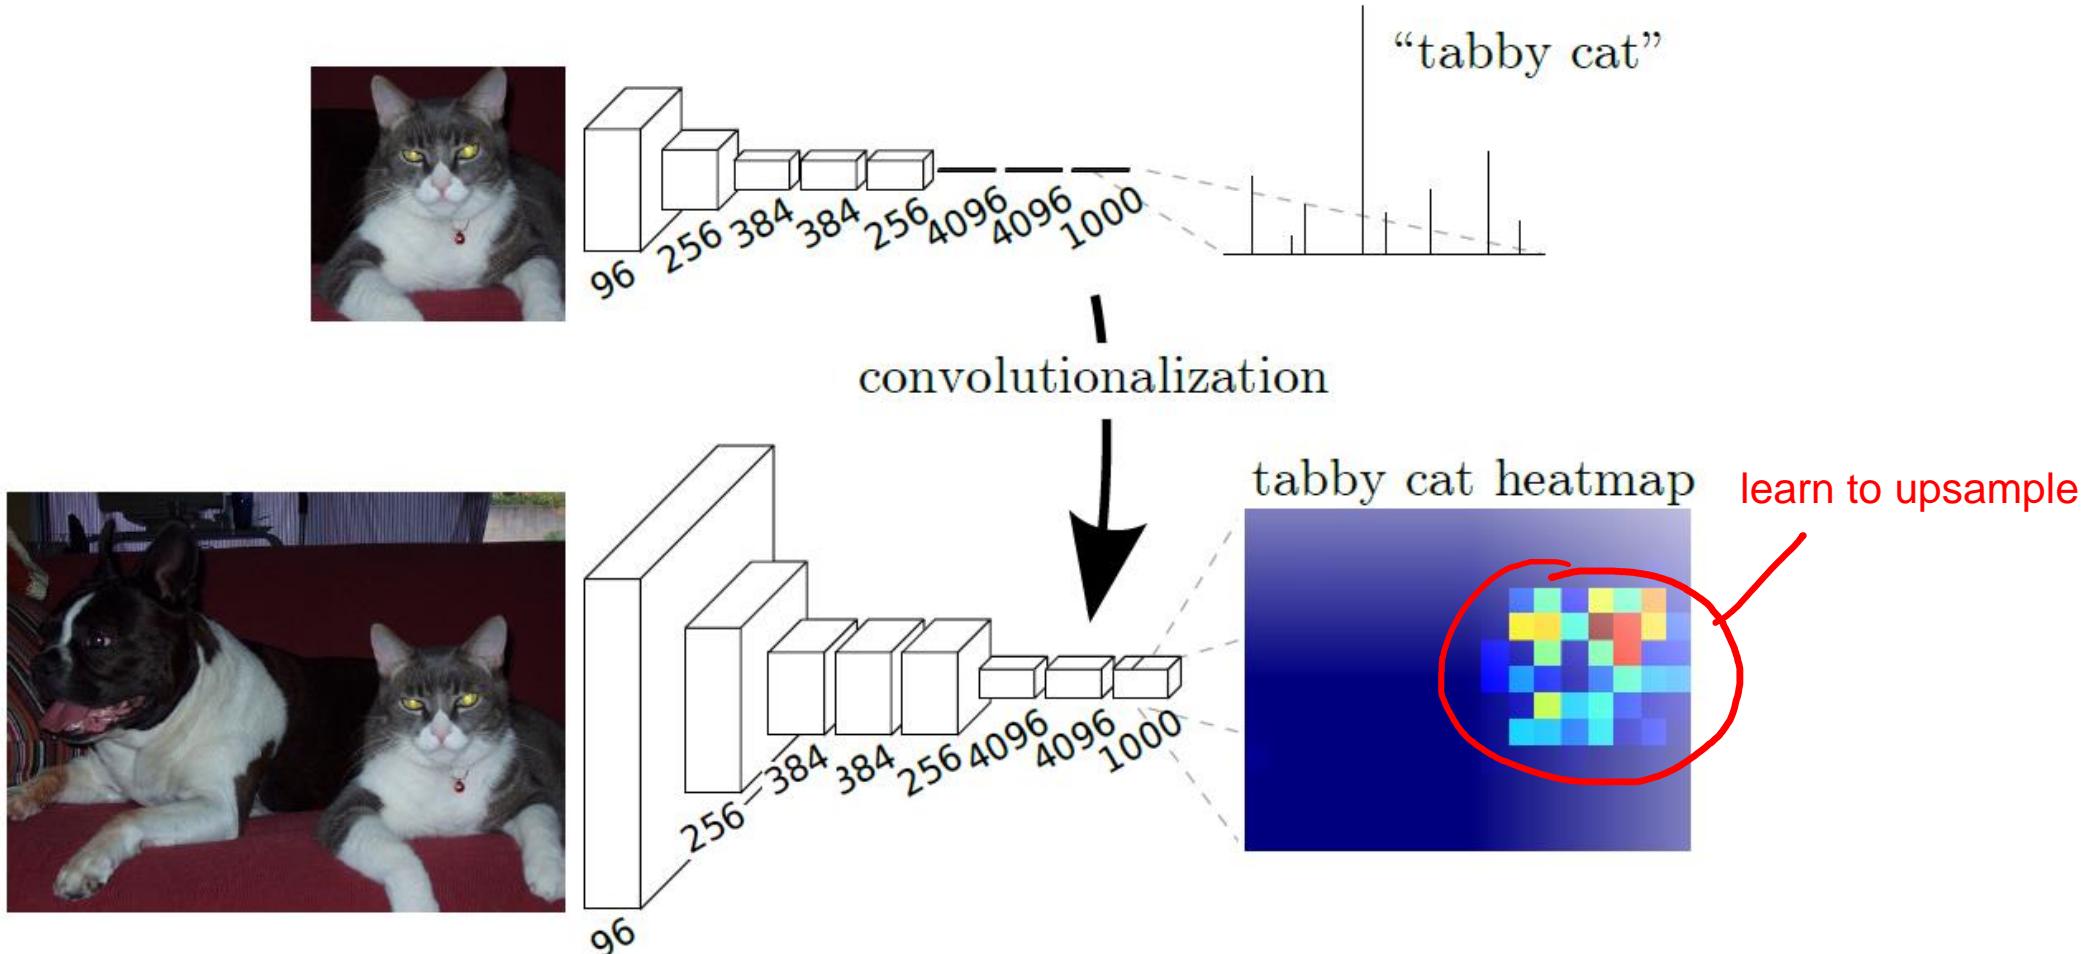


classification map
for label 4

Beyond classification



Beyond classification



Up-sampling

- Unpooling

1	2
3	4



1	1	2	2
1	1	2	2
3	3	4	4
3	3	4	4

average (un)pooling

1	2
3	4

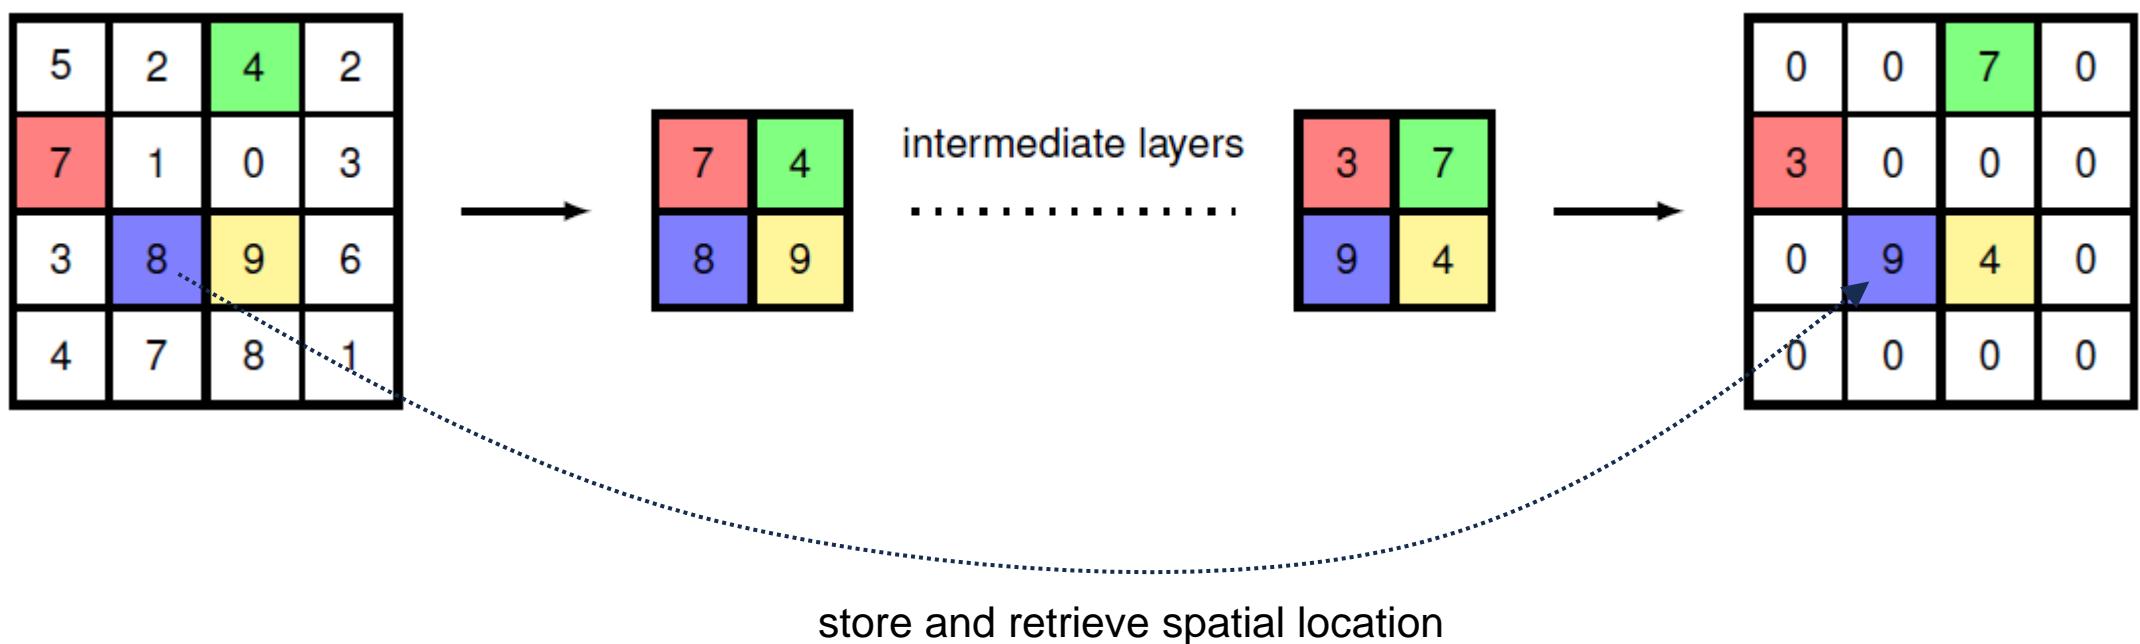


1	0	2	0
0	0	0	0
3	0	4	0
0	0	0	0

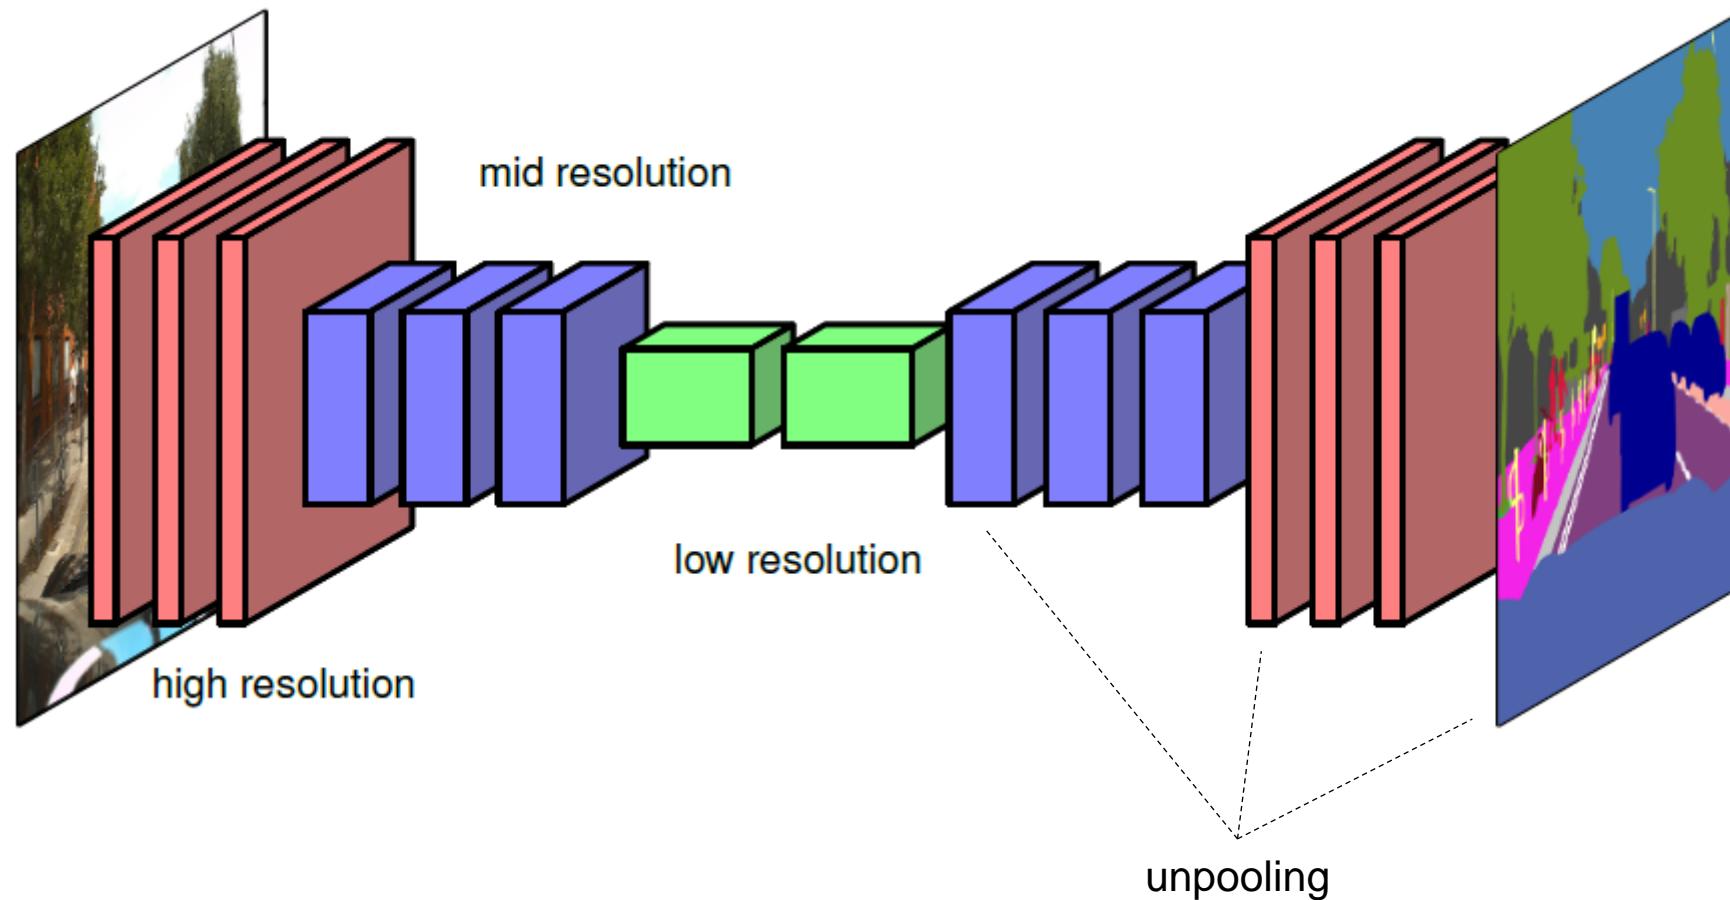
max (un)pooling

Up-sampling

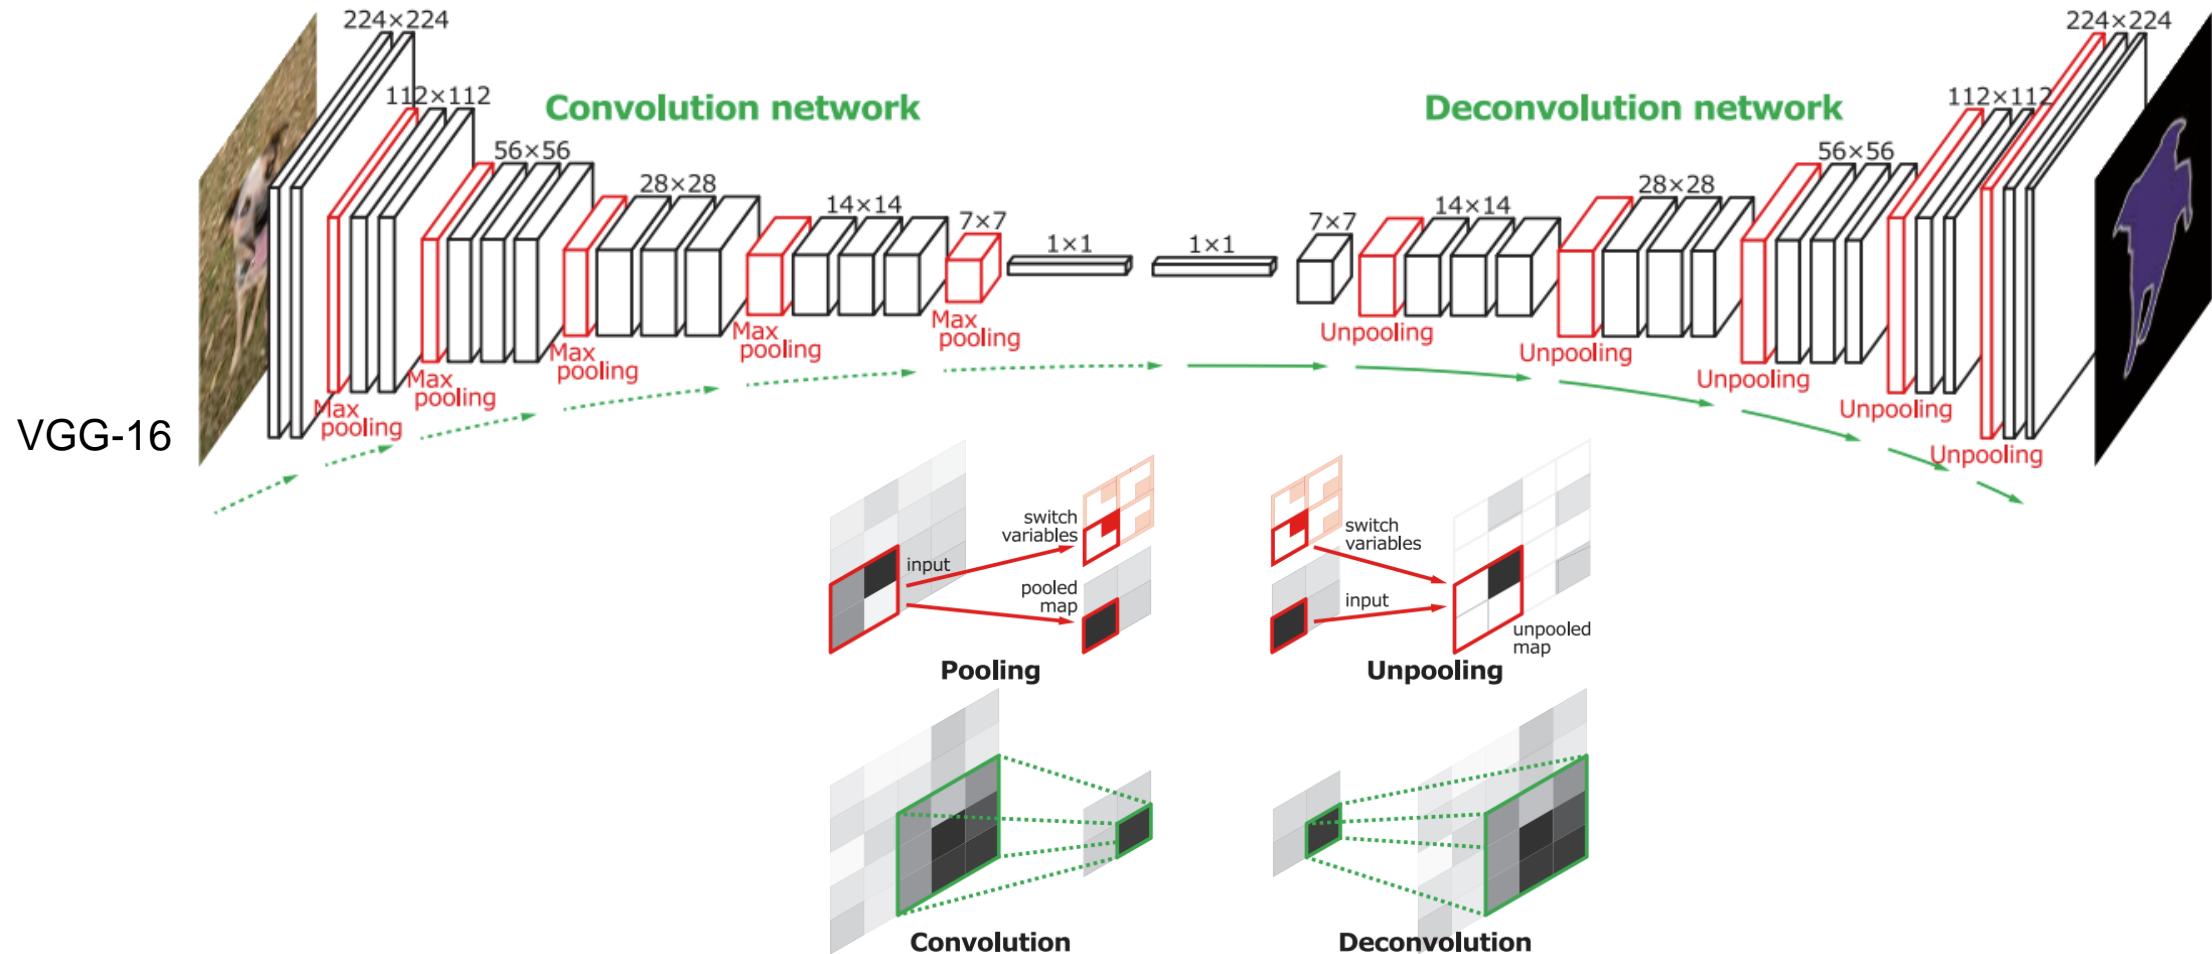
- Unpooling with spatial information



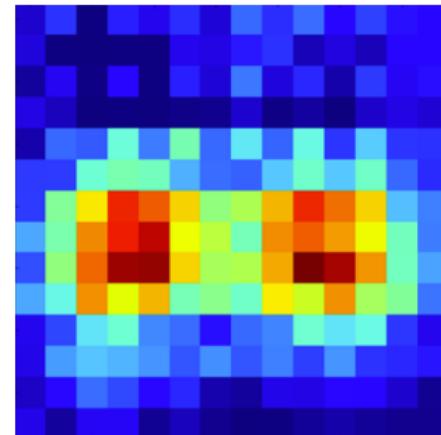
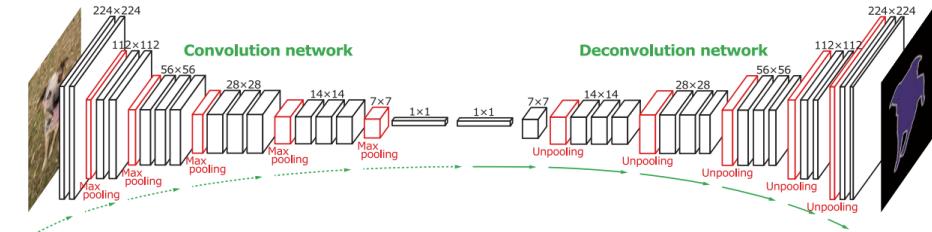
Encoder-decoder network



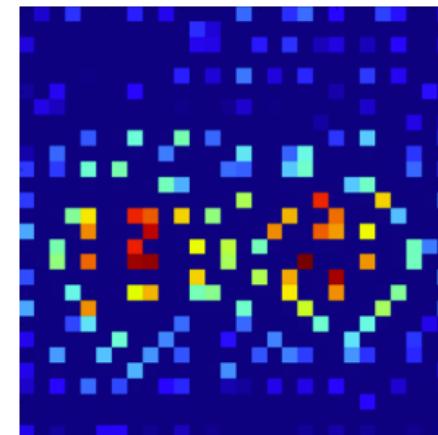
Encoder-decoder network



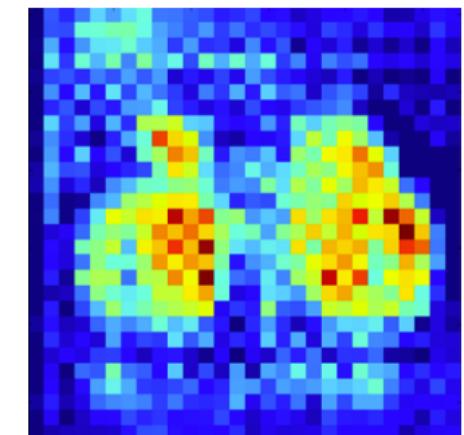
Encoder-decoder network



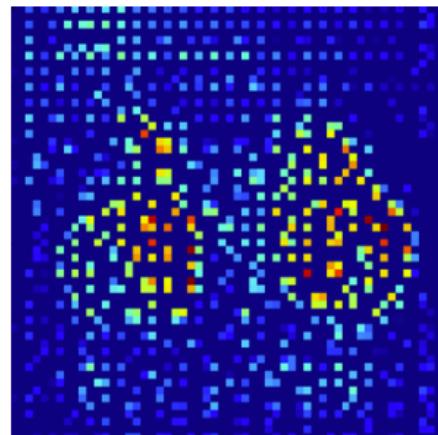
14×14 deconv



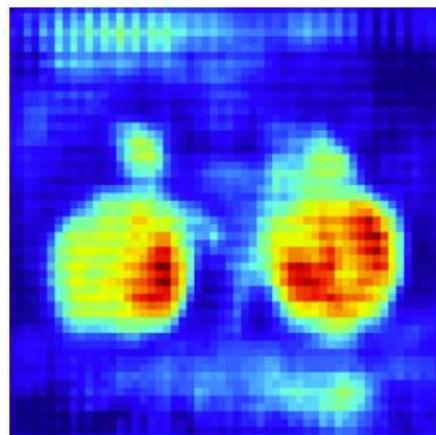
28×28 unpool



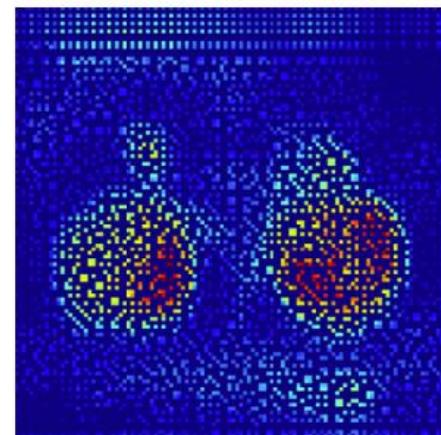
28×28 deconv



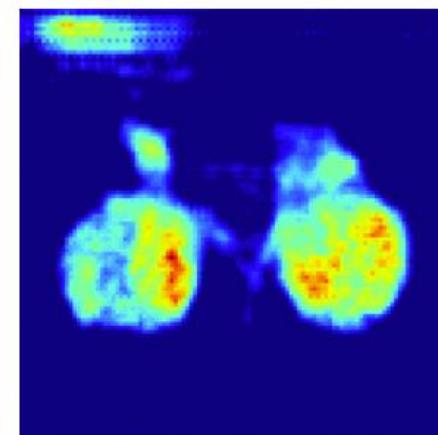
56×56 unpool



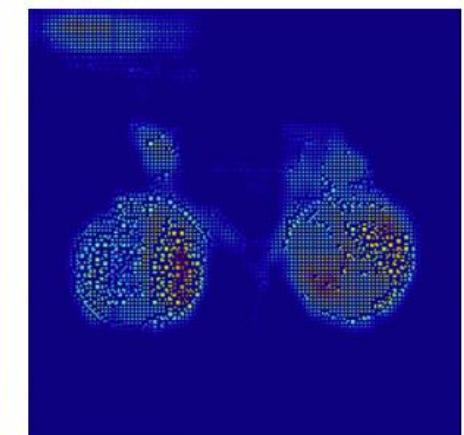
56×56 deconv



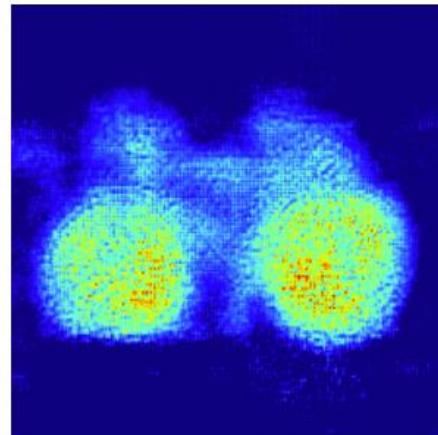
112×112 unpool



112×112 deconv

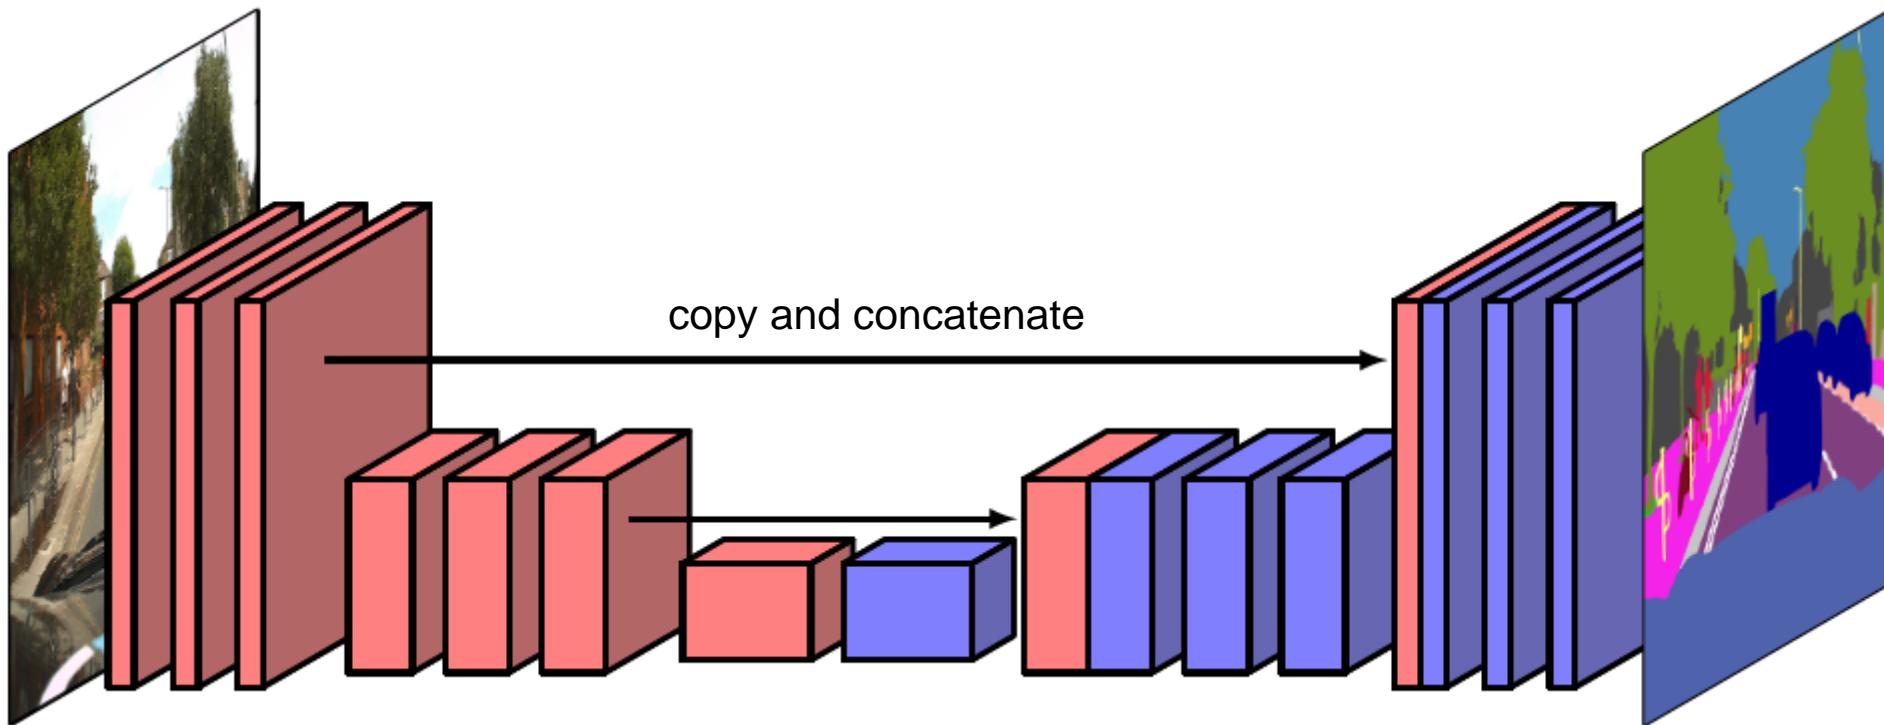


224×224 unpool

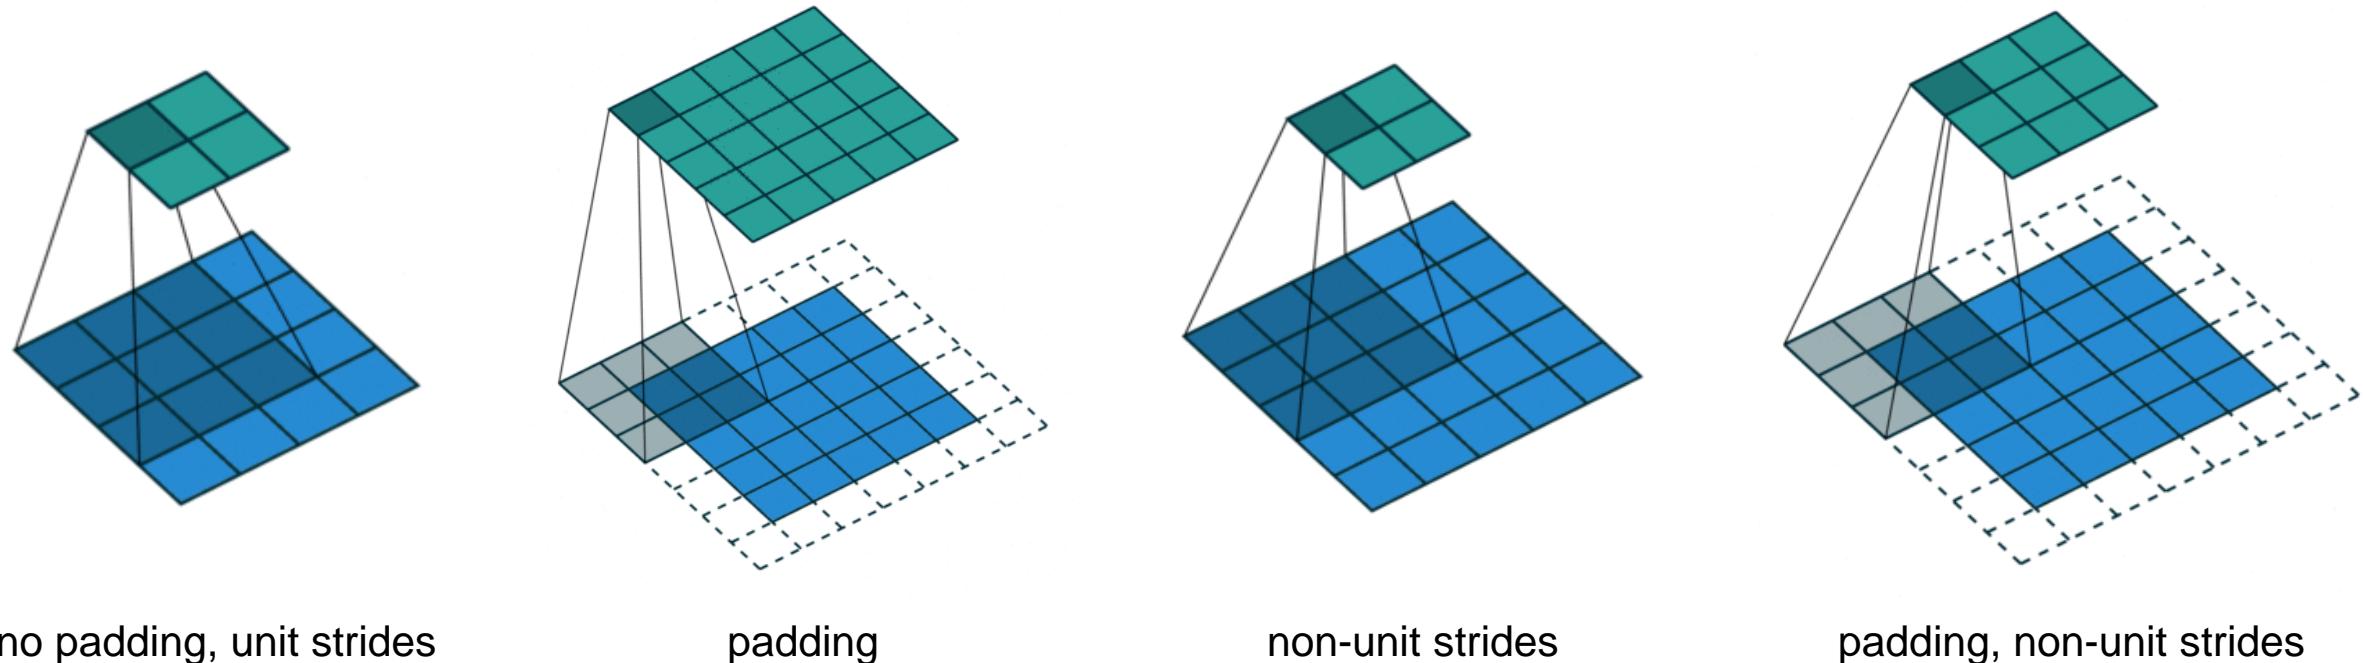


224×224 deconv

U-Net architecture

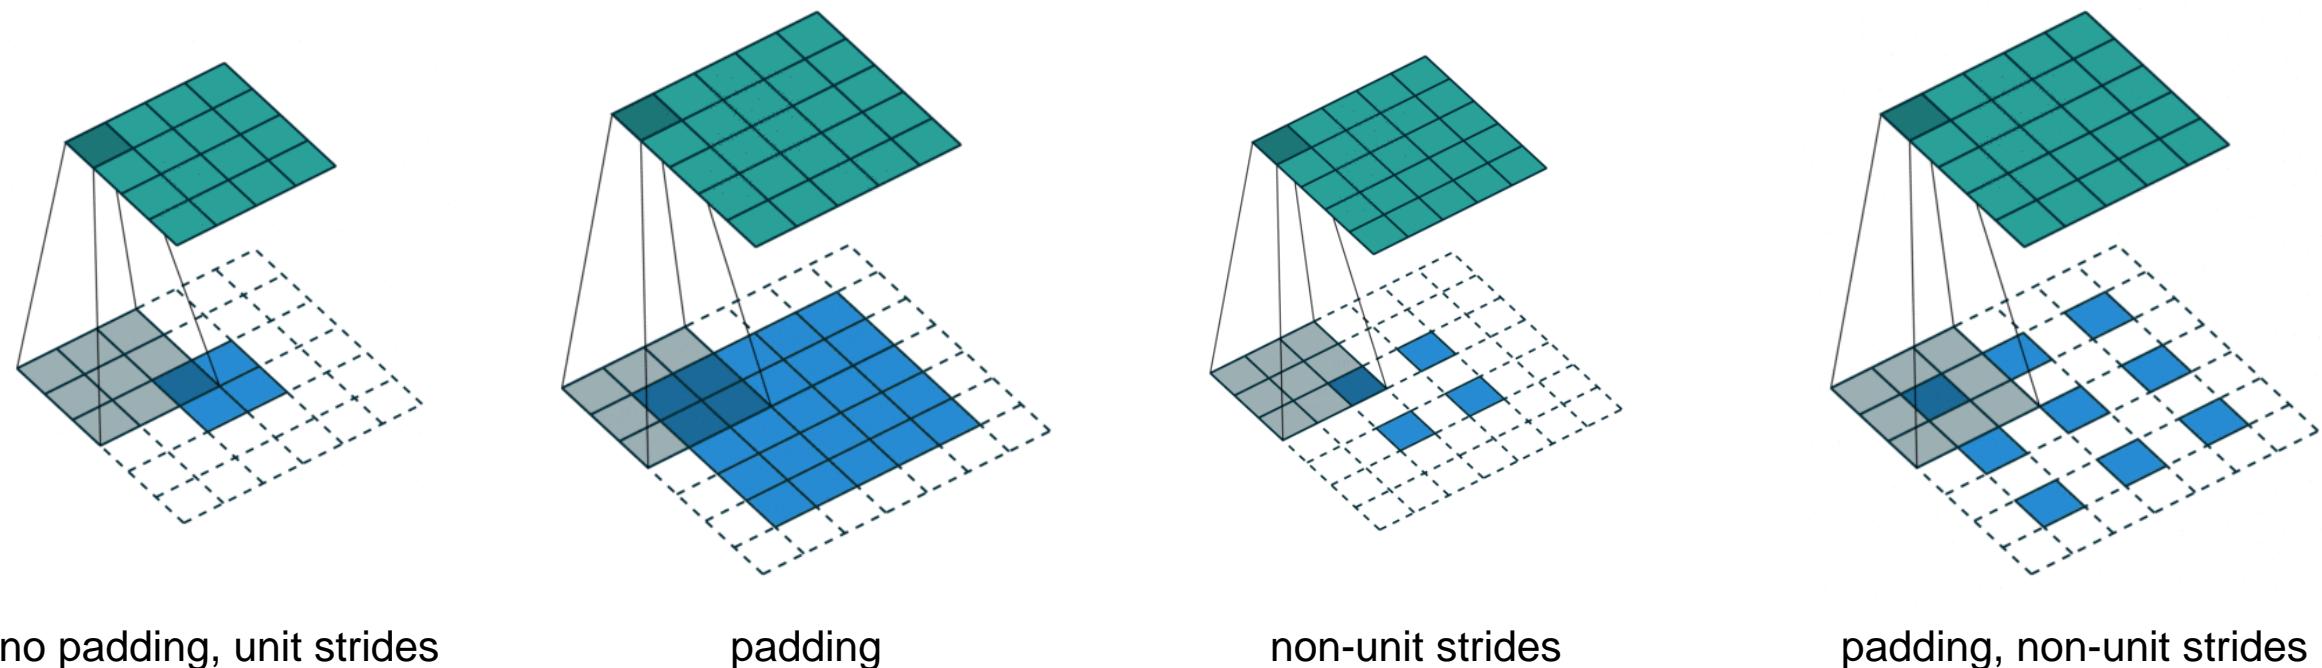


Convolutions



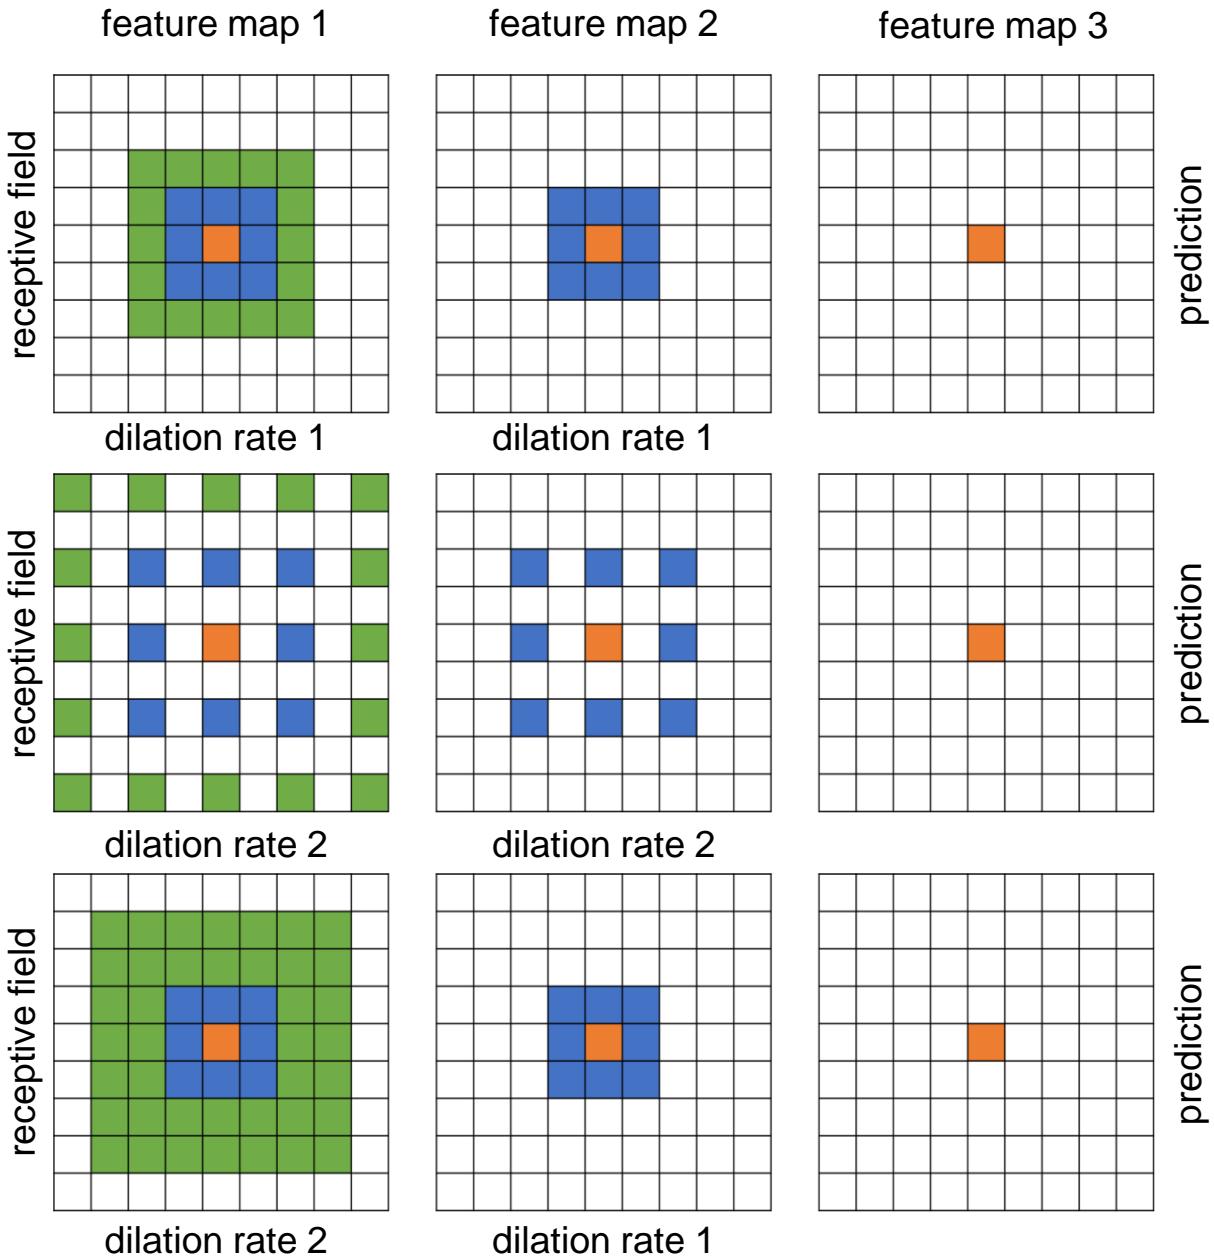
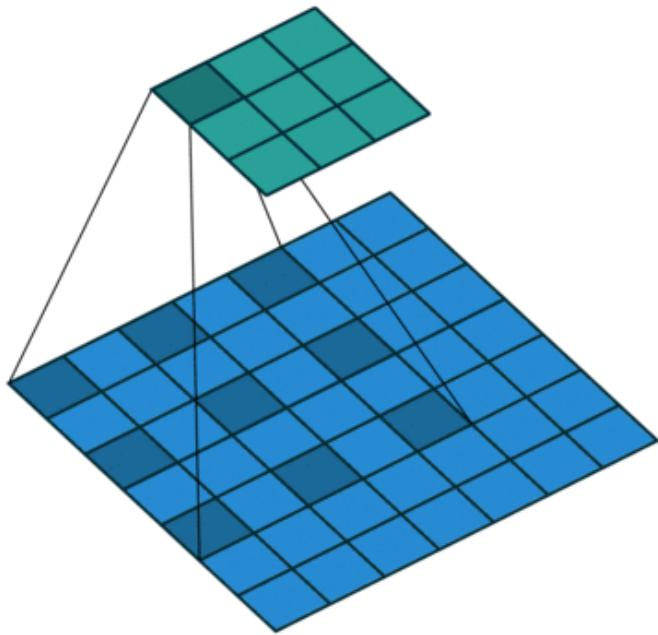
Convolutions can be used for 'learned' down-sampling

Transpose convolutions

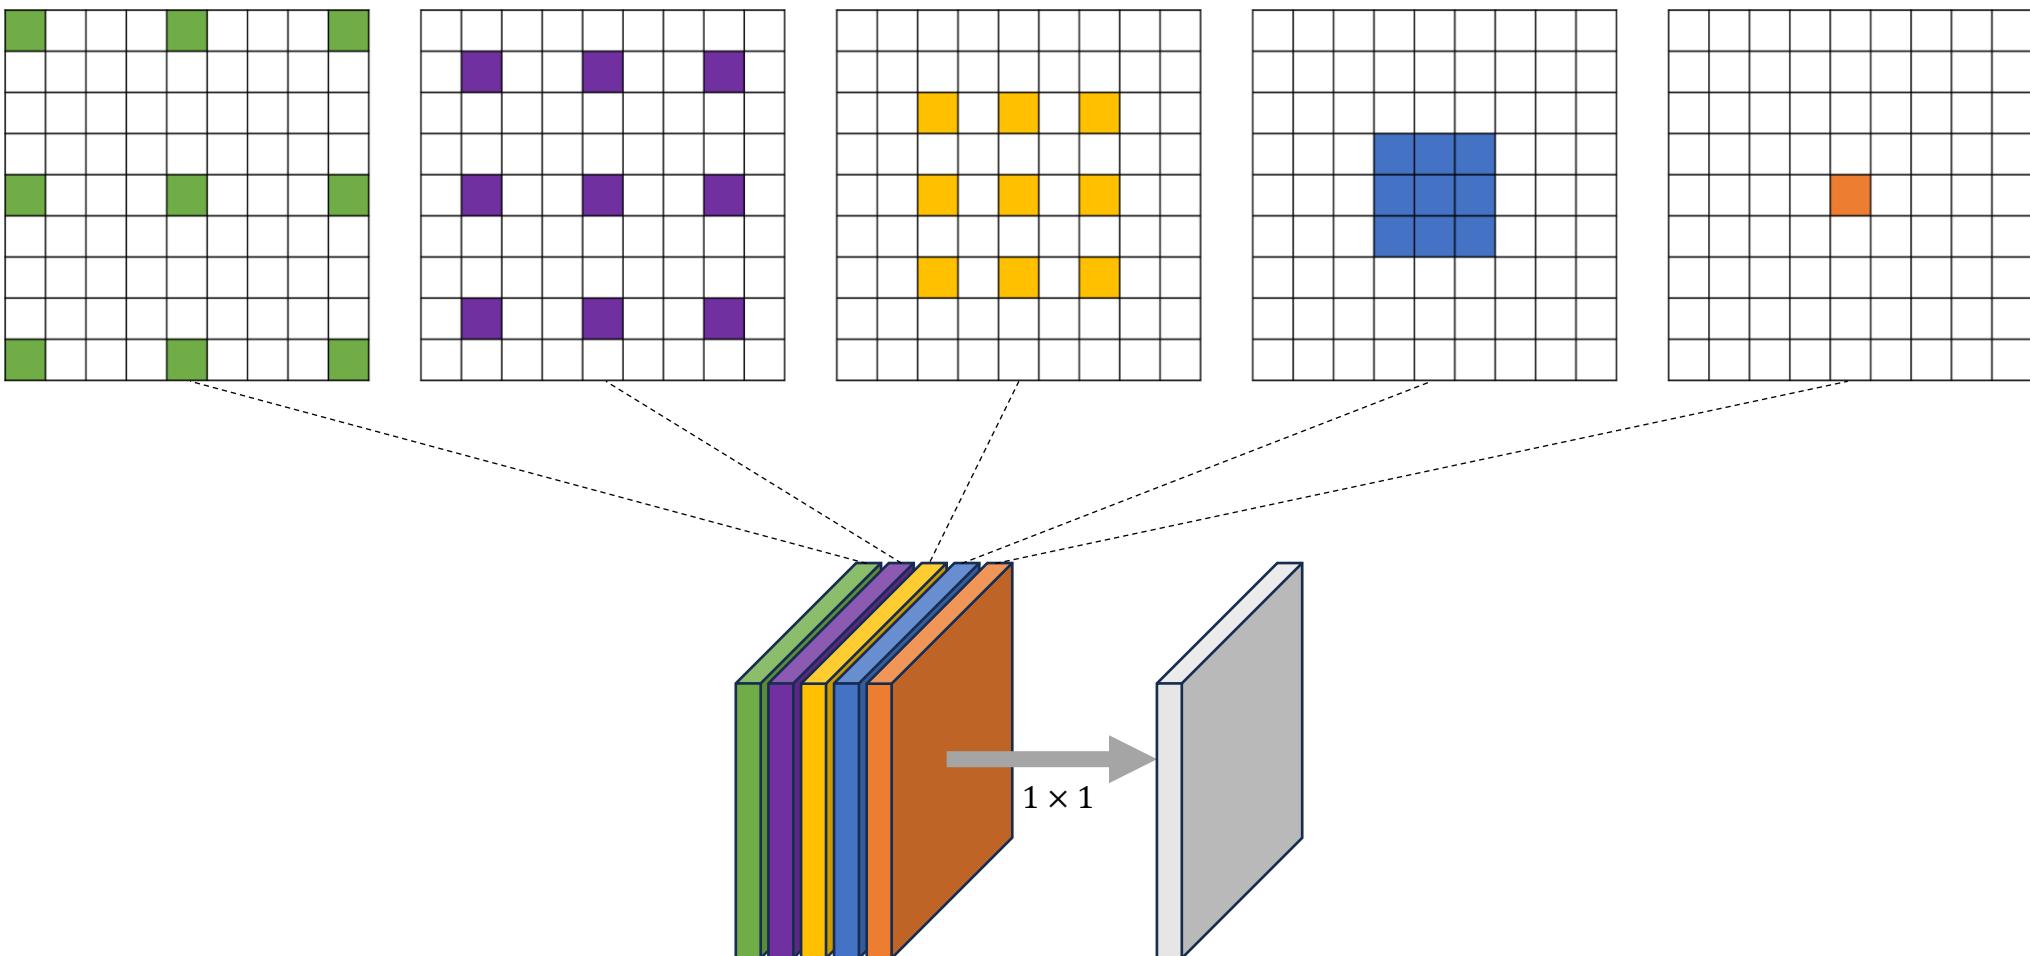


Transpose convolutions can be used for ‘learned’ up-sampling

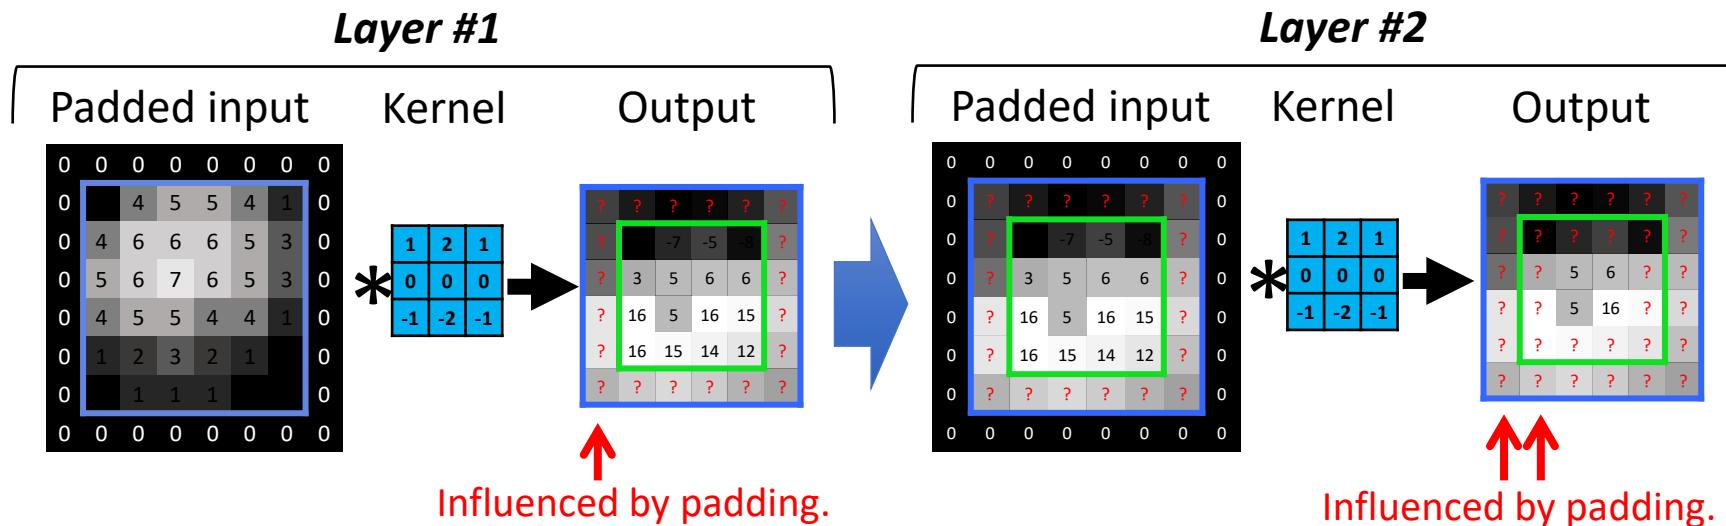
Dilated convolutions



Atrous spatial pyramid pooling

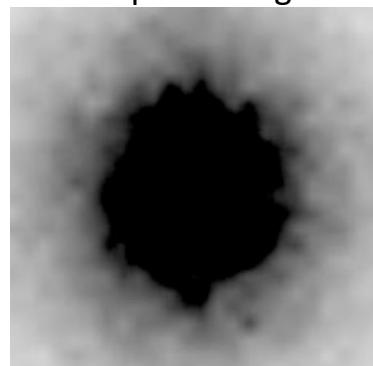


Padding effects

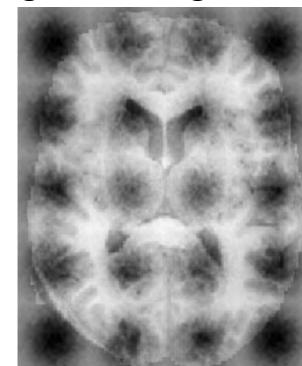
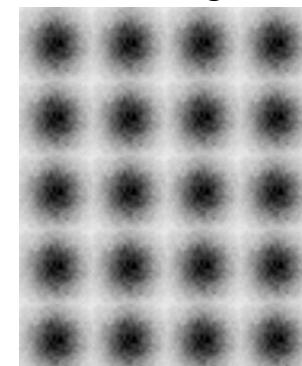


Artifacts that may show up due to padding (exaggerated for clarity):

If processing whole image at testing.

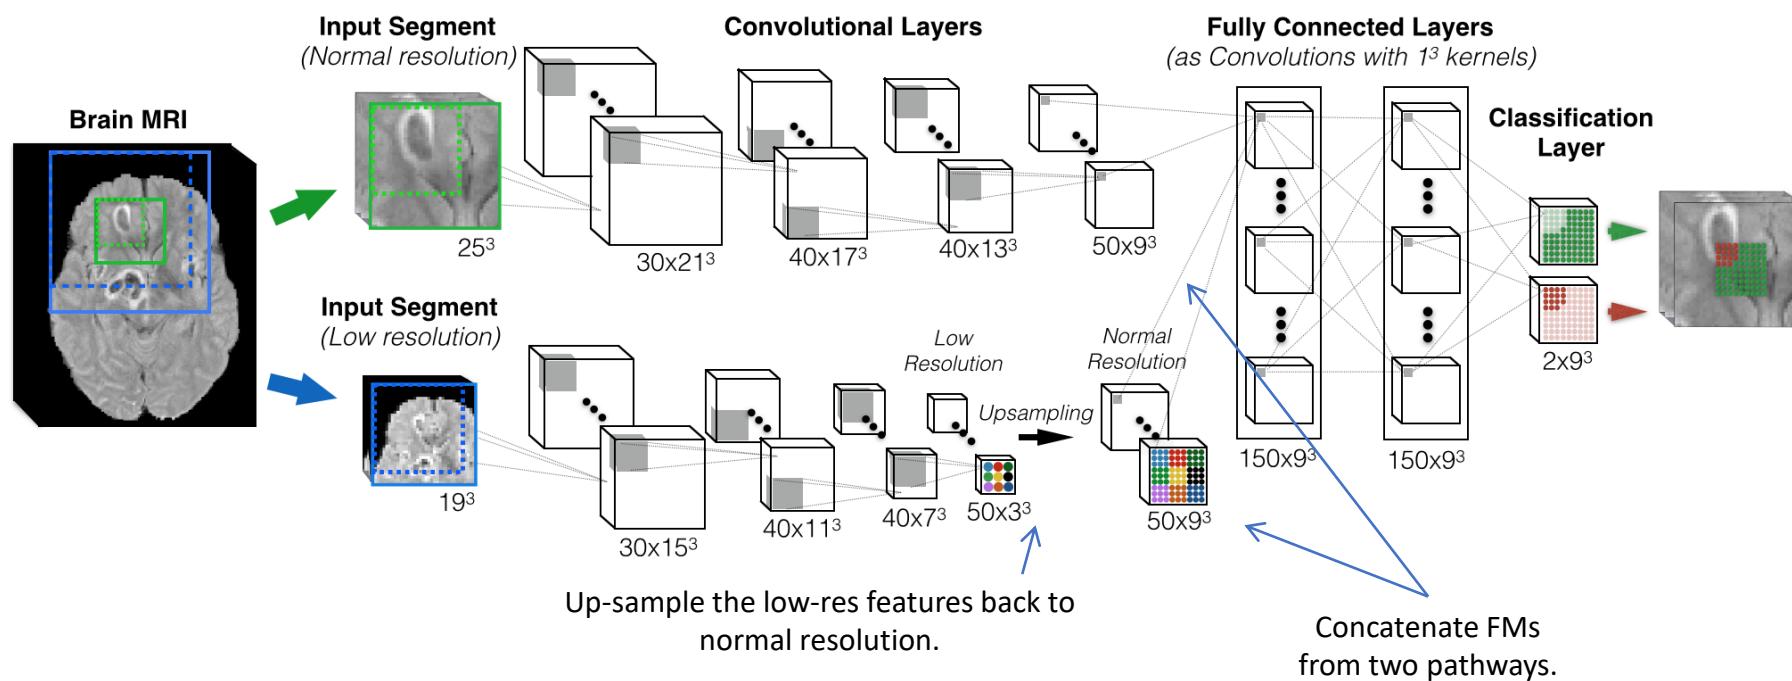


If tiling the image at testing.

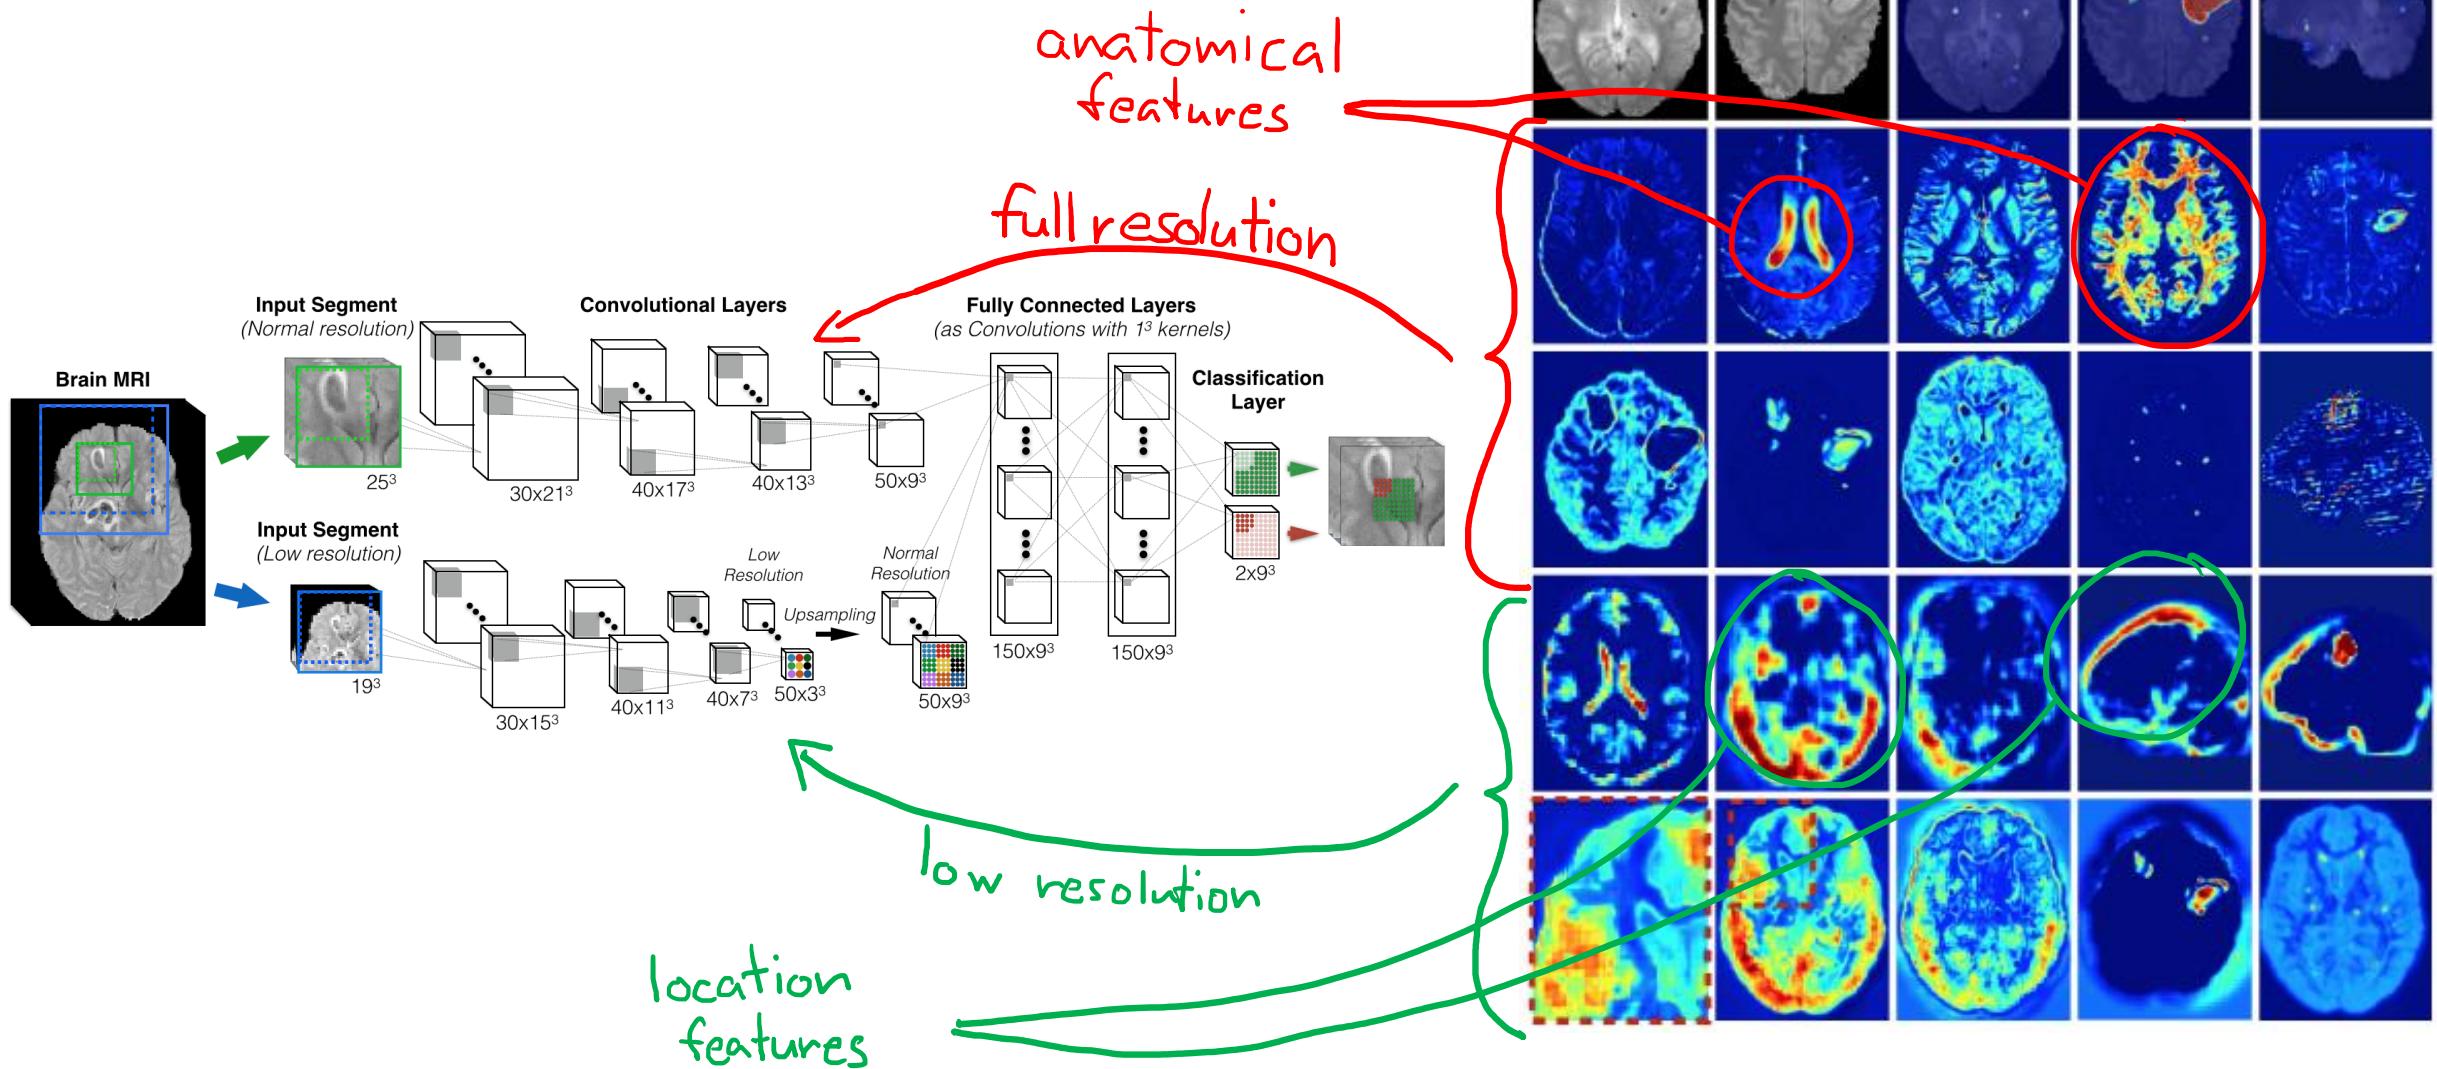


Multi-scale processing

- What else can we do to make the network “seeing” more context?
- **Idea:** Add pathways which process down-sampled images

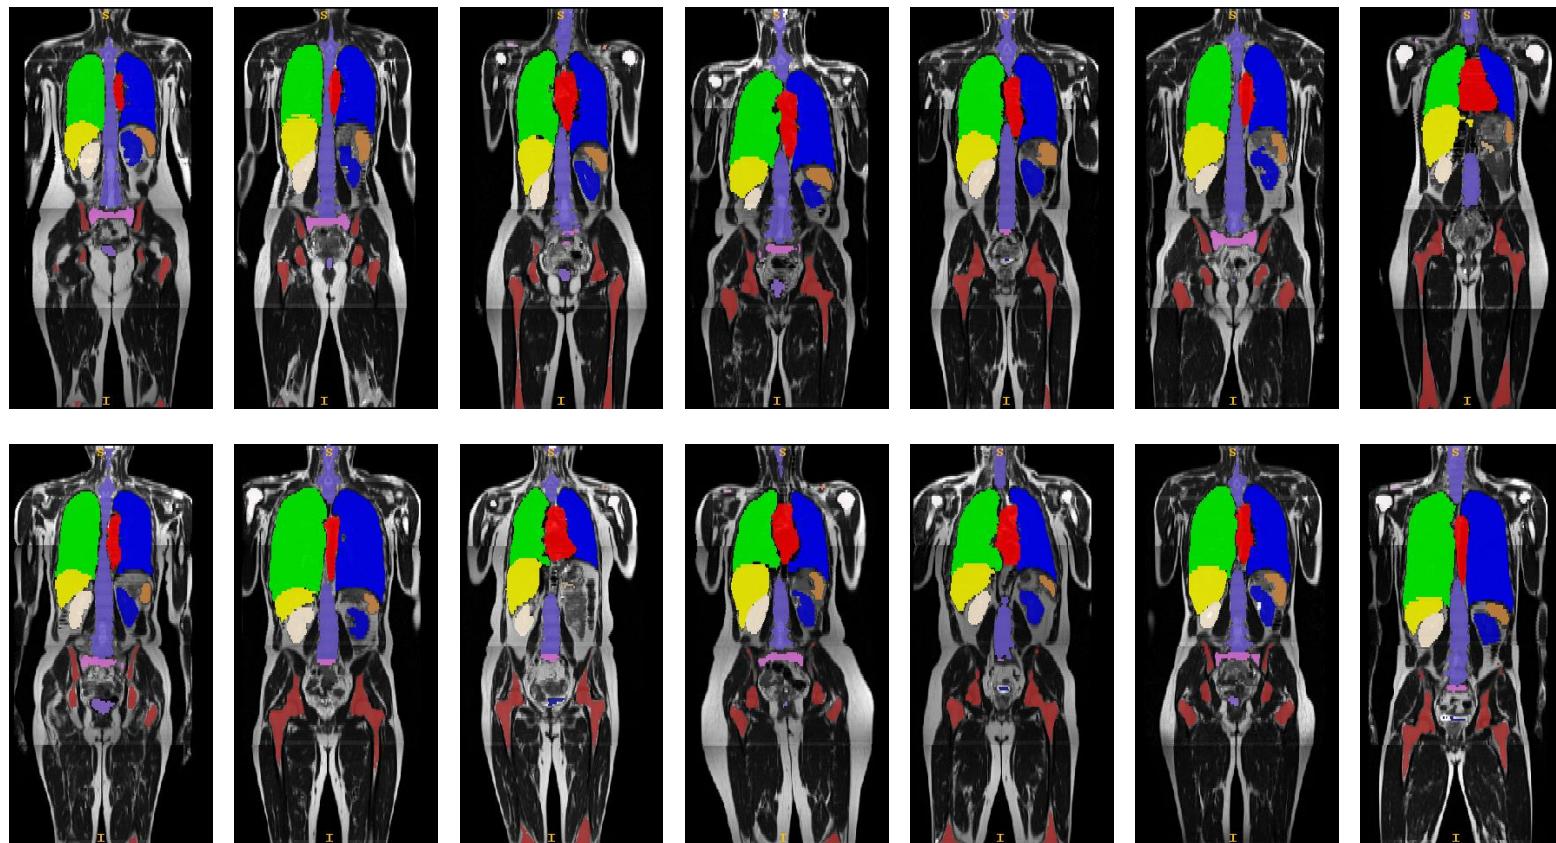


Multi-scale feature maps



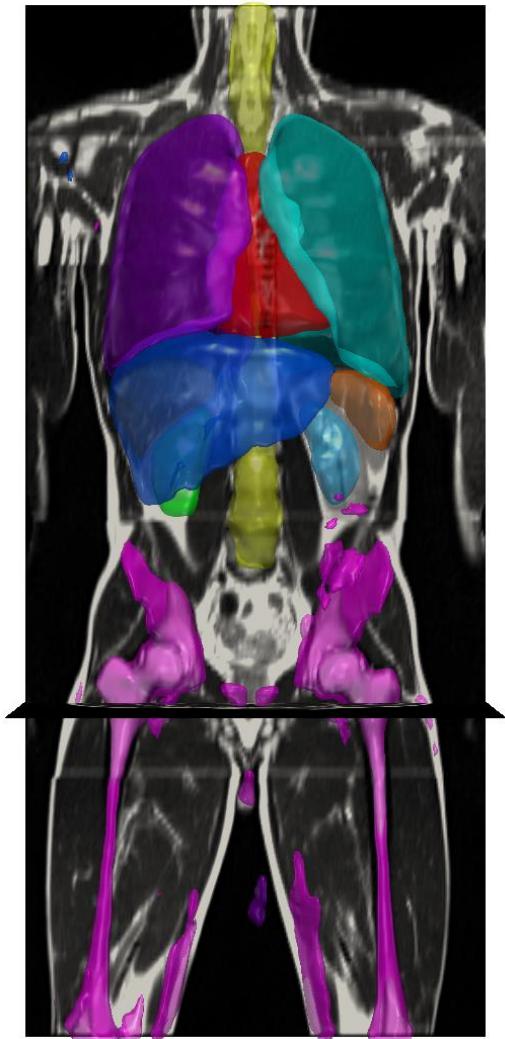
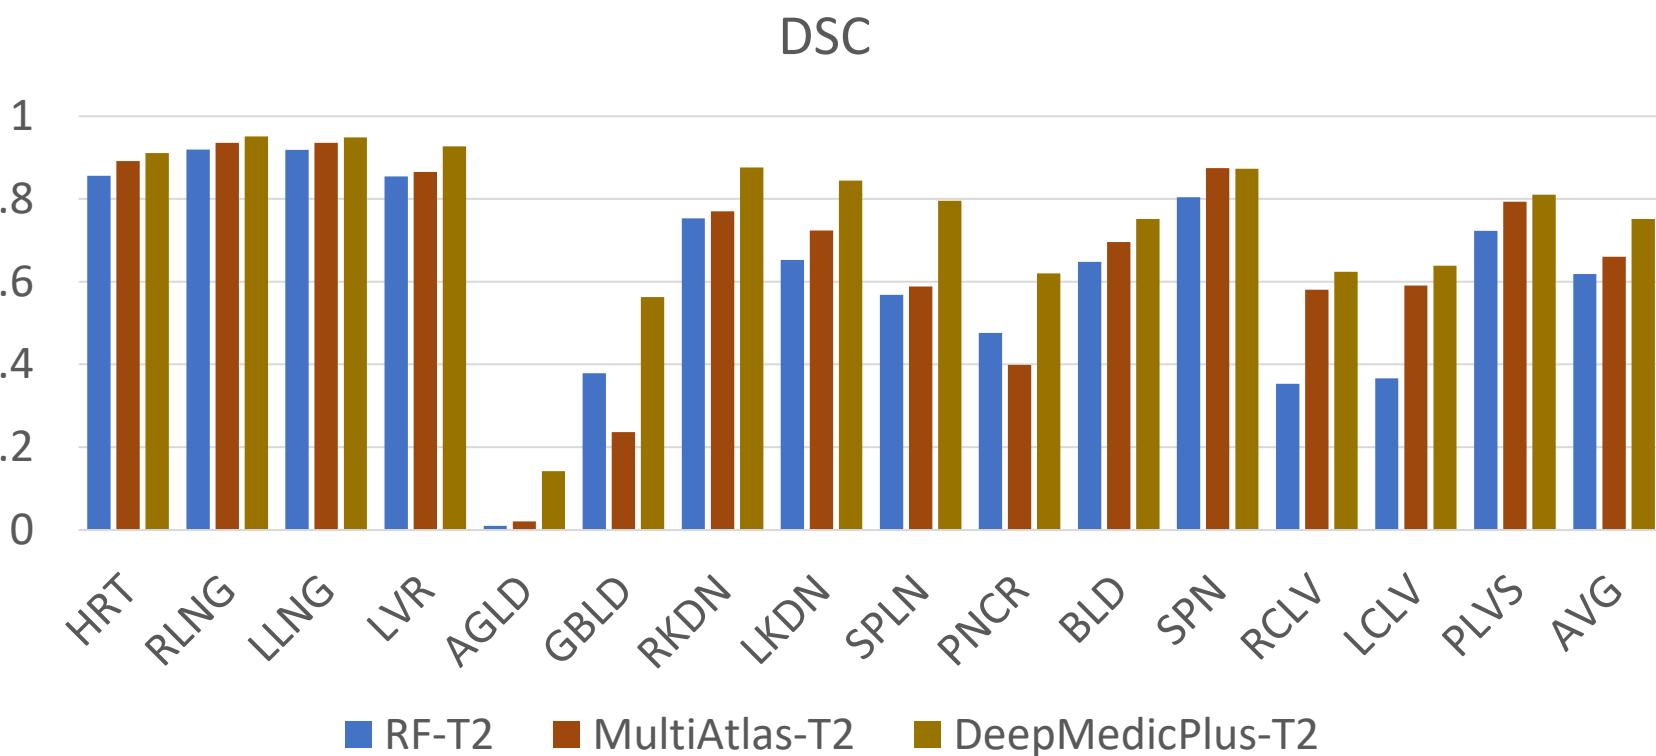
Example: Multi-organ segmentation

- Multi-atlas vs random forests vs CNNs



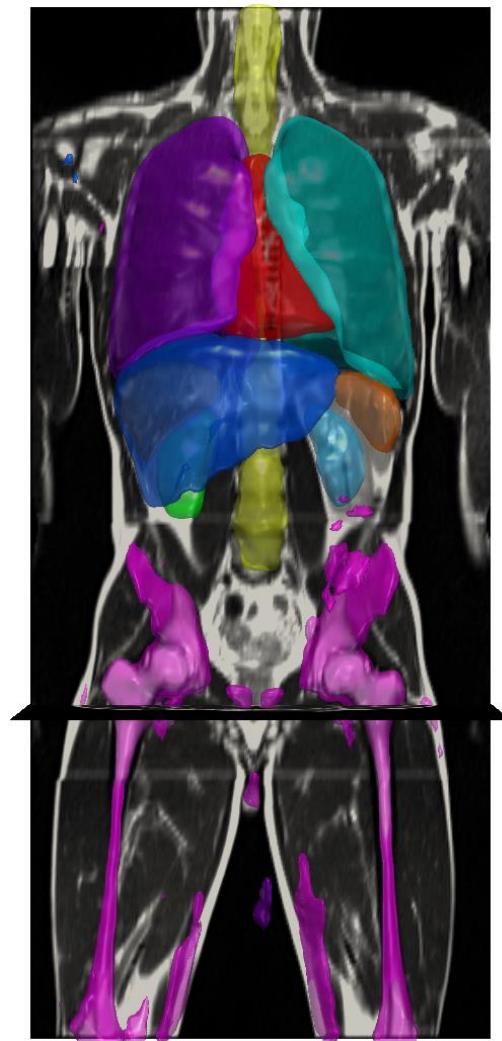
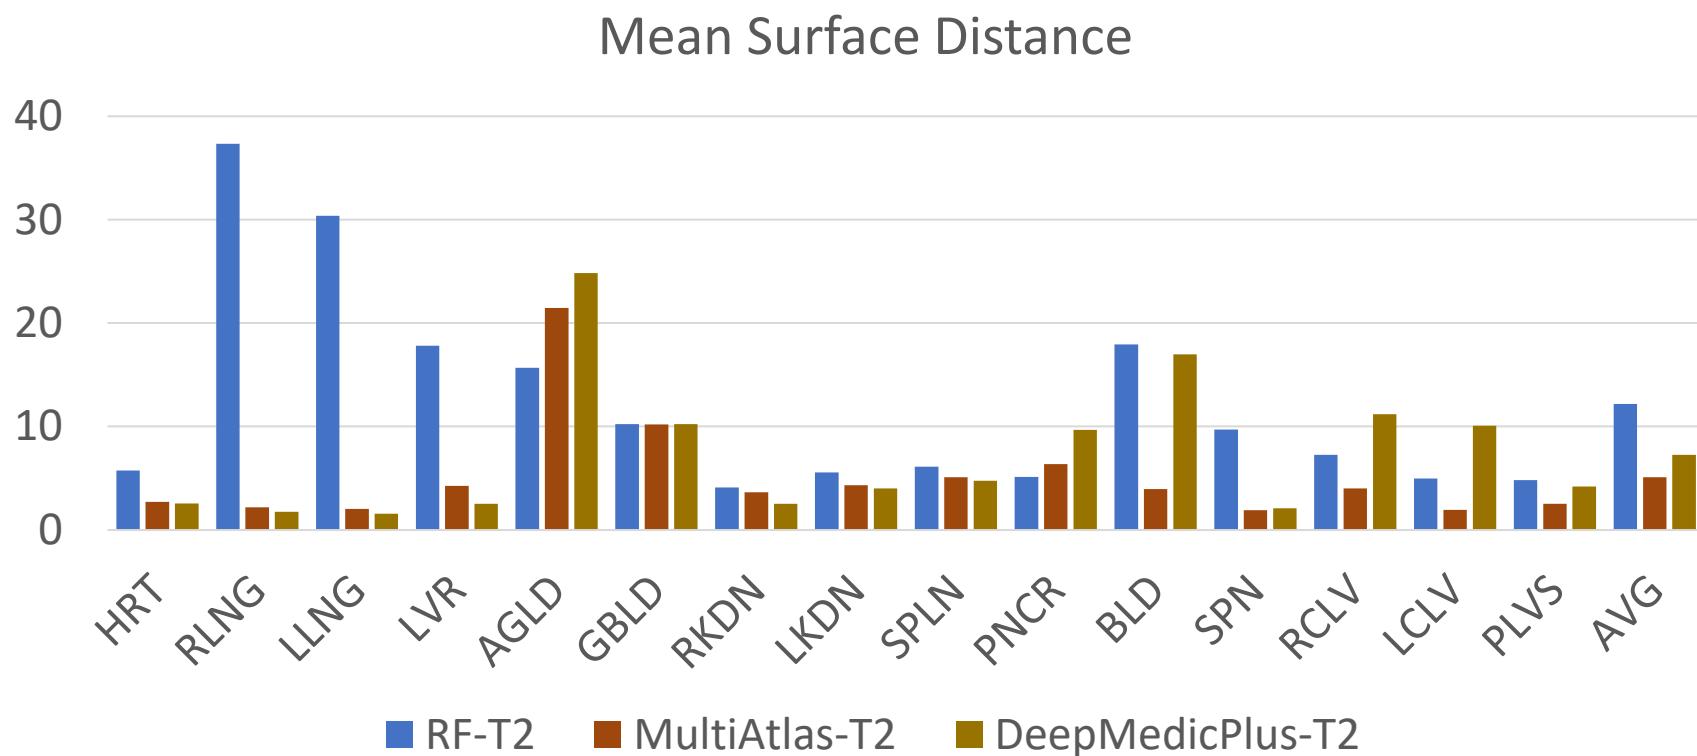
Example: Multi-organ segmentation

- Multi-atlas vs random forests vs CNNs

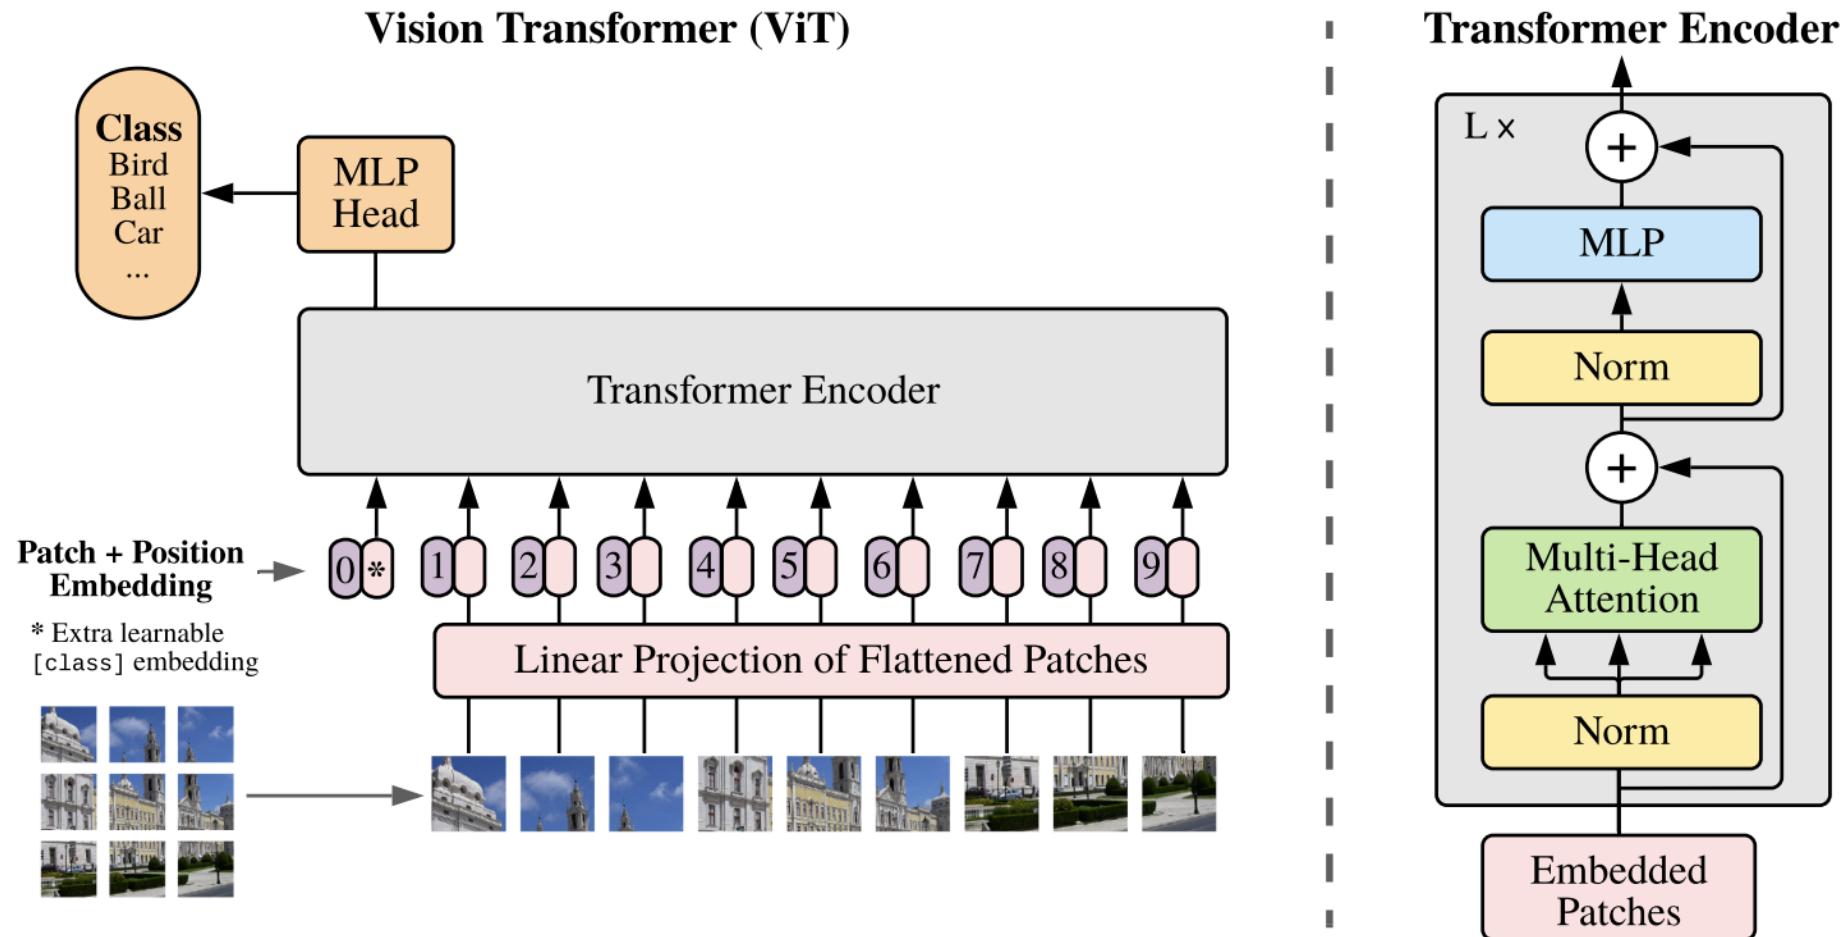


Example: Multi-organ segmentation

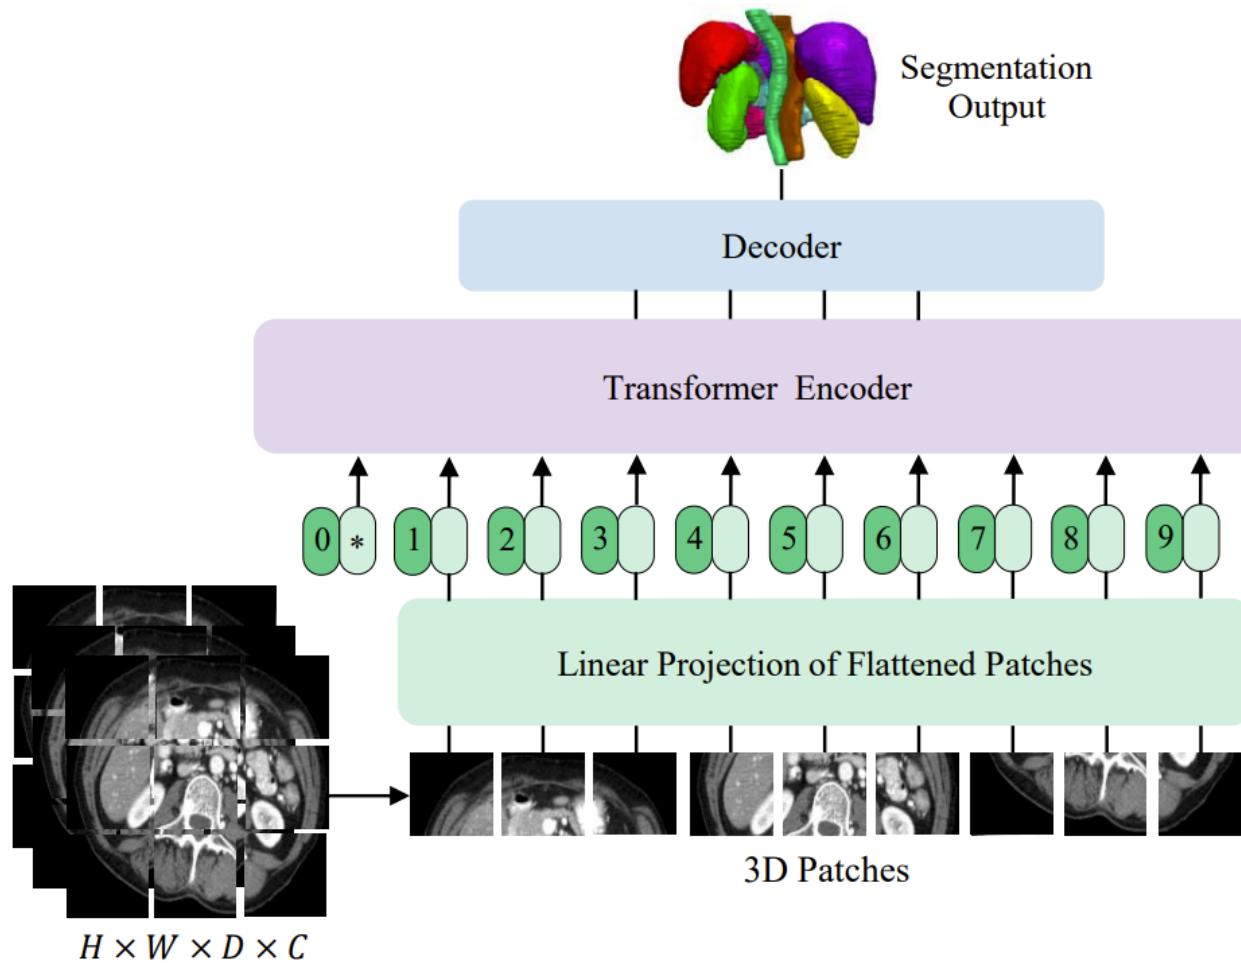
- Multi-atlas vs random forests vs CNNs



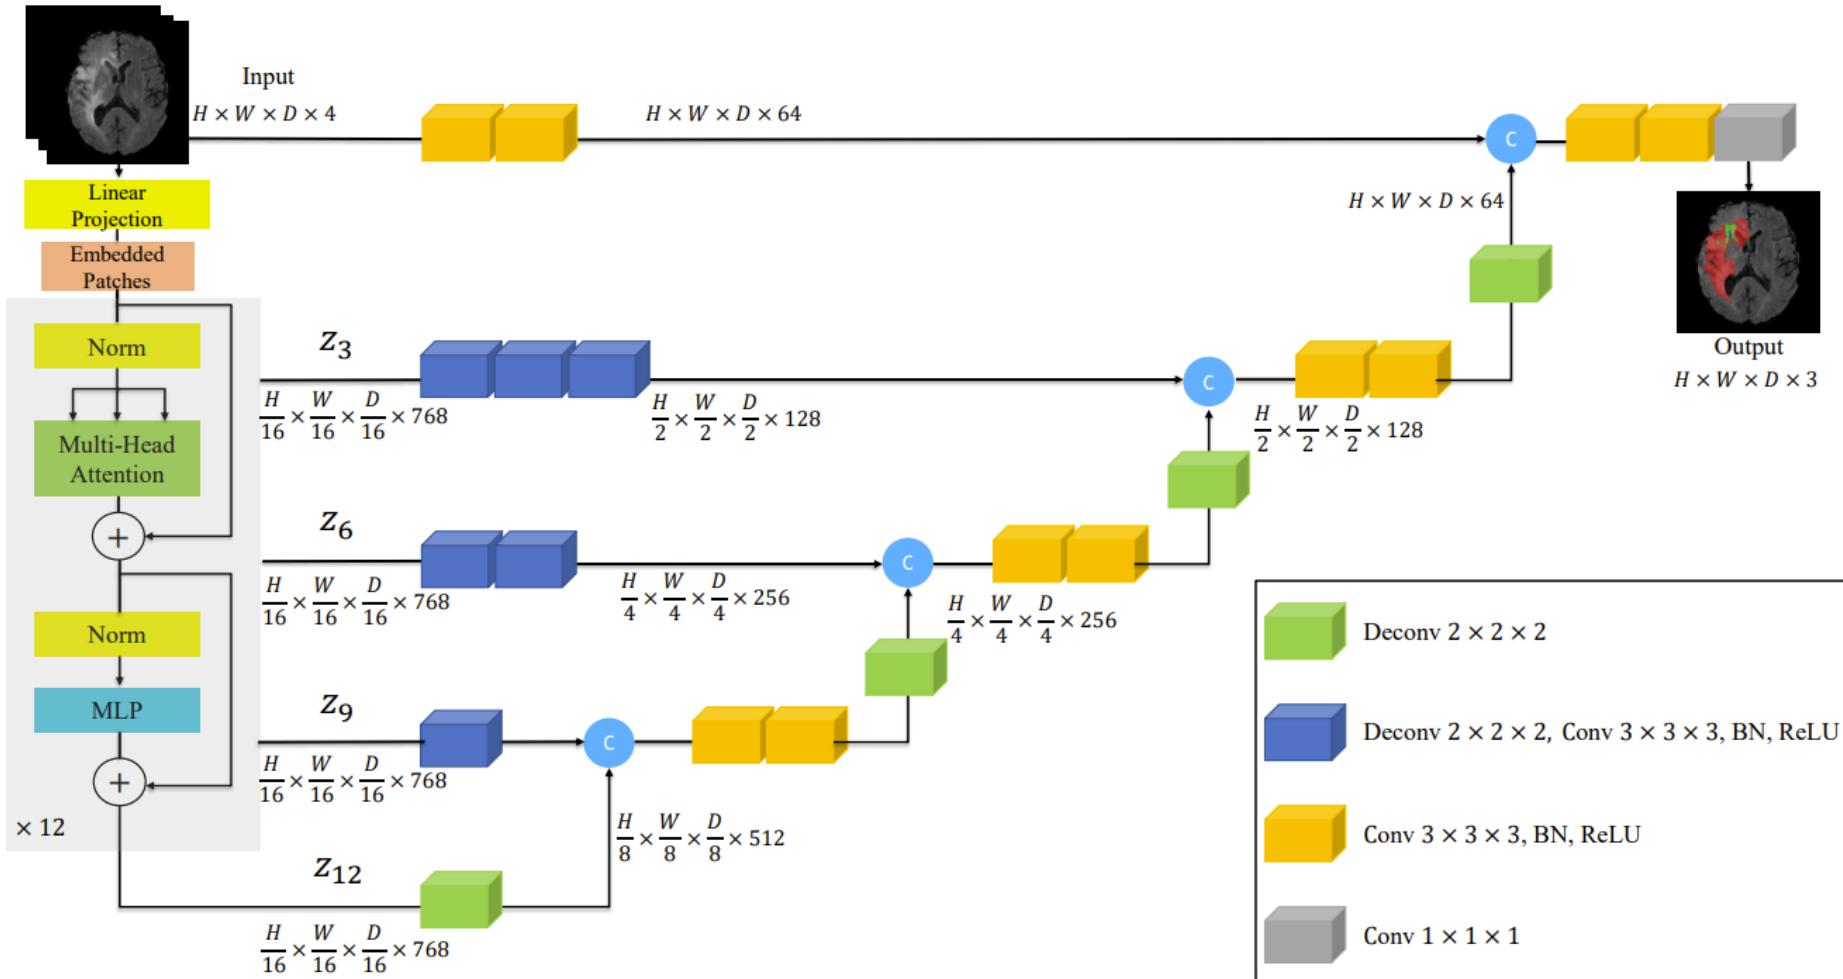
Vision transformers



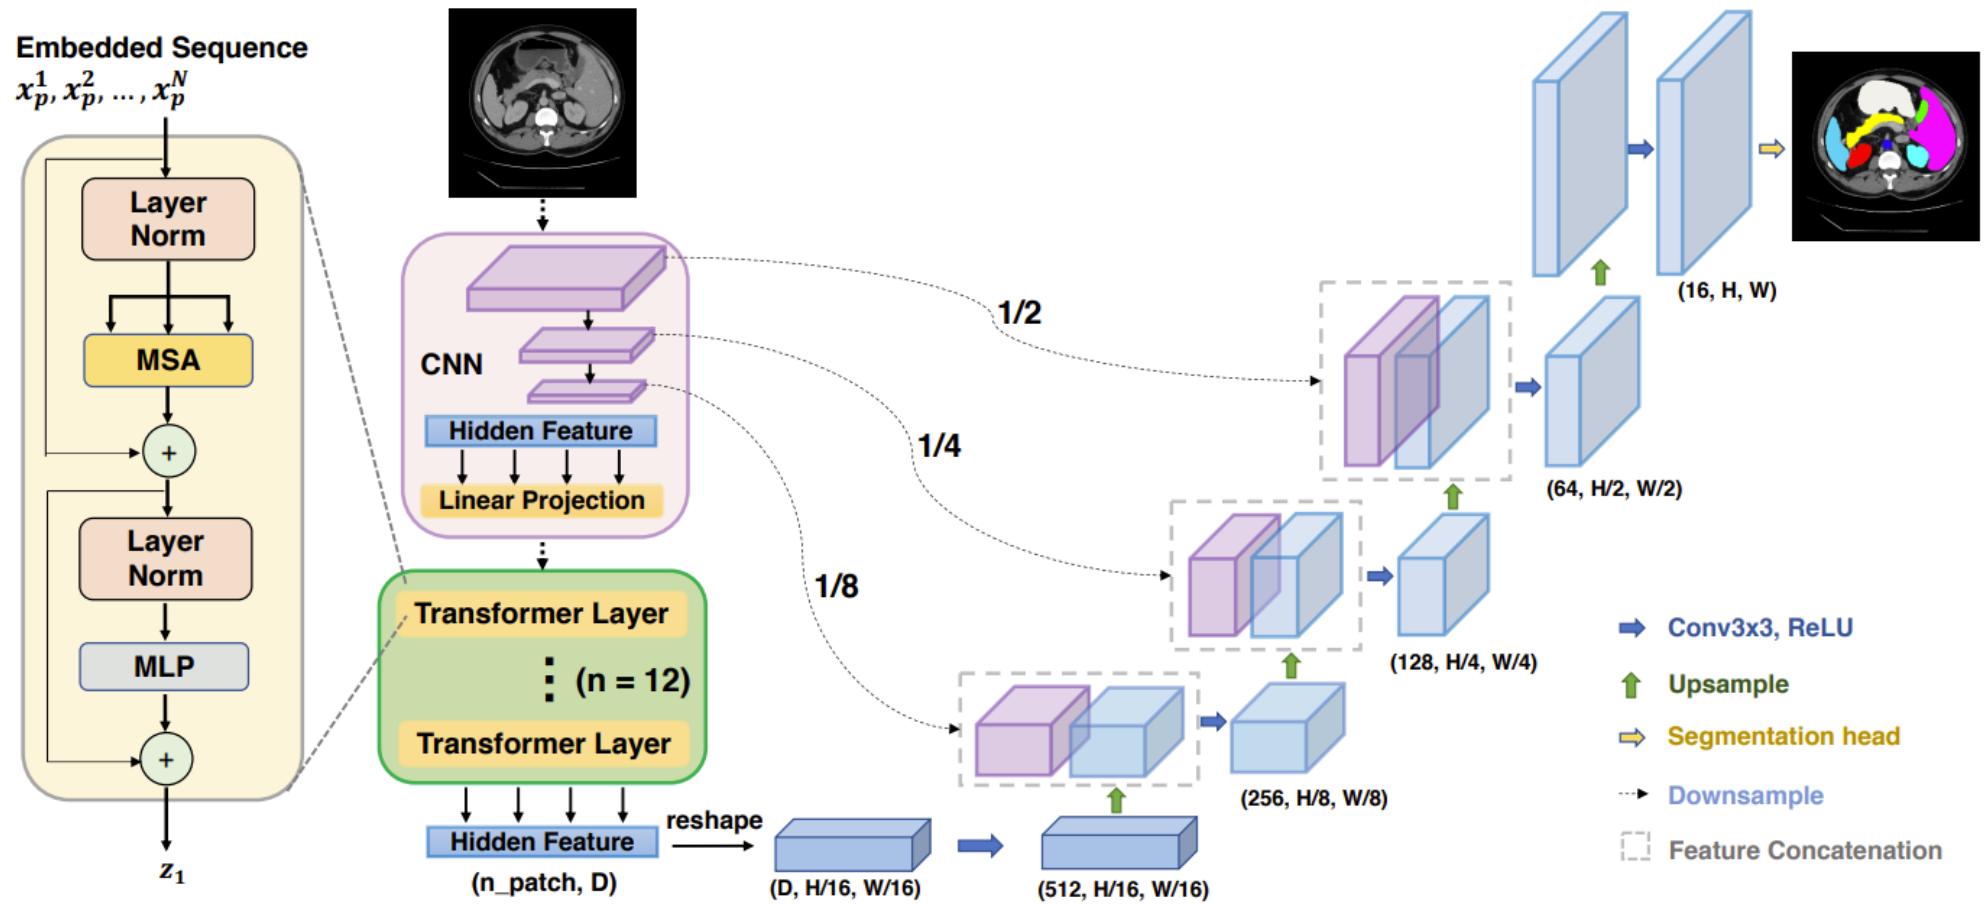
Transformers for image segmentation



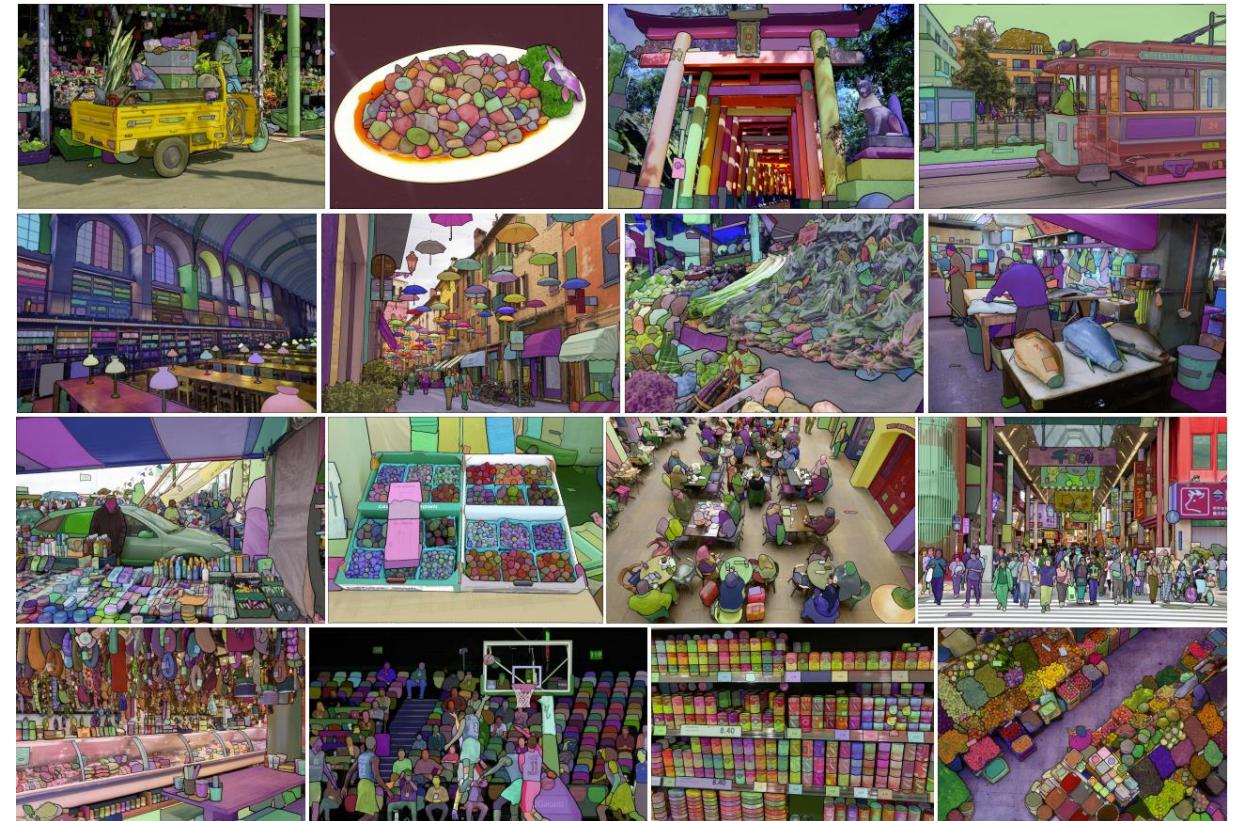
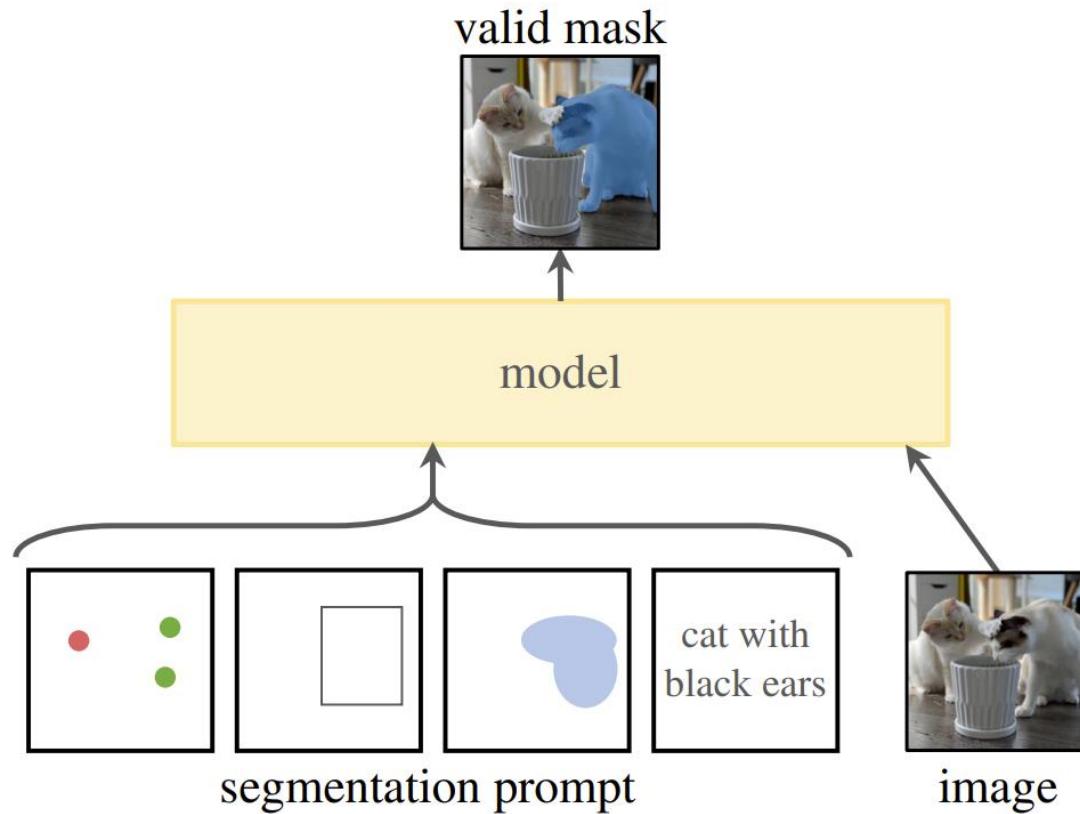
Transformers for image segmentation



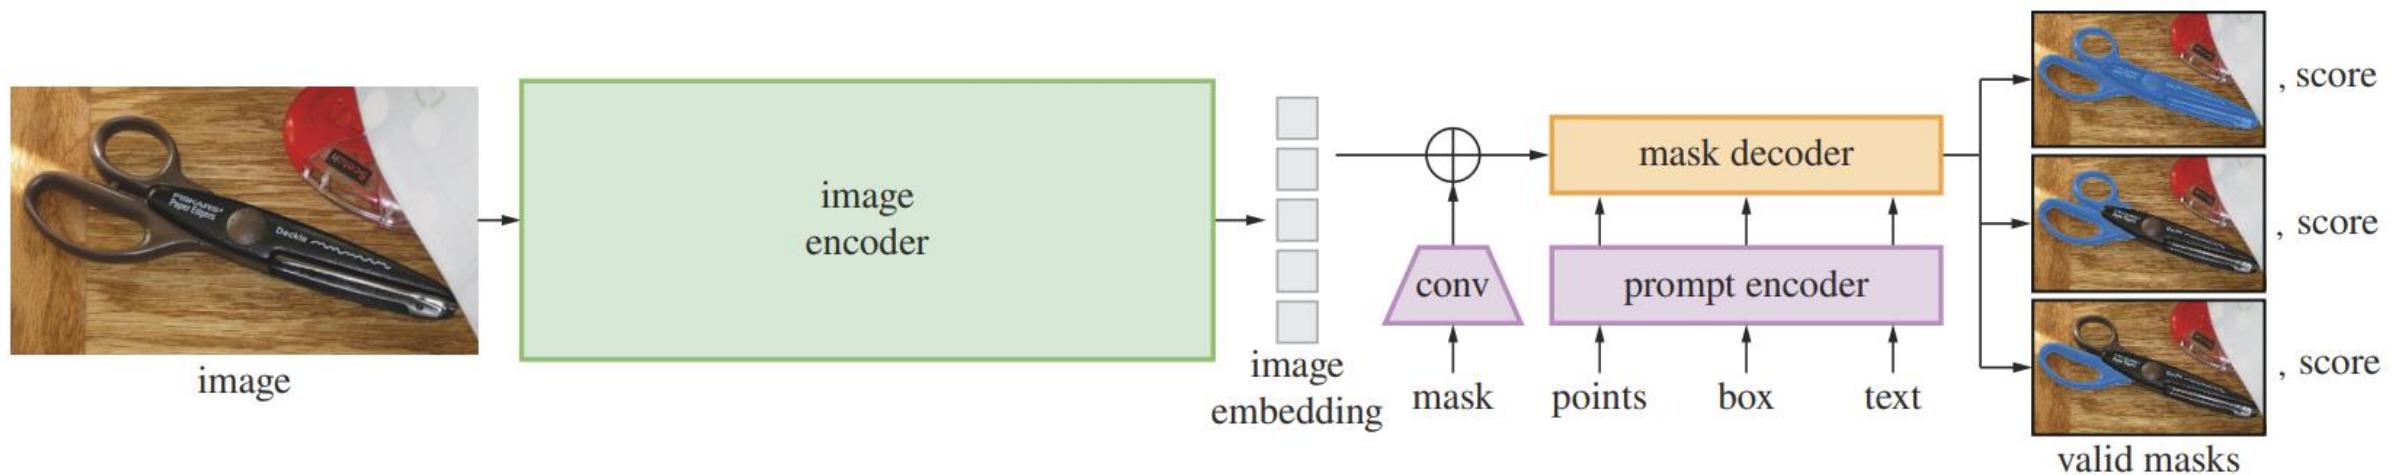
Transformers for image segmentation



Transformers for image segmentation



Segment anything model



Tutorial 2 - Segmentation