

# Machine Learning for Imaging Image Registration

Daniel Rueckert  
Department of Computing  
Imperial College London, UK

## Inverse problems in imaging

- Observe:

$$y = Ax + n$$

## Inverse problems in imaging

- Observe:

$$y = Ax + n$$

- Goal:

Recover  $x$  from  $y$

# Inverse problems in imaging

- Observe:

$$y = Ax + n$$

- Goal:

Recover  $x$  from  $y$

- Inpainting
- Deblurring
- Denoising
- Super-resolution
- Image Reconstruction  
(Medical imaging)

$y$



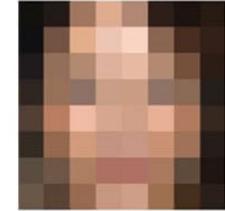
$x$



Inpainting



Deblurring



Super-resolution

## Classical Approach

- Example: deblurring



## Classical Approach

- Example: deblurring
- Least squares problem:

$$\arg \min_x \mathcal{D}(Ax, y)$$
$$\mathcal{D}(Ax, y) = \frac{1}{2} \|Ax - y\|_2^2$$

Operator  $A$  describes  
the forward map (here a  
linear operation)



## Classical Approach

- Example: deblurring
- Least squares problem:

$$\arg \min_x \mathcal{D}(Ax, y)$$

$$\mathcal{D}(Ax, y) = \frac{1}{2} \|Ax - y\|_2^2$$

- Least squares solution:

$$\hat{x} = (A^T A)^{-1} A^T y$$



## Classical Approach

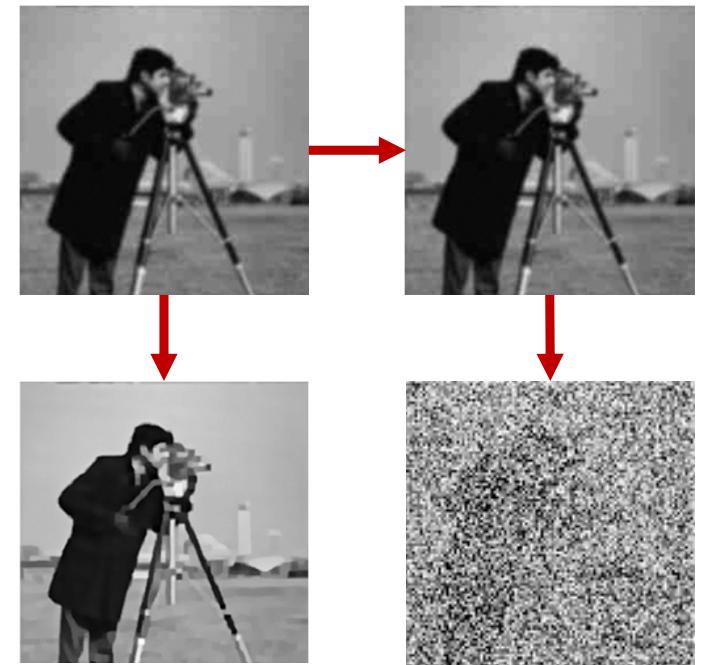
- Example: deblurring
- Least squares problem:

$$\arg \min_x \mathcal{D}(Ax, y)$$

$$\mathcal{D}(Ax, y) = \frac{1}{2} \|Ax - y\|_2^2$$

- Least squares solution:

$$\hat{x} = (A^T A)^{-1} A^T y$$



Wildly different solutions

## Classical Approach

- Example: deblurring
- Least squares problem:

$$\arg \min_x \mathcal{D}(Ax, y) + \lambda \mathcal{R}(x)$$

 Regularisation

$$\mathcal{D}(Ax, y) = \frac{1}{2} \|Ax - y\|_2^2 \quad \text{and} \quad \mathcal{R}(x) = \frac{1}{2} \|x\|_2^2$$

## Classical Approach

- Example: deblurring
- Least squares problem:

$$\arg \min_x \mathcal{D}(Ax, y) + \lambda \mathcal{R}(x)$$

 Regularisation

$$\mathcal{D}(Ax, y) = \frac{1}{2} \|Ax - y\|_2^2 \quad \text{and} \quad \mathcal{R}(x) = \frac{1}{2} \|x\|_2^2$$

- Least squares solution with Tikhonov regularisation:

$$\hat{x} = \underbrace{(A^T A + \lambda I)^{-1} A^T y}_{\text{Better conditioned (and noise suppression)}}$$

## General approach for images

$$\arg \min_x \mathcal{D}(Ax, y) + \lambda \mathcal{R}(x)$$

$\mathcal{D}(Ax, y)$  Data consistency term

$\mathcal{R}(x)$  Regularisation term (encoding prior knowledge on  $x$ )

$\lambda$  Regularisation parameter

## General approach for images

$$\arg \min_x \mathcal{D}(Ax, y) + \lambda \mathcal{R}(x)$$

$\mathcal{D}(Ax, y)$	Data consistency term
$\mathcal{R}(x)$	Regularisation term (encoding prior knowledge on $x$ )
$\lambda$	Regularisation parameter

Common regularisers:

$$\mathcal{R}(x) = \|\nabla x\|_2^2 \quad \text{Tikhonov regularisation}$$

$$\mathcal{R}(x) = \|\nabla x\|_{2,1} = \sum_{i=0}^n \sqrt{\sum_d \left( |\nabla x|_i^{(d)} \right)^2} \quad \text{Total variation regularisation}$$

$$\mathcal{R}(x) = \|\Psi x\|_1 \quad \text{Wavelet regularisation}$$

## General approaches for images

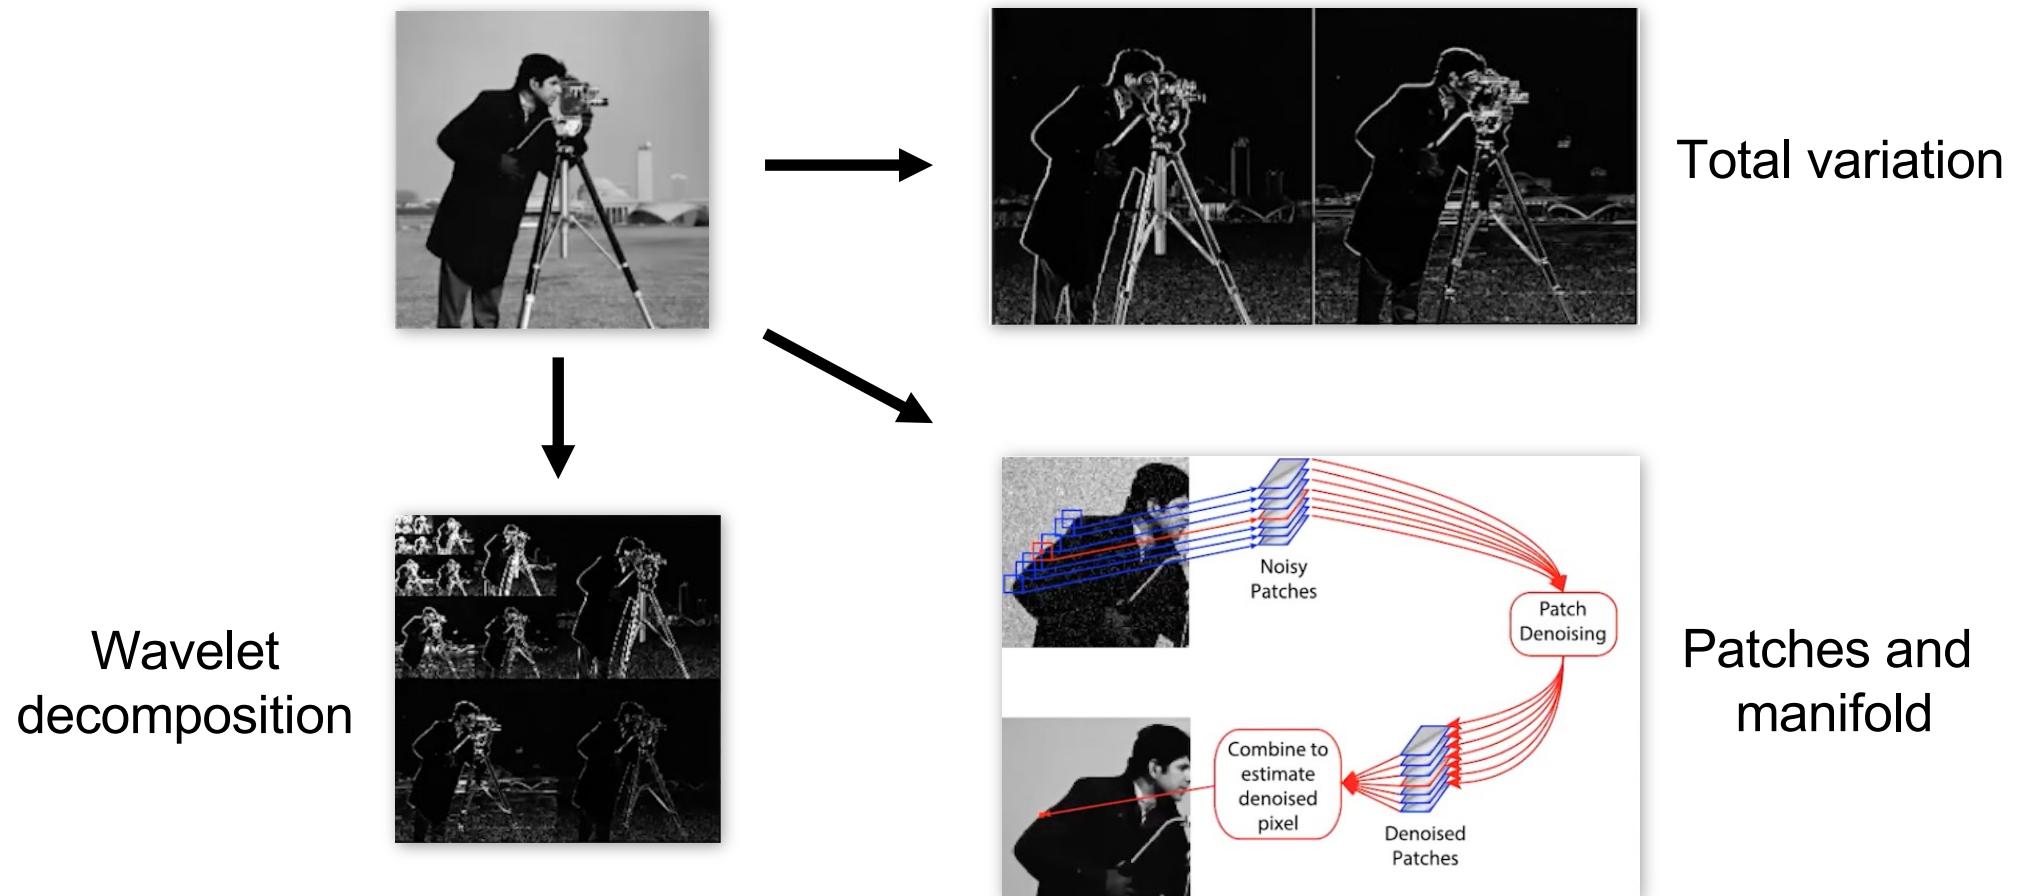
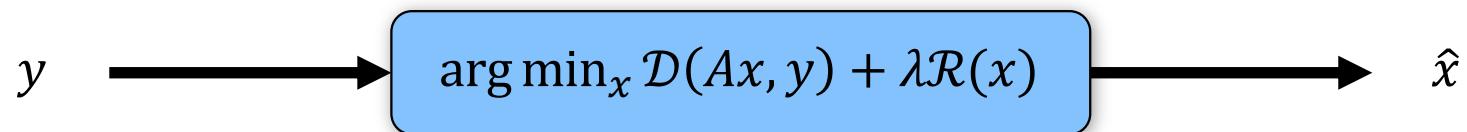


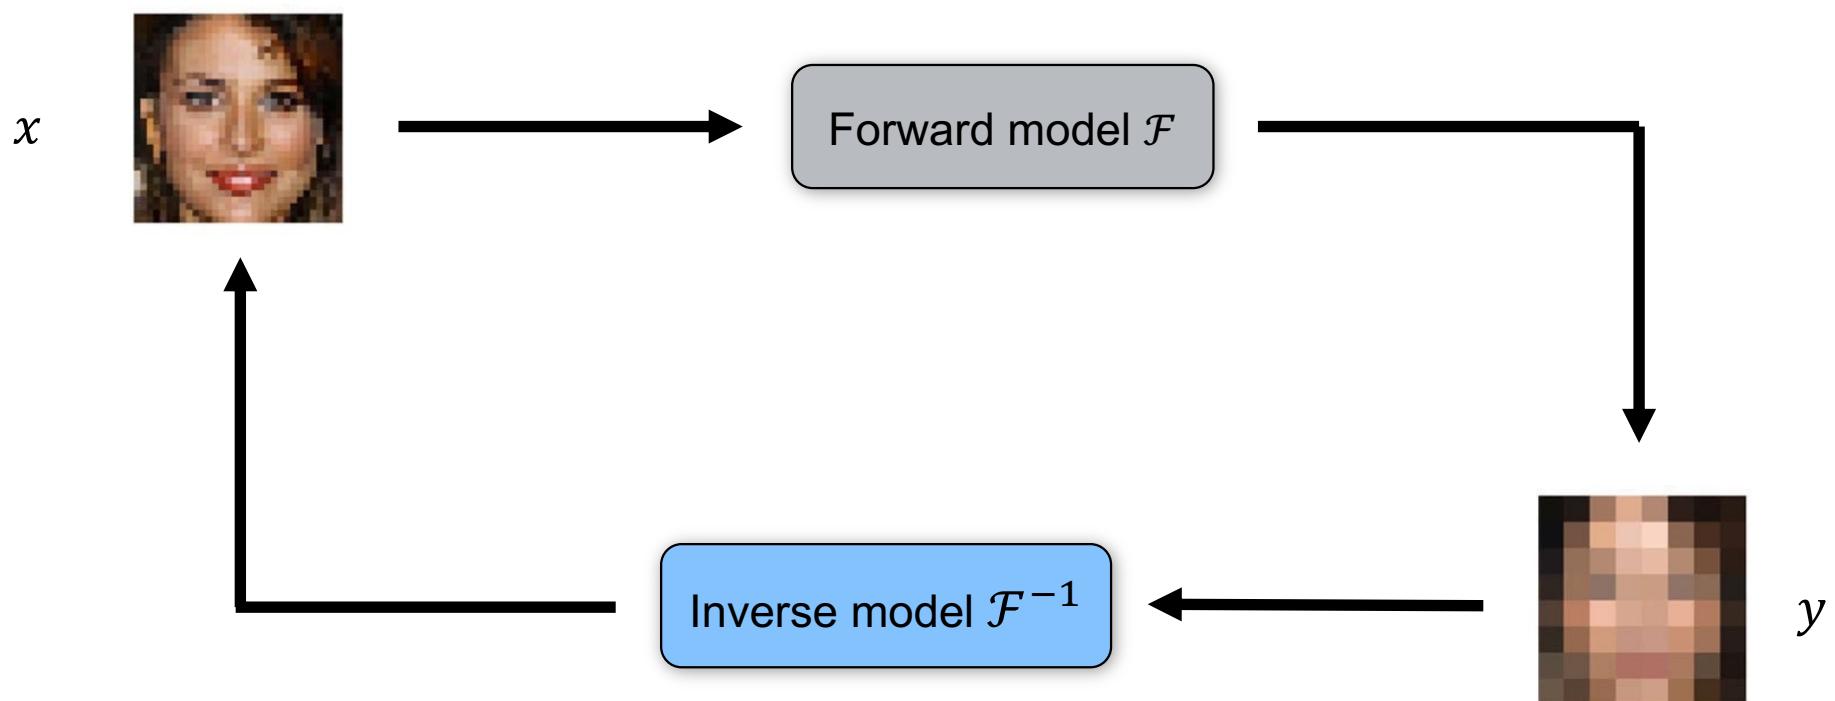
Figure adapted from Rebecca Willett

## Regularisation in inverse problems

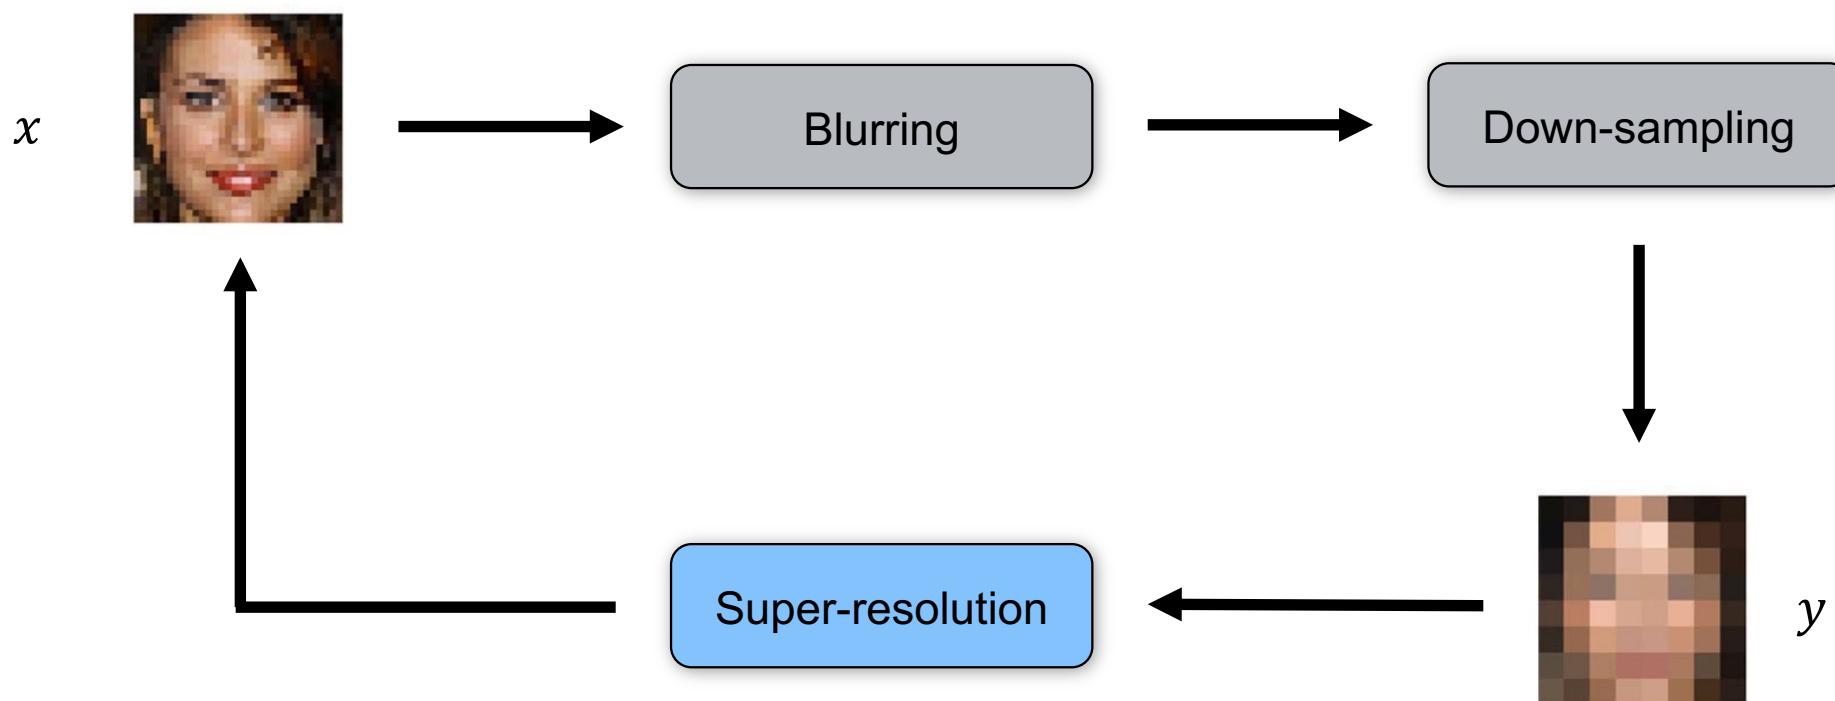


Instead of choosing  $\mathcal{R}$  a-priori based on a simple (geometric or other) model of image, learn  $\mathcal{R}$  from training data

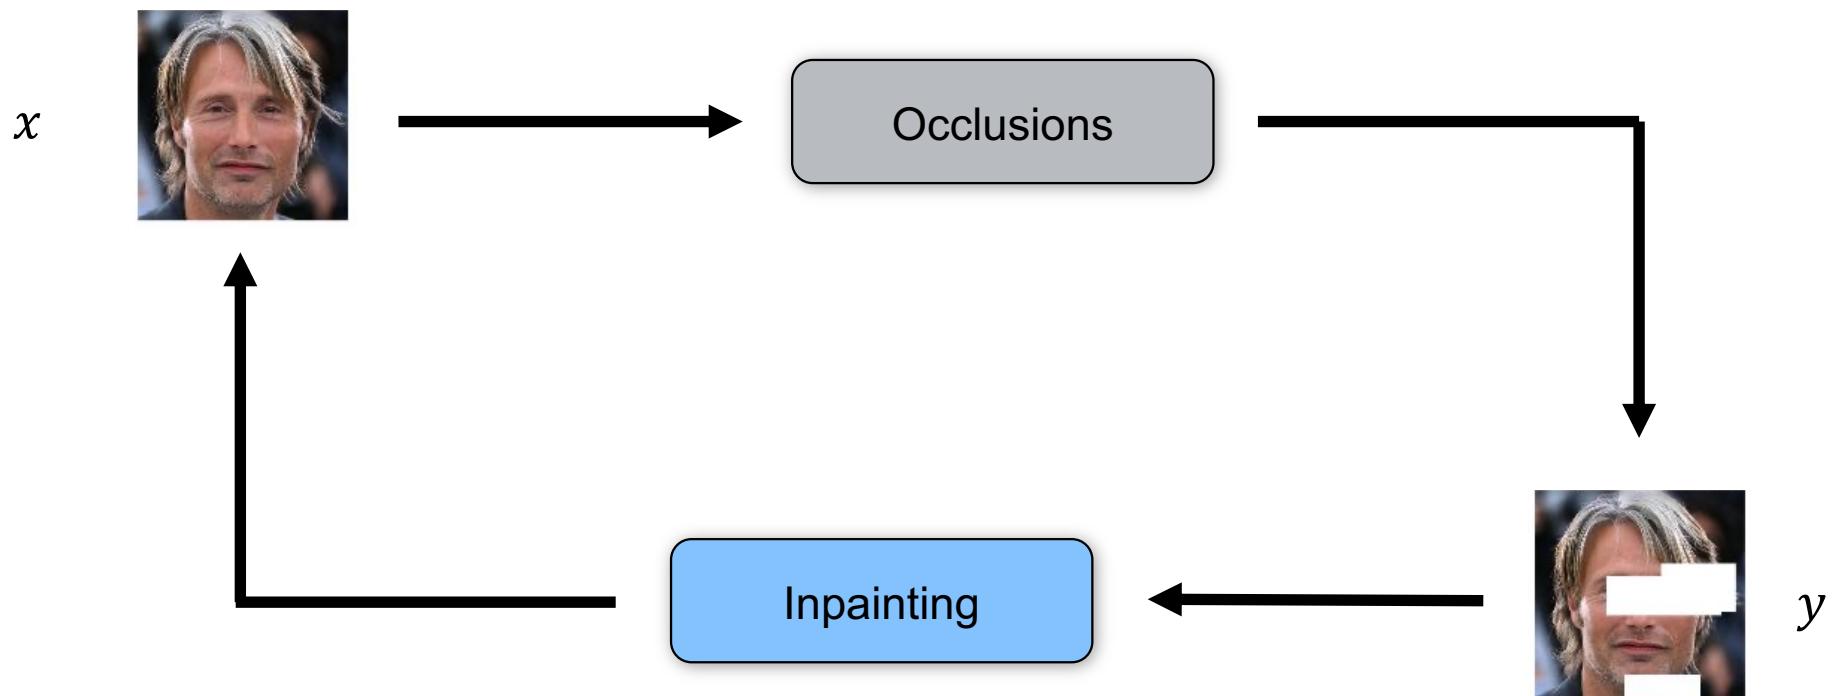
## Forward model/Inverse model



## Forward model/Inverse model



## Forward model/Inverse model



# Solving Inverse Problems with Deep Learning

Daniel Rueckert  
Department of Computing  
Imperial College London, UK

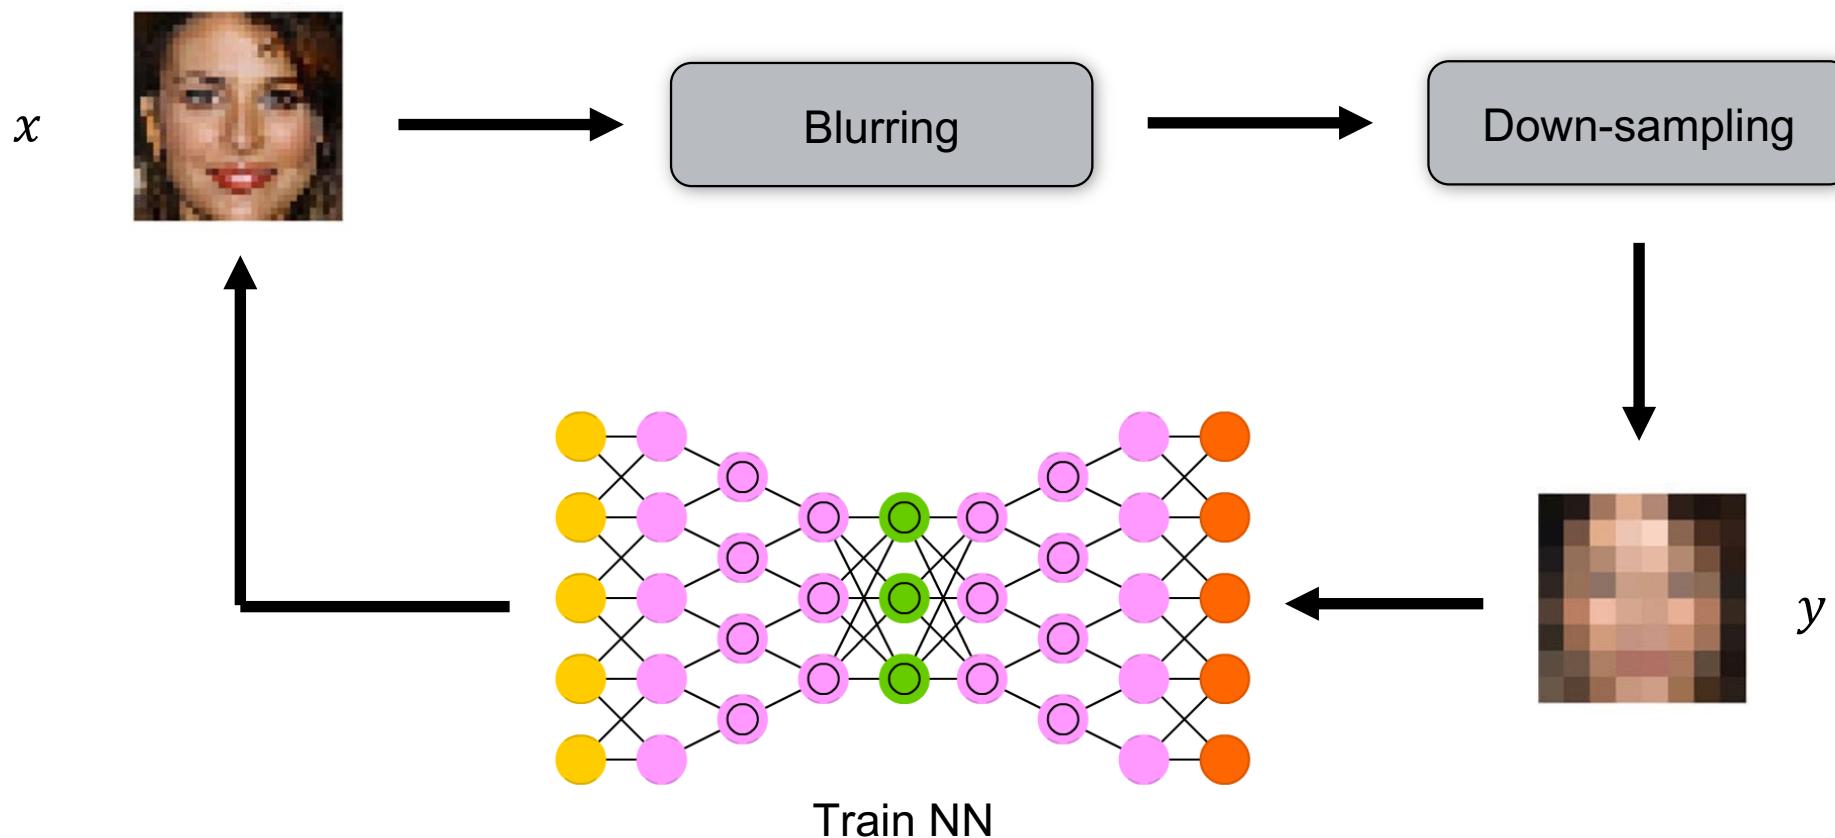
## Classes of methods

Model agnostic  
(ignores forward  
model)

Decoupled  
(First learn, then  
reconstruct)

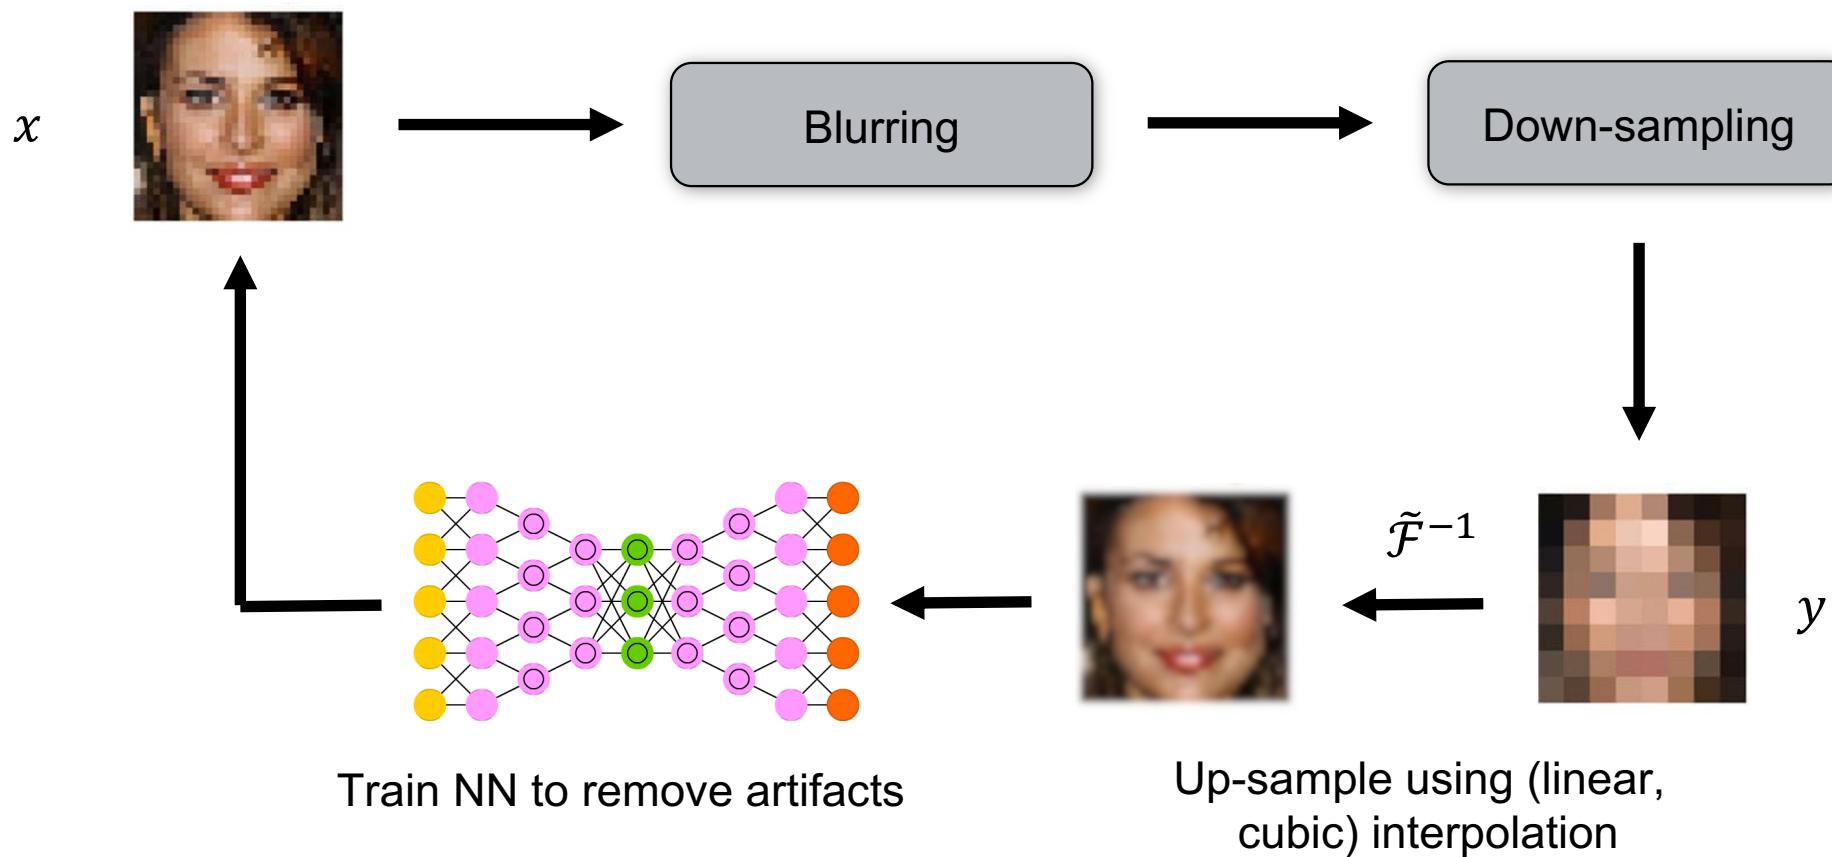
Unrolled  
optimisation

# Solving the inverse problem



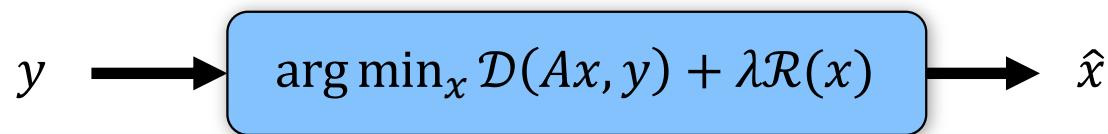
(partly)  
model  
agnostic

## Solving the inverse problem



Decoupled  
(First learn, then  
reconstruct)

## Solving the inverse problem: Deep proximal gradient



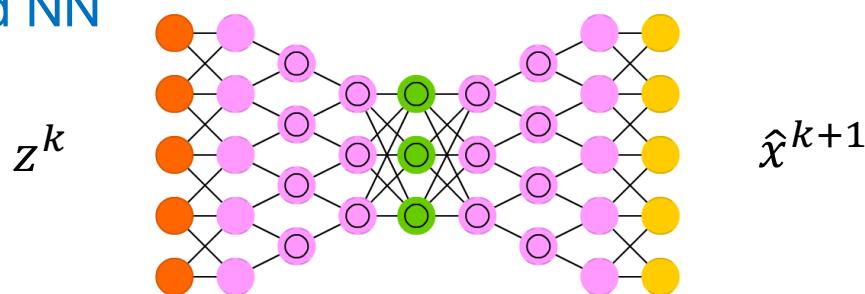
set  $\hat{x}^{(1)}$  and stepsize  $\eta$

for  $k = 1, 2, \dots$

$$z^k = \hat{x}^{(k)} + \eta A^T(y - A\hat{x}^{(k)}) \quad \text{gradient descent (DC)}$$

$$\hat{x}^{(k+1)} = \boxed{\arg \min_x \|z^k - x\|_2^2 + \eta \lambda \mathcal{R}(x)} \quad \text{denoising}$$

replace with learned NN



Decoupled  
(First learn, then  
reconstruct)

## Solving the inverse problem: GANs

$$y \rightarrow \arg \min_x \mathcal{D}(Ax, y) + \lambda \mathcal{R}(x) \rightarrow \hat{x}$$

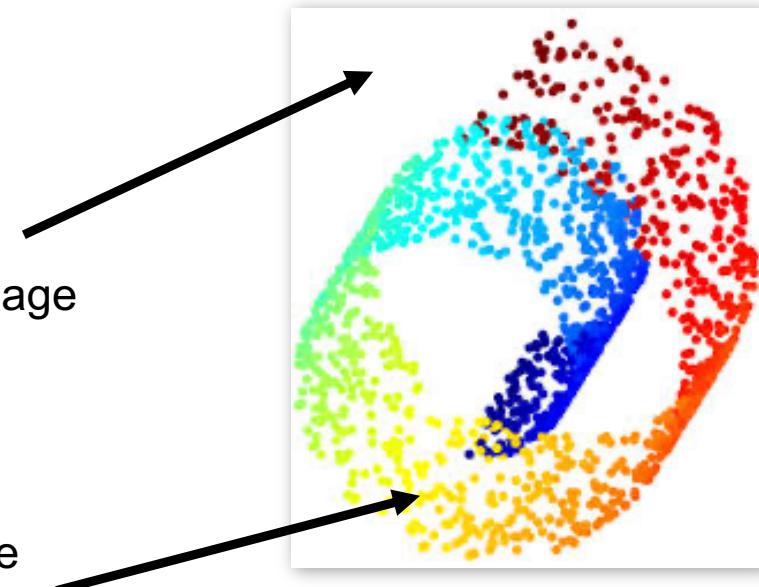
$$\mathcal{R}(x) = \begin{cases} 0 & x \text{ on image manifold} \\ \infty & \text{otherwise} \end{cases}$$



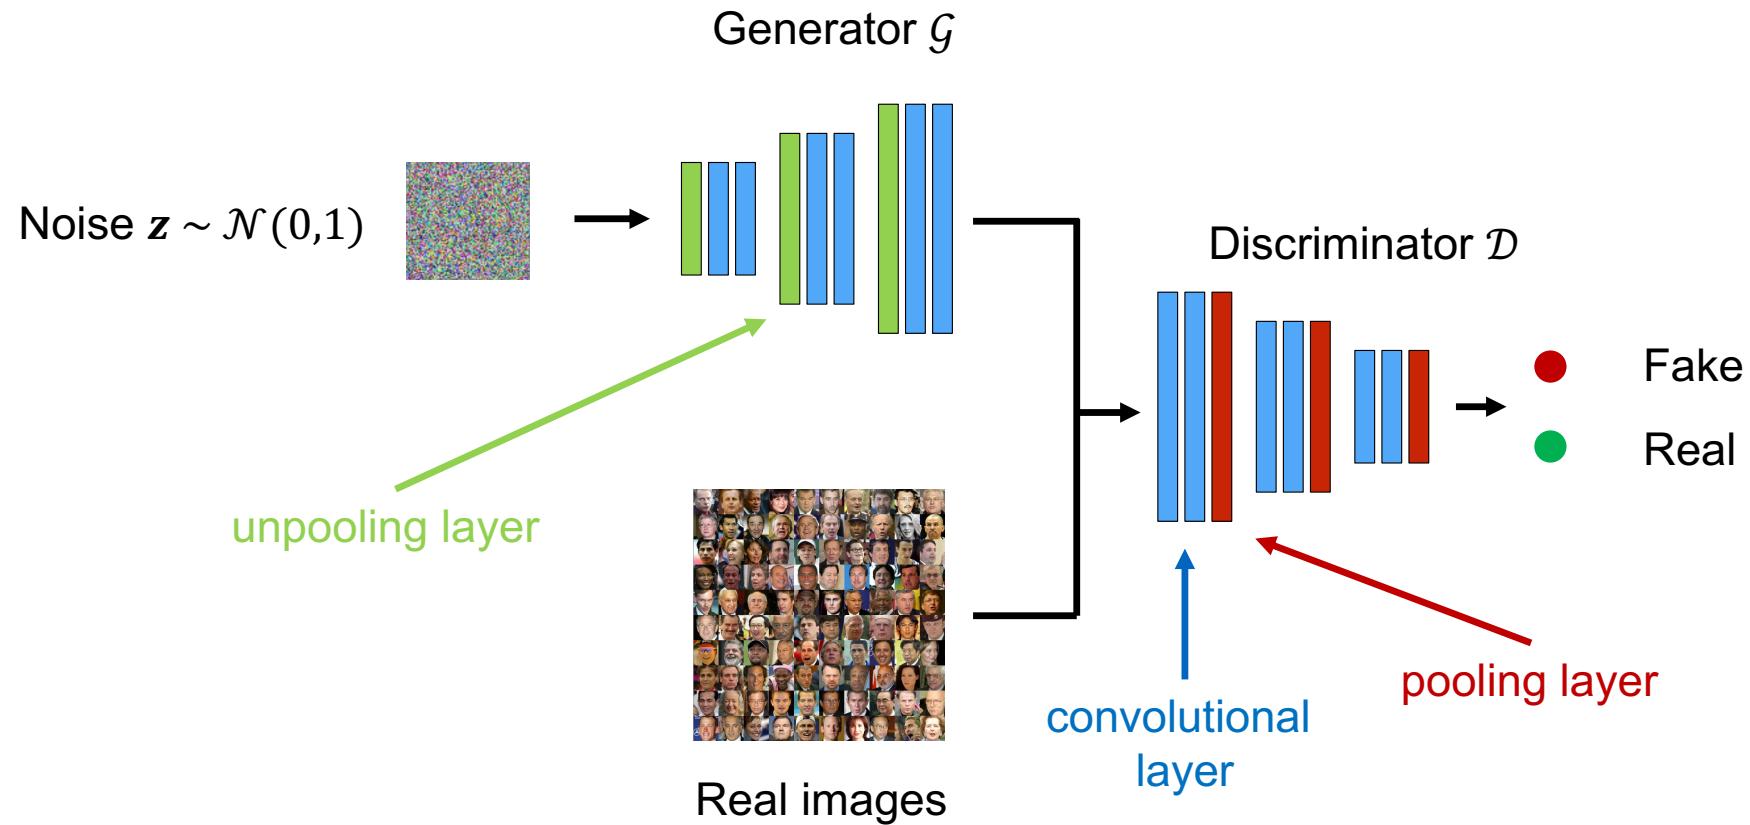
Non-natural image  
(bad)



Natural image  
(good)



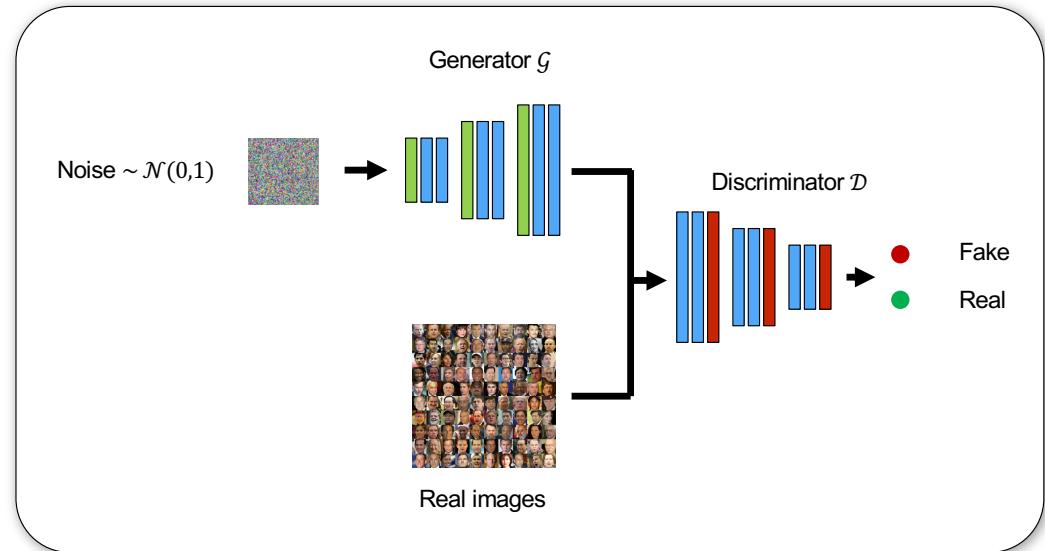
# Generative adversarial networks (GANs)



I. Goodfellow, NIPS, 2014

# Generative adversarial networks (GANs)

- $\mathcal{D}(x)$  represents the probability that  $x$  came from the real data rather than from  $\mathcal{G}$ .
- $\mathcal{D}$  is trained to maximize the probability of assigning the correct label to both real examples and samples from  $\mathcal{G}$ .
- Simultaneously train  $\mathcal{G}$  to minimize  $\log(1 - \mathcal{D}(\mathcal{G}(\mathbf{z}))$ , e.g.  $\mathcal{D}$  and  $\mathcal{G}$  play a two-player minimax game with value function  $V(\mathcal{G}, \mathcal{D})$ :



$$\min_{\mathcal{G}} \max_{\mathcal{D}} V(\mathcal{G}, \mathcal{D}) = \underbrace{\mathbb{E}_{x \sim p_{data}(x)} [\log \mathcal{D}(x)]}_{\text{real}} + \underbrace{\mathbb{E}_{z \sim p_z(z)} [\log(1 - \mathcal{D}(\mathcal{G}(z)))]}_{\text{synthetic (aka "fake")}}$$

# Solving the inverse problem: GANs

Decoupled  
(First learn, then  
reconstruct)

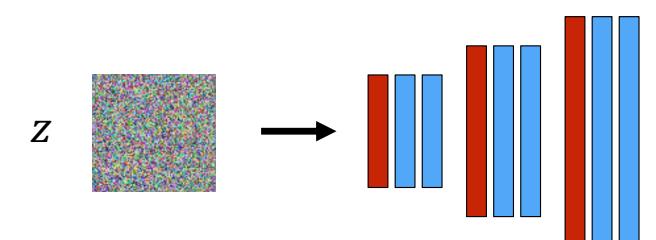
$$y \rightarrow \boxed{\arg \min_x \mathcal{D}(Ax, y) + \lambda \mathcal{R}(x)} \rightarrow \hat{x}$$

$$\mathcal{R}(x) = \begin{cases} 0 & x \text{ on image manifold} \\ \infty & \text{otherwise} \end{cases}$$

- Learn generator  $\mathcal{G}$  that outputs  $x \in \mathbb{R}^d$  given  $z \in \mathbb{R}^{d'}$  for  $d' < d$

$$\mathcal{R}(x) = \begin{cases} 0 & x \in \text{range}(\mathcal{G}) \\ \infty & \text{otherwise} \end{cases}$$

Generator  $\mathcal{G}$



## Solving the inverse problem: GANs

Decoupled  
(First learn, then reconstruct)

$$y \rightarrow \arg \min_x \mathcal{D}(Ax, y) + \lambda \mathcal{R}(x) \rightarrow \hat{x}$$

$$\mathcal{R}(x) = \begin{cases} 0 & x \text{ on image manifold} \\ \infty & \text{otherwise} \end{cases}$$

- Learn generator  $\mathcal{G}$  that outputs  $x \in \mathbb{R}^d$  given  $z \in \mathbb{R}^{d'}$  for  $d' < d$

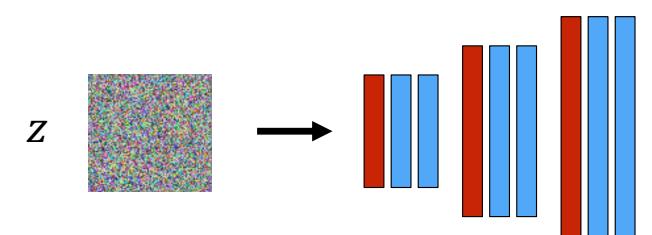
$$\mathcal{R}(x) = \begin{cases} 0 & x \in \text{range}(\mathcal{G}) \\ \infty & \text{otherwise} \end{cases}$$

Generator  $\mathcal{G}$

- Choose  $x \in \text{range}(\mathcal{G})$  that best fits the data

$$\hat{x} = \arg \min_{x \in \text{range}(\mathcal{G})} \|y - Ax\|_2^2 = \mathcal{G}(z)$$

$$\hat{z} = \arg \min_z \|y - A\mathcal{G}(z)\|_2^2$$



## Solving the inverse problem: Gradient descent networks

Let's assume that  $\mathcal{R}(x)$  is differentiable

$$\arg \min_x \mathcal{D}(Ax, y) + \lambda \mathcal{R}(x)$$

## Solving the inverse problem: Gradient descent networks

Let's assume that  $\mathcal{R}(x)$  is differentiable

$$\arg \min_x \|y - Ax\|_2^2 + \mathcal{R}(x)$$

set  $\hat{x}^{(1)}$  and stepsize  $\eta$

for  $k = 1, 2, \dots$

$$\hat{x}^{(k+1)} = \hat{x}^{(k)} + \eta A^T (y - A\hat{x}^{(k)}) - \boxed{\eta \nabla \mathcal{R}(\hat{x}^{(k)})}$$

replace with learned NN

Unrolled  
optimisation

## Solving the inverse problem: Gradient descent networks

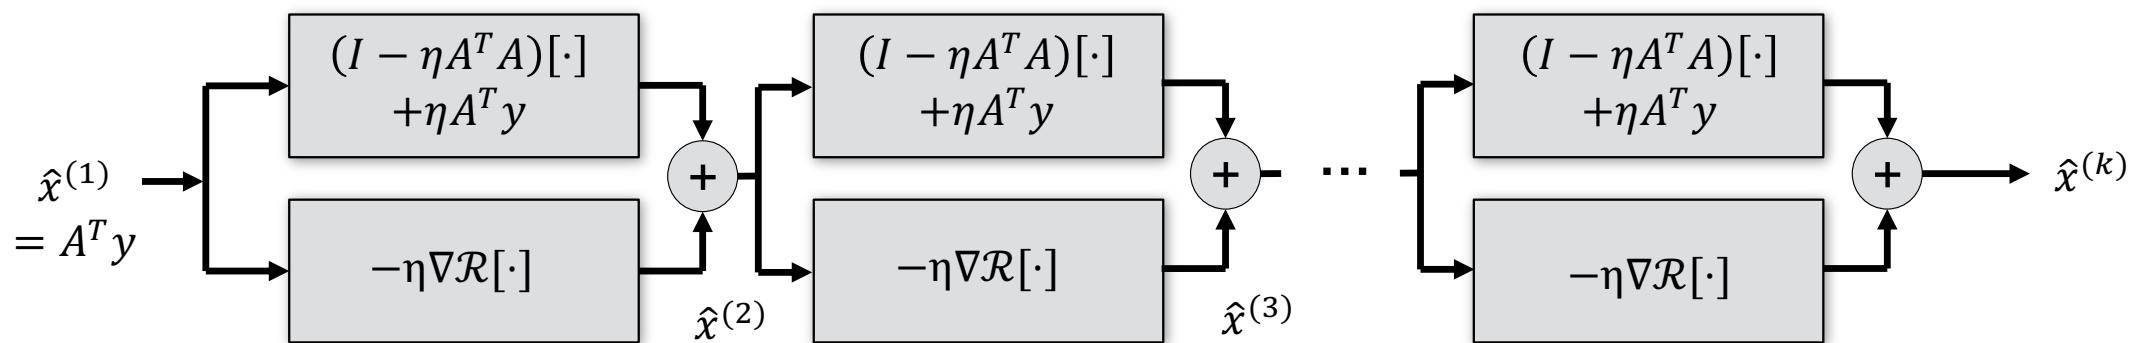
Let's assume that  $\mathcal{R}(x)$  is differentiable

$$\arg \min_x \|y - Ax\|_2^2 + \mathcal{R}(x)$$

set  $\hat{x}^{(1)}$  and stepsize  $\eta$

for  $k = 1, 2, \dots$

$$\hat{x}^{(k+1)} = \hat{x}^{(k)} + \eta A^T (y - A\hat{x}^{(k)}) - \boxed{\eta \nabla \mathcal{R}(\hat{x}^{(k)})} \quad \text{replace with learned NN}$$



# Solving the inverse problem: Gradient descent networks

Unrolled  
optimisation

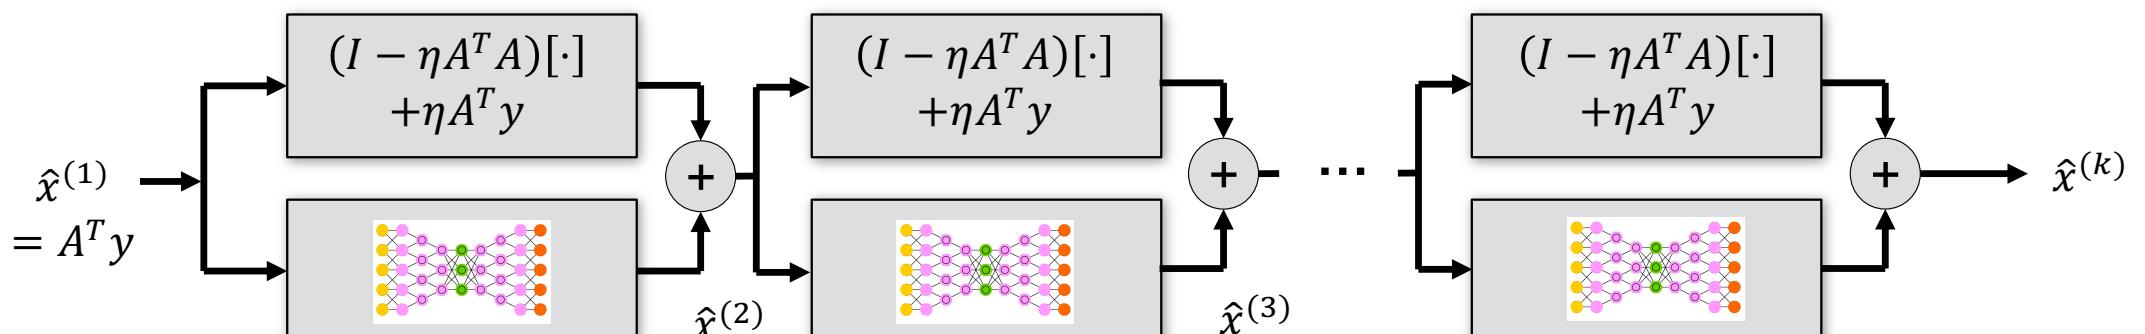
Let's assume that  $\mathcal{R}(x)$  is differentiable

$$\arg \min_x \|y - Ax\|_2^2 + \mathcal{R}(x)$$

set  $\hat{x}^{(1)}$  and stepsize  $\eta$

for  $k = 1, 2, \dots$

$$\hat{x}^{(k+1)} = \hat{x}^{(k)} + \eta A^T (y - A\hat{x}^{(k)}) - \boxed{\eta \nabla \mathcal{R}(\hat{x}^{(k)})} \quad \text{replace with learned NN}$$



Unrolled optimization framework **trained end-to-end**

Imperial College  
London

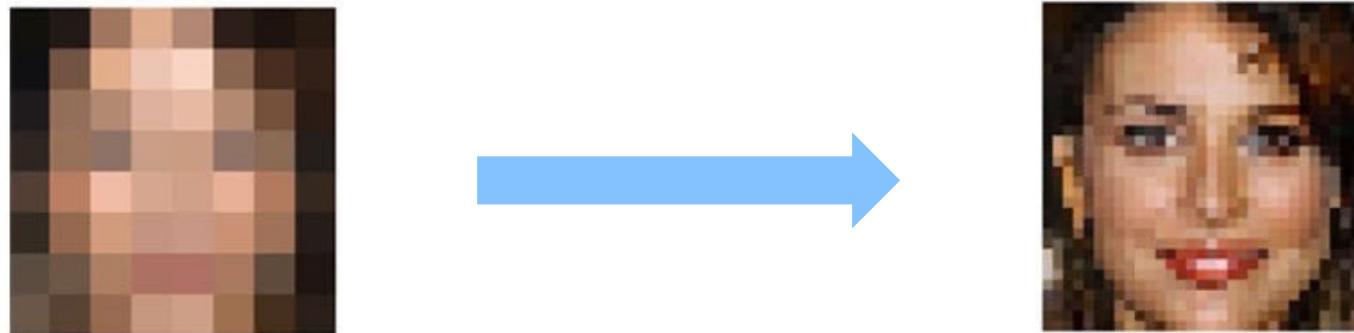
Imperial College  
London

## Deep Learning for Image Super-Resolution

Daniel Rueckert  
Department of Computing  
Imperial College London, UK

## Problem formulation

- Upsample low-resolution (LR) image to high-resolution (HR or SR) image



- Forward model (going from high- to low-resolution) is straightforward and involves some image degradation followed by downsampling
- Inverse model: e.g. interpolation-based models

## Super-resolution framework: Post-upampling super-resolution

- Directly upsamples LR image into SR
- Requires learnable upsampling layers

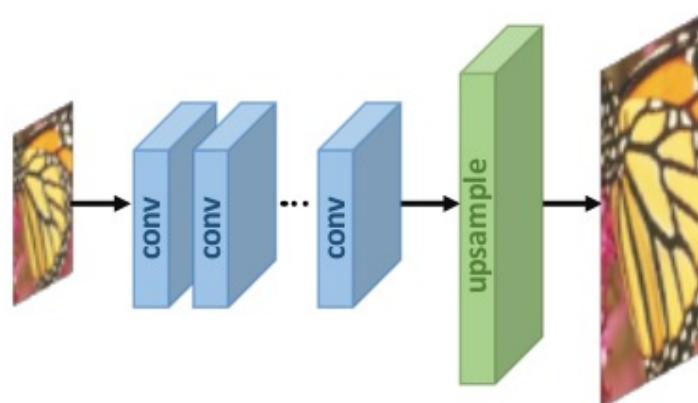


Figure from: Deep Learning for Image Super-resolution: A Survey  
<https://arxiv.org/abs/1902.06068>

- Advantages:
  - Fast and low memory requirements
- Disadvantages:
  - Network has to learn entire upsampling pipeline
  - Network typically limited to a specific up-sampling factor

# Super-resolution framework: Pre-upsampling super-resolution

- Two stage process:
  1. First use traditional upsampling algorithm (e.g. linear interpolation) to obtain SR images
  2. Then refining upsampled using a deep neural network (usually a CNN)

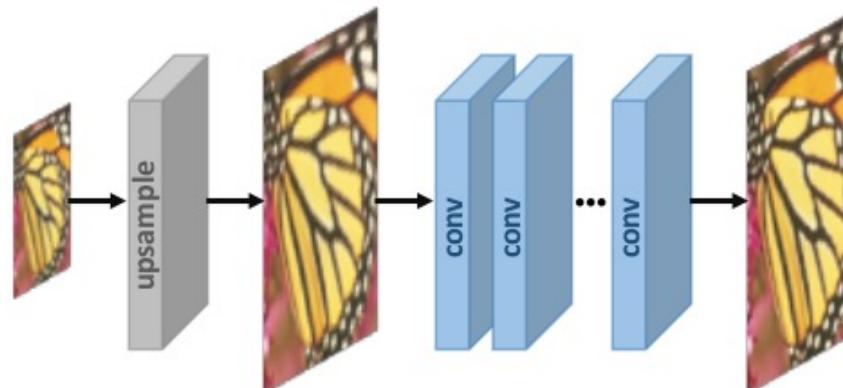


Figure from: Deep Learning for Image Super-resolution: A Survey  
<https://arxiv.org/abs/1902.06068>

- Advantages:
  - Upsampling operation is performed using interpolation, then correct smaller details
  - Can be applied to a range of upscaling factors and image sizes
- Disadvantages:
  - Operates on SR image, thus requires more computation and memory

# Super-resolution framework: Progressive upsampling super-resolution

- Multi-stage process:
  - Use a cascade of CNNs to progressively reconstruct higher-resolution images.
  - At each stage, the images are upsampled to higher resolution and refined by CNNs

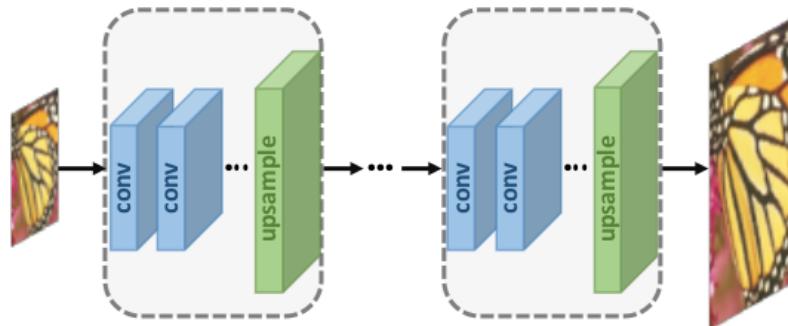


Figure from: Deep Learning for Image Super-resolution: A Survey  
<https://arxiv.org/abs/1902.06068>

- Advantages:
  - Decomposes complex task into simple tasks
  - Reasonable efficiency
- Disadvantages:
  - Sometimes difficult to train very deep models

# Super-resolution framework: Iterative up-and-down sampling super-resolution

- Approach:
  - Alternate between upsampling and downsampling (back-projection) operations
  - Mutually connected up- and down-sampling stages

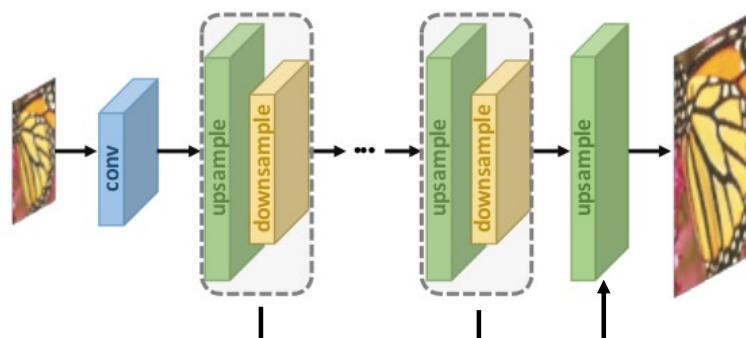
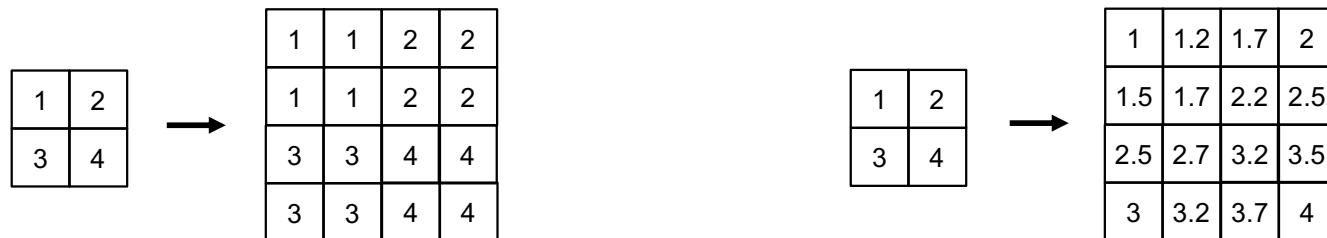


Figure from: Deep Learning for Image Super-resolution: A Survey  
<https://arxiv.org/abs/1902.06068>

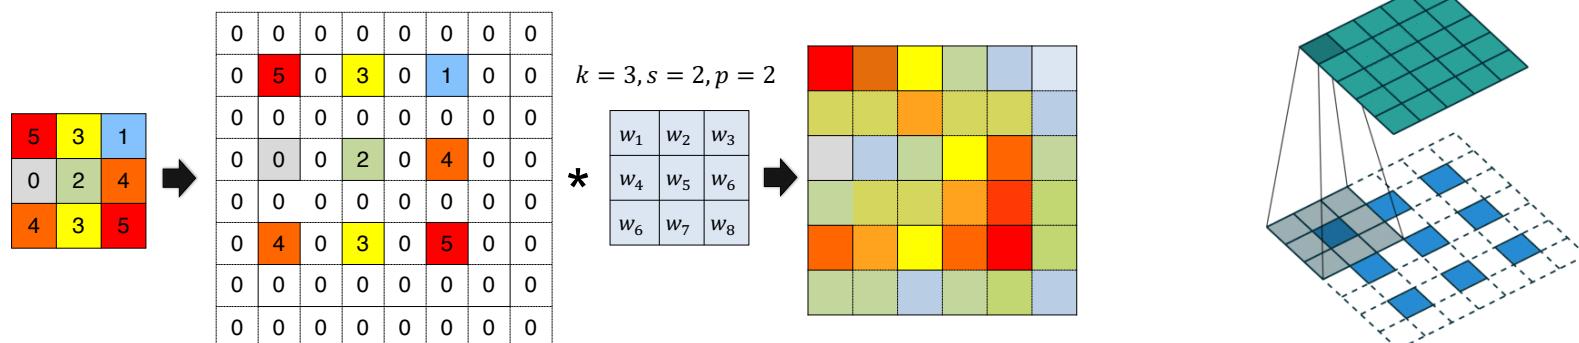
- Advantages:
  - Has shown superior performance as it allows error feedback
  - Easier training of deep networks

## How to implement upsampling?

- Traditional interpolation-based upsampling (NN, bi-linear) can be implemented as a convolutional layer with fixed weights:



- Or using learnable weights as transpose convolution:



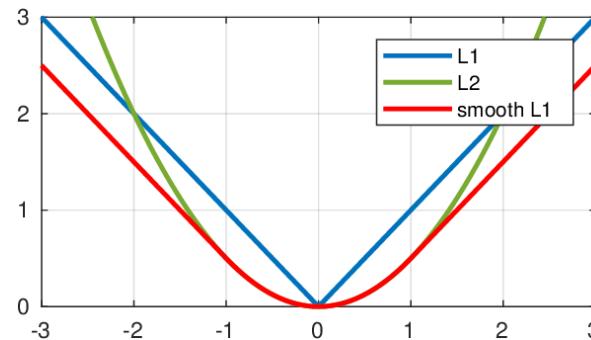
## Loss functions for super-resolution

- Pixel-wise loss function (either L1 or L2)

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N |x_i - \hat{x}_i|^p$$

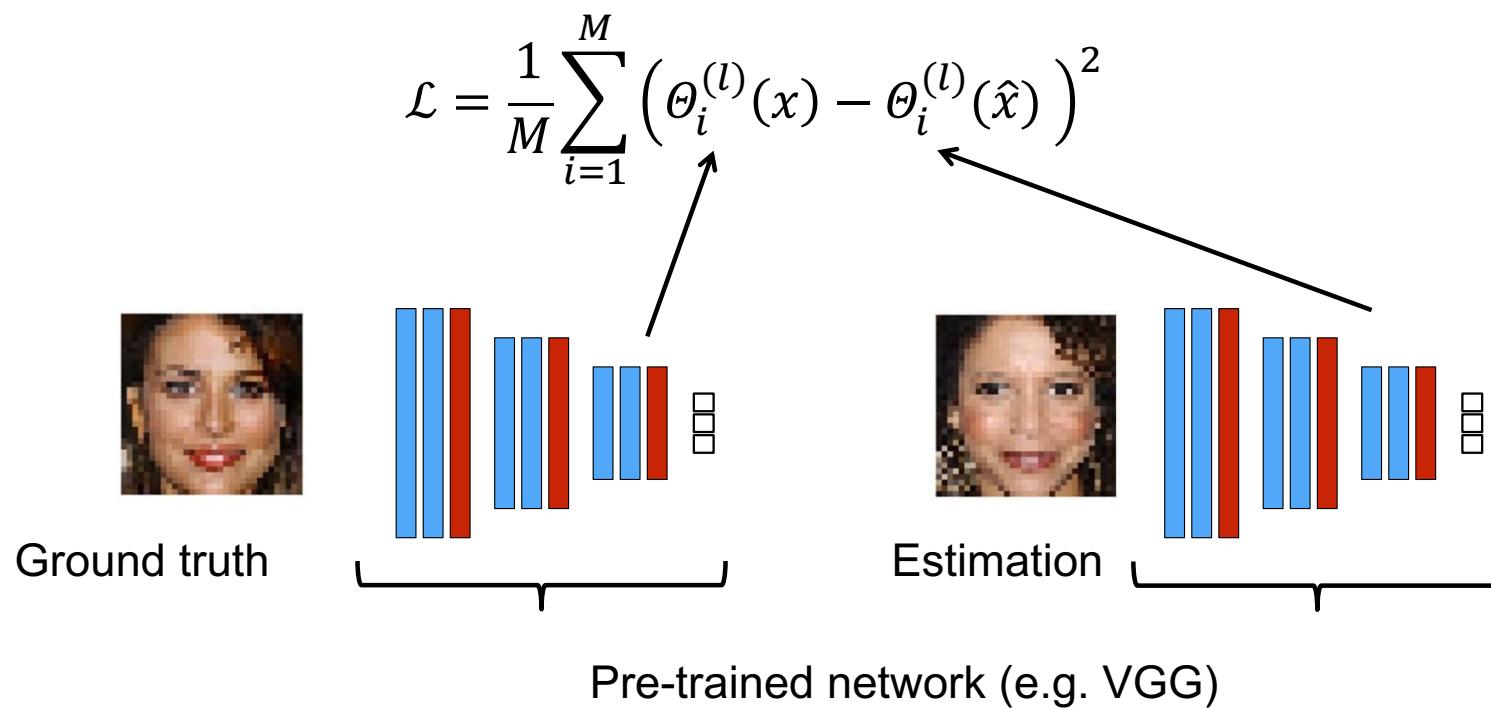
- Alternative: Huber loss function

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N \delta(x_i - \hat{x}_i) \quad \delta(a) = \begin{cases} 0.5a^2 & \text{for } |a| \leq 1 \\ |a| - 0.5 & \text{otherwise} \end{cases}$$



## Loss functions for super-resolution

- Perceptual loss: Computes loss on the output  $\theta$  of an intermediate layer  $l$  of a pre-trained network:



## Loss functions for super-resolution

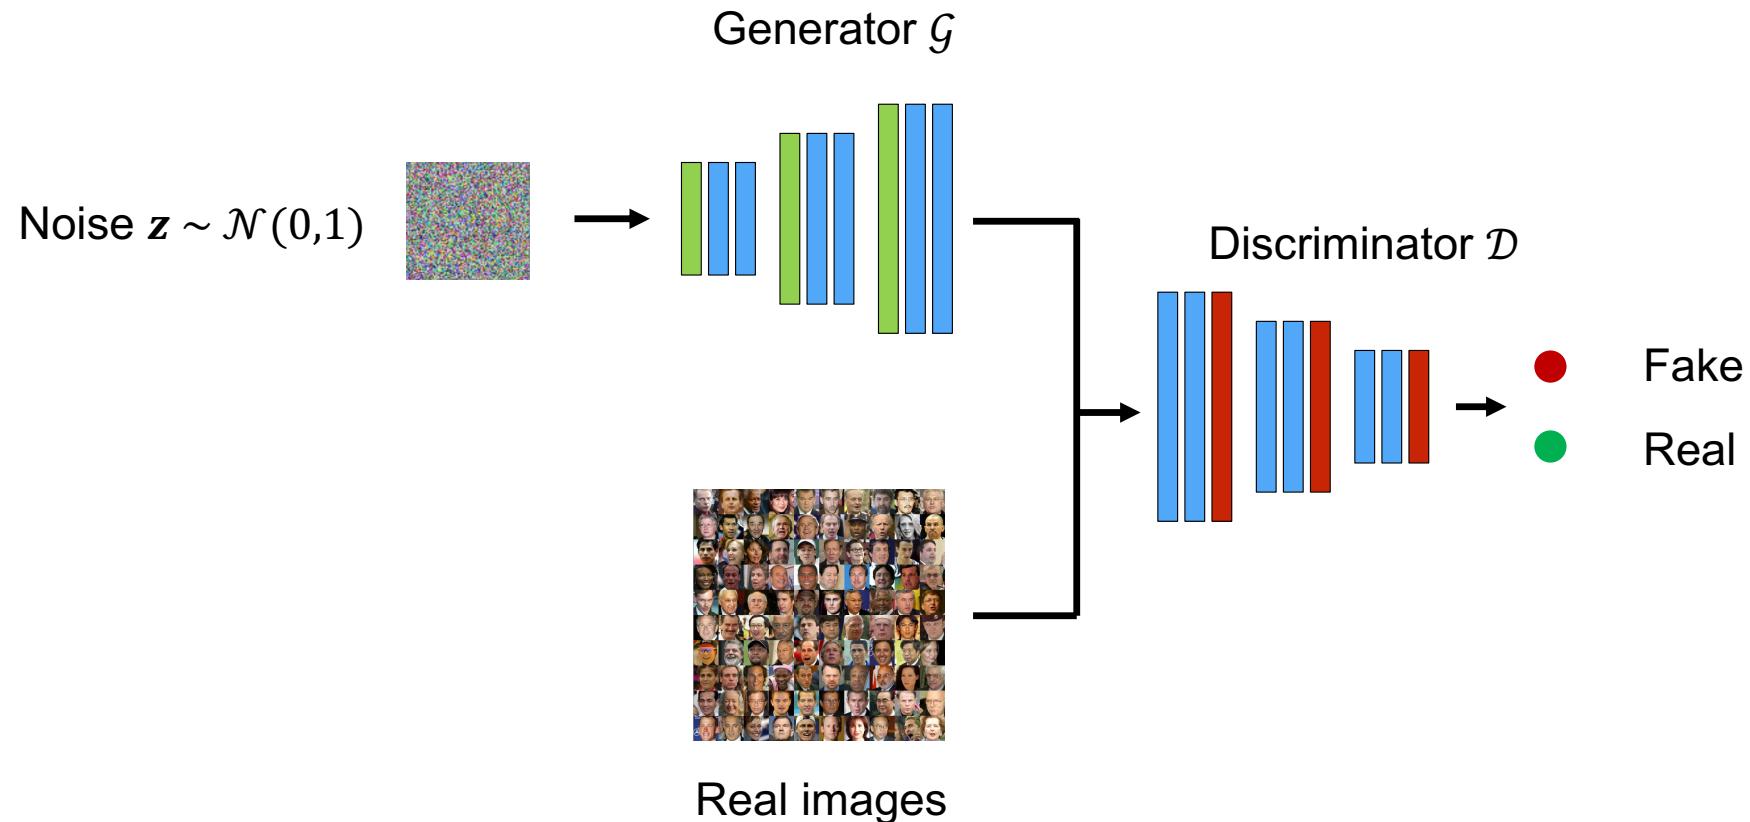
- Total variation

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N \sqrt{\sum_d \left( |\nabla x|_i^{(d)} \right)^2}$$

- Assumption: Absolute value of gradient of the image is low, e.g. image is piecewise constant
- Implementation for 2D image (using forward differences to approximate the gradient):

$$\mathcal{L} = \frac{1}{n_1 n_2} \sum_{i=1}^{n_1-1} \sum_{j=1}^{n_2-1} \sqrt{|x_{i-1,j} - x_{i,j}|^2 + |x_{i,j-1} - x_{i,j}|^2}$$

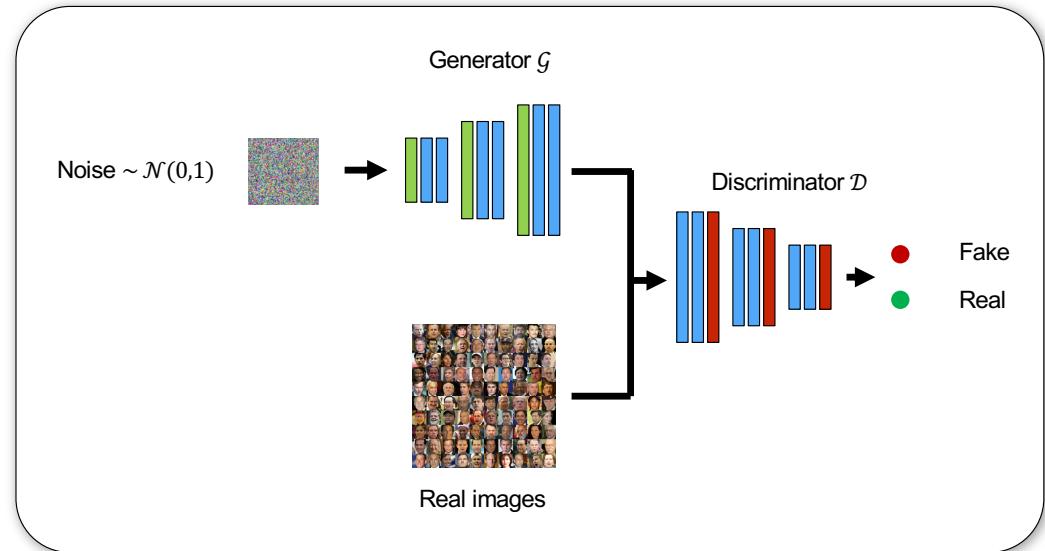
# Generative adversarial networks (GANs)



I. Goodfellow, NIPS, 2014

# Generative adversarial networks (GANs)

- $\mathcal{D}(x)$  represents the probability that  $x$  came from the real data rather than from  $\mathcal{G}$ .
- $\mathcal{D}$  is trained to maximize the probability of assigning the correct label to both real examples and samples from  $\mathcal{G}$ .
- Simultaneously train  $\mathcal{G}$  to minimize  $\log(1 - \mathcal{D}(\mathcal{G}(\mathbf{z}))$ , e.g.  $\mathcal{D}$  and  $\mathcal{G}$  play a two-player minimax game with value function  $V(\mathcal{G}, \mathcal{D})$ :



$$\min_{\mathcal{G}} \max_{\mathcal{D}} V(\mathcal{G}, \mathcal{D}) = \underbrace{\mathbb{E}_{x \sim p_{data}(x)} [\log \mathcal{D}(x)]}_{\text{real}} + \underbrace{\mathbb{E}_{z \sim p_z(z)} [\log(1 - \mathcal{D}(\mathcal{G}(z)))]}_{\text{synthetic (aka "fake")}}$$

## Loss functions for super-resolution based on GANs

- Minimax game for super-resolution:

$$\min_{\mathcal{G}} \max_{\mathcal{D}} \mathbb{E}_{I_{HR} \sim p_{train}} [\log \mathcal{D}(I_{HR})] + \mathbb{E}_{I_{LR} \sim P_{\mathcal{G}}(I_{LR})} [\log(1 - \mathcal{D}(\mathcal{G}(I_{LR})))]$$

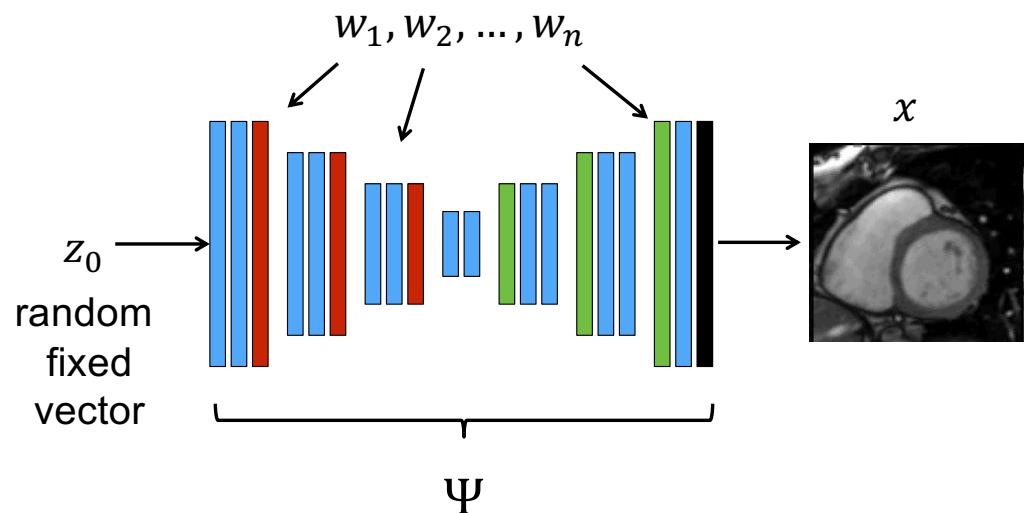
- Aim: To detect
  - Which images are (real) high-resolution?
  - Which images are super-resolved?
- Use discriminator output as loss function:

$$\mathcal{L} = -\log \mathcal{D}(\mathcal{G}(I_{LR}))$$

# Parameterizing images via generative networks

- Consider a generator network  $\Psi$  with a fixed input  $z_0$
- The network parameters  $w$  can be thought as a ***parameterization*** of images

$$w \mapsto x = \Psi(z_0; w)$$



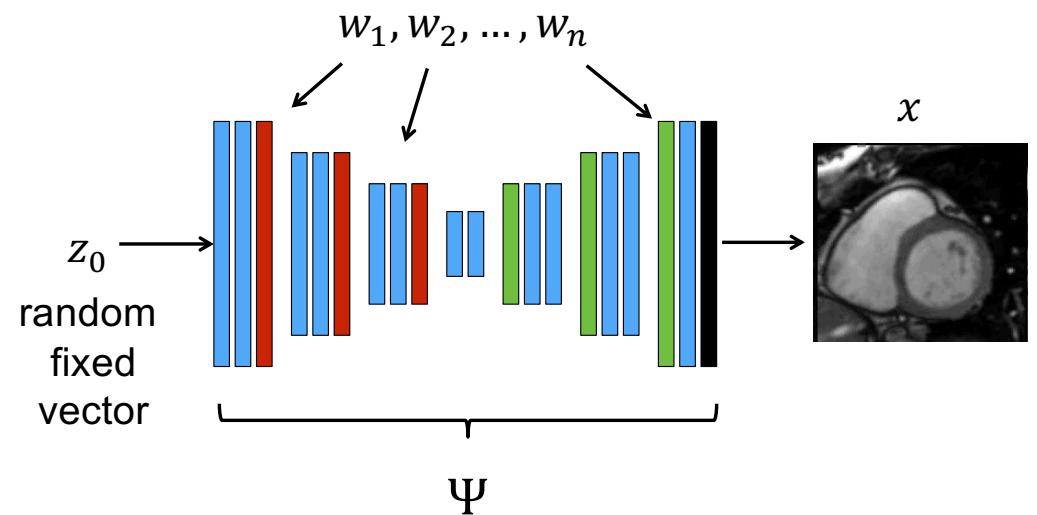
# Parameterizing images via generative networks

Fit a network to a single image:

- Start with a random-initialised network
- Given an image  $x$ , its parameters  $w$  are recovered by solving the following optimization problem:

$$\min_w \|x - \Psi(z_0; w)\|^2$$

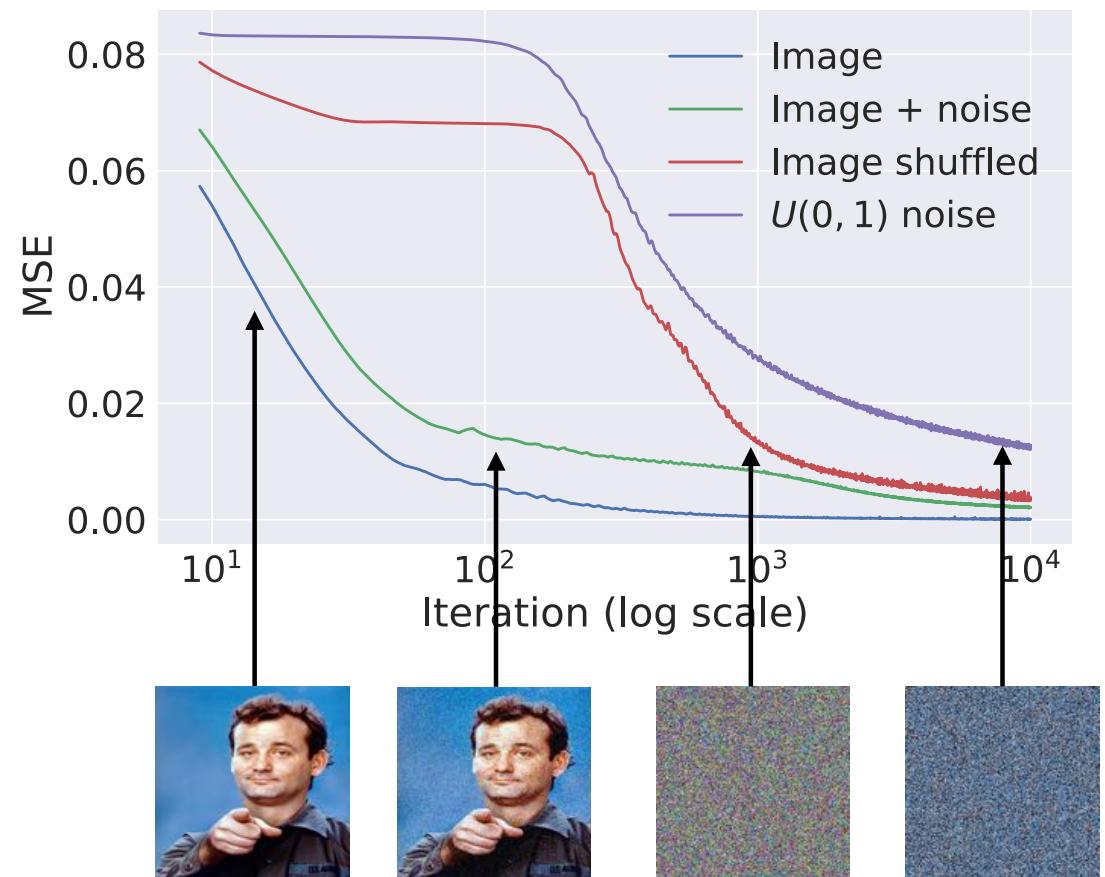
- This is similar to learning a network from a single image



# Deep image prior

For most generative networks fitting naturally looking images is easier/faster than fitting others

Deep Image Prior. Ulyanov, Vedaldi & Lempitsky, CVPR 2018



# Deep image prior: Application to inpainting

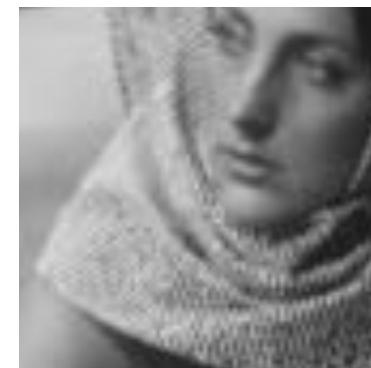
For inpainting we only reconstruct the visible pixels, implicitly inferring the pixels that are masked out by mask  $m$ :

$$\min_w \|m \odot (x - \Psi(z_0; w))\|^2$$

Conventional  
inpainting



Inpainting  
using deep  
image prior

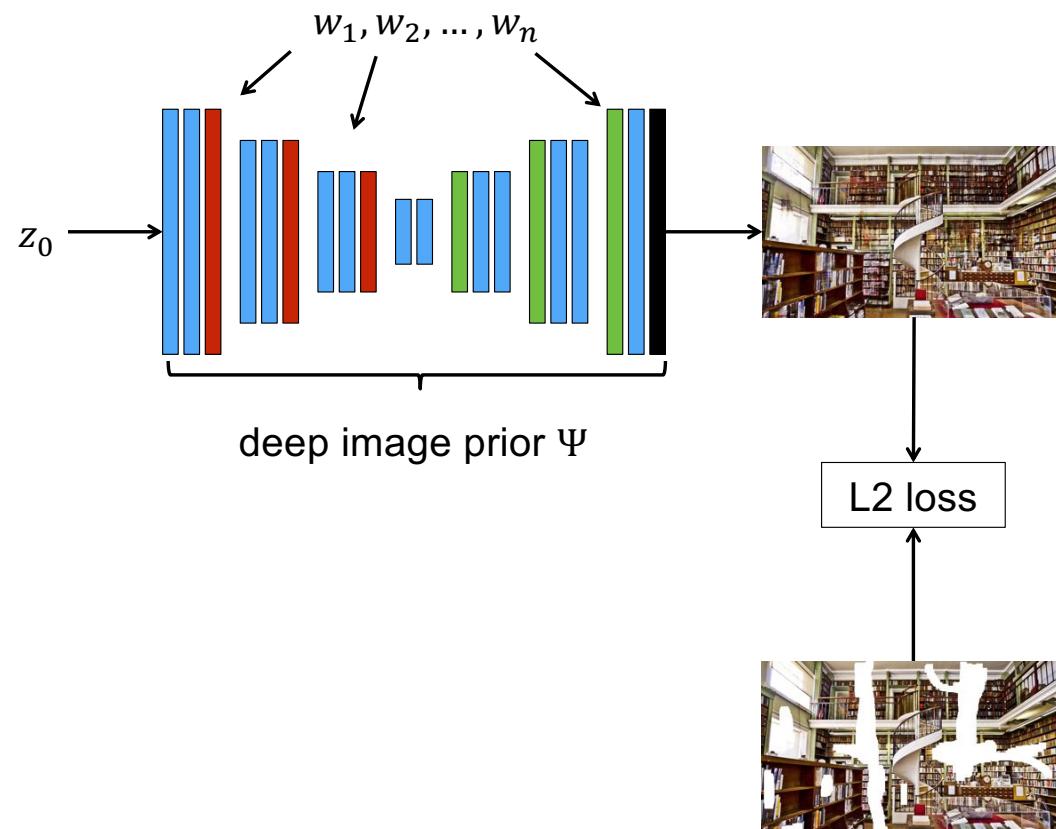


# Deep image prior: Application to inpainting



Deep Image Prior. Ulyanov, Vedaldi & Lempitsky, CVPR 2018

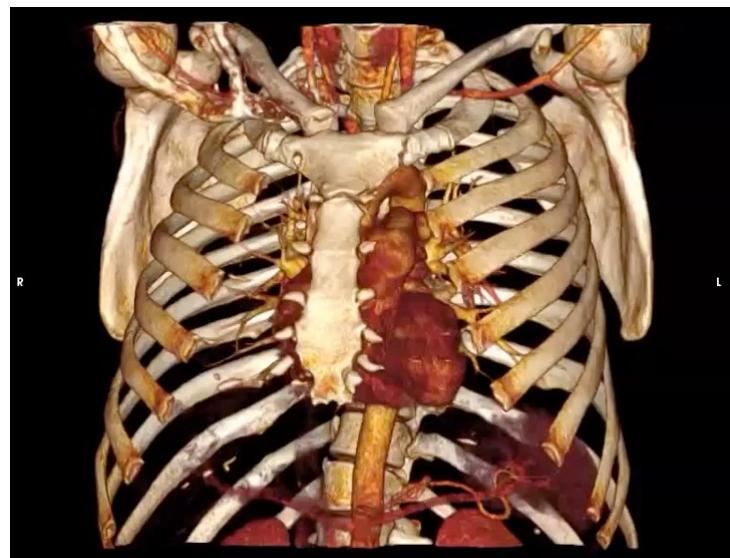
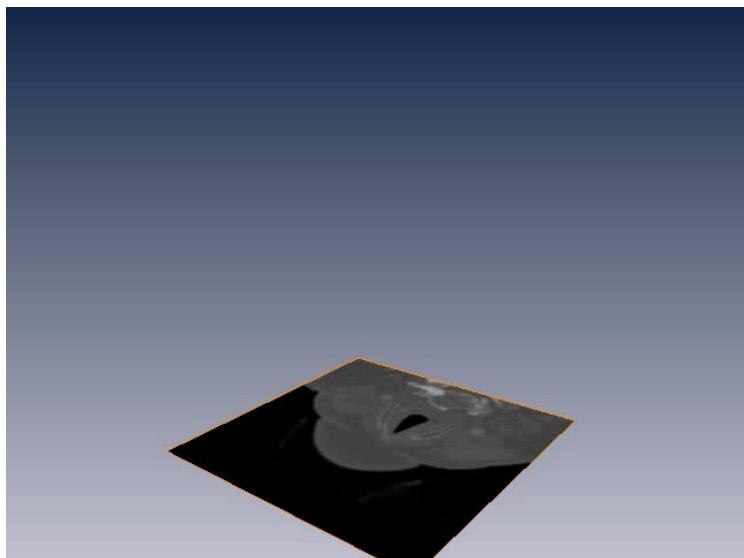
# Inpainting via deep image priors



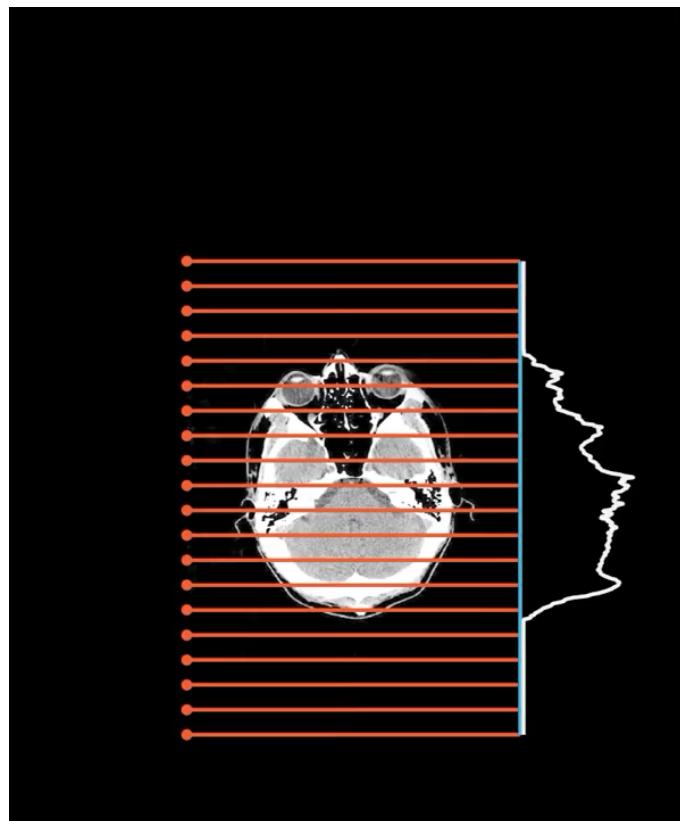
## Deep Learning for Image Reconstruction

Daniel Rueckert  
Department of Computing  
Imperial College London, UK

# Medical Imaging: X-ray computed tomography (CT)



# Medical Imaging: X-ray computed tomography (CT)

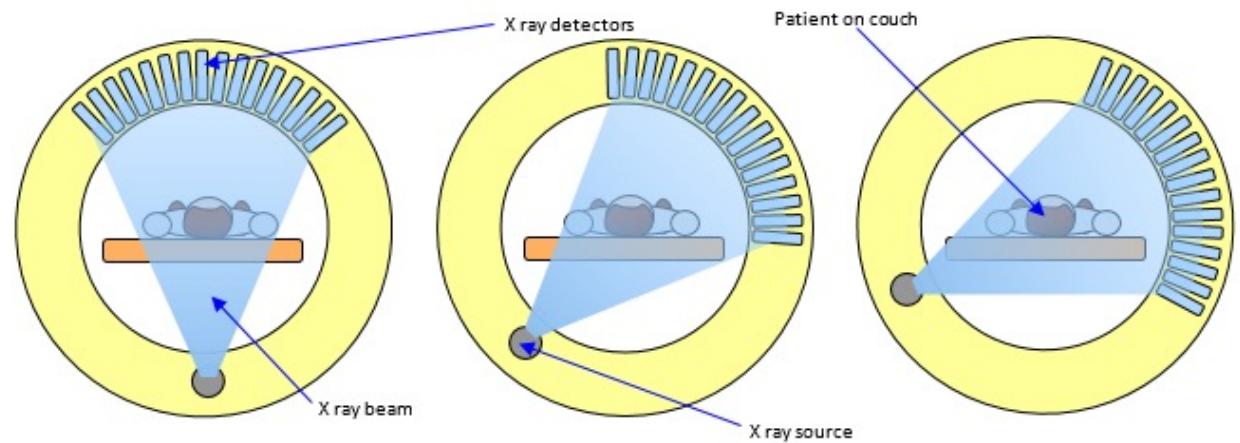


source: [https://youtu.be/q7Rt\\_OY\\_7tU](https://youtu.be/q7Rt_OY_7tU)

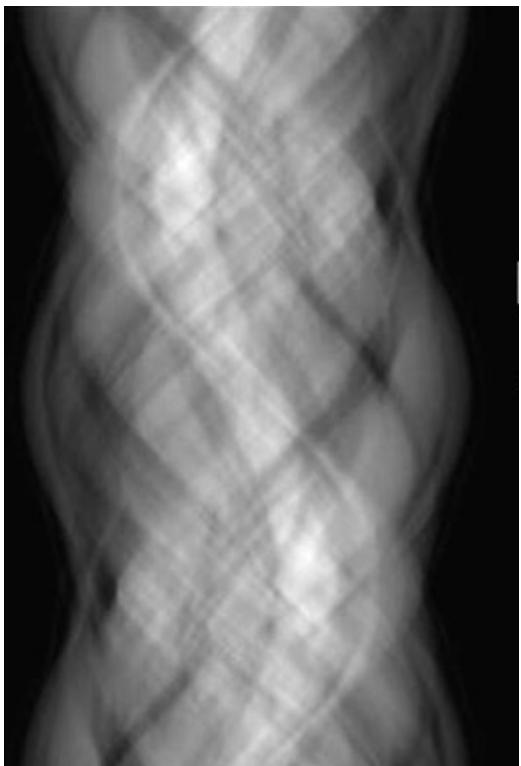
# Medical Imaging: X-ray computed tomography (CT)

- CT imaging

- high contrast
- high spatial resolution
- fast acquisition
- but ionising radiation



# Medical Imaging: X-ray computed tomography (CT)



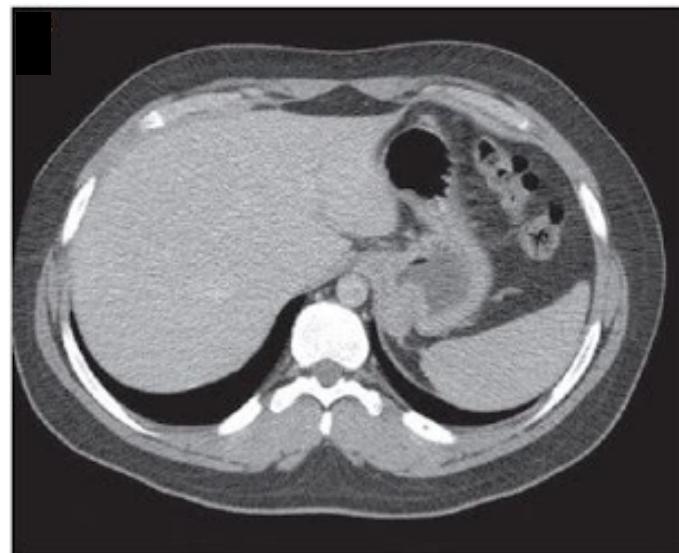
Sinogram

Forward problem  
(Radon transform)



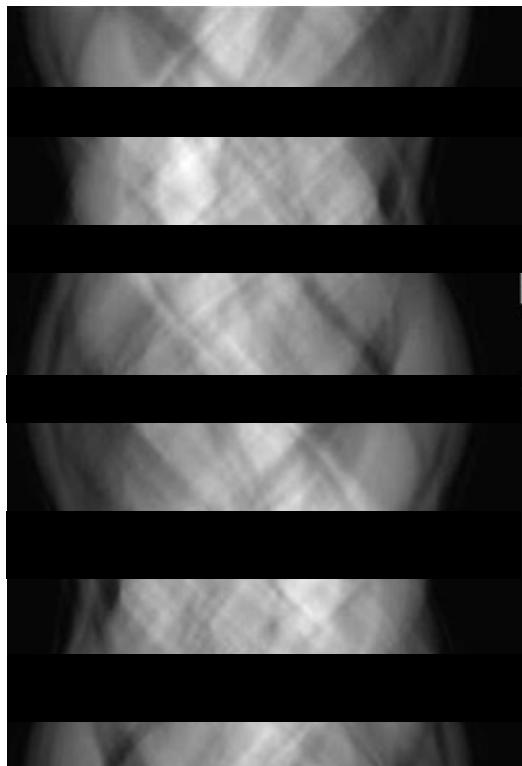
Inverse problem:

Fully sampled  
reconstruction



CT image

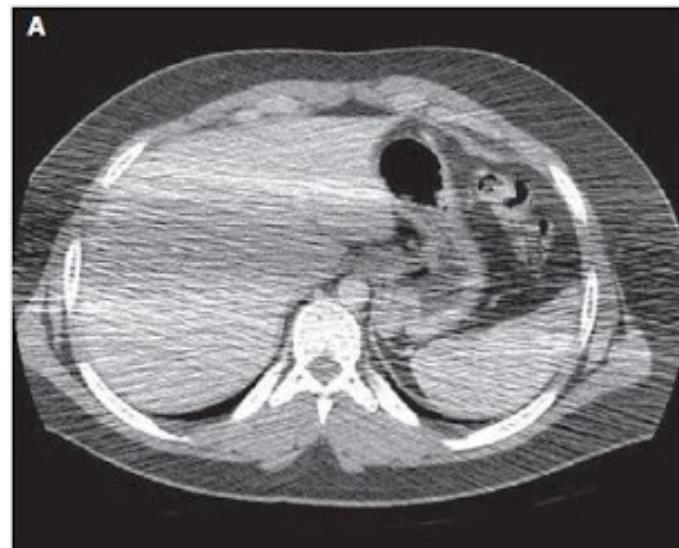
# Medical Imaging: X-ray computed tomography (CT)



Sinogram

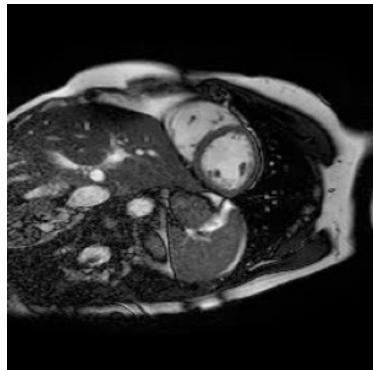


Under-sampled  
reconstruction

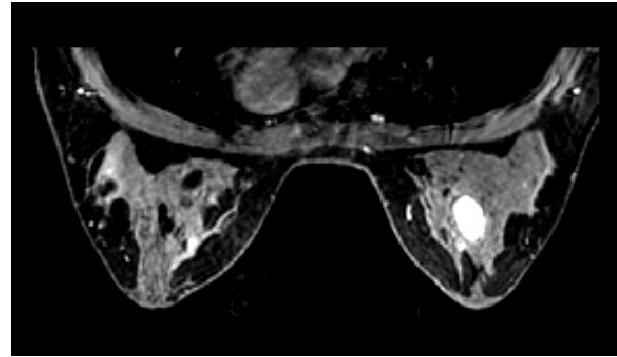


CT image

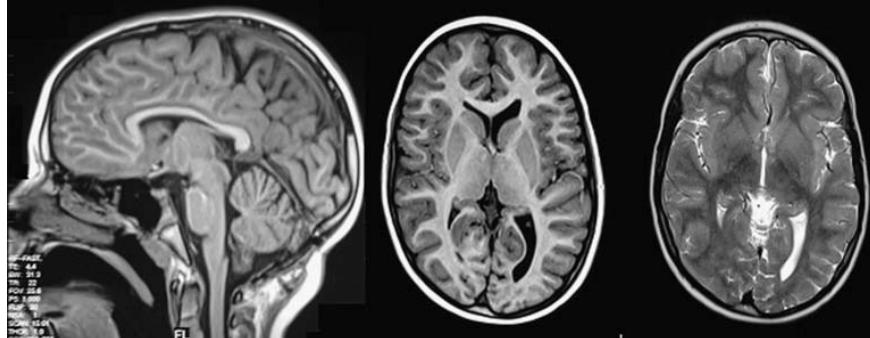
# Medical Imaging: Magnetic Resonance (MR)



heart



cancer



brain



lung

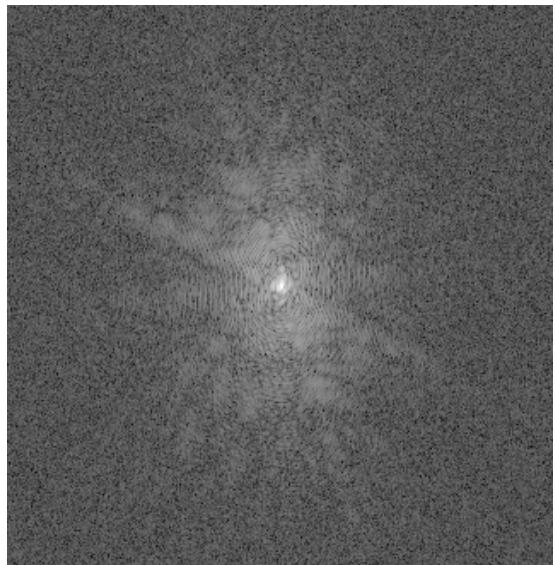


Whole body

# Medical Imaging: Magnetic Resonance (MR)

- MR imaging:
    - high contrast
    - high spatial resolution
    - no ionising radiation
    - but slow acquisition process
  - Slow acquisition is
    - ok for static objects (e.g. brain, bones, etc)
    - problematic for moving objects (e.g. heart, liver, fetus)
  - Options for MR acquisition:
    - real-time MR: fast, but 2D and relatively poor image quality
    - gated MRI fine for period motion, e.g. respiration or cardiac motion but requires gating (ECG or navigators) leading to long acquisition times (30-90 min).
- 

# Medical Imaging: Magnetic Resonance (MR)



K-Space

Forward problem  
(Fourier transform)



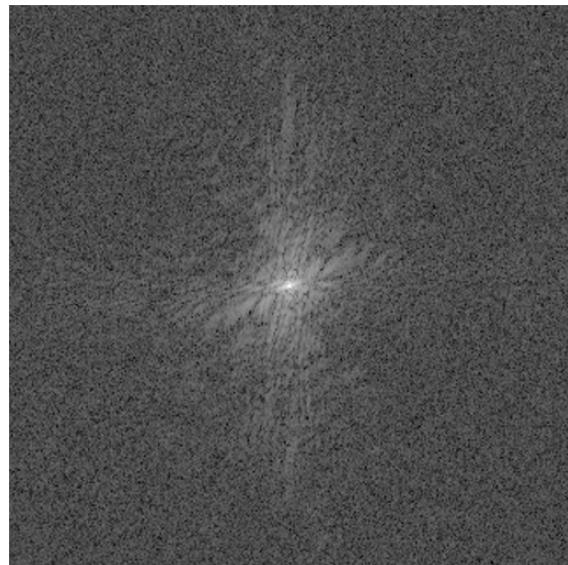
Inverse problem:

Fully sampled  
reconstruction



MR image

# Medical Imaging: Magnetic Resonance (MR)



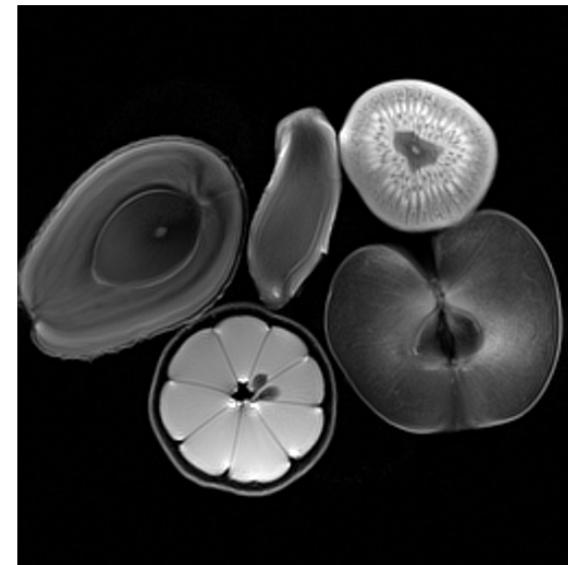
K-Space

Forward problem  
(Fourier transform)



Inverse problem:

Fully sampled  
reconstruction

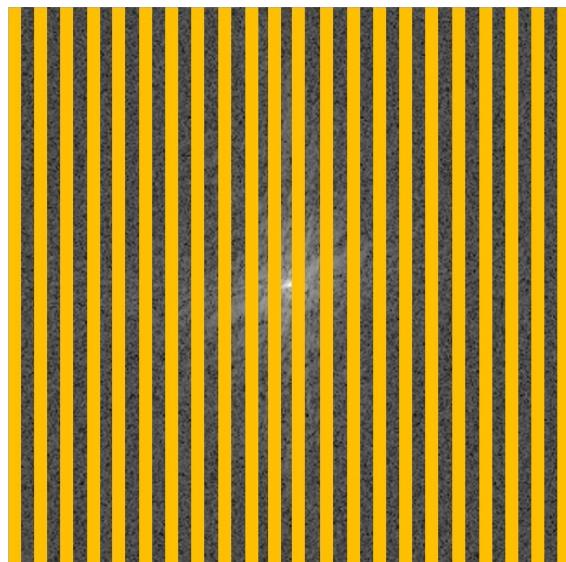


MR image

Recover an image  $x \in \mathbb{K}^{N_x}$  from a set of observations  $y \in \mathbb{K}^{N_y}$  which are corrupted by noise  $n \in \mathbb{K}^{N_y}$ , and  $A: \mathbb{K}^{N_x} \rightarrow \mathbb{K}^{N_y}$  is a linear operator

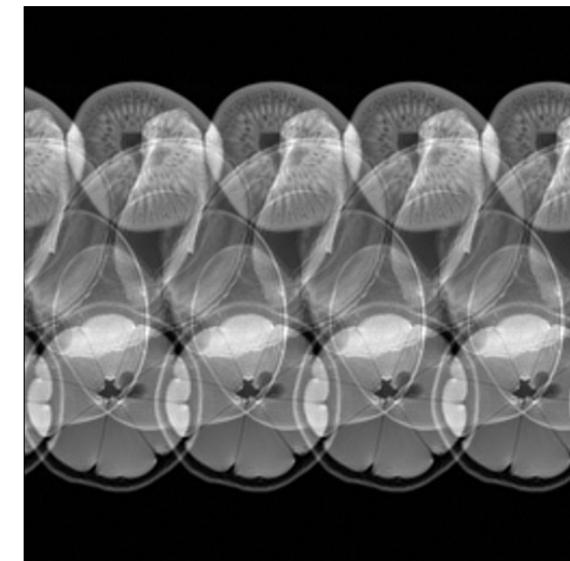
$$y = Ax + n$$

# Medical Imaging: Magnetic Resonance (MR)



Undersampled K-Space

Inverse problem

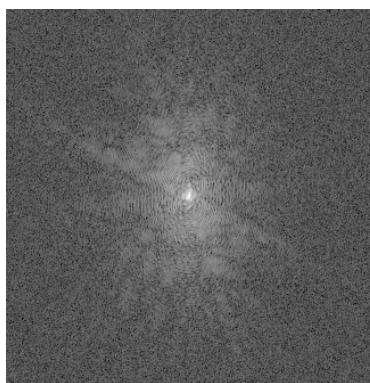


Aliased MR image

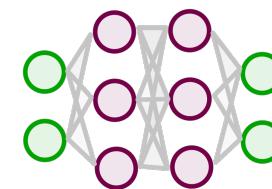
Recover an image  $x \in \mathbb{K}^{N_x}$  from a set of observations  $y \in \mathbb{K}^{N_y}$  which are corrupted by noise  $n \in \mathbb{K}^{N_y}$ , and  $A: \mathbb{K}^{N_x} \rightarrow \mathbb{K}^{N_y}$  is a linear operator

$$y = Ax + n$$

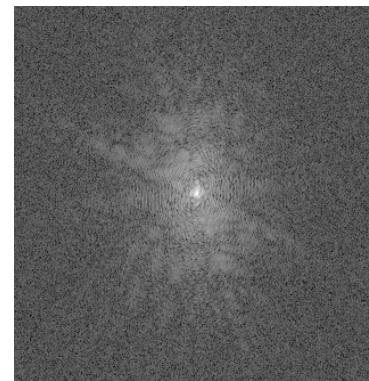
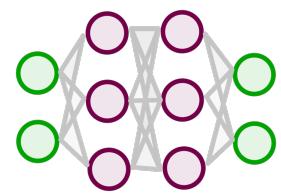
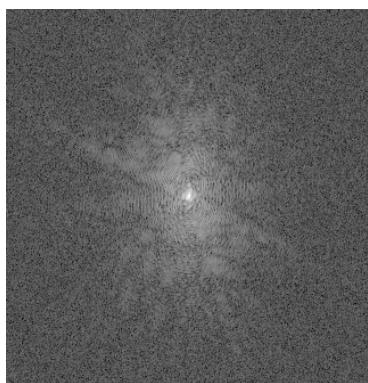
## Possible reconstruction approaches: Image domain



$$\xrightarrow{\mathcal{F}}$$



## Possible reconstruction approaches: k-space domain

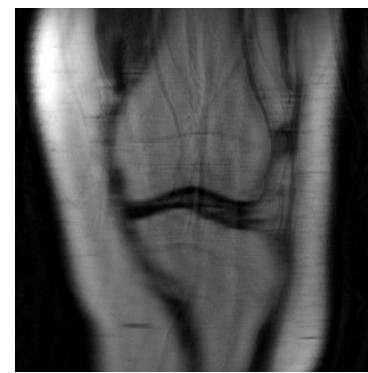
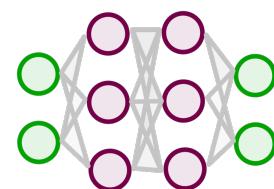
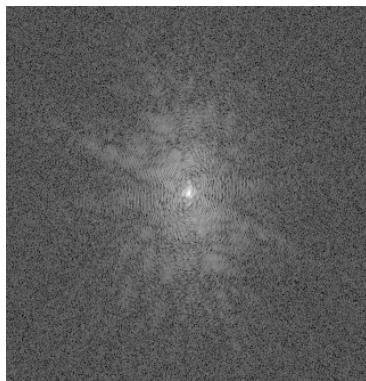


$$\xrightarrow{\mathcal{F}}$$



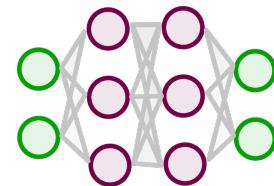
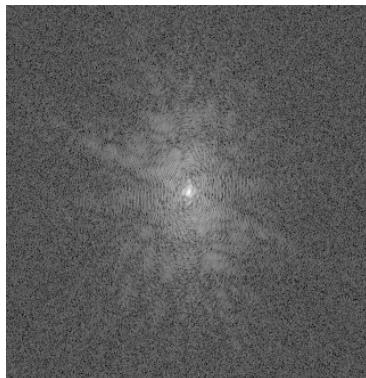
## Training Phase

- Optimize  $\theta$  of the model  $f_\theta$  using the training set



## Training Phase

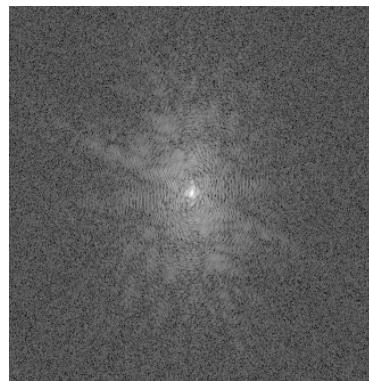
- Optimize  $\theta$  of the model  $f_\theta$  using the training set



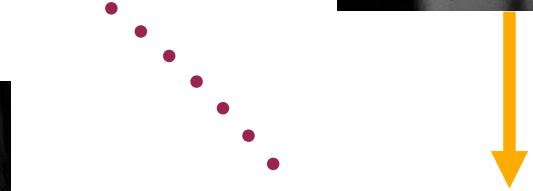
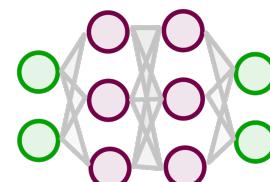
**Cost Function**

## Training Phase

- Optimize  $\theta$  of the model  $f_\theta$  using the training set



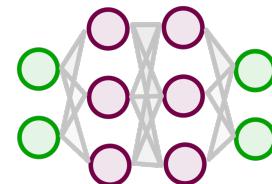
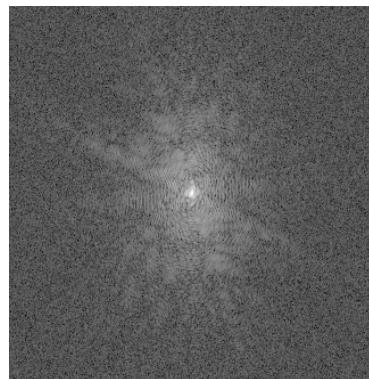
Error Propagation



Cost Function

## Training Phase

- Optimize  $\theta$  of the model  $f_\theta$  using the training set



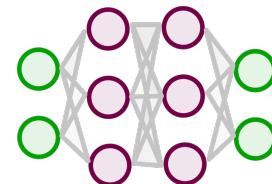
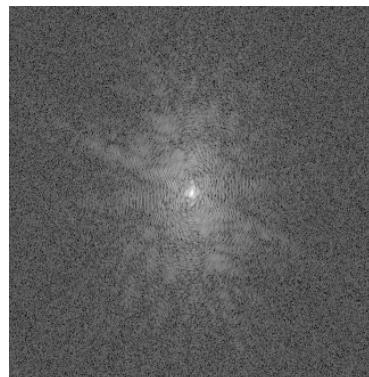
Error Propagation



Cost Function

## Training Phase

- Optimize  $\theta$  of the model  $f_\theta$  using the training set



Error Propagation



Cost Function

## Cost Function

**Synonym: Loss function, error function, similarity measure**

- Quantifies how well the model prediction matches targets
- Needs to be selected according to the underlying task
- Is optimized during training → needs to be differentiable!

## Cost Function

**Synonym: Loss function, error function, similarity measure**

- Quantifies how well the model prediction matches targets
- Needs to be selected according to the underlying task
- Is optimized during training → needs to be differentiable!

**Example for image reconstruction: Mean Squared Error**

$$MSE(y, \hat{y}) = \frac{1}{N_S} \sum_{n=1}^{N_S} \| y_n - \hat{y}_n \|_2^2$$

## Training Phase

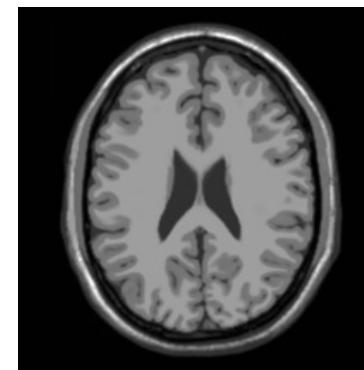
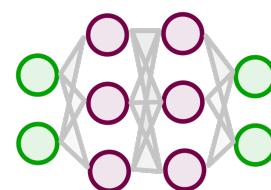
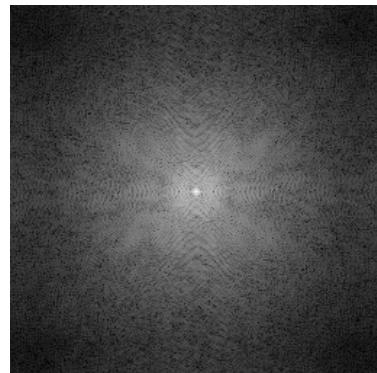
**Learning model:**  $\hat{y}_n = f_\theta(x_n)$

- The parameters  $\theta$  of the mapping function  $f_\theta$  are optimized under a cost function
- The cost function quantifies how well  $\hat{y}_n = f_\theta(x_n)$  is predicted given  $x_n$
- The parameters  $\theta$  by minimizing the cost function  $\mathcal{L}$  with learning rate  $\tau$

$$\theta^{(k+1)} = \theta^{(k)} - \tau \frac{\partial \mathcal{L}(\theta, x)}{\partial \theta} \Big|_{\theta=\theta^{(k)}}$$

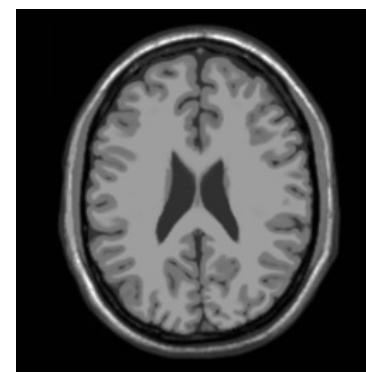
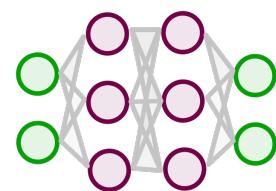
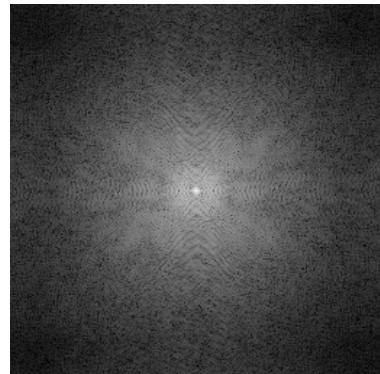
## Testing Phase (Inference)

- Apply  $f_\theta$  using the optimized  $\theta$  to the test set



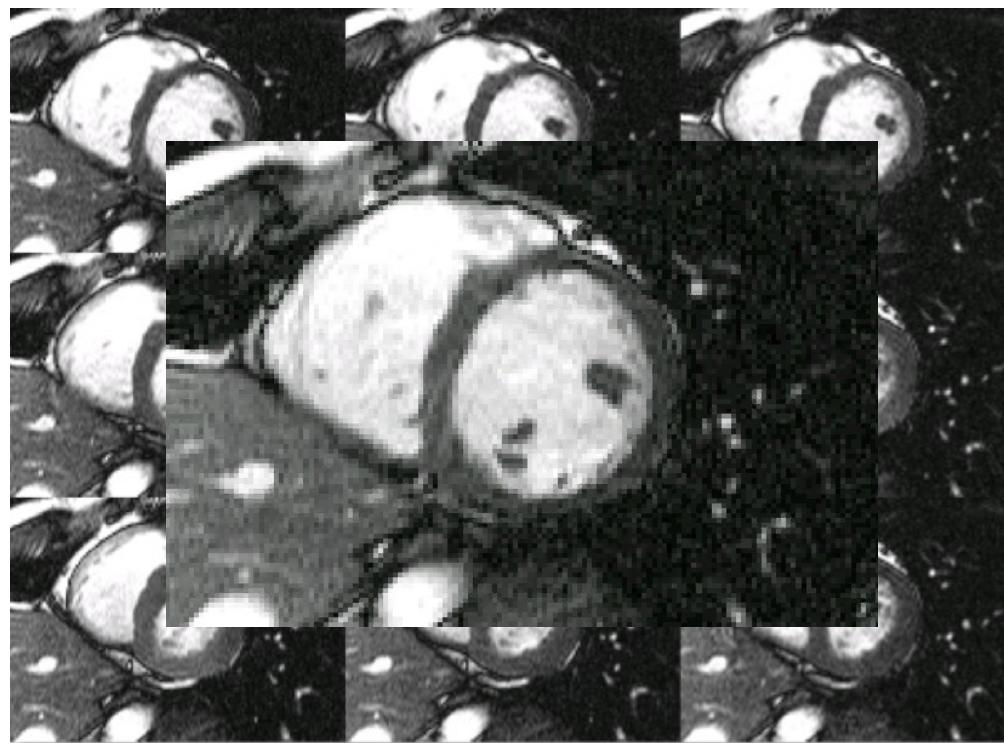
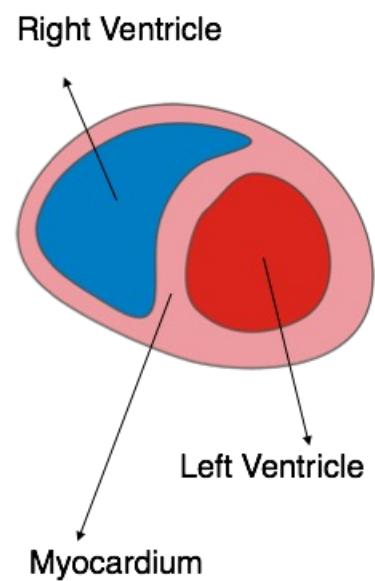
## Testing Phase (Inference)

- Apply  $f_{\theta}$  using the optimized  $\theta$  to the test set

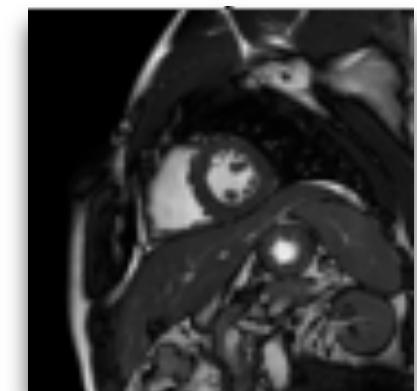
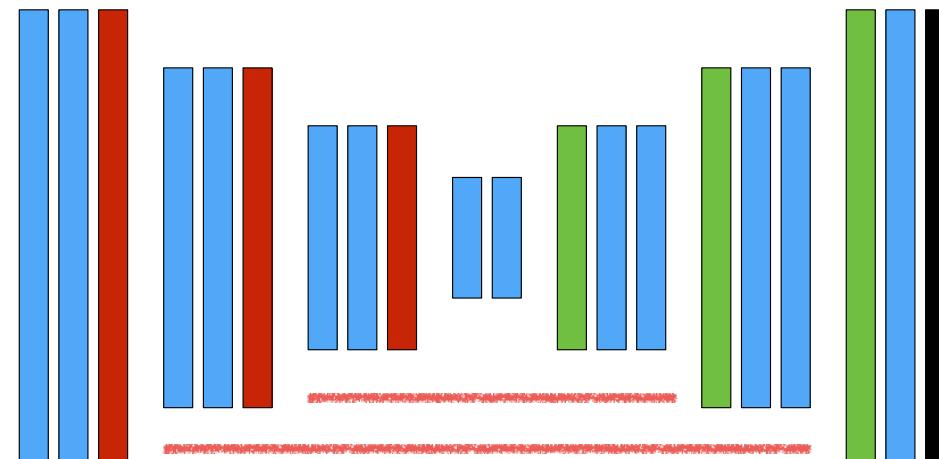
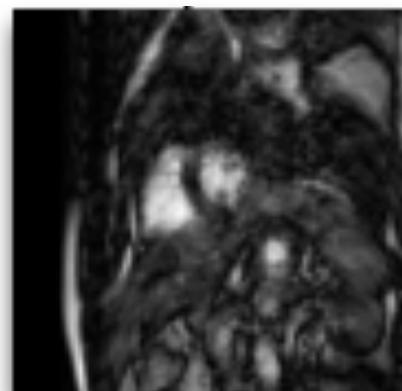


Generalization: Ability to correctly predict unseen examples

# Magnetic Resonance (MR)



# Deep learning for image reconstruction



■ Convolution + RELU

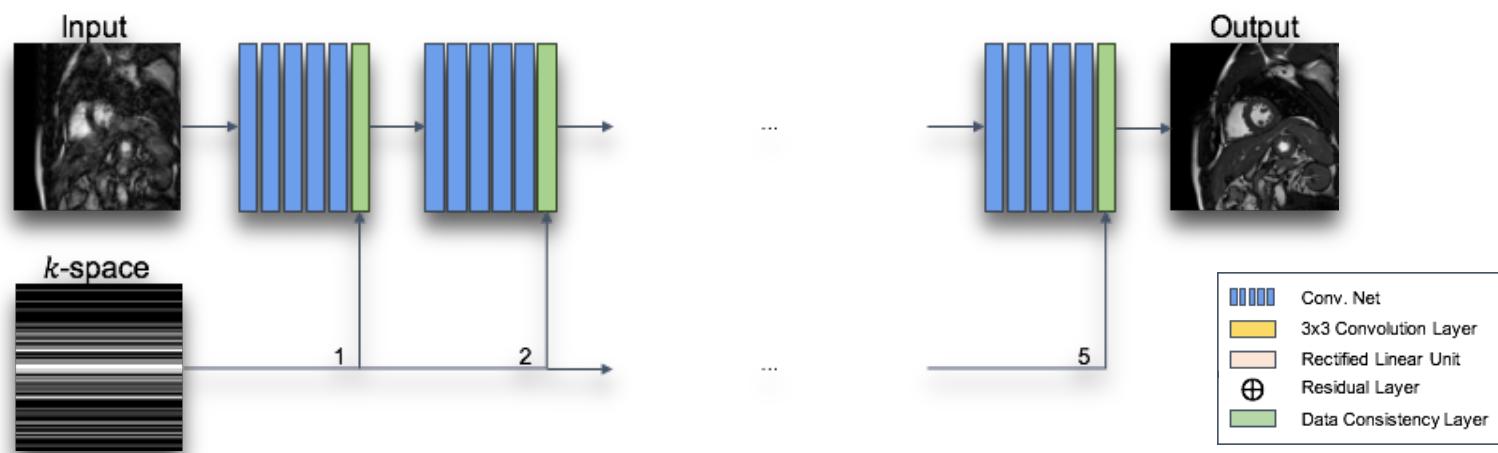
■ Max pooling

■ Transposed convolution

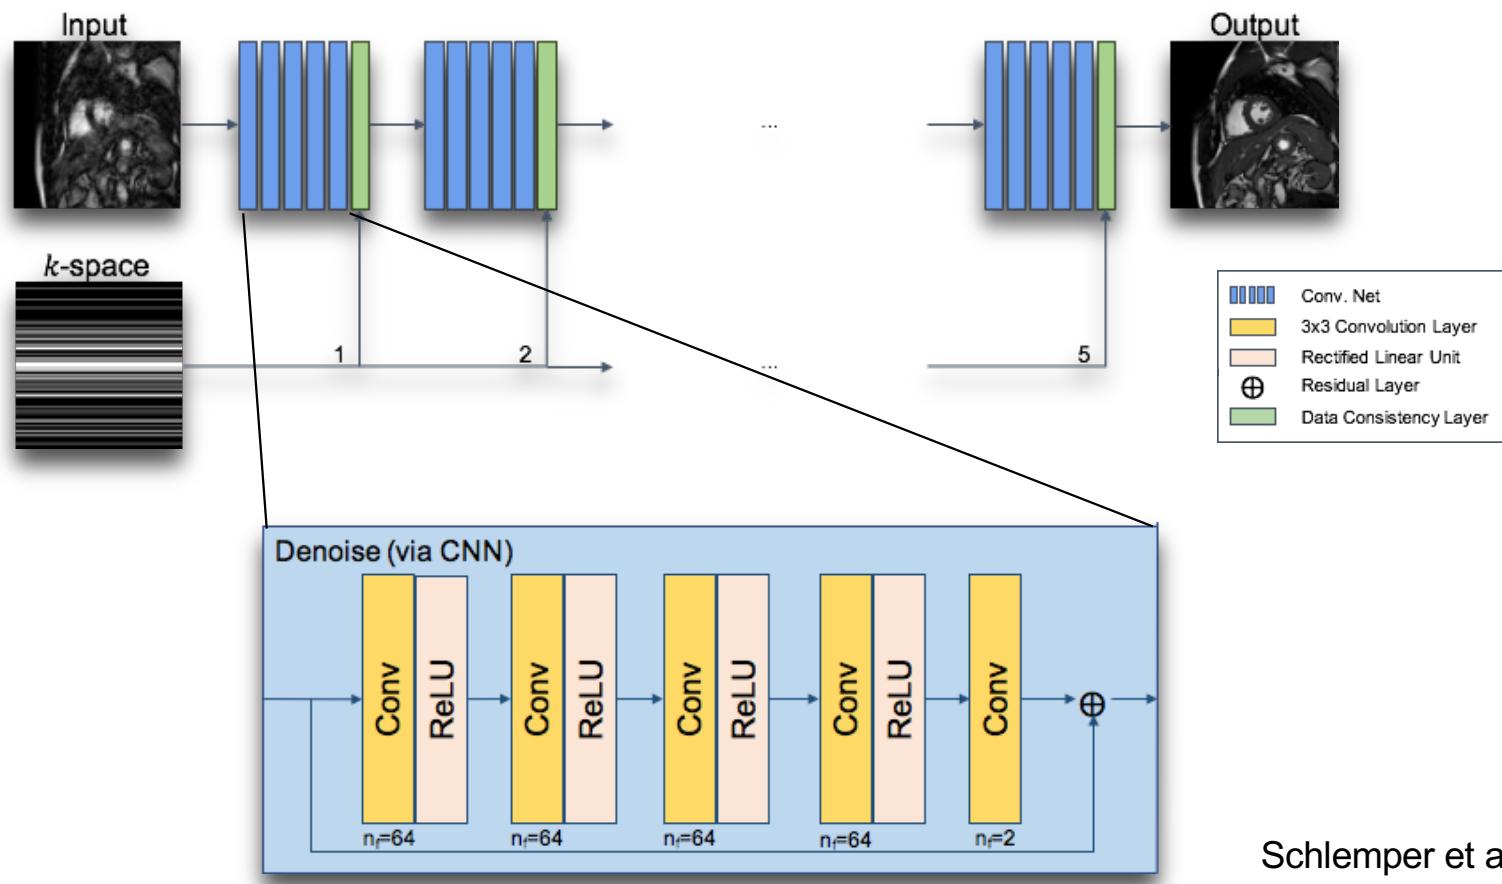
■ Softmax

— Skip layers

# Deep learning for image reconstruction

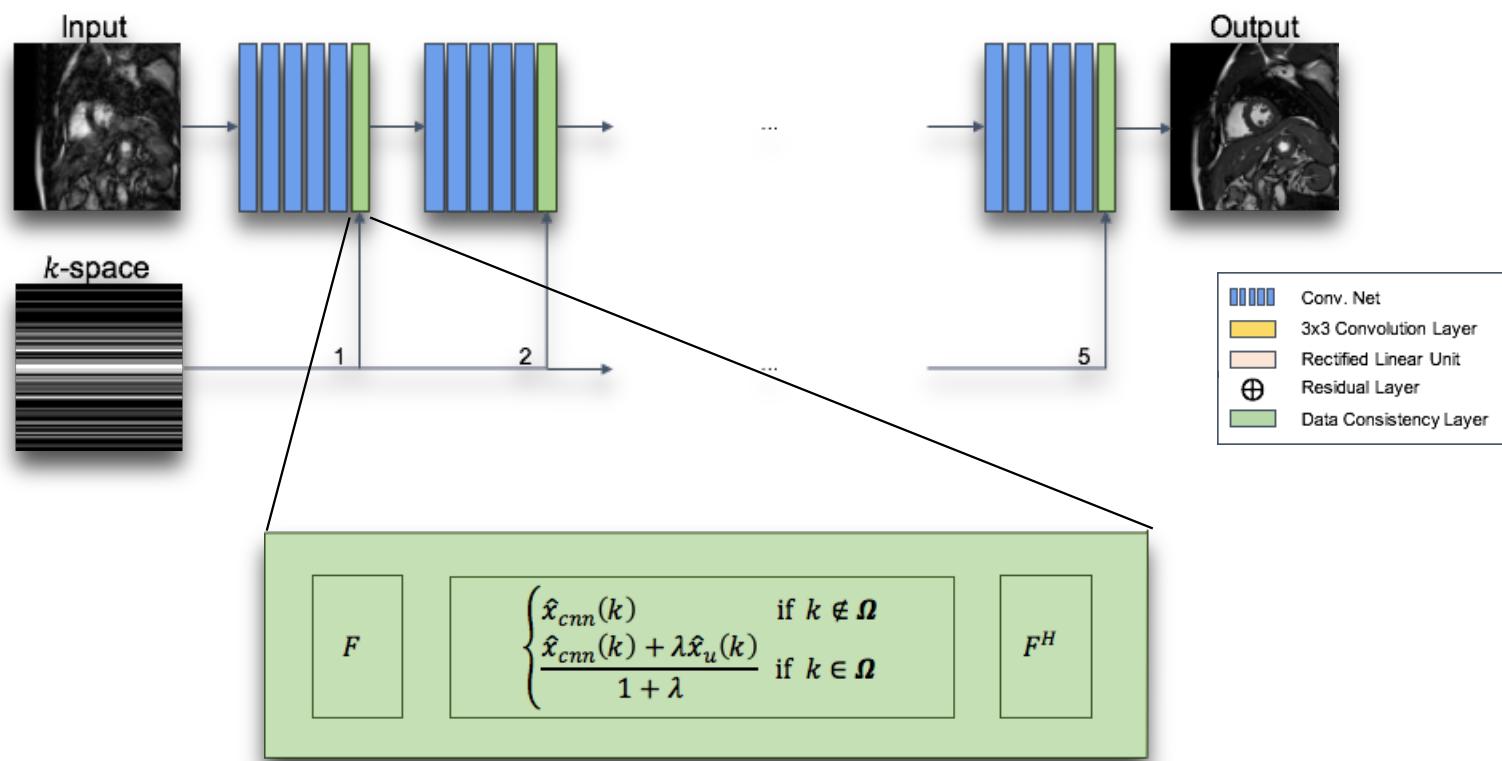


# Deep learning for image reconstruction

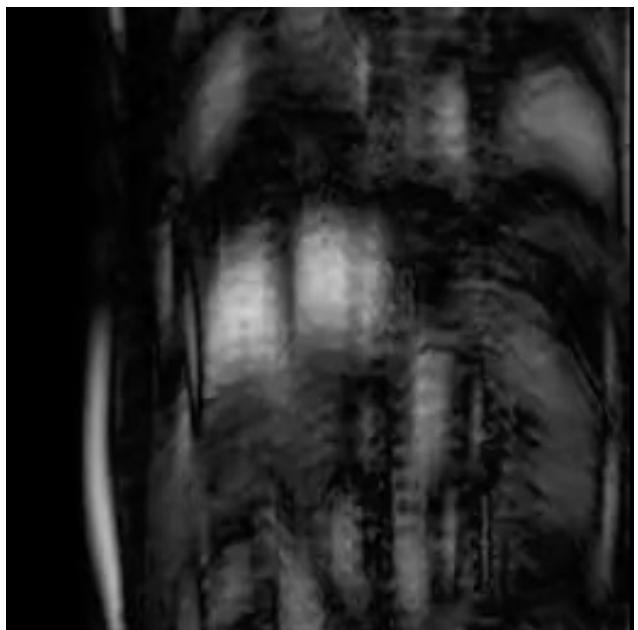


Schlemper et al. IEEE TMI 2017

# Deep learning for image reconstruction



## Magnitude reconstruction (6-fold)



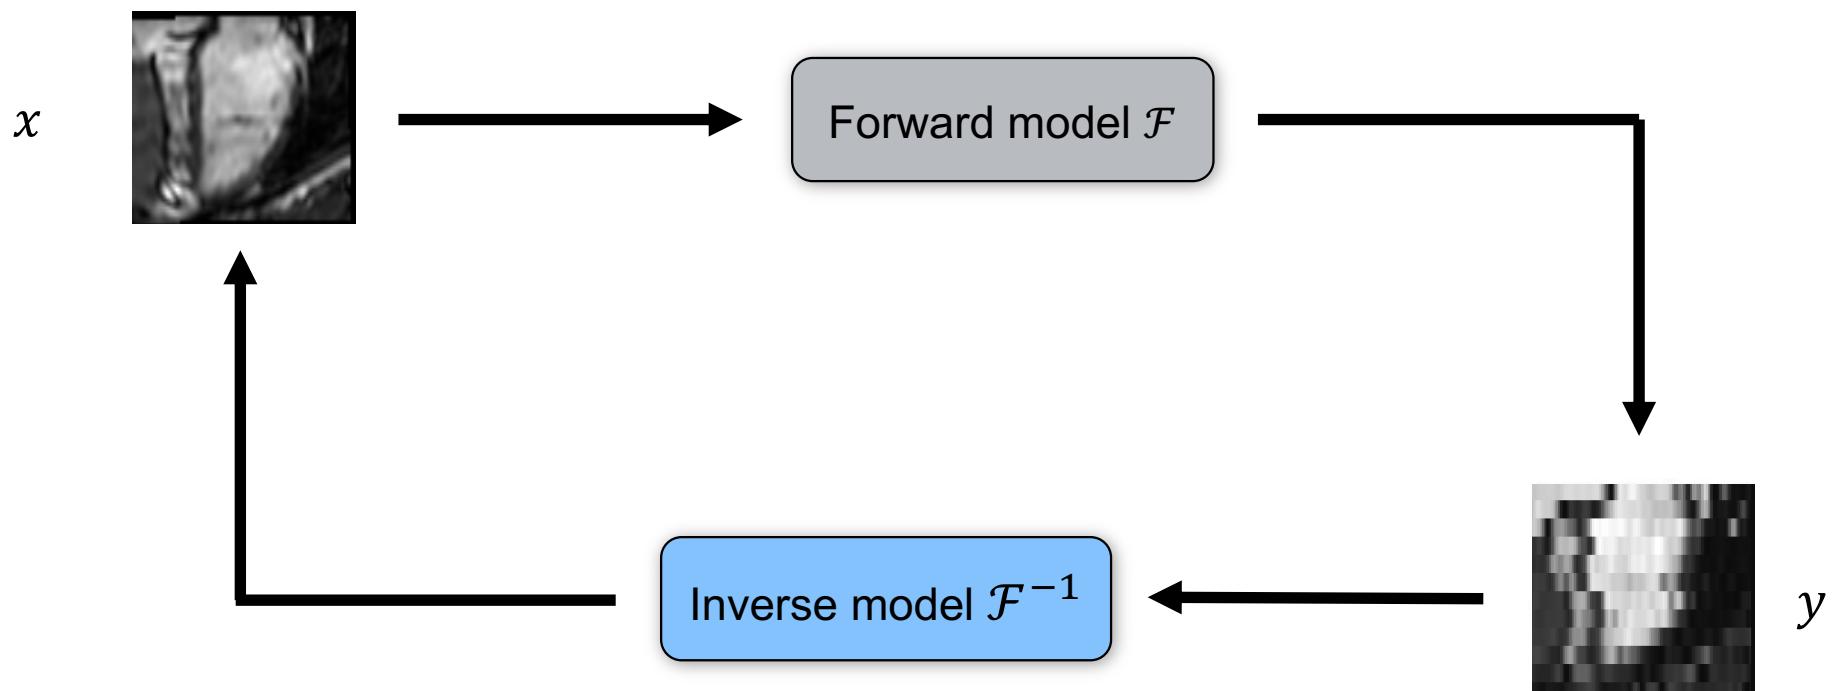
(a) 6x Undersampled

## Magnitude reconstruction (11-fold)

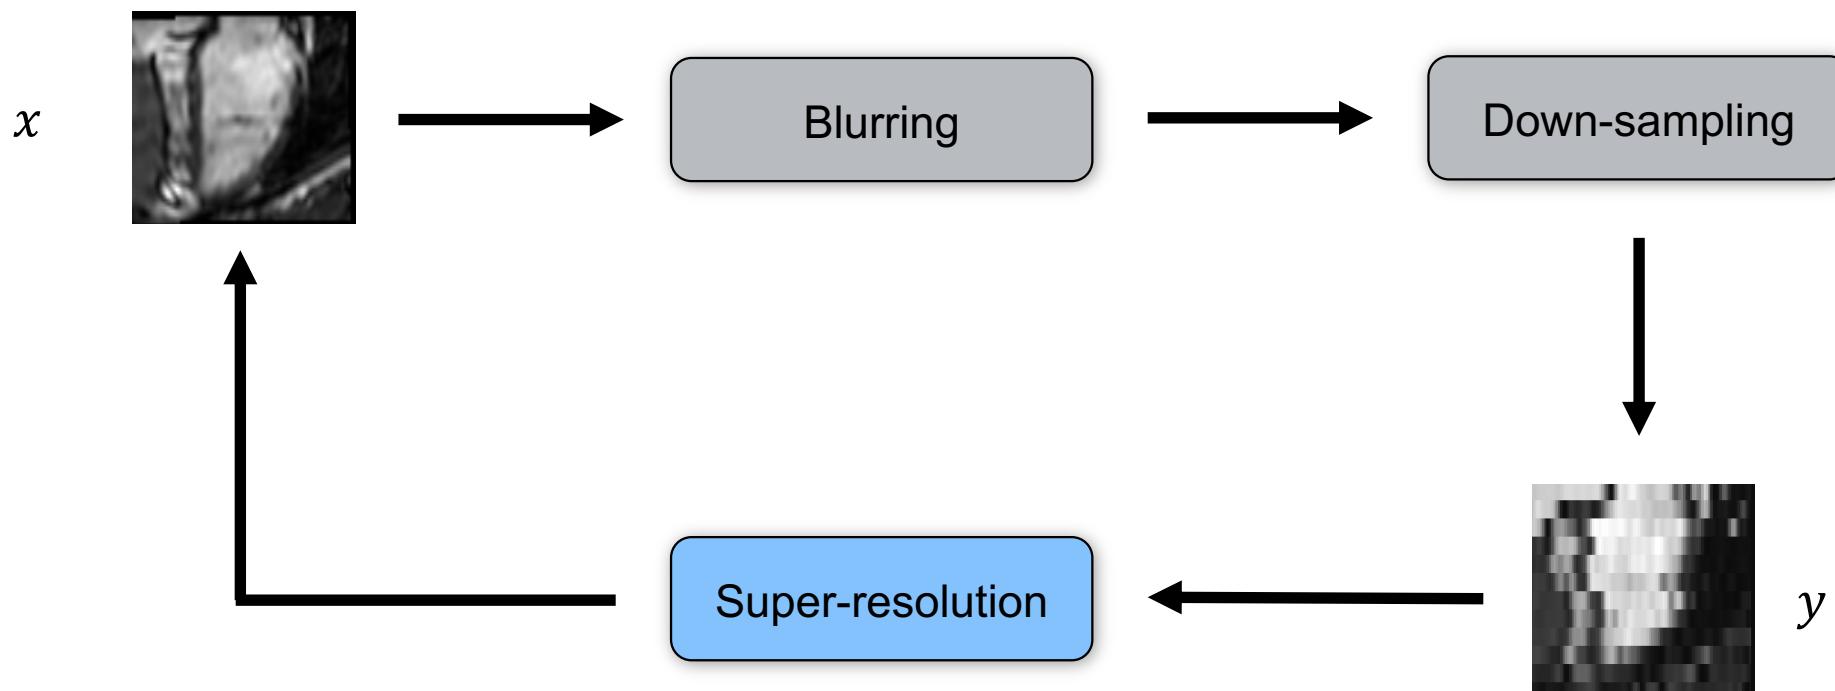


(a) 11x Undersampled

## Forward model/Inverse model

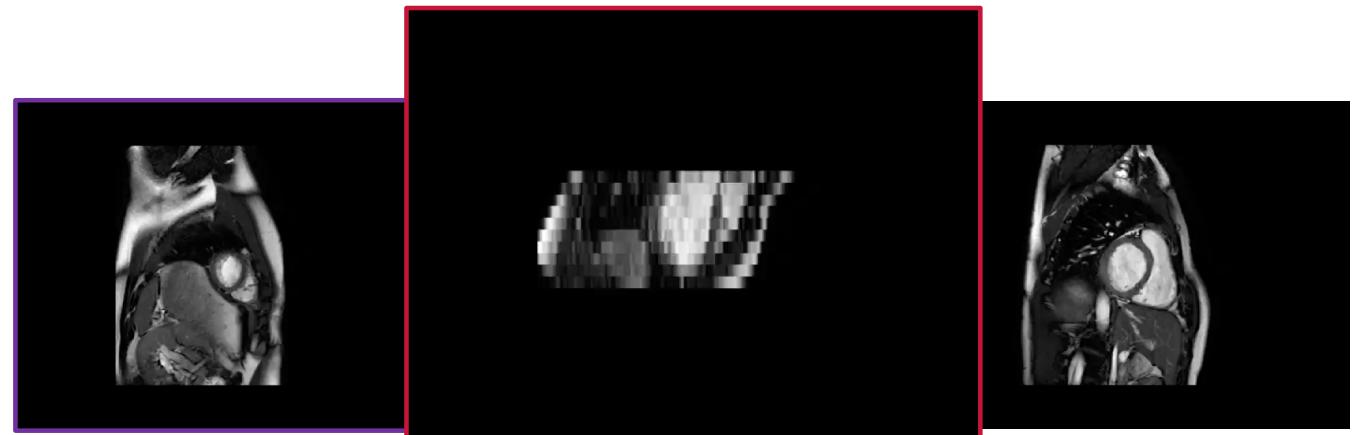
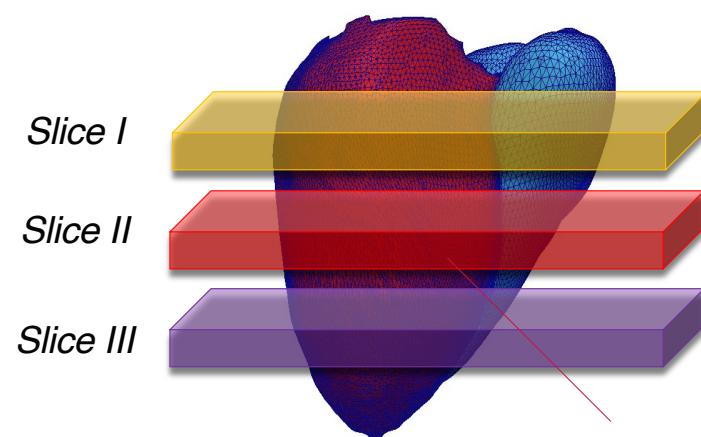


## Forward model/Inverse model

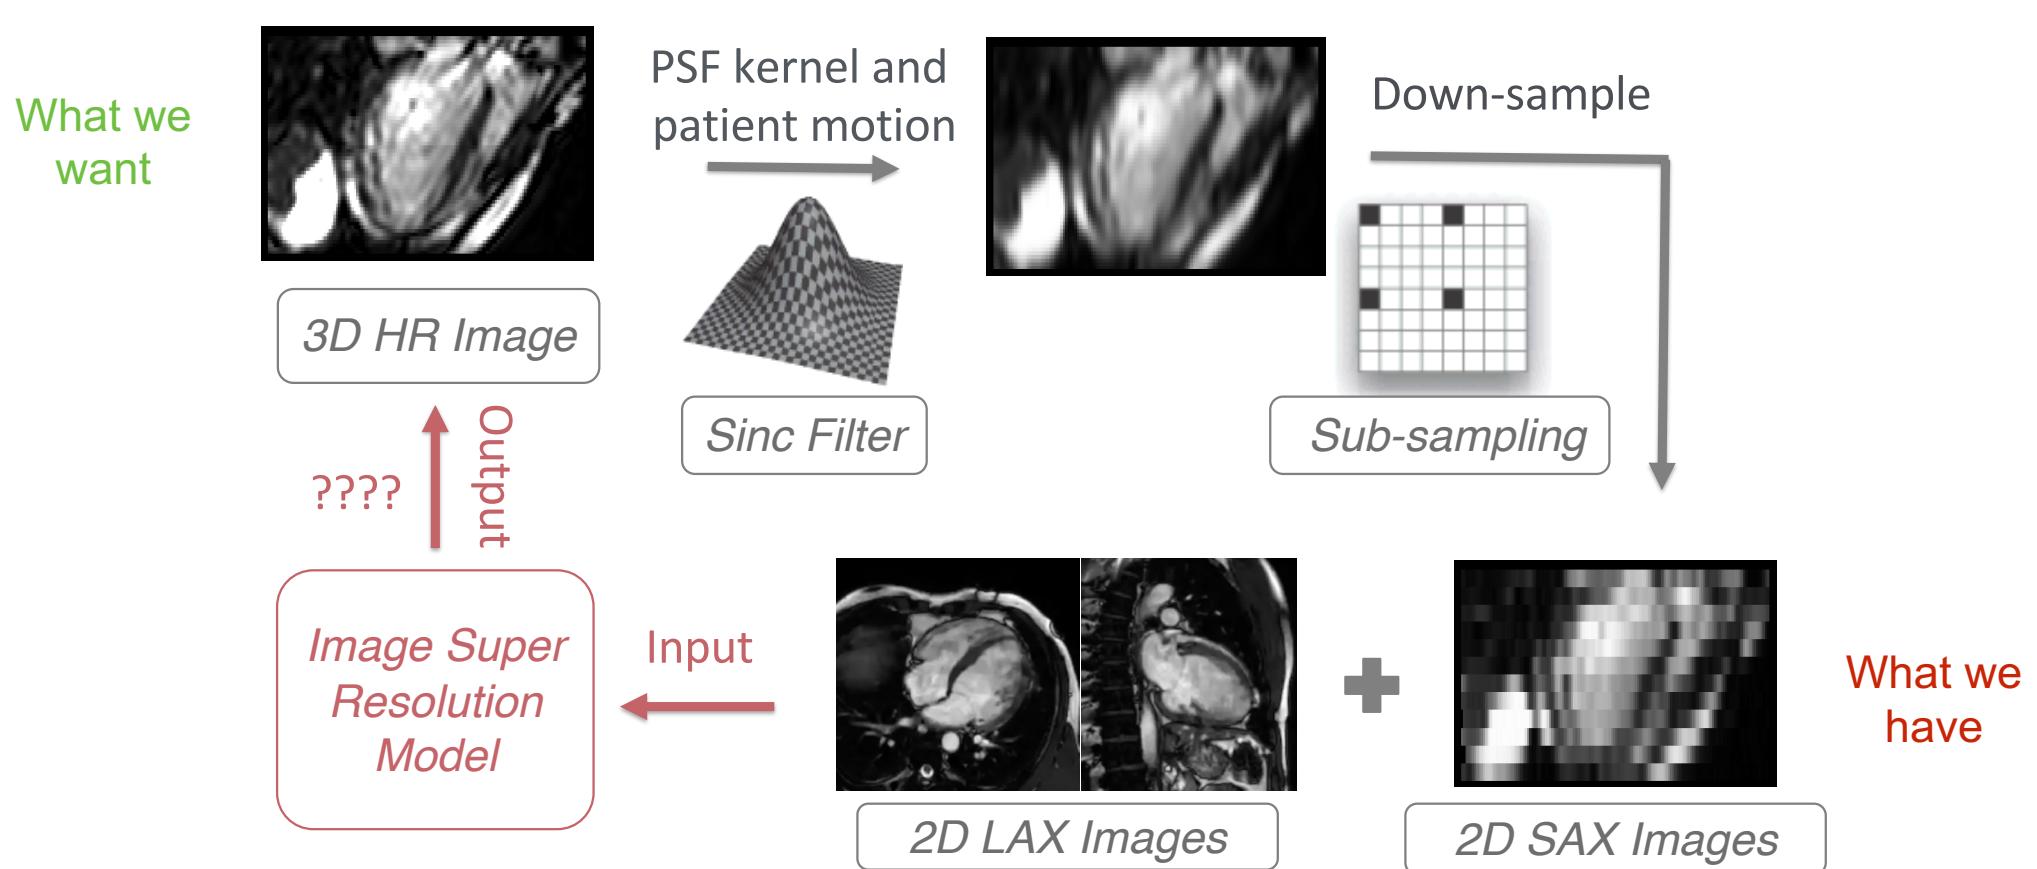


# AI-enabled image super-resolution

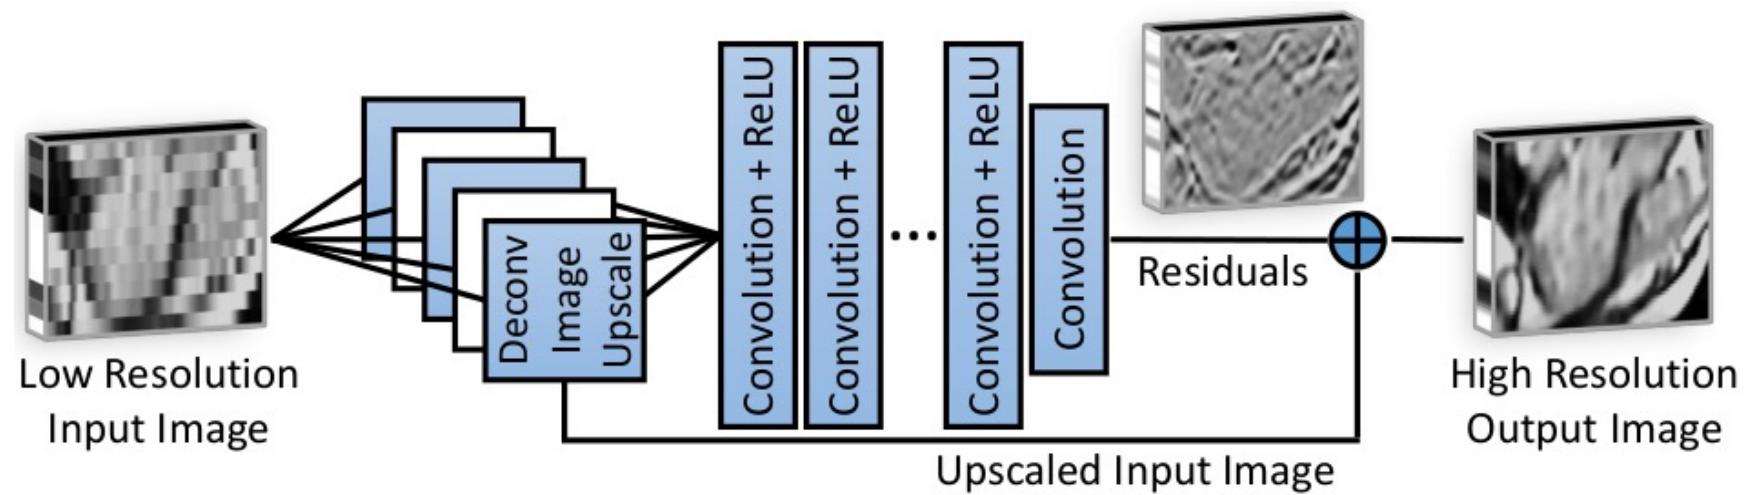
- Acquisition of cardiac MRI typically consists of 2D multi-slice data due to
  - constraints on SNR
  - breath-hold time
  - total acquisition time
- This leads to thick slice data (thickness 8-10 mm per slice)



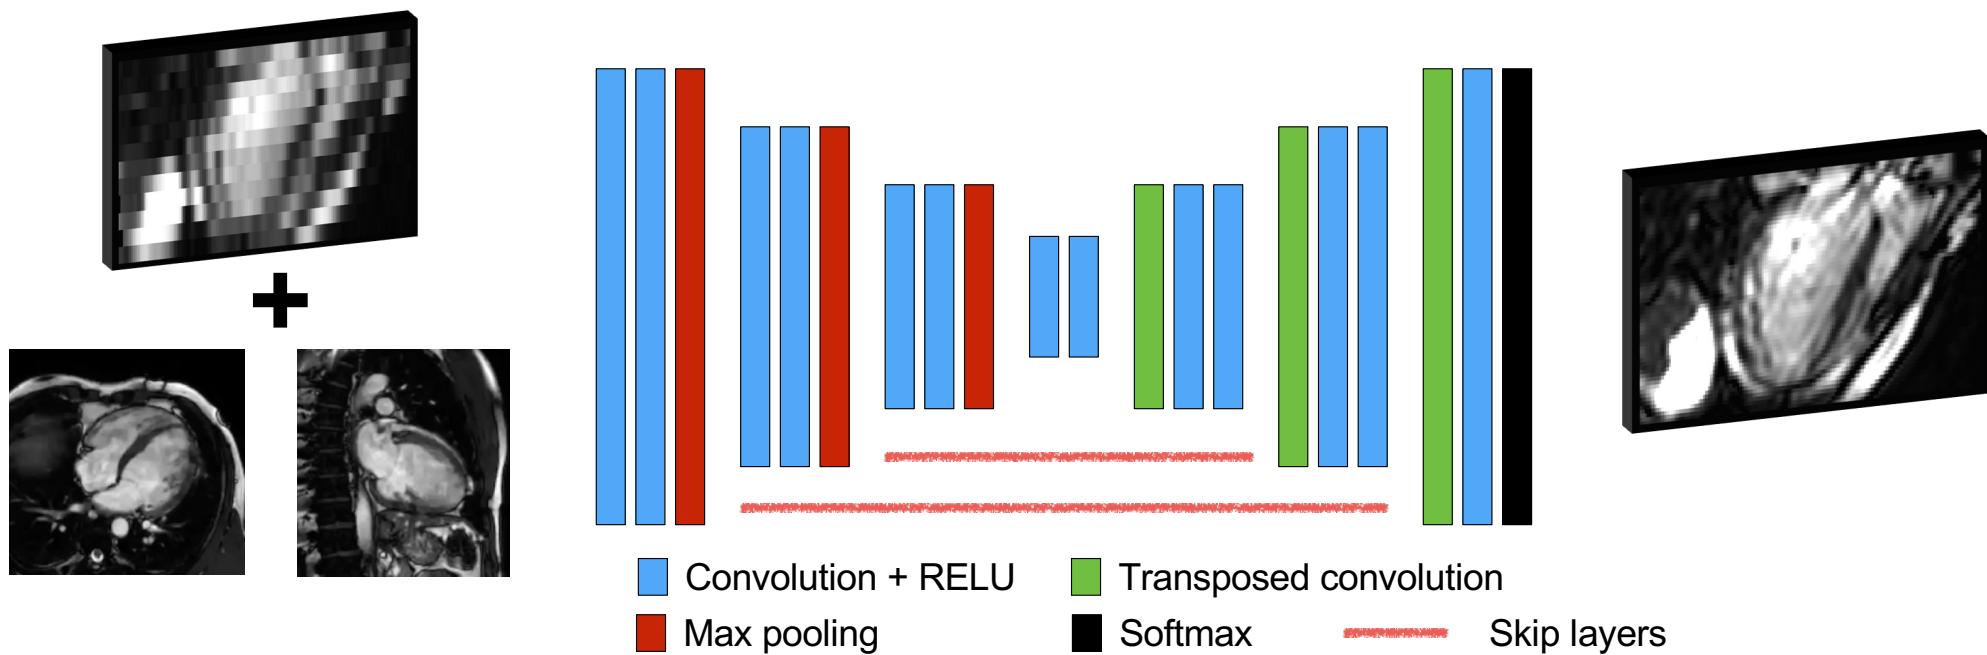
# Deep learning for image super-resolution



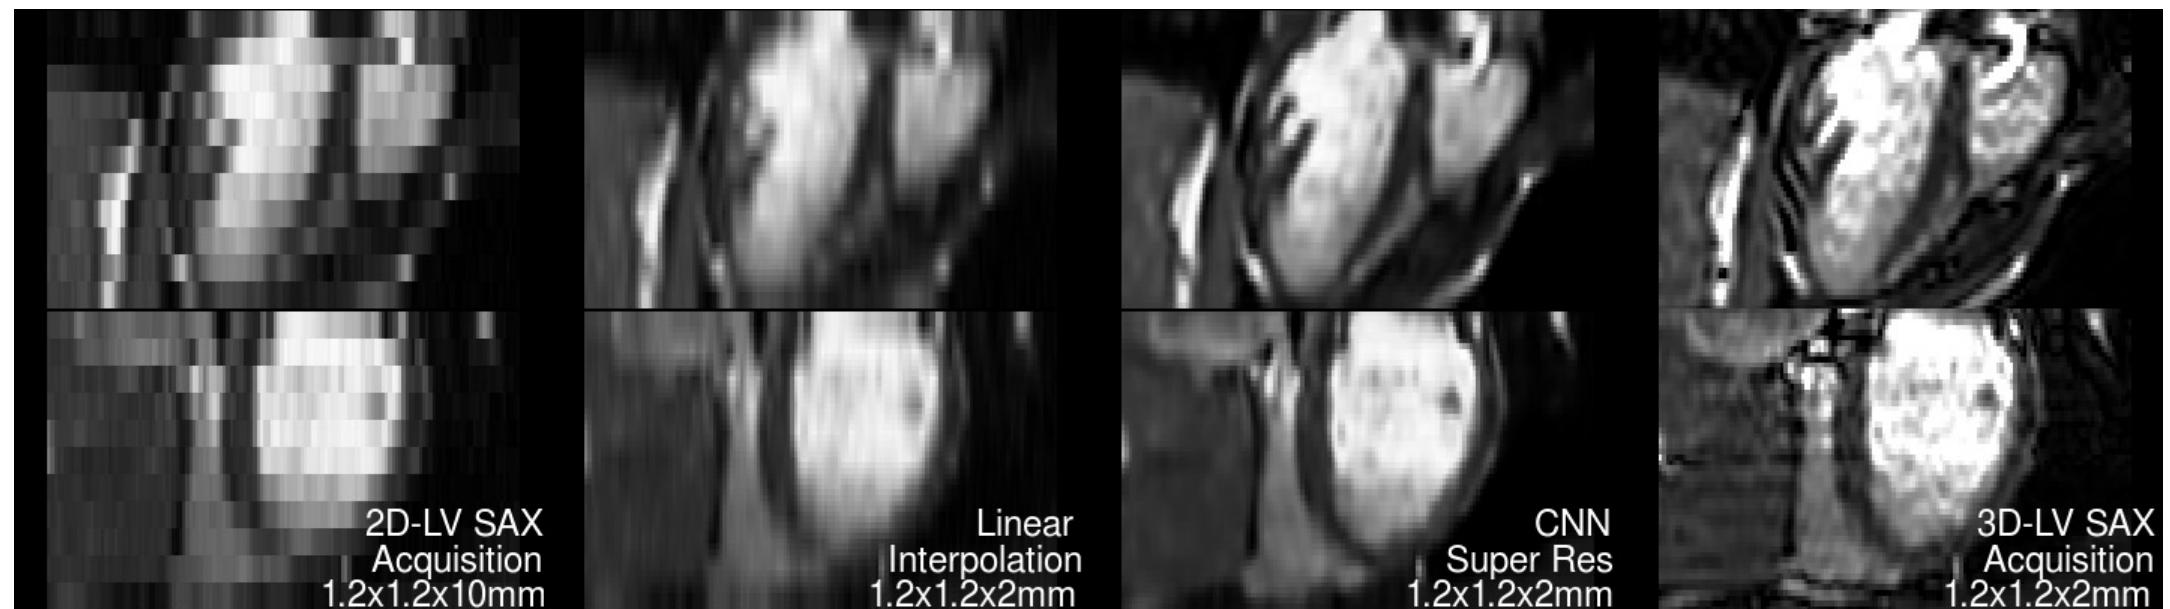
# Deep learning for image super-resolution



# AI-enabled image super-resolution



## AI-enabled image super-resolution



O. Oktay et al. IEEE TMI 2018

That's all for now...

## References

- Deep Learning for Image Super-resolution: A Survey
  - <https://arxiv.org/abs/1902.06068>
- Deep Learning Techniques for Inverse Problems in Imaging
  - <https://arxiv.org/abs/2005.06001>